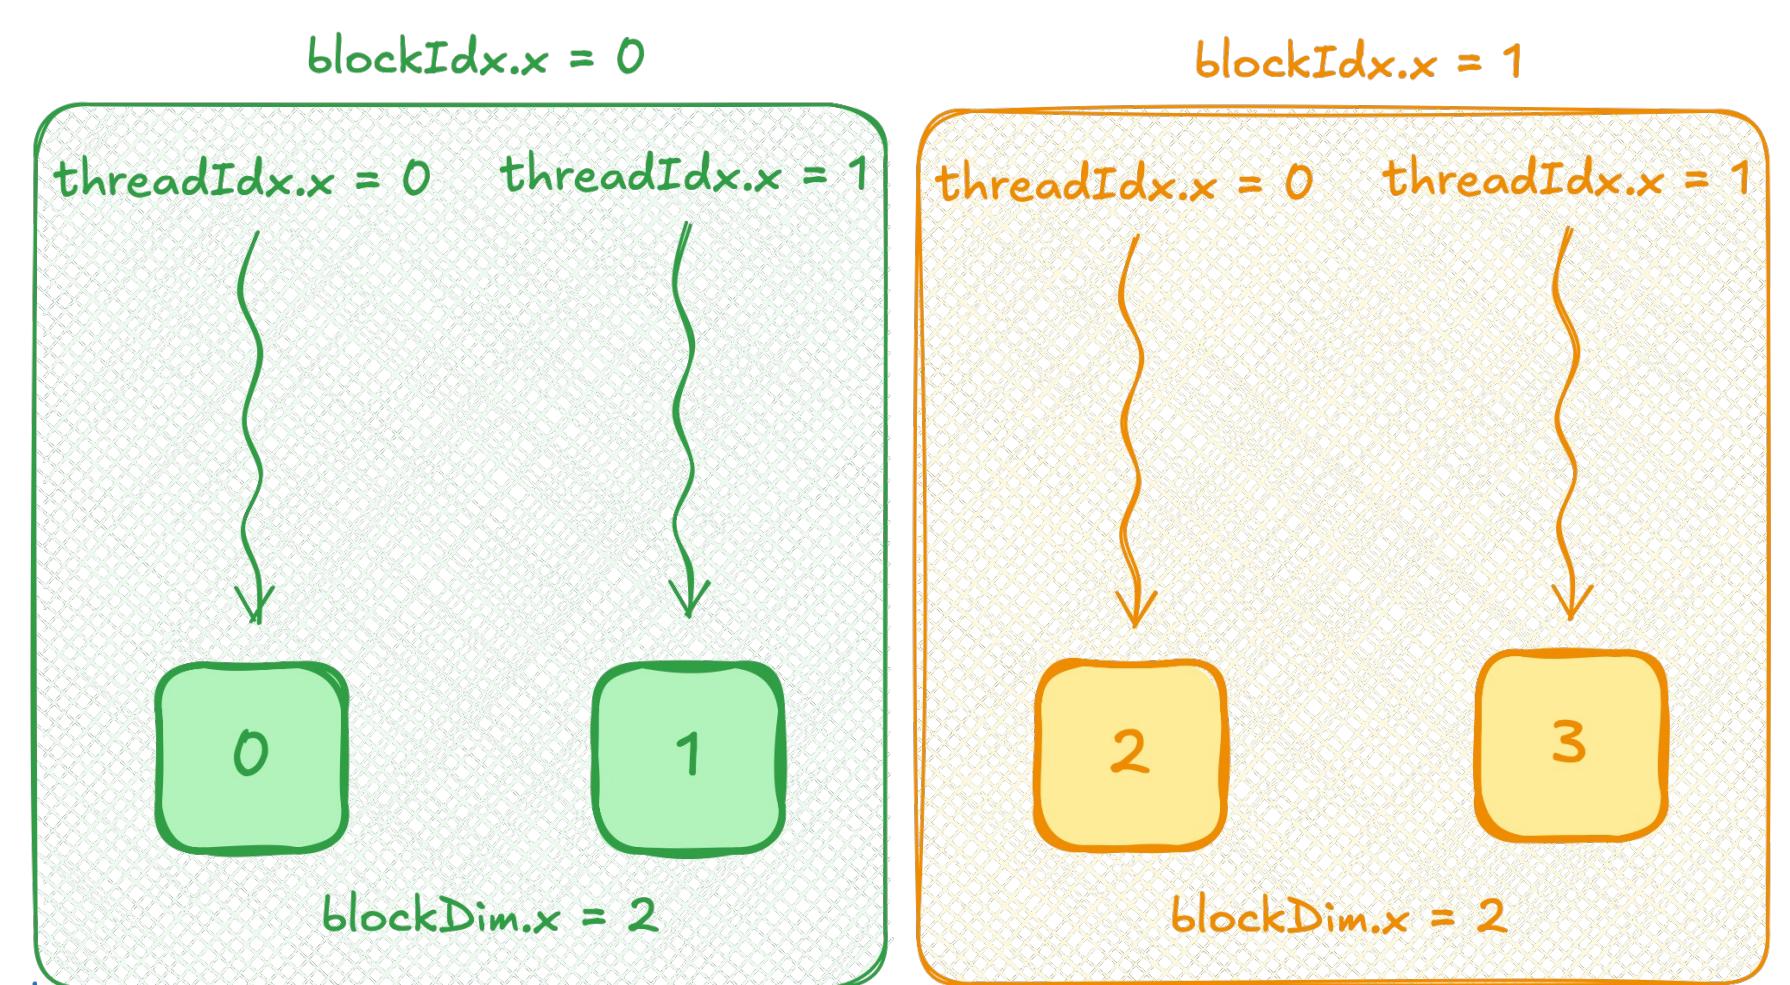
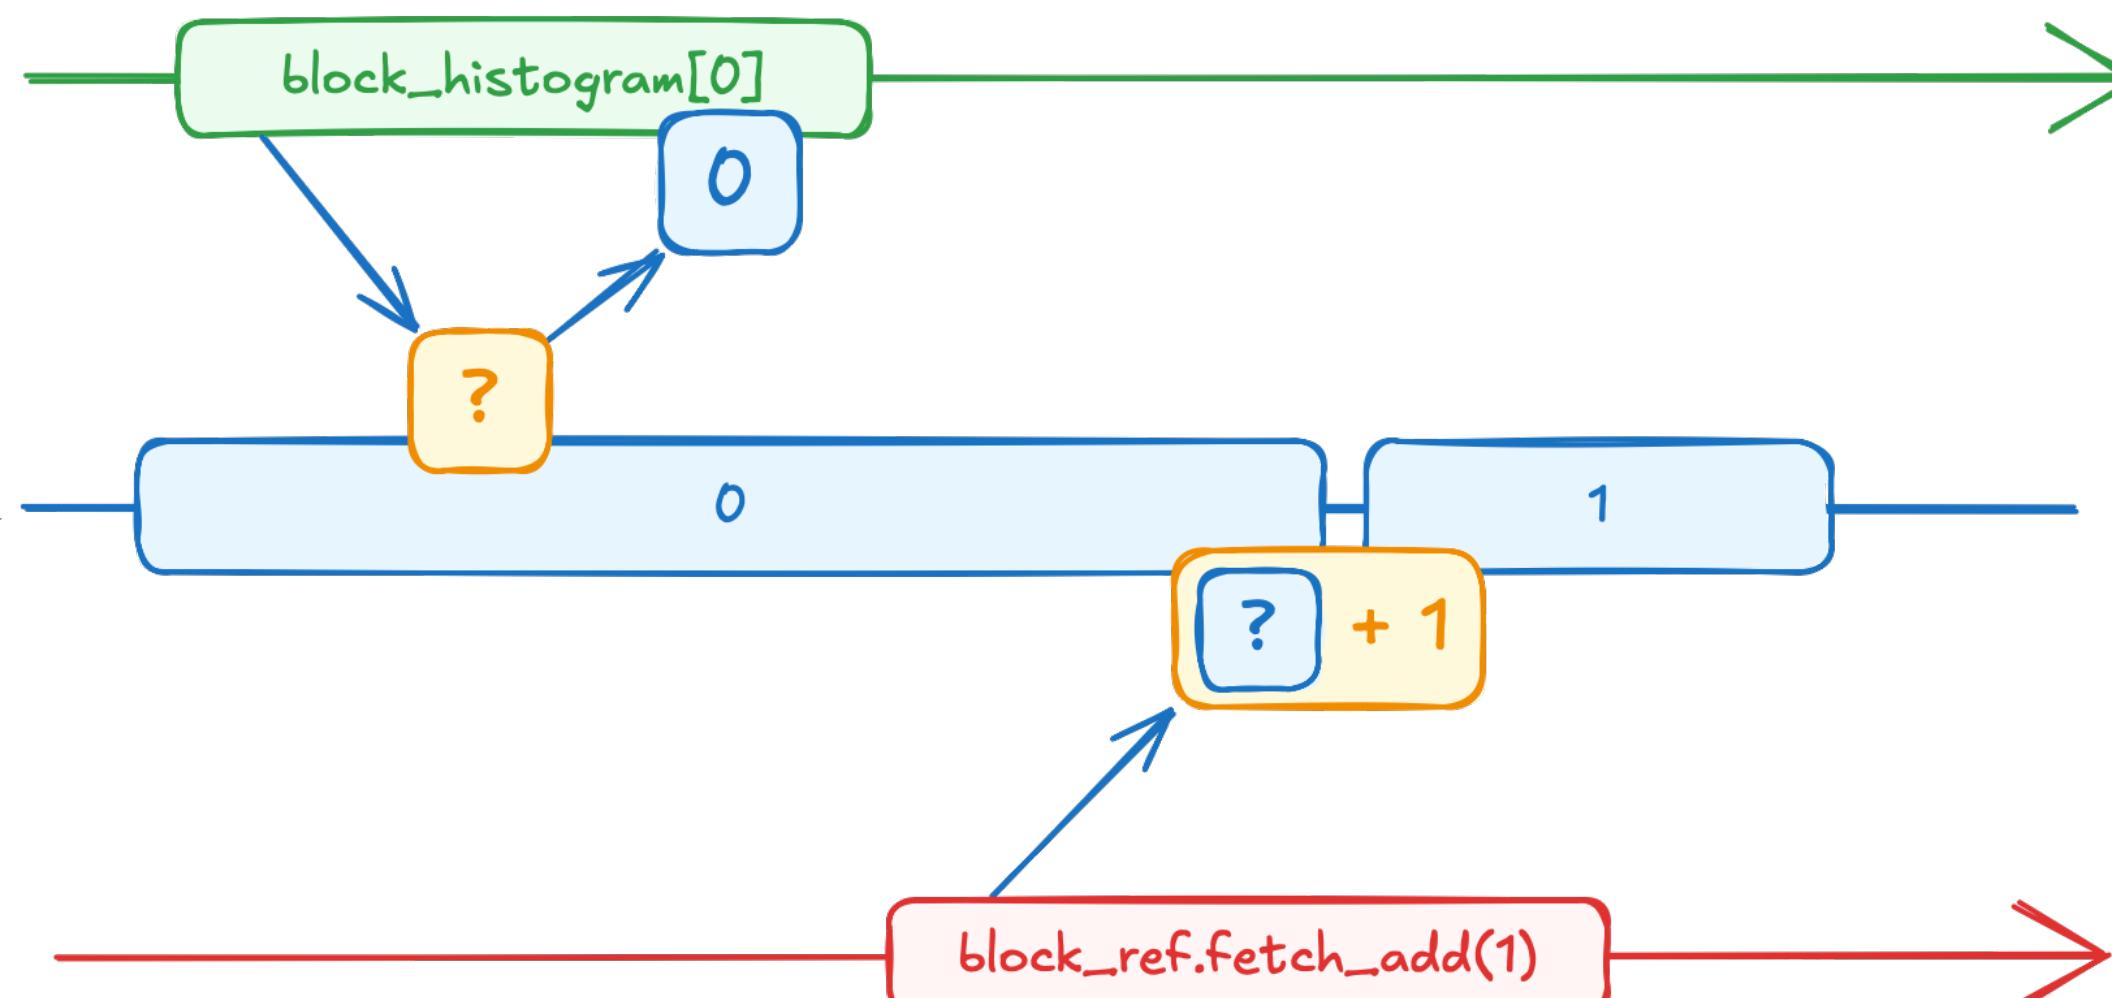


# Implementing New Algorithms with CUDA Kernels

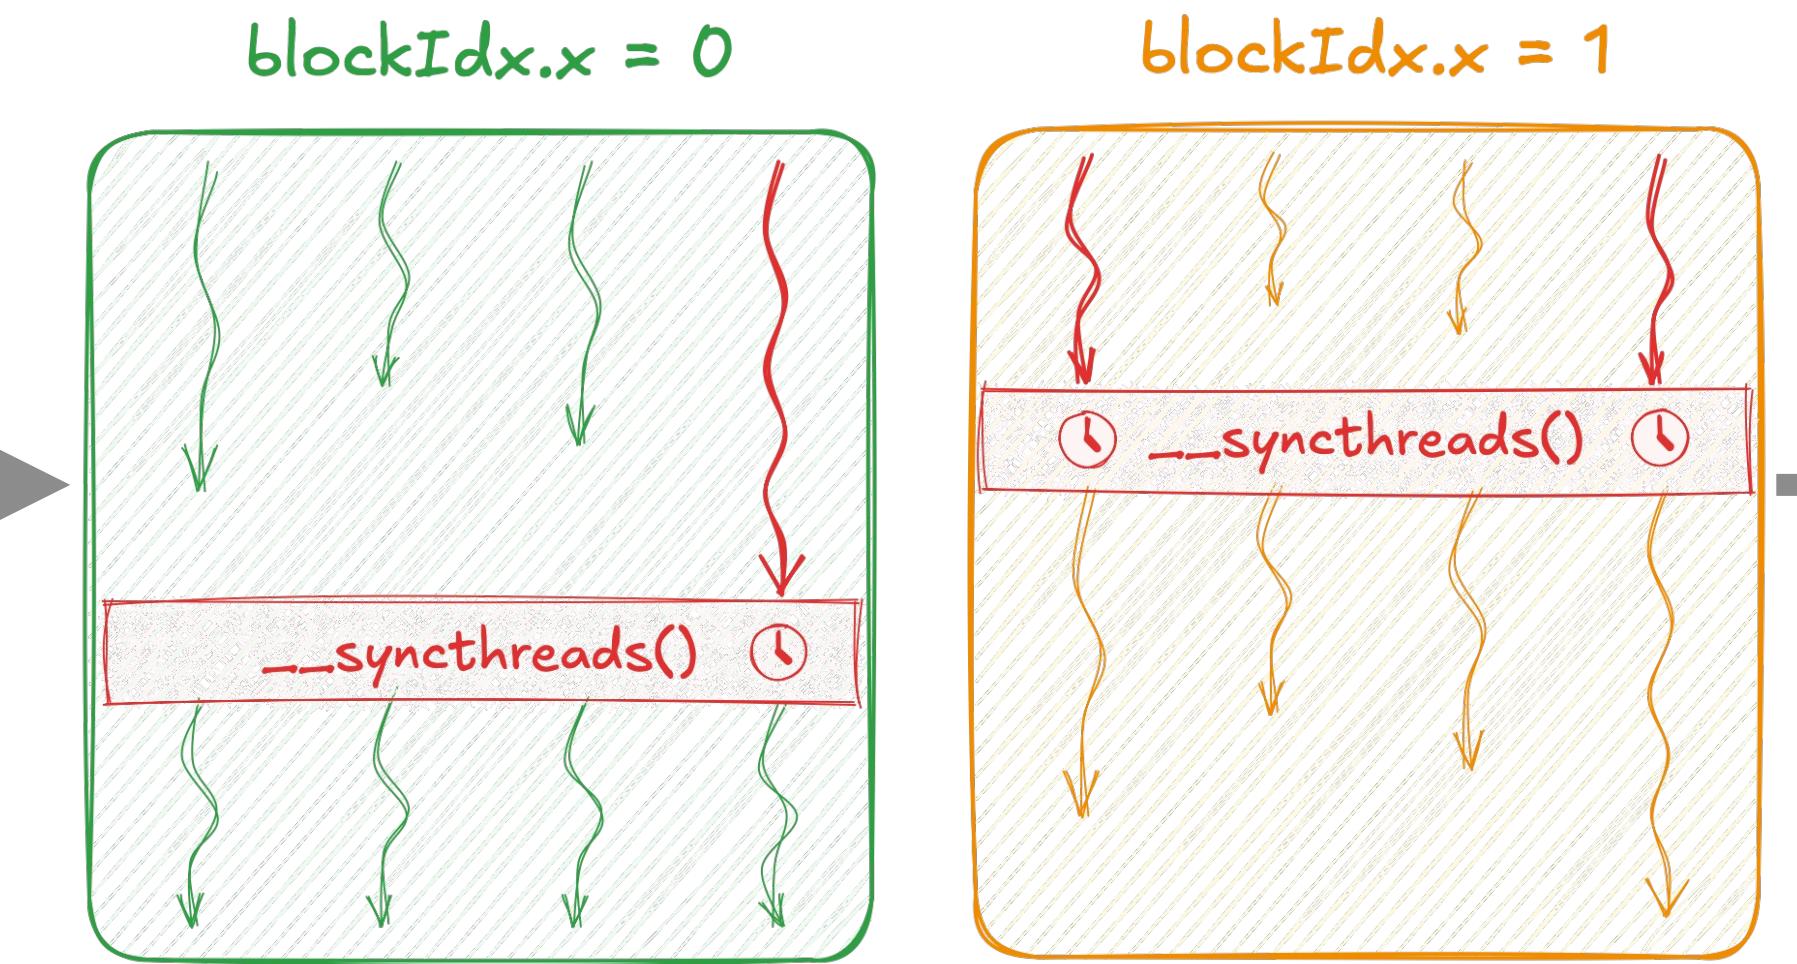
# Section at a Glance



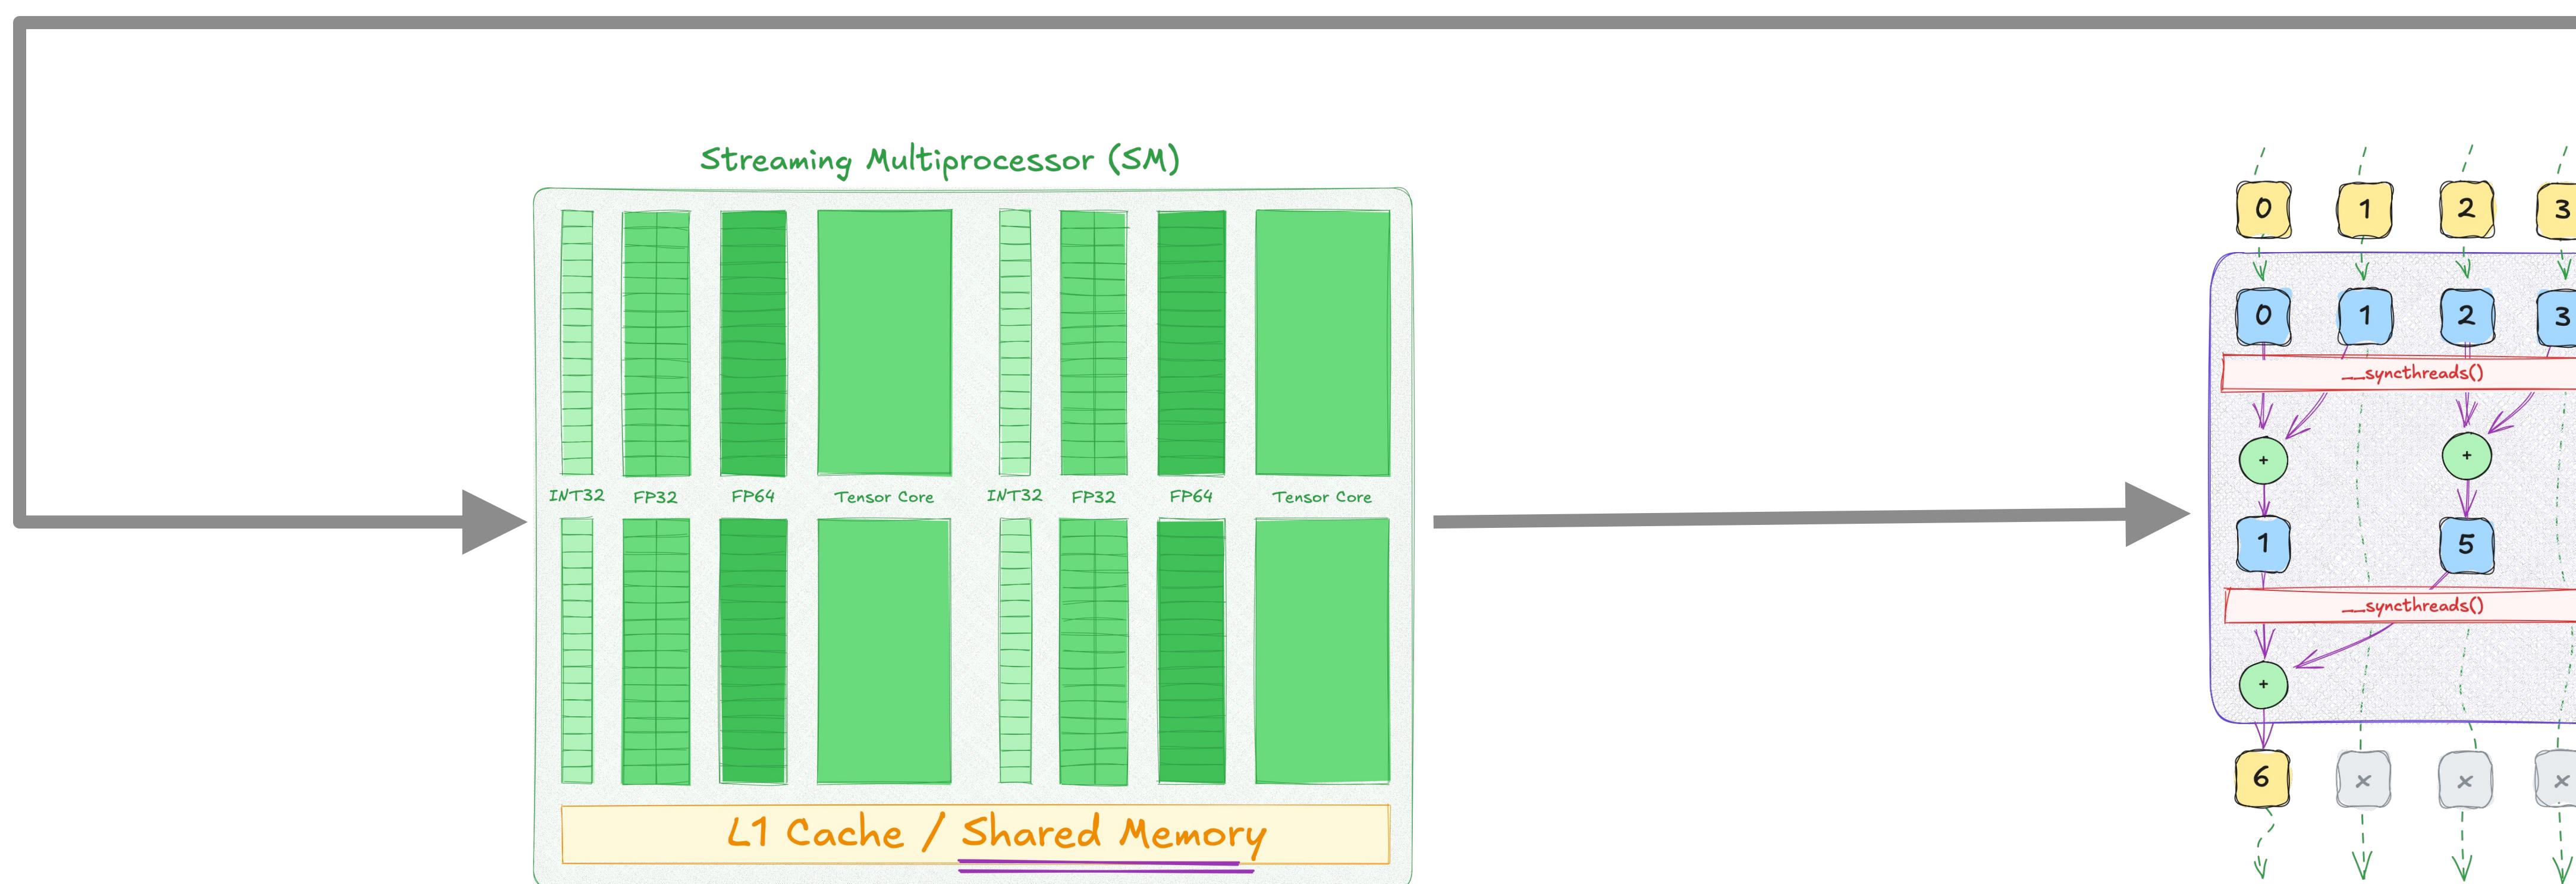
Learn the notion of 3.2 **CUDA Kernels** as we implement our first CUDA algorithm



Learn about 3.3 **Atomics**  
Memory operations



Learn the notion of thread block  
3.4 **Synchronization**

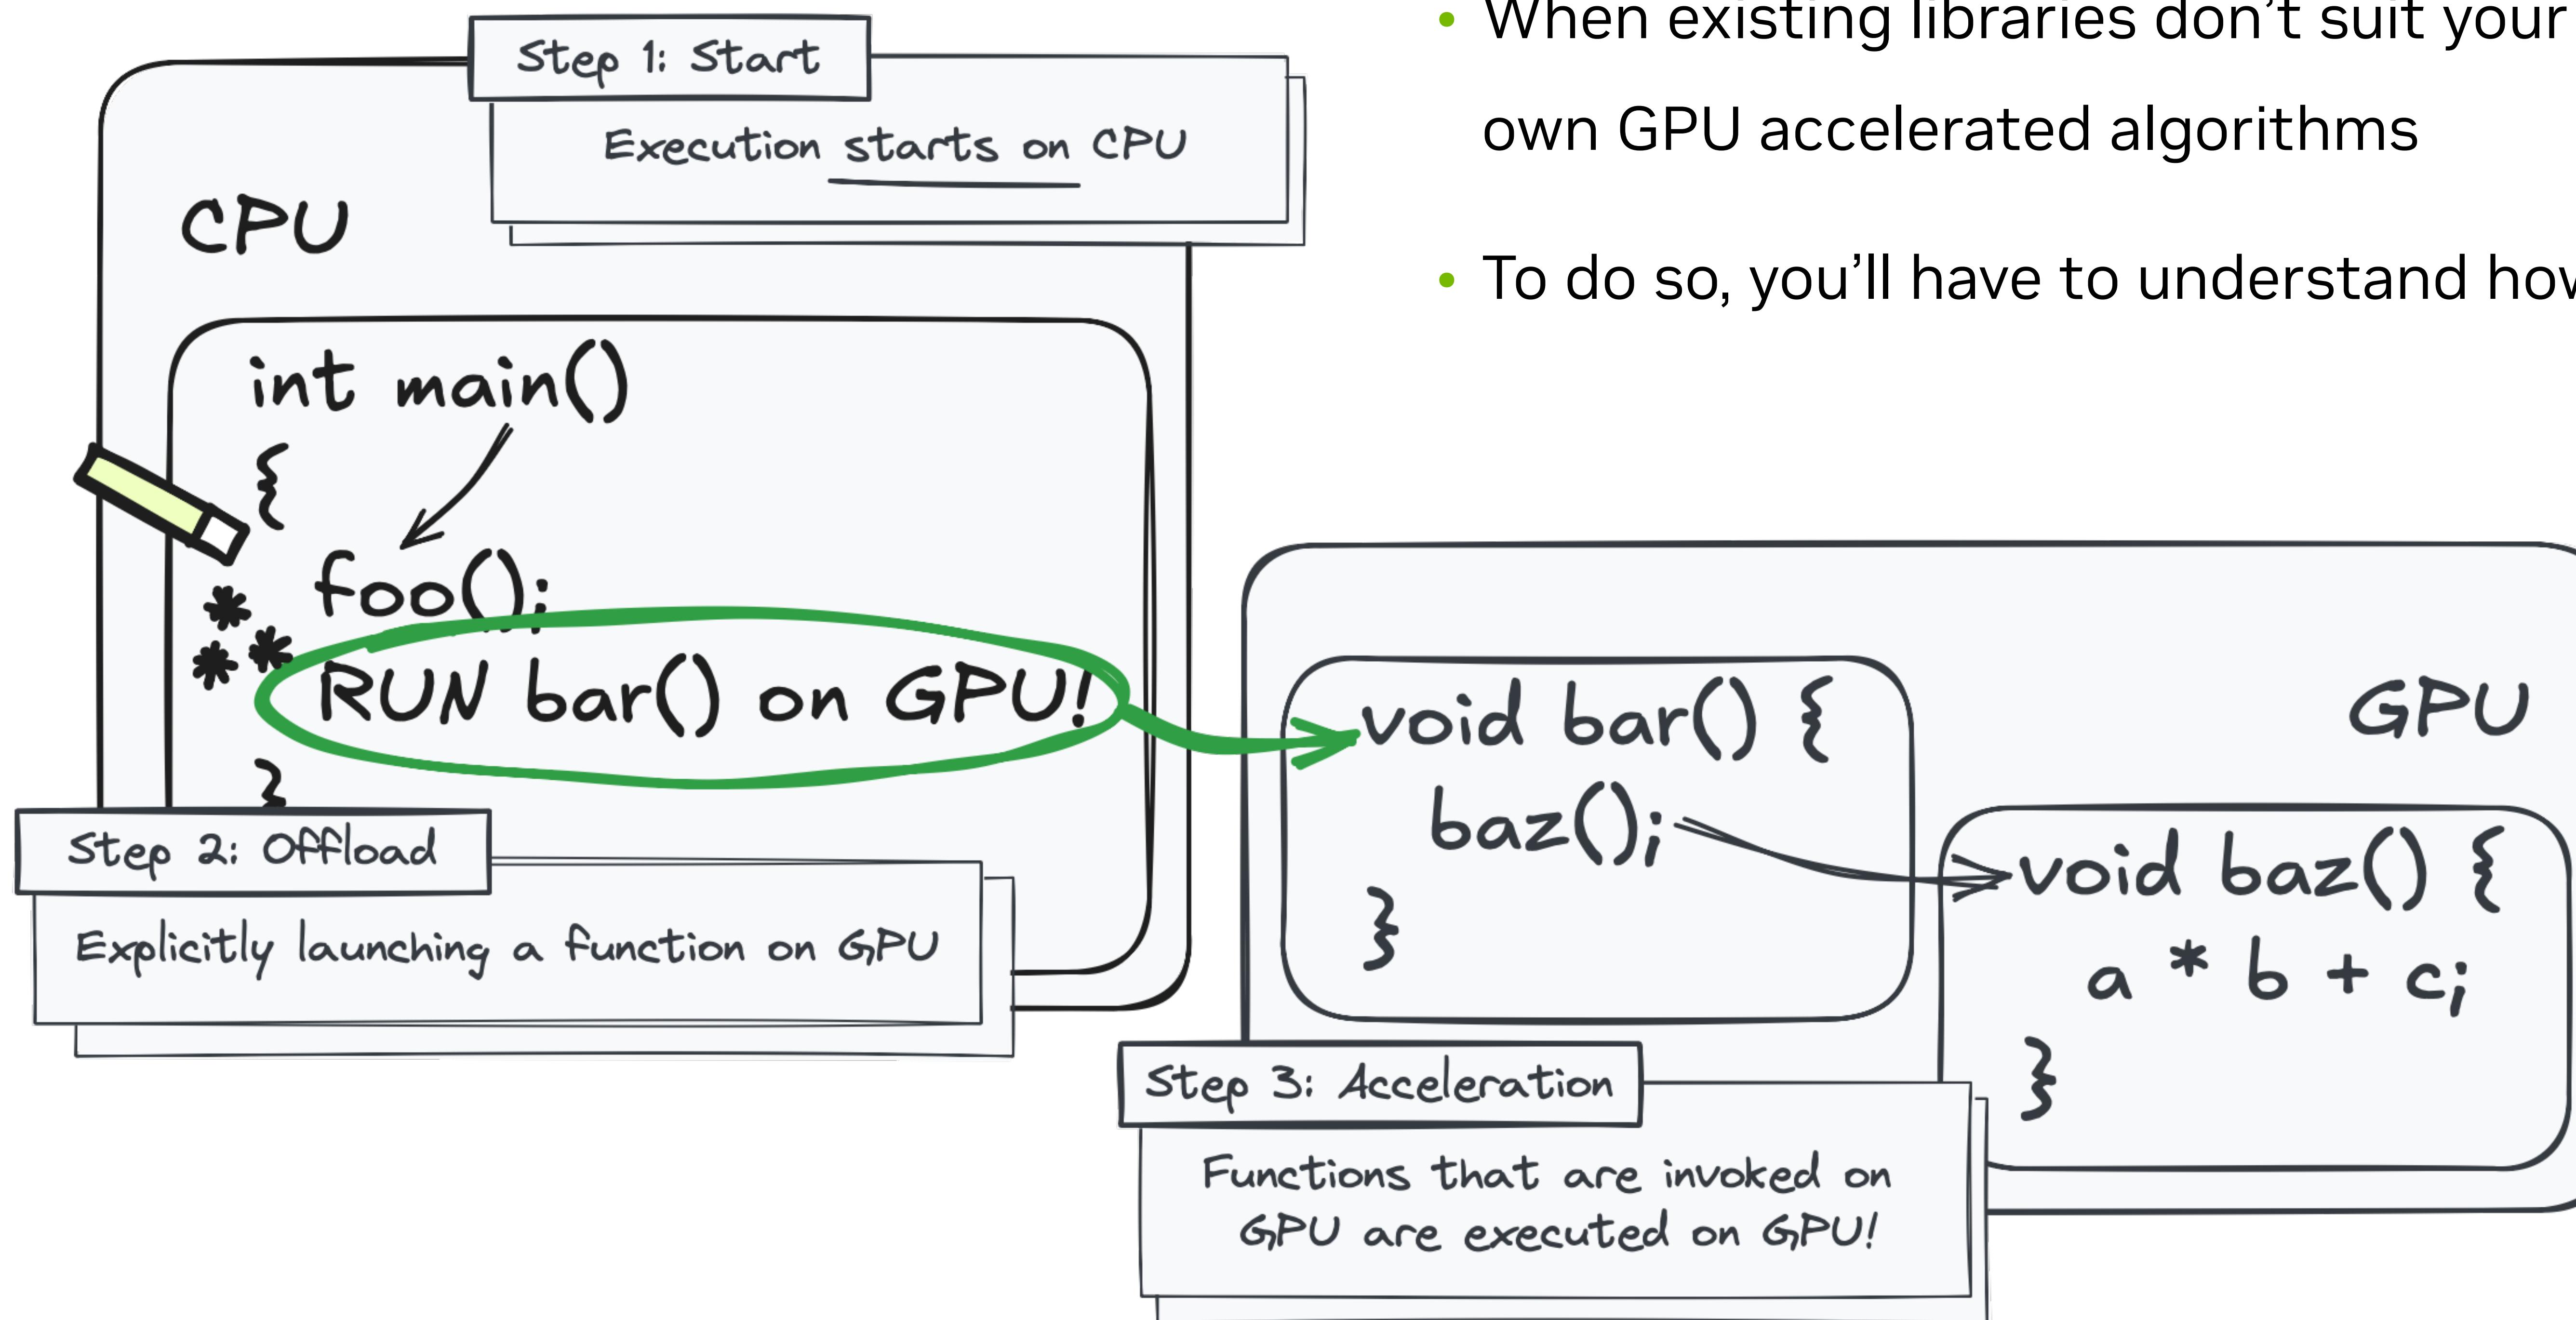


Learn the notion of 3.5 **Shared Memory** as accelerate out kernel and achieve 4x speedup

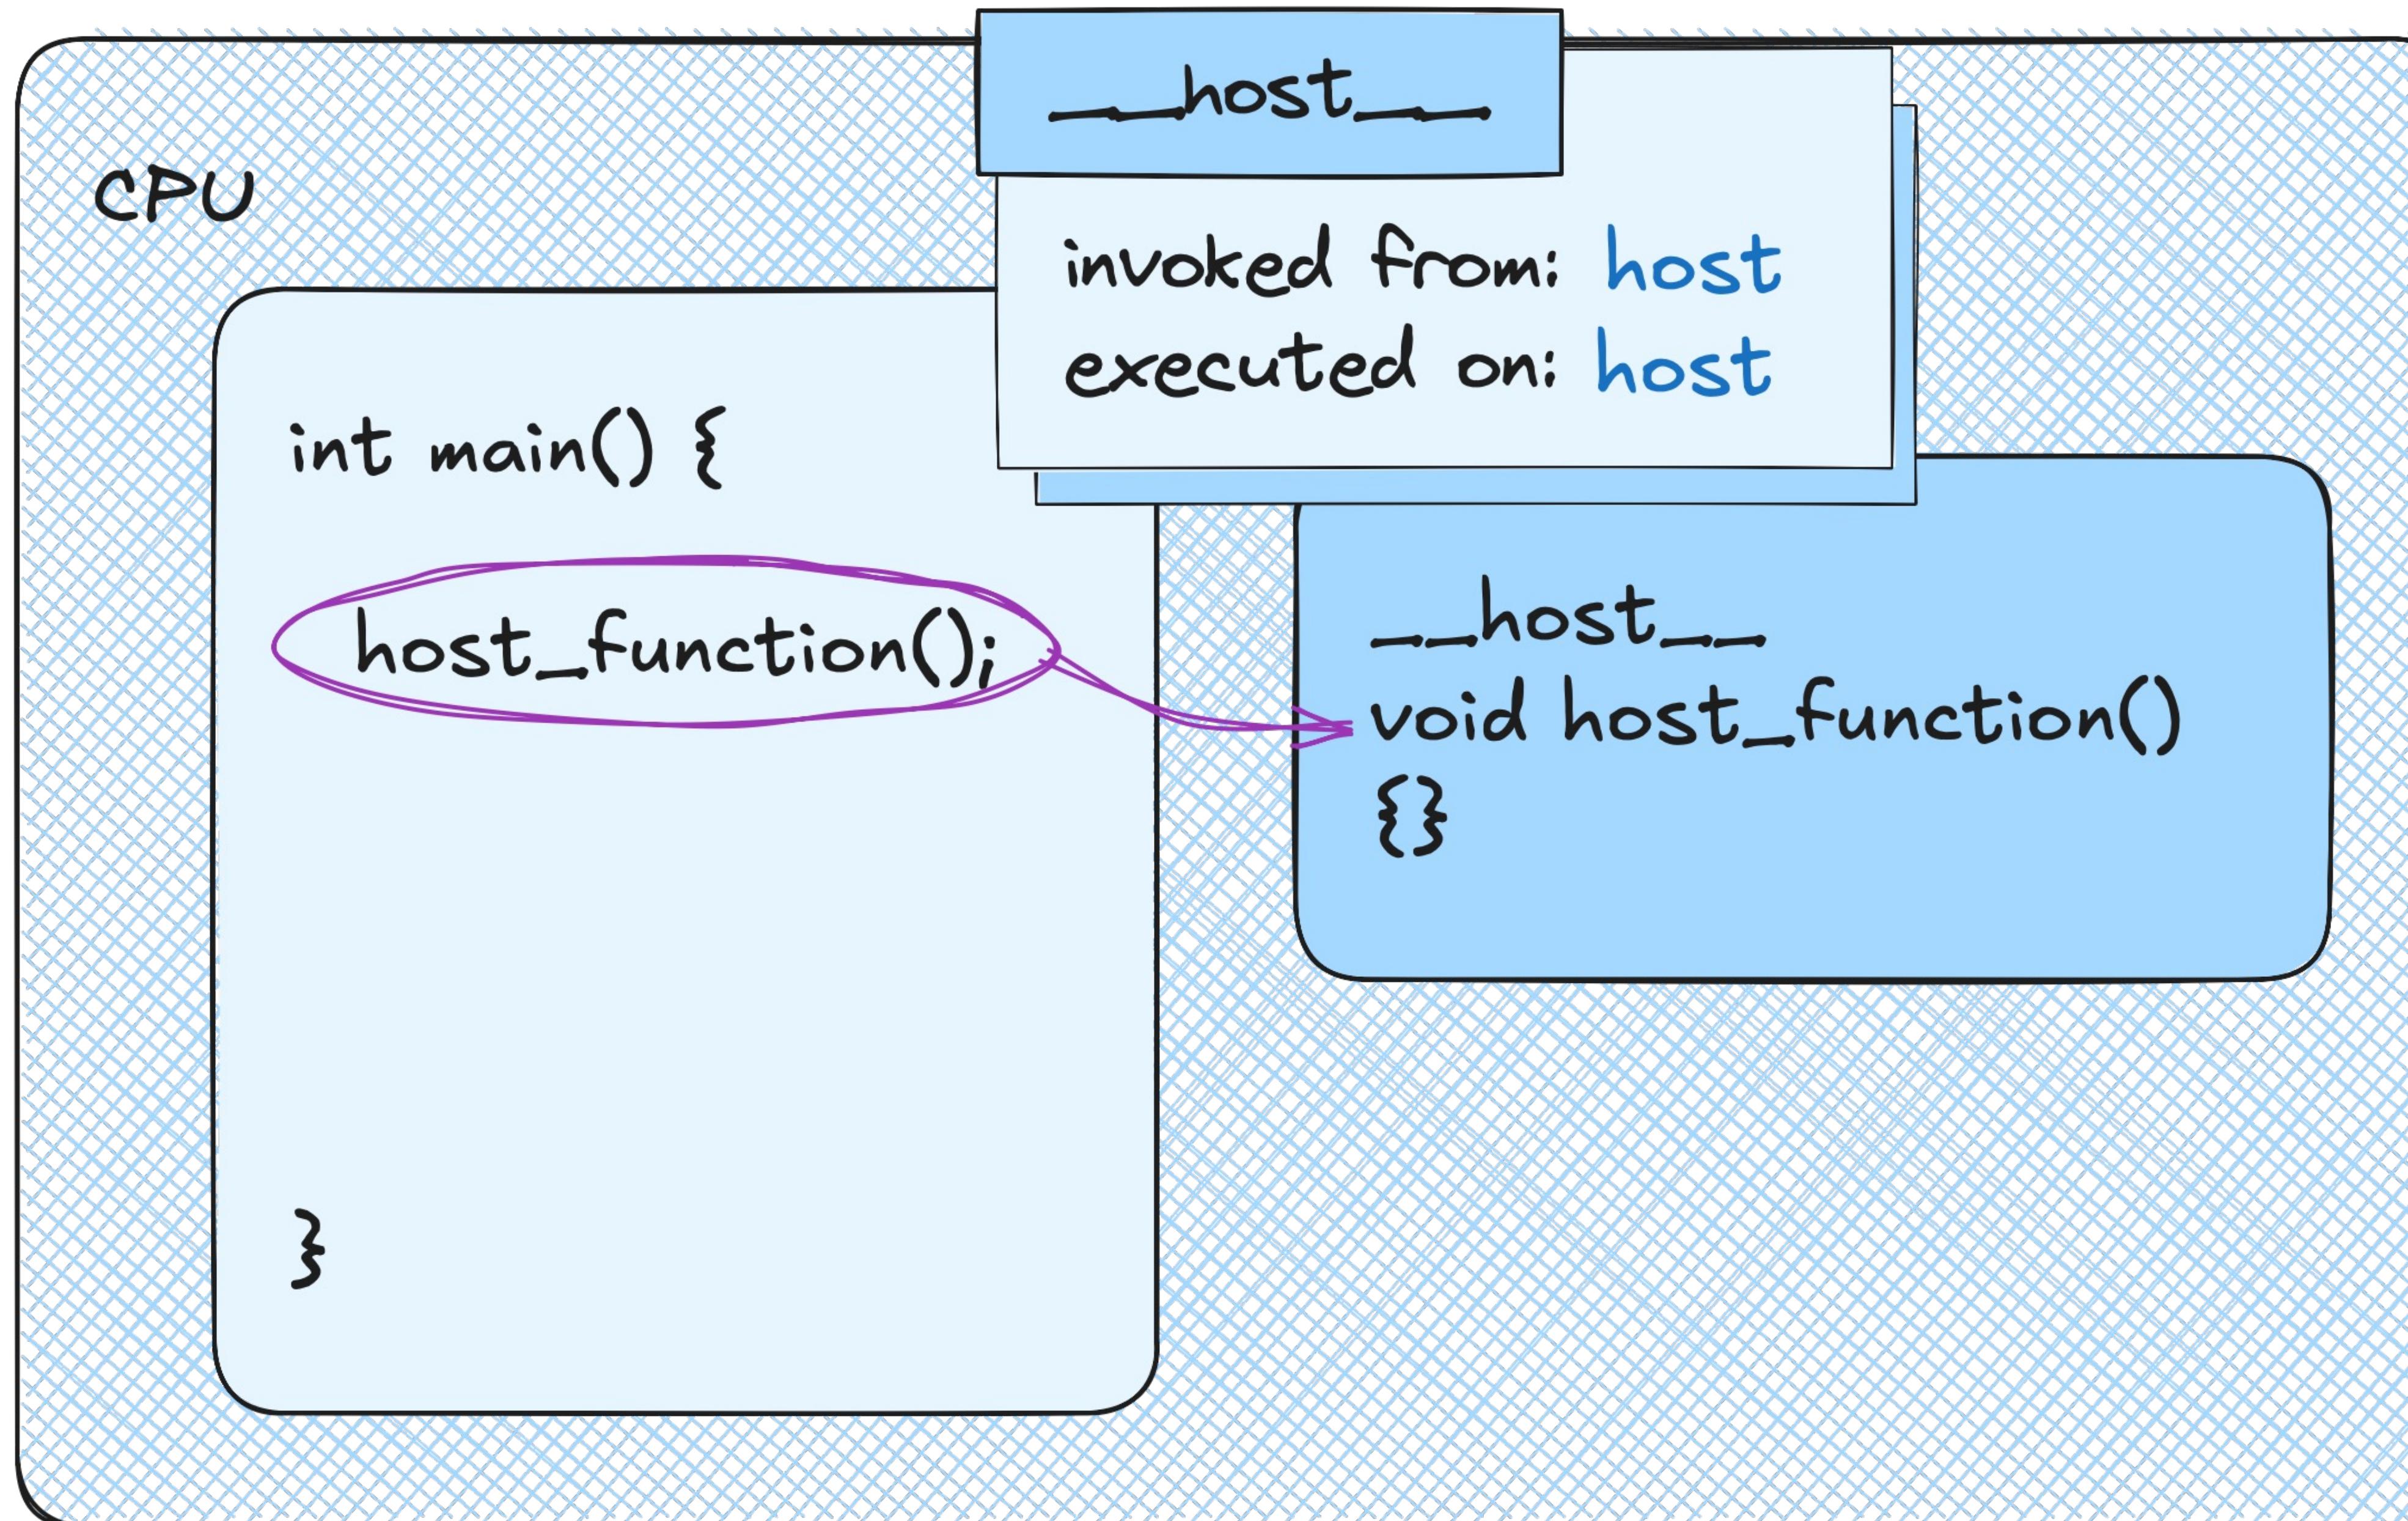
Learn the notion of 3.6 **Cooperative Algorithms** as we coarse the grid and accelerate IO

# CUDA Kernels

- CUDA libraries do not have magic powers to invoke code on GPU
- When existing libraries don't suit your use case, you'll have to write your own GPU accelerated algorithms
- To do so, you'll have to understand how to launch work on GPU directly

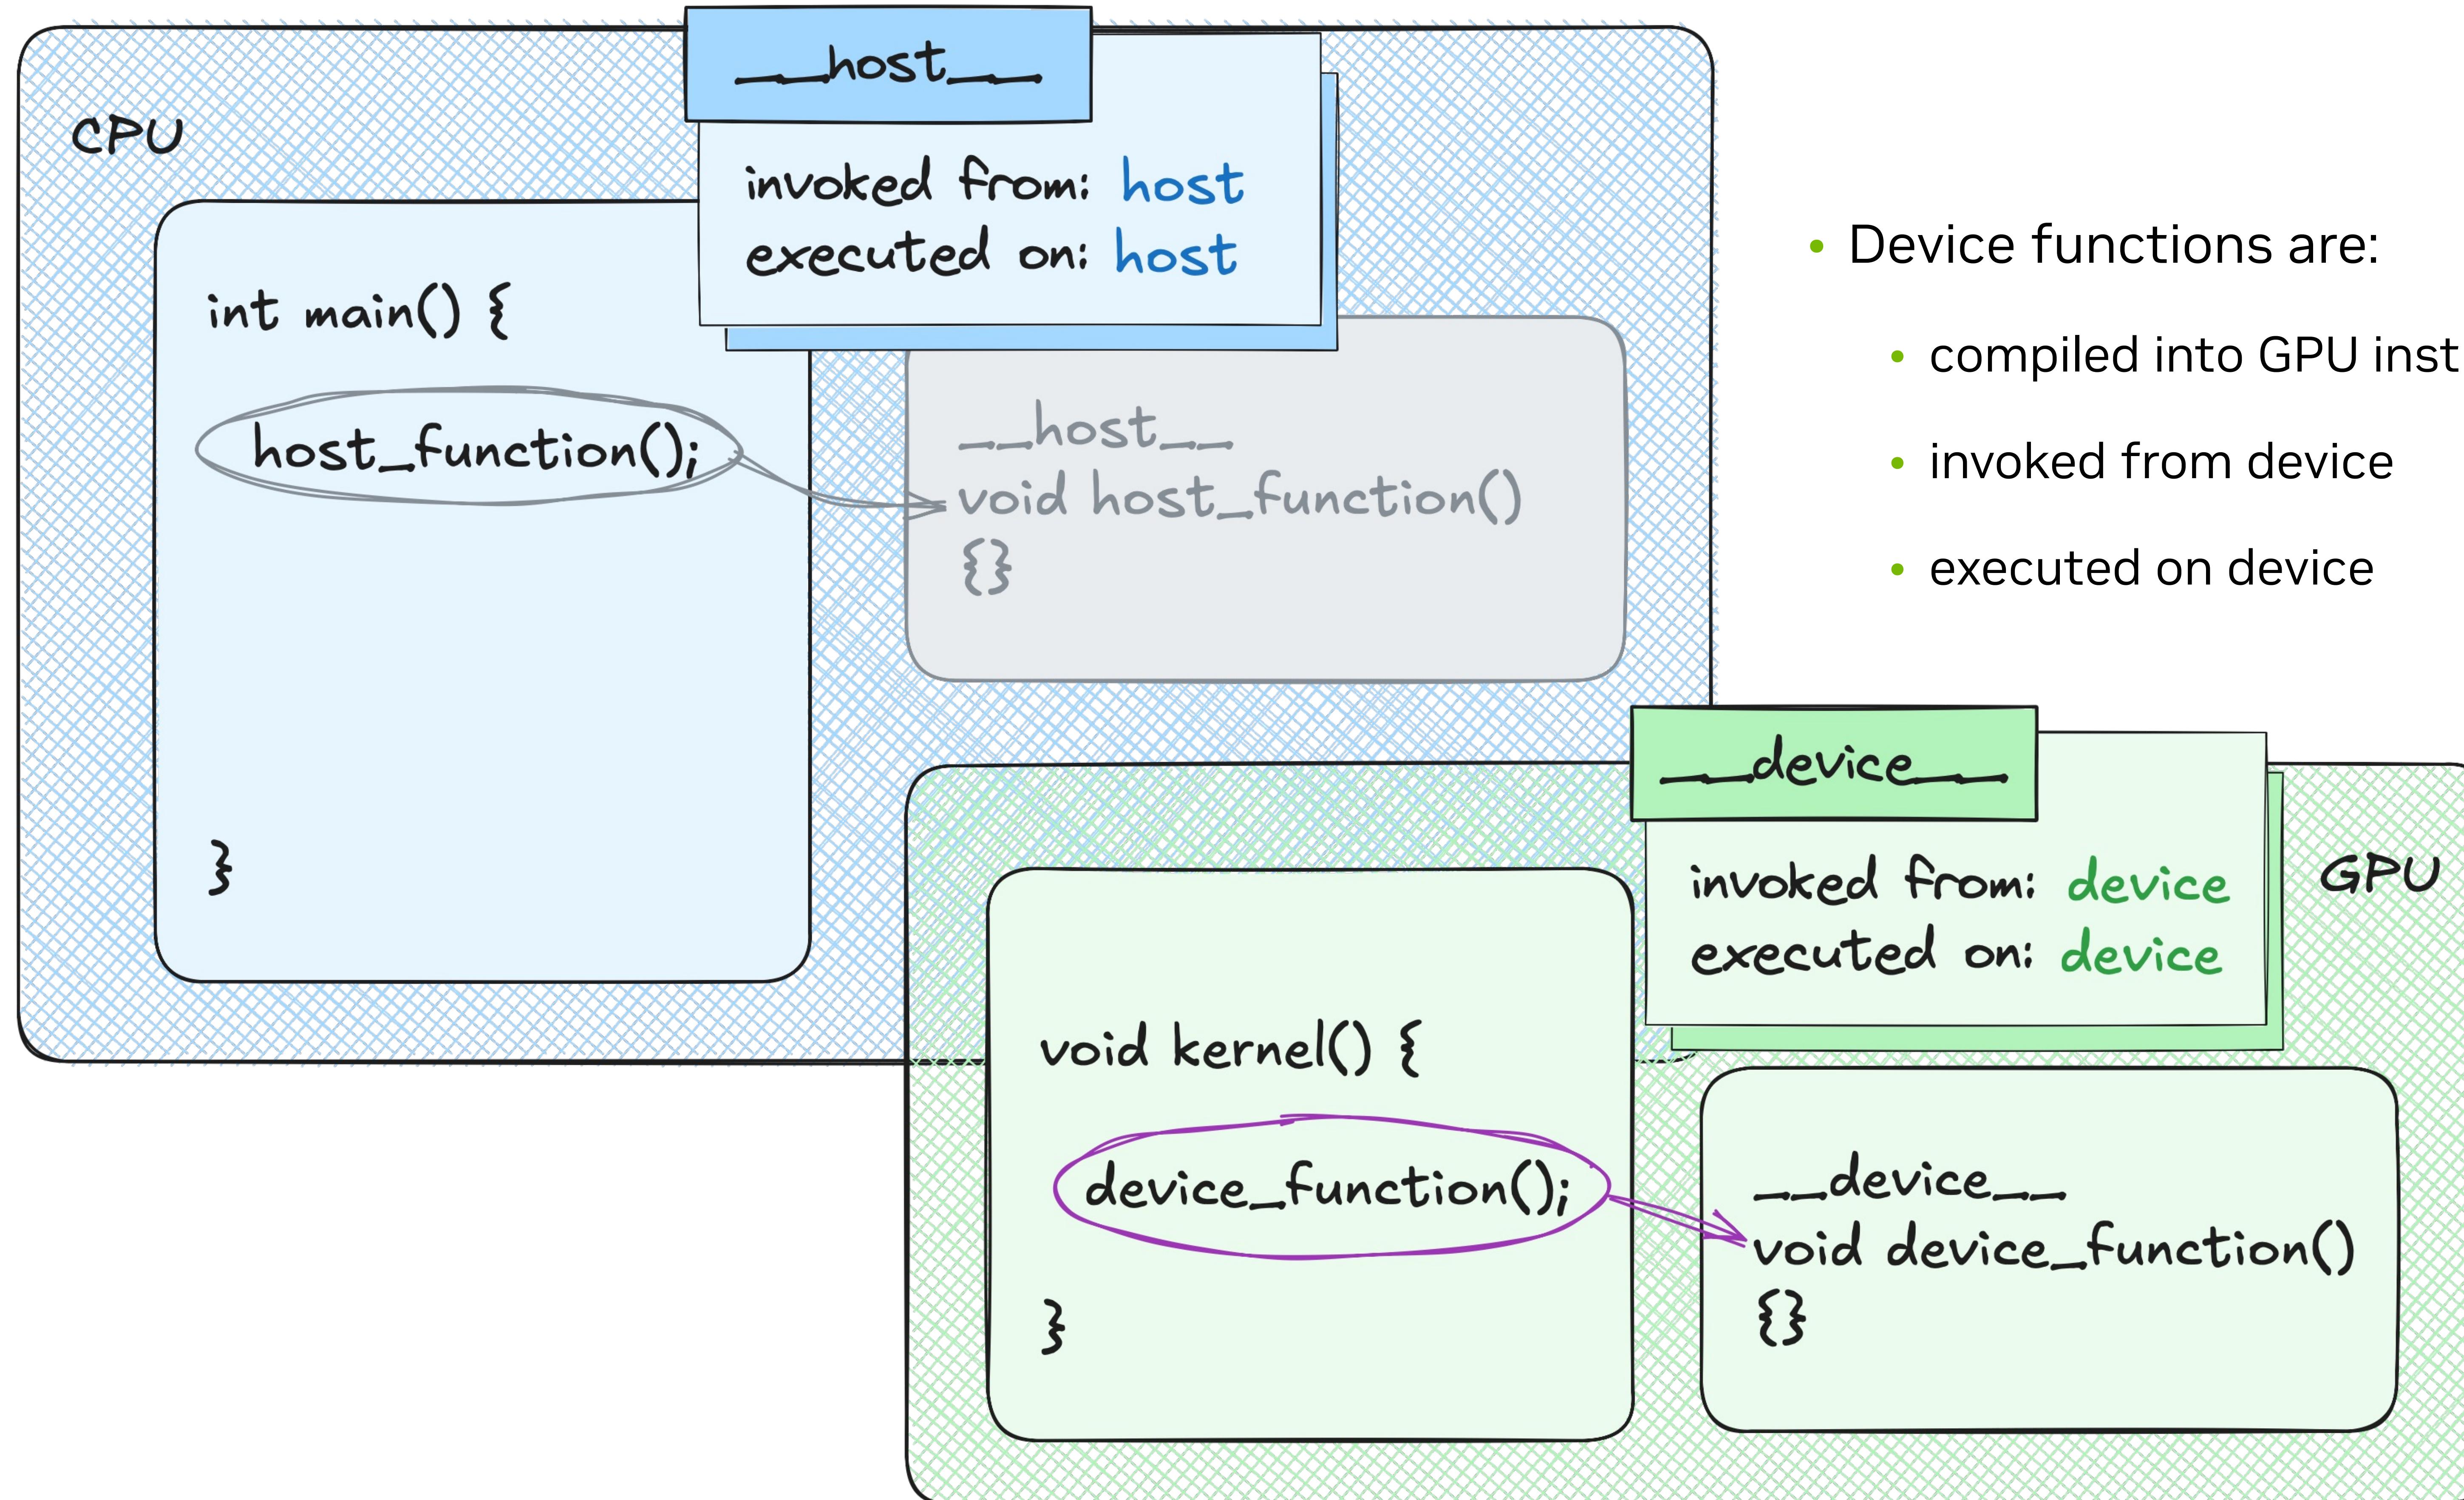


# CUDA Kernels

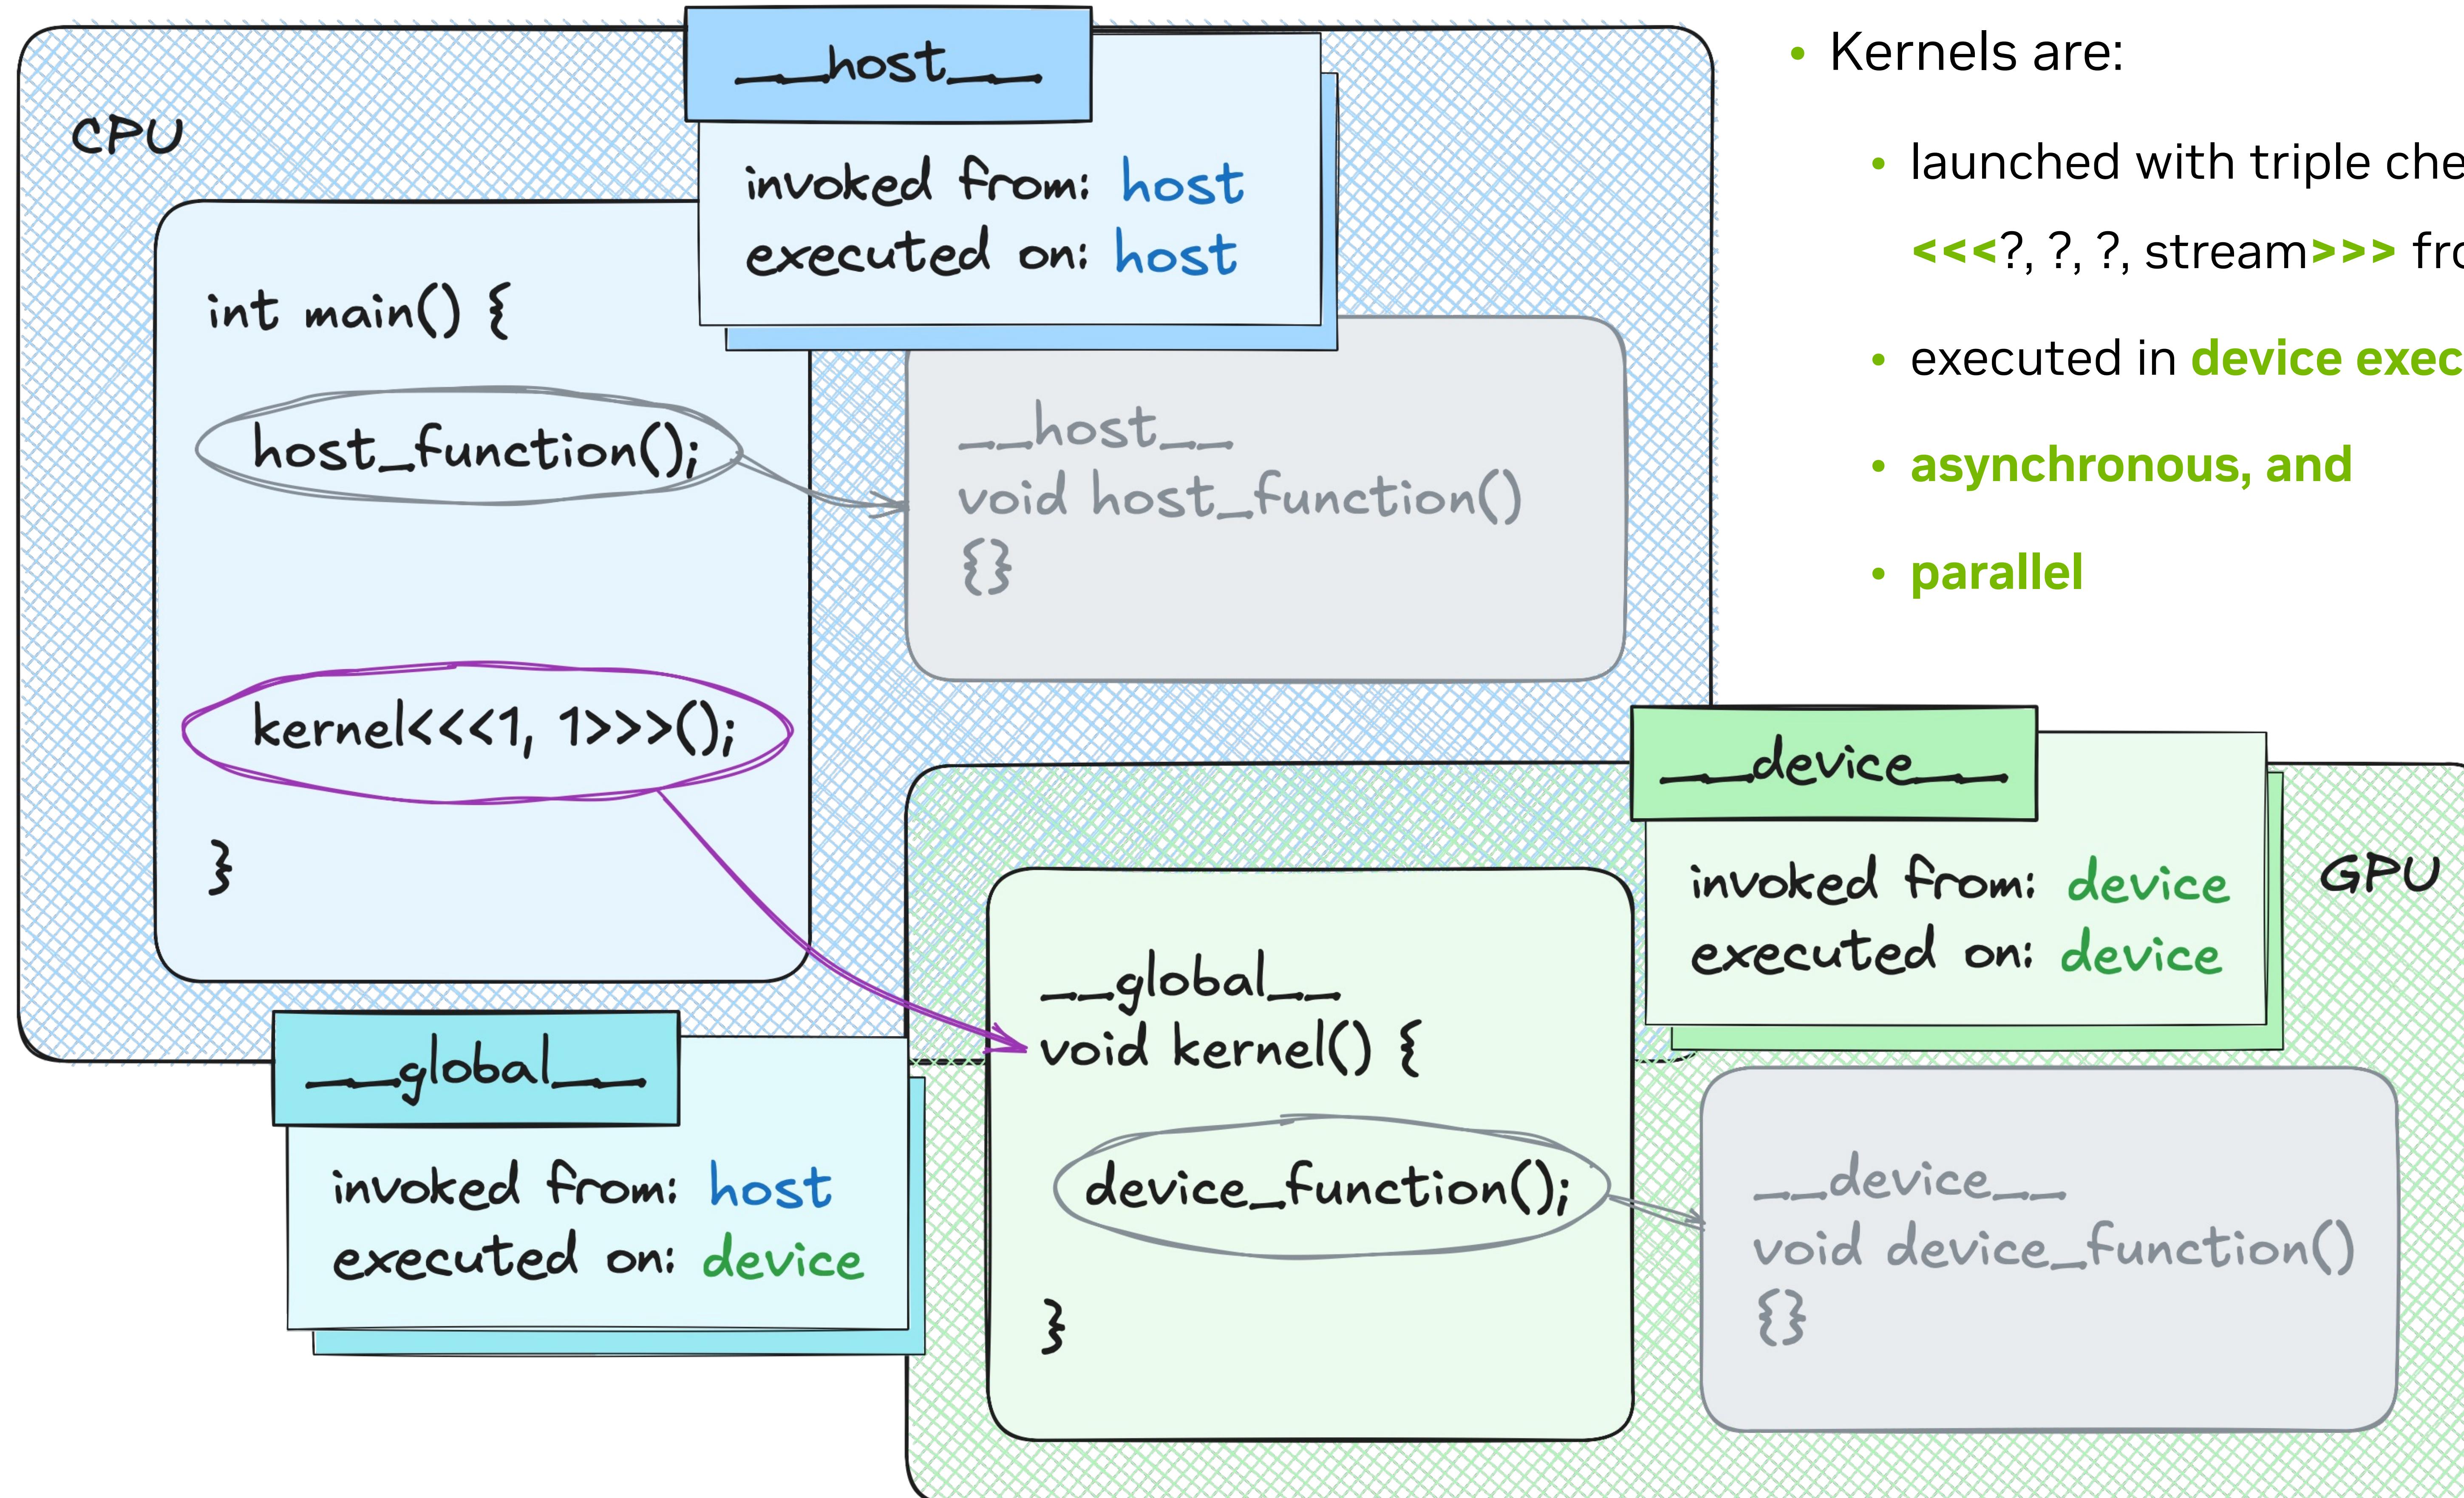


- All functions are implicitly `__host__`
- Host functions are:
  - compiled into CPU instructions
  - invoked from host
  - executed on host

# CUDA Kernels



# CUDA Kernels



# Current Code

```
using temperature_grid_f =
    cuda::std::mdspan<float, cuda::std::extents<int, 2>>;
void simulate(temperature_grid_f in,
              float *out,
              cudaStream_t stream)
{
    auto ids = thrust::make_counting_iterator(0);
    cub::DeviceTransform::Transform(
        ids, out, in.size(),
        [in] __host__ __device__(int cell_id) {
            return ach::compute(cell_id, in);
        }, stream);
}
```

- Currently, the simulator is:
  - using `cuda::std::mdspan` to view data as 2D
  - Using CUB to transform cell indices into the next temperatures
- Under the hood, CUB:
  - invokes kernels on provided stream
  - doesn't synchronize the stream,
  - which makes CUB asynchronous
- Let's try implementing simulator without CUB!

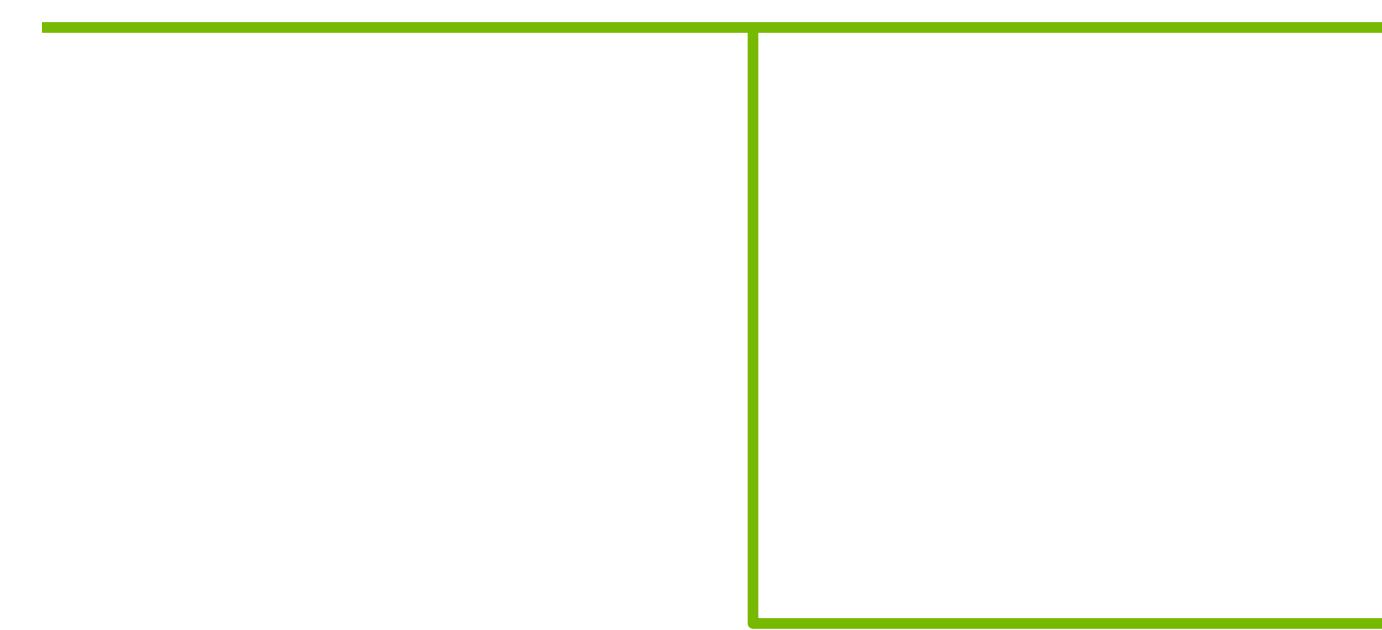
# CUDA Kernels

```
__global__
void single_thread_kernel(ach::temperature_grid_f in, float *out)
{
    for (int id = 0; id < in.size(); id++)
    {
        out[id] = ach::compute(id, in);
    }
}
```

Triple chevron accepts four parameters:

- You know what stream is at this point
- we'll go over remaining parameters one at a time

```
void simulate(ach::temperature_grid_f in,
              float *out,
              cudaStream_t stream)
{
    single_thread_kernel<<<1, 1, 0, stream>>>(in, out);
```



# CUDA Kernels

```
__global__
void single_thread_kernel(ach::temperature_grid_f in, float *out)
{
    for (int id = 0; id < in.size(); id++)
    {
        out[id] = ach::compute(id, in);
    }
}
```

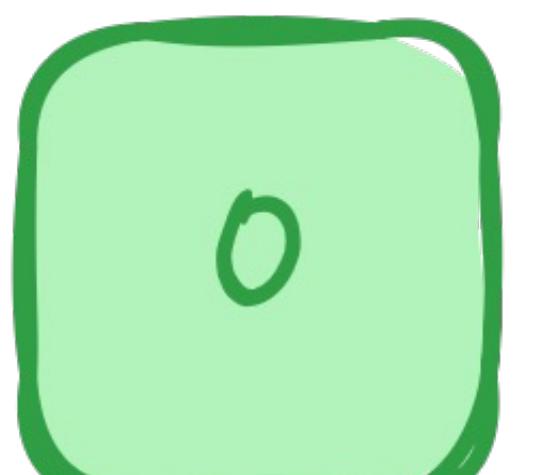
As we saw previously, GPUs do not automatically parallelize your code:

- Single-threaded code remains serial
- This makes our kernel incredibly slow

10 compute steps take:

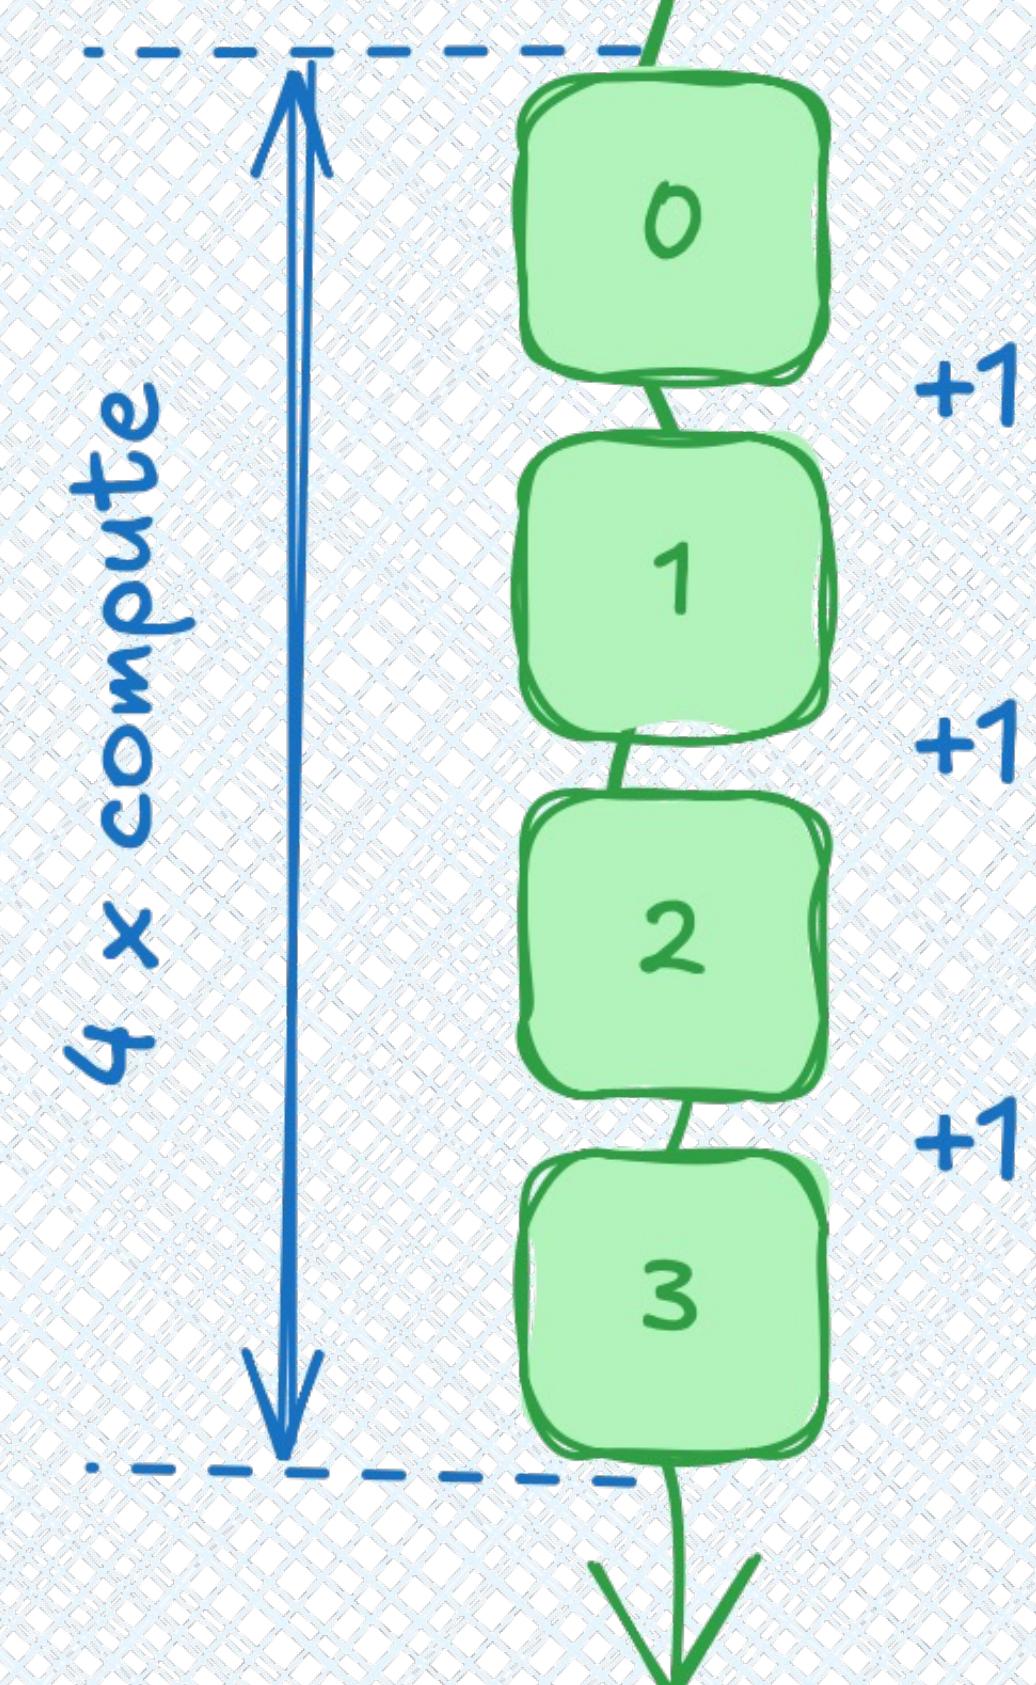
- cub: 0.000043
- 1 thread: 4.15

thread



time view

thread  
id



# CUDA Kernels

```
const int number_of_threads = 2;

__global__
void block_kernel(ach::temperature_grid_f in, float *out)
{
    int thread_index = threadIdx.x;

    for (int id = thread_index; id < in.size(); id += number_of_threads)
    {
        out[id] = ach::compute(id, in);
    }
}

void simulate(ach::temperature_grid_f in,
             float *out,
             cudaStream_t stream)
{
    block_kernel<<<1, number_of_threads, 0, stream>>>(in, out);
}
```

10 compute steps take:

- cub: 0.000043
- 1 thread: 4.15
- 2 threads: 2.6

# CUDA Kernels

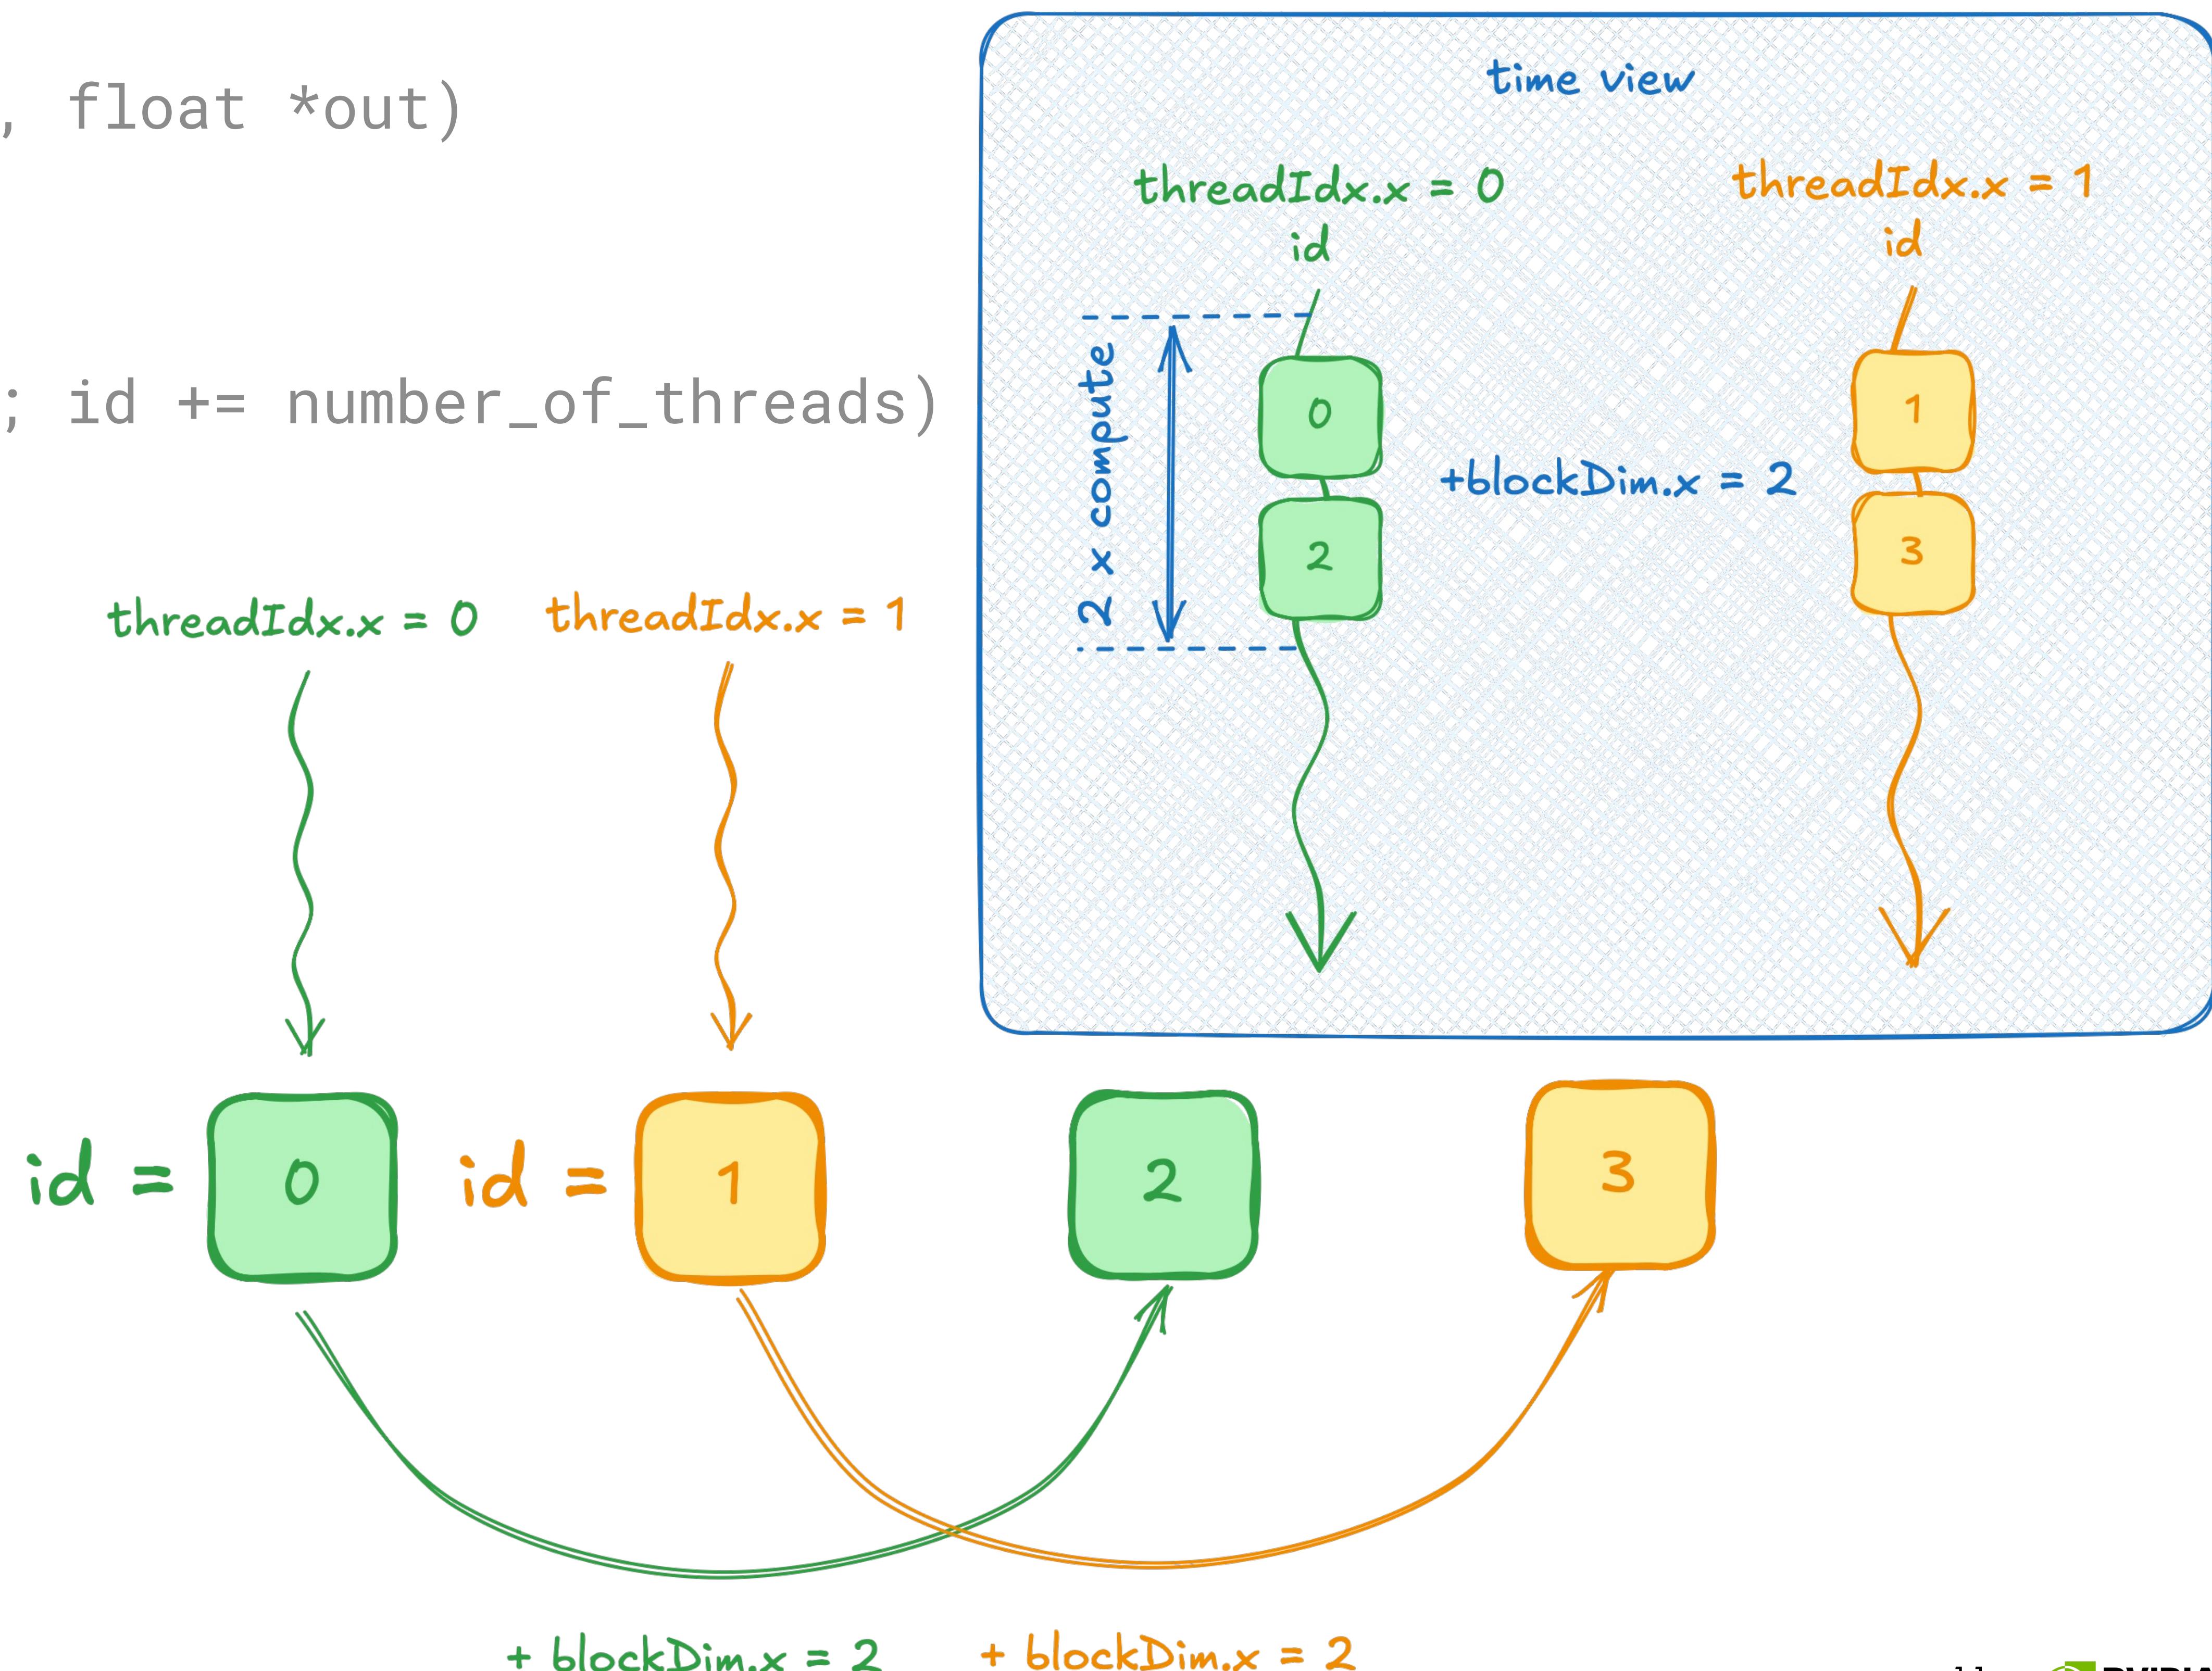
```
const int number_of_threads = 2;

__global__
void block_kernel(ach::temperature_grid_f in, float *out)
{
    int thread_index = threadIdx.x;

    for (int id = thread_index; id < in.size(); id += number_of_threads)
    {
        out[id] = ach::compute(id, in);
    }
}
```

CUDA extends device code with a few variables:

- `threadIdx.x` stores index of the current thread



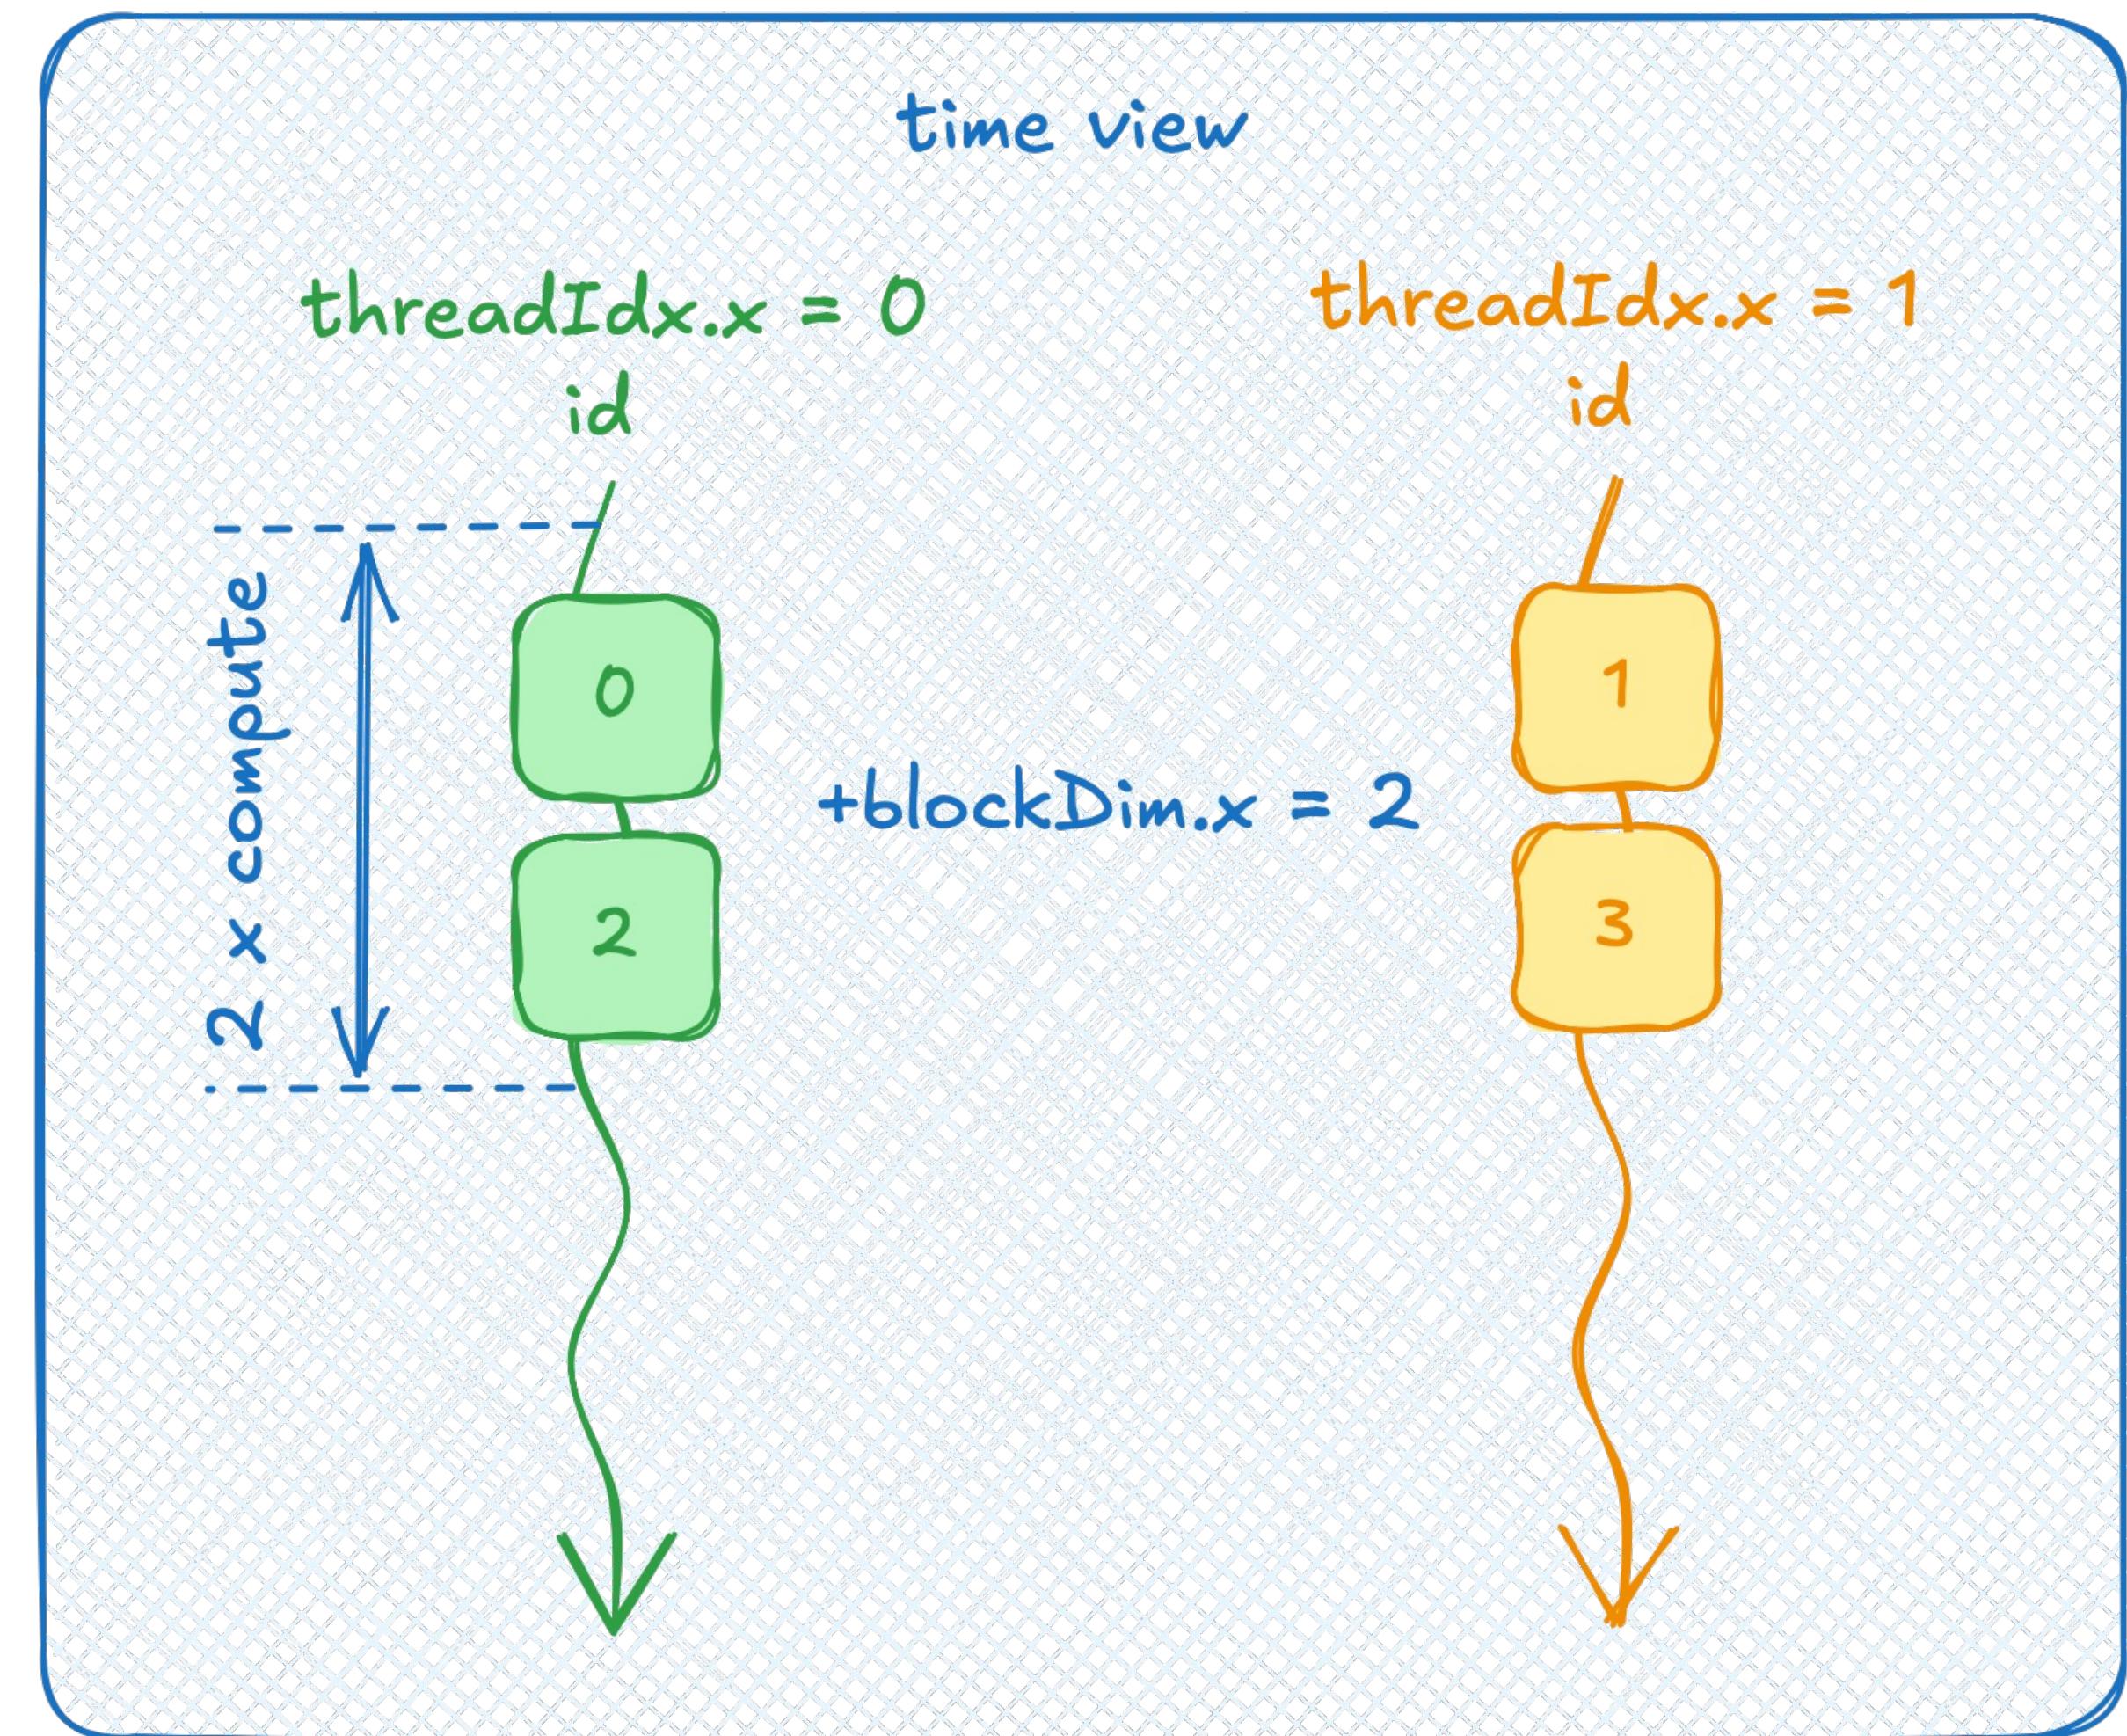
# CUDA Kernels

```
const int number_of_threads = 2;

__global__
void block_kernel(ach::temperature_grid_f in, float *out)
{
    int thread_index = threadIdx.x;

    for (int id = thread_index; id < in.size(); id += number_of_threads)
    {
        out[id] = ach::compute(id, in);
    }
}
```

Same kernel is executed by multiple GPU threads



`threadIdx.x = 0`

```
int thread_index = 0;
```

```
out[0] = ach::compute(0, in);  
out[2] = ach::compute(2, in);
```

`threadIdx.x = 1`

```
int thread_index = 1;
```

```
out[1] = ach::compute(1, in);  
out[3] = ach::compute(3, in);
```

# CUDA Kernels

```
const int number_of_threads = 2;

__global__
void block_kernel(ach::temperature_grid_f in, float *out)
{
    int thread_index = threadIdx.x;

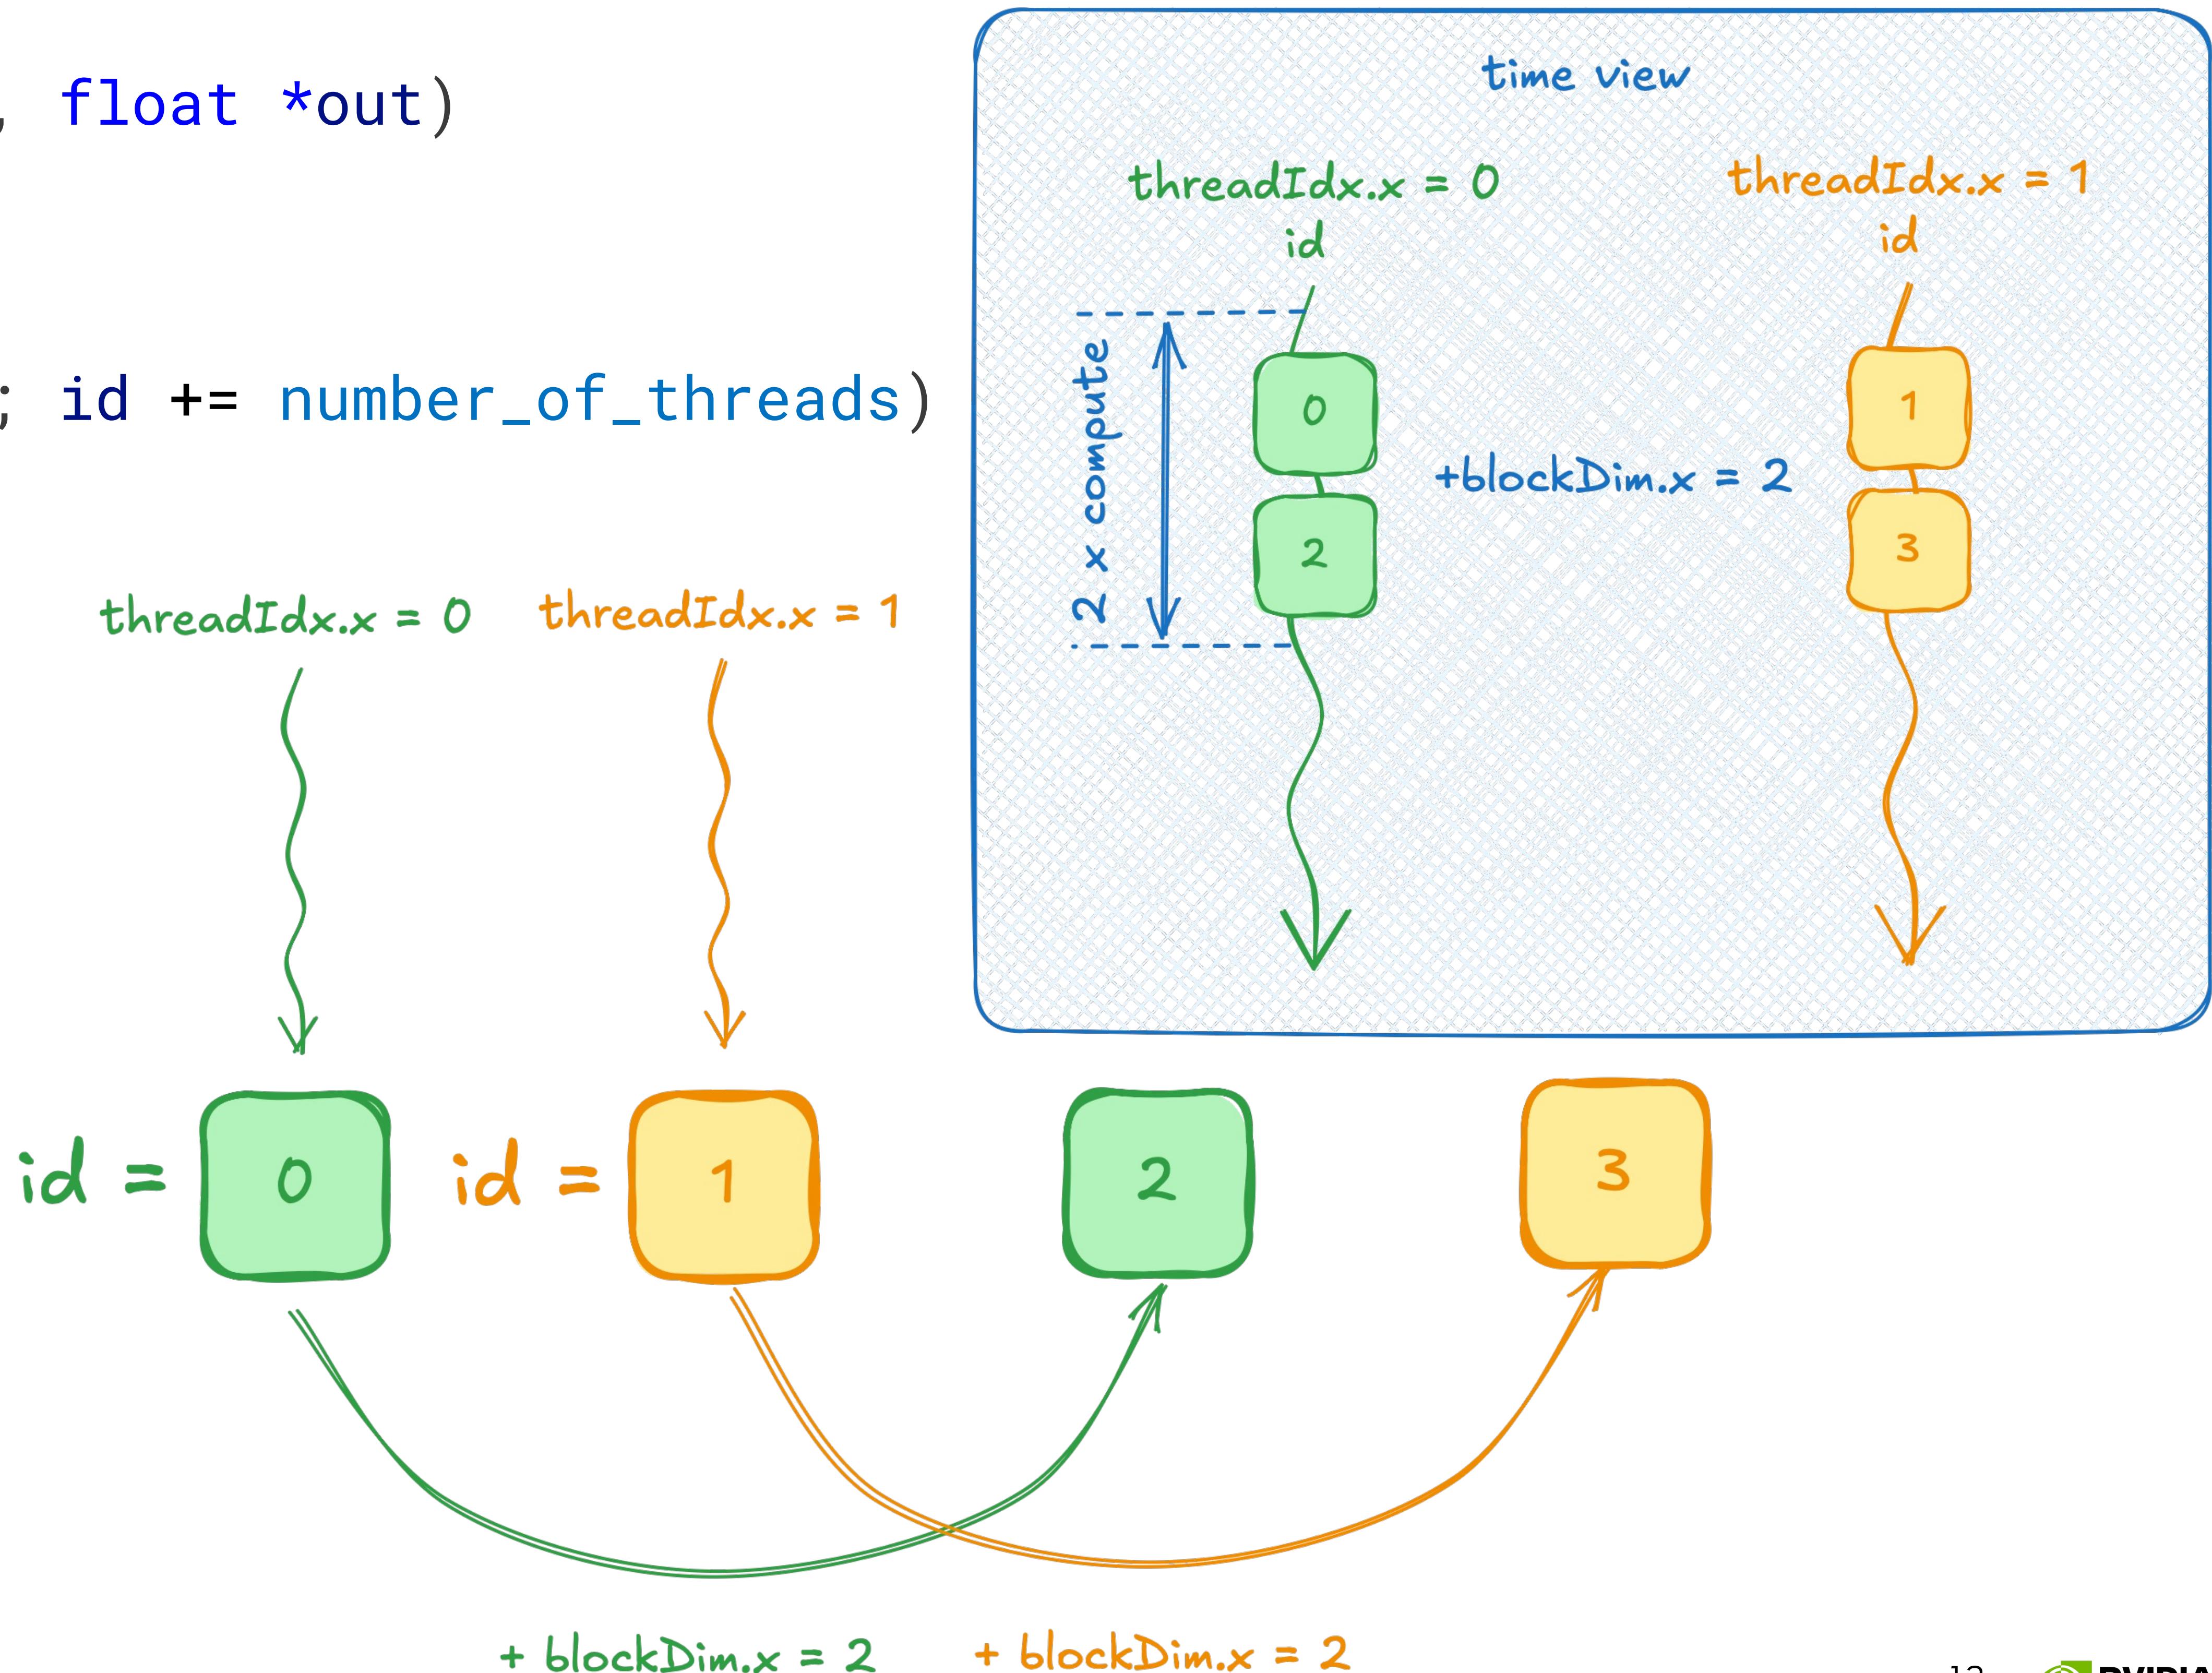
    for (int id = thread_index; id < in.size(); id += number_of_threads)
    {
        out[id] = ach::compute(id, in);
    }
}
```

CUDA extends device code with a few variables:

- `threadIdx.x` stores index of the current thread

10 compute steps take:

- cub: 0.000043
- 1 thread: 4.15
- 2 threads: 2.6



# CUDA Kernels

```
const int number_of_threads = 256;

__global__
void block_kernel(ach::temperature_grid_f in, float *out)
{
    int thread_index = threadIdx.x;

    for (int id = thread_index; id < in.size(); id += number_of_threads)
    {
        out[id] = ach::compute(id, in);
    }
}

void simulate(ach::temperature_grid_f in,
             float *out,
             cudaStream_t stream)
{
    block_kernel<<<1, number_of_threads, 0, stream>>>(in, out);
}
```

- Scheme works!
- More threads process input faster
- Let's just keep adding them

10 compute steps take:

- |                |          |
|----------------|----------|
| • cub:         | 0.000043 |
| • 1 thread:    | 4.15     |
| • 2 threads:   | 2.6      |
| • 256 threads: | 0.037    |

# CUDA Kernels

```
const int number_of_threads = 2048;

__global__
void block_kernel(ach::temperature_grid_f in, float *out)
{
    int thread_index = threadIdx.x;

    for (int id = thread_index; id < in.size(); id += number_of_threads)
    {
        out[id] = ach::compute(id, in);
    }
}

void simulate(ach::temperature_grid_f in,
             float *out,
             cudaStream_t stream)
{
    block_kernel<<<1, number_of_threads, 0, stream>>>(in, out);
}
```

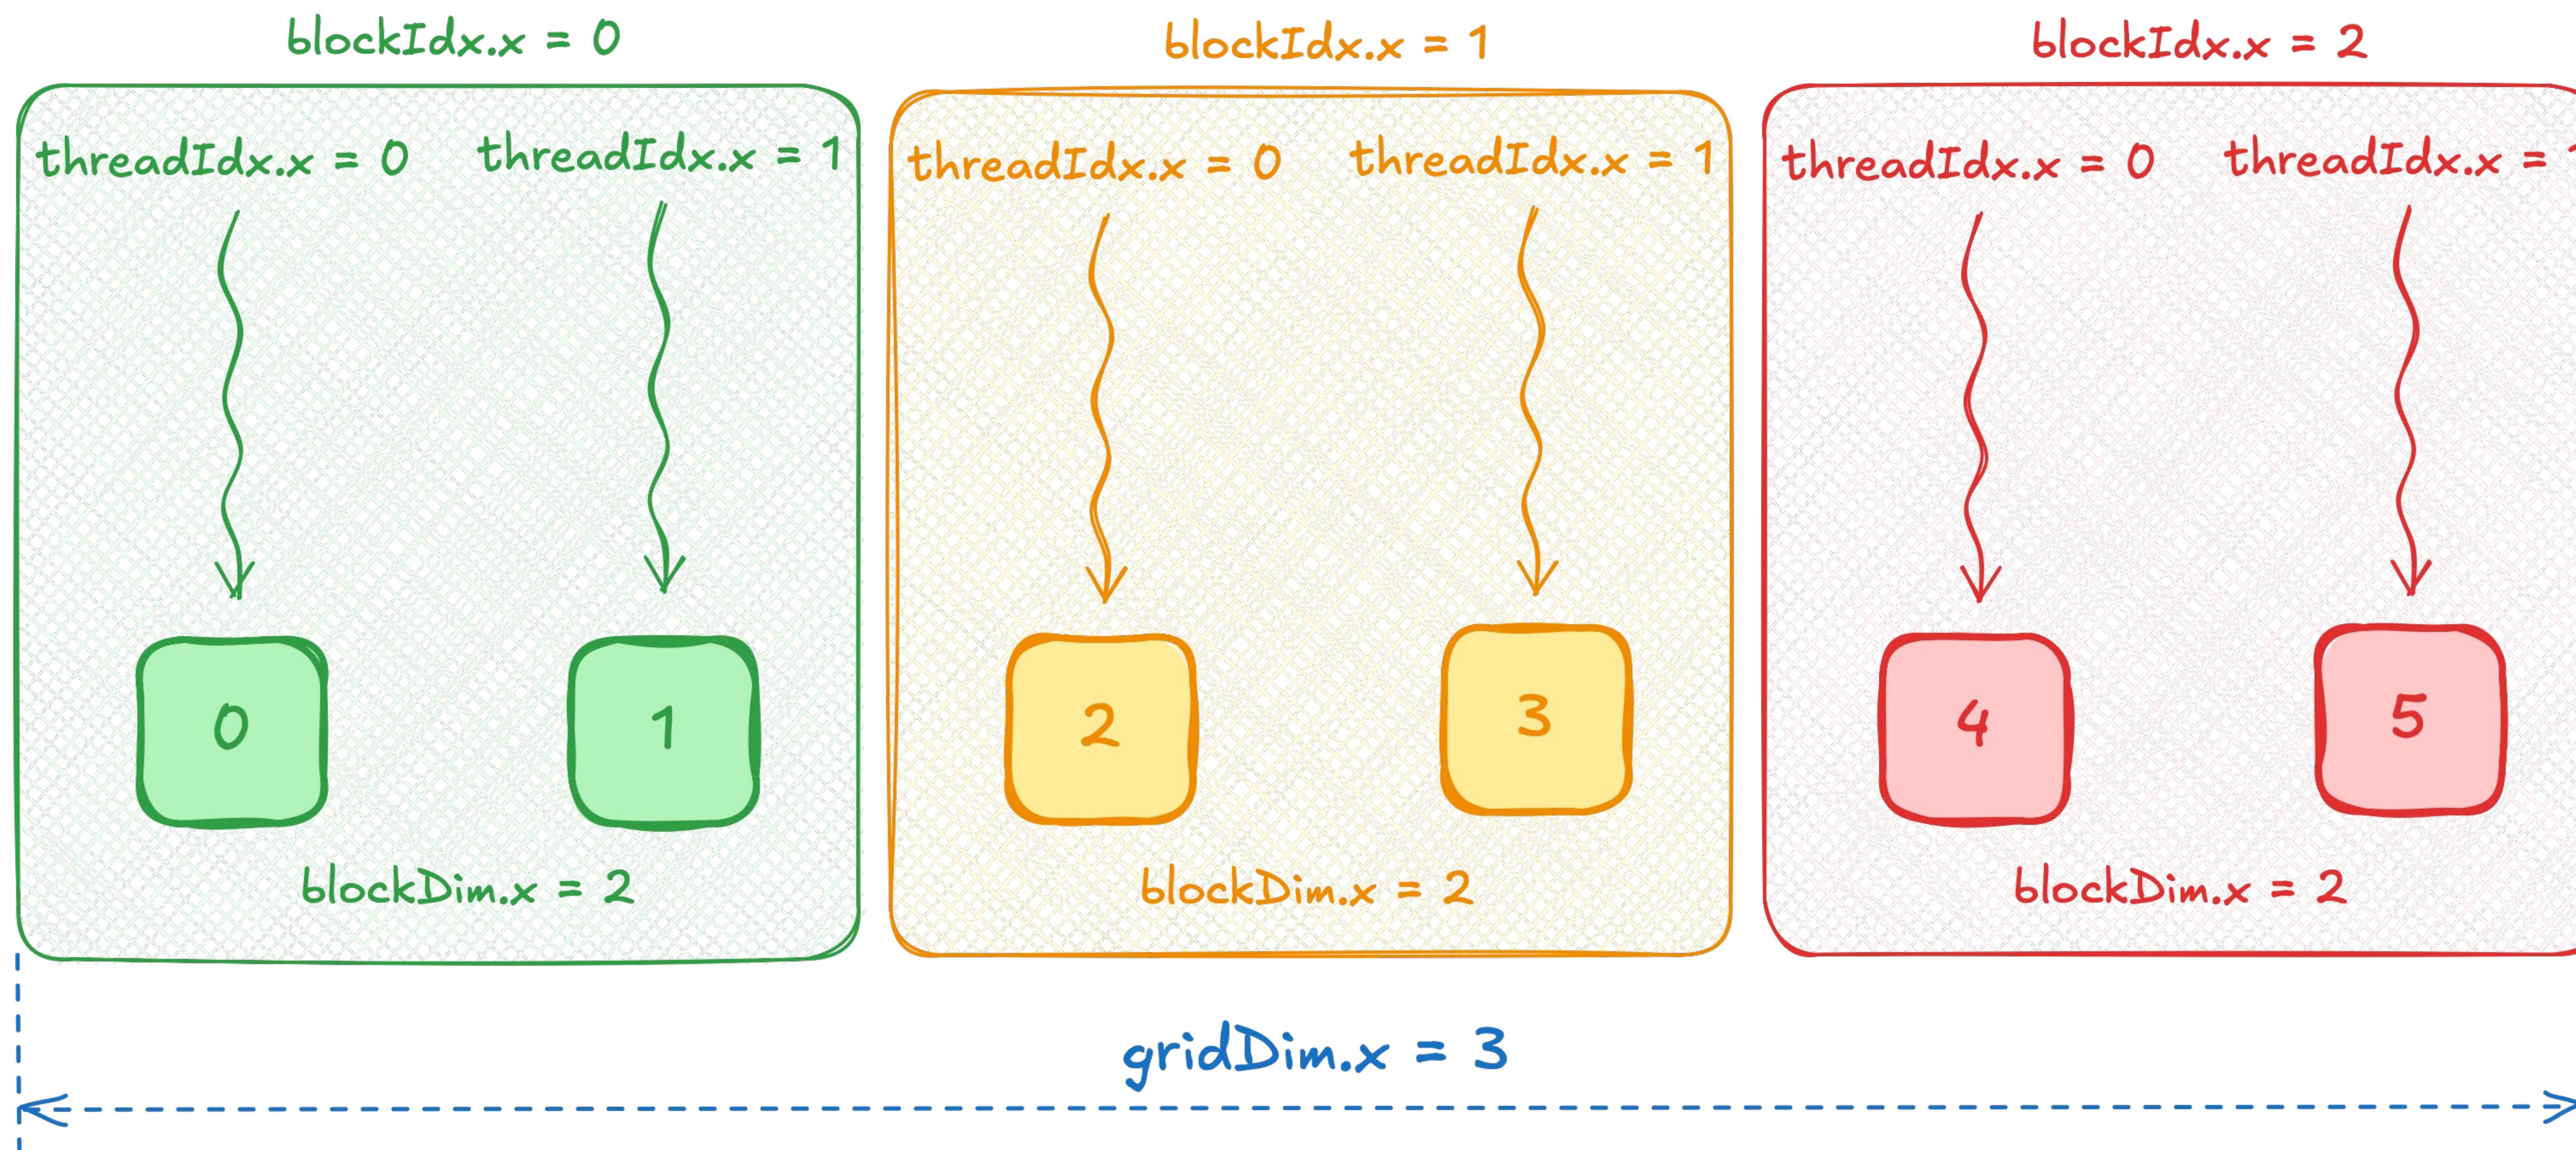
The error is expected:

- You cannot launch more than 1024 threads
- ... in a thread block, which is a HW limit

terminate called after throwing an instance of  
'std::runtime\_error'  
what(): invalid configuration argument

# Missing Context

- threads are grouped in blocks
- all blocks are of the same size (max 1024 threads)
- thread indexing is local within a thread block
- collection of blocks is called grid
- `blockIdx.x` stores index of current block within
- `blockDim.x` stores the number of blocks in the grid
- `gridDim.x` stores the number of blocks in the grid

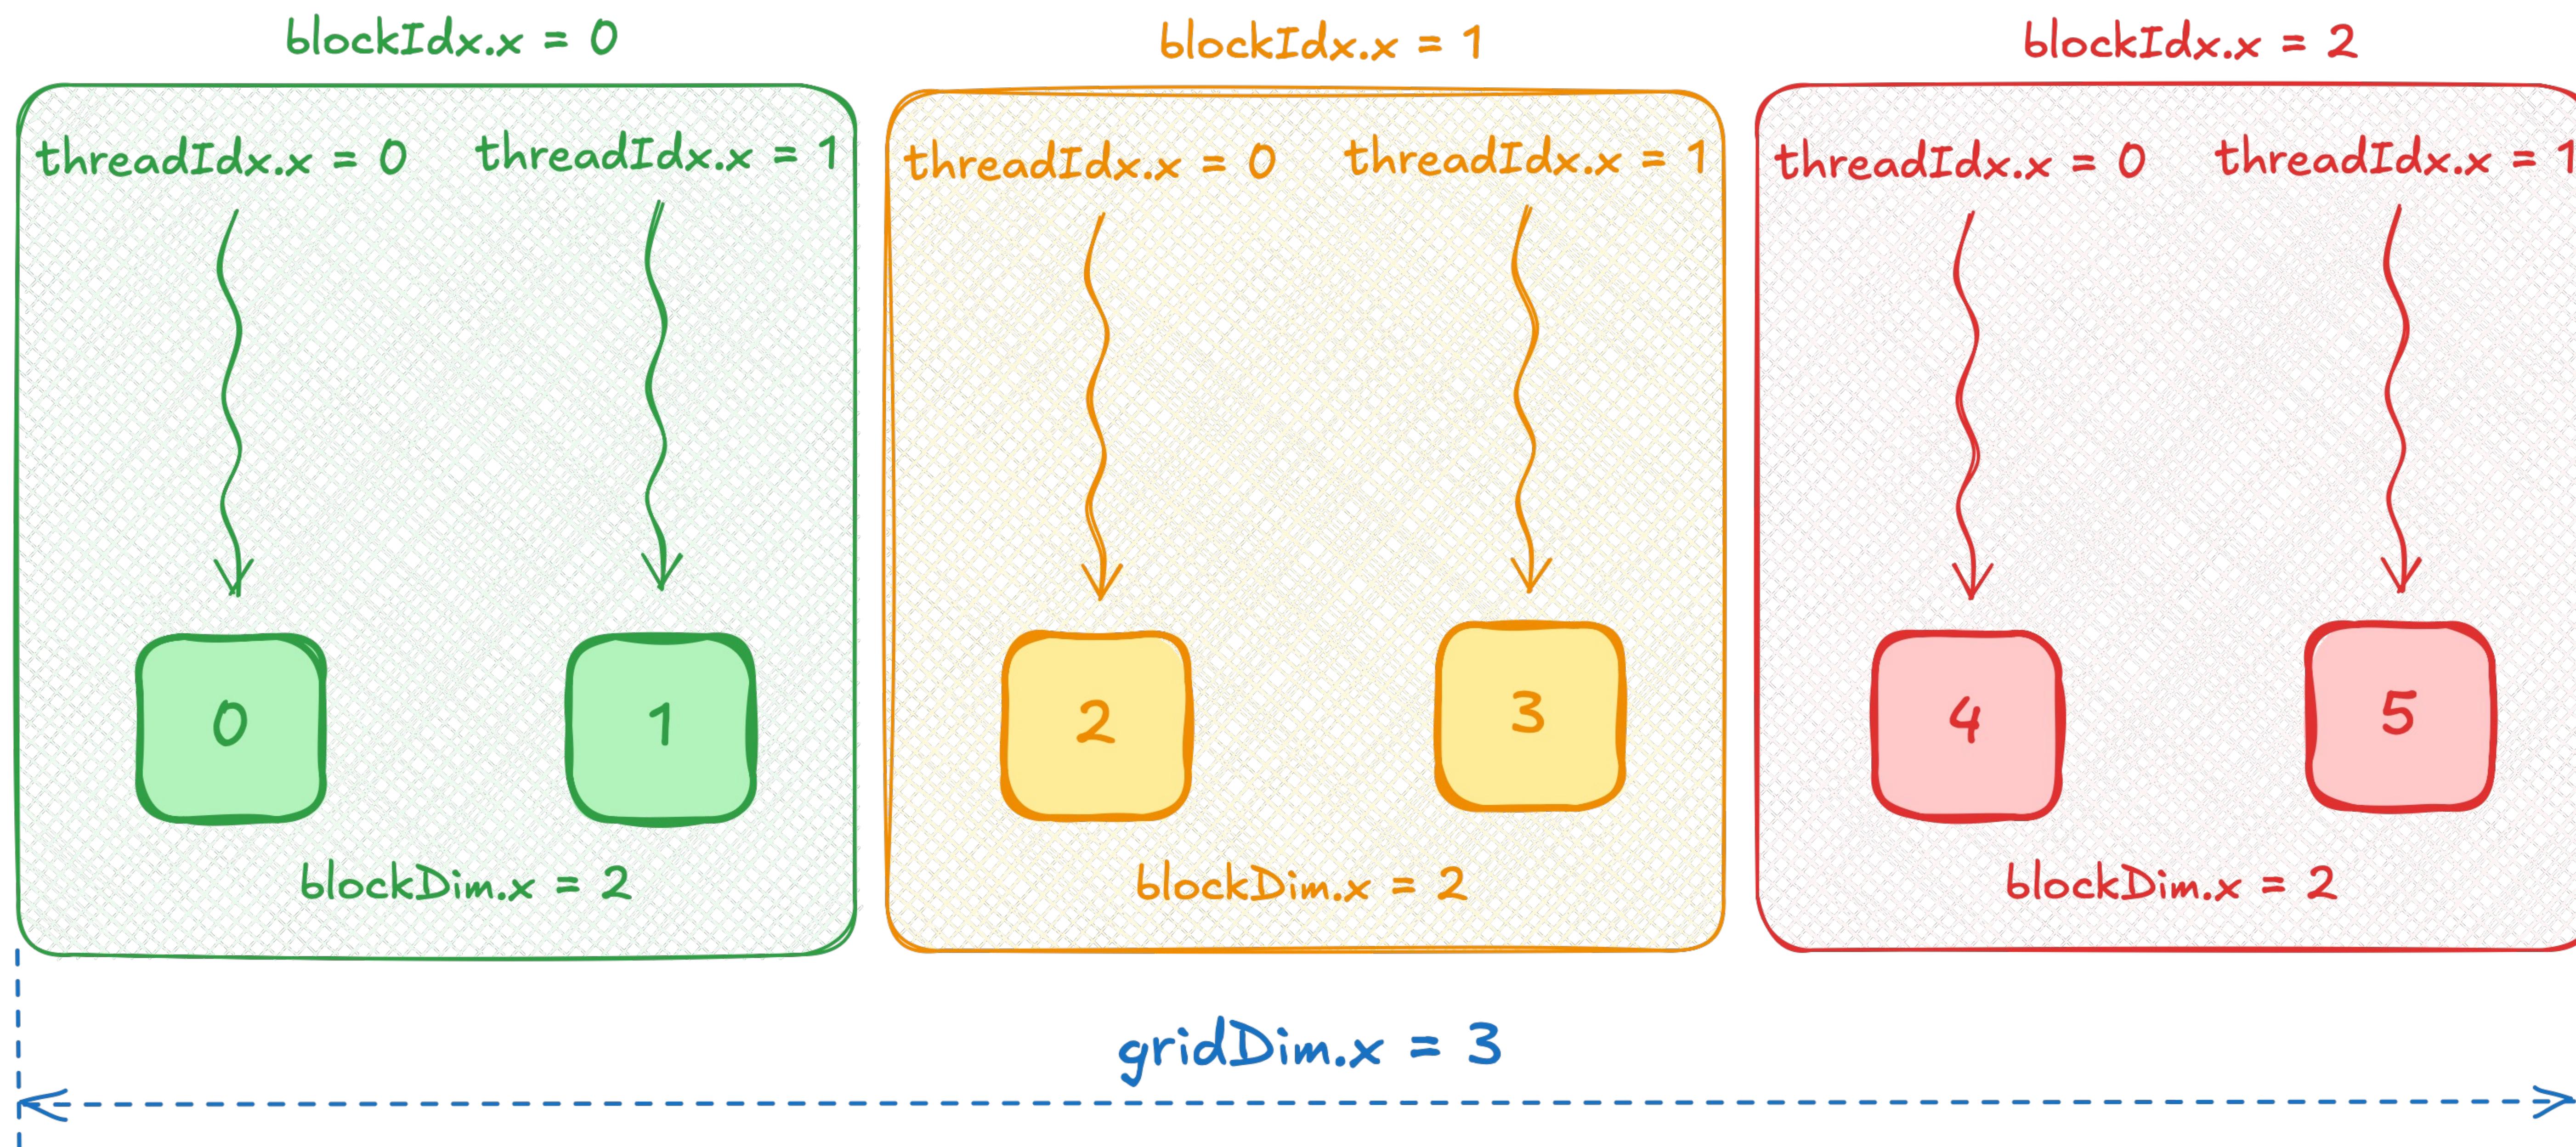


# CUDA Kernels

```
__global__
void grid_kernel(...)

{
    int number_of_threads = blockDim.x * gridDim.x;
    ...
}
```

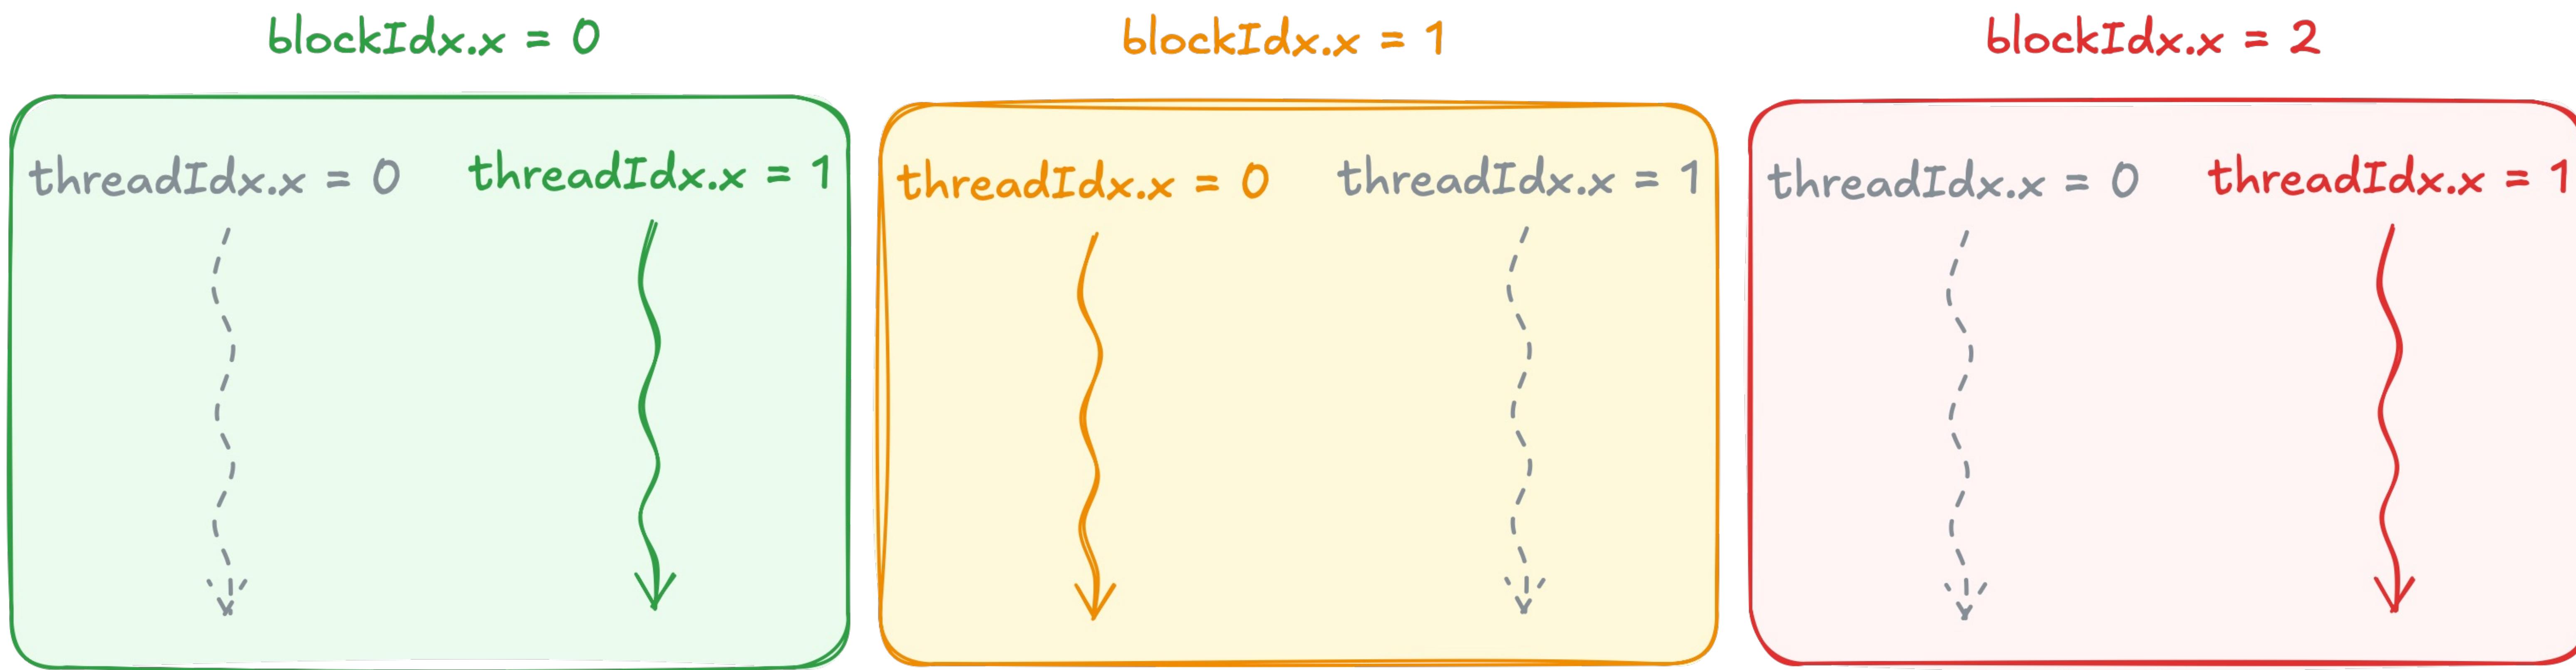
- collection of blocks is called grid
- `blockIdx.x` stores index of current block within grid
- `blockDim.x` stores the number of threads in the block
- `gridDim.x` stores the number of blocks in the grid



# CUDA Kernels

```
__global__  
void grid_kernel(...)  
{  
    int thread_index = blockDim.x * blockIdx.x  
        + threadIdx.x;  
    ...  
}
```

- collection of blocks is called grid
- `blockIdx.x` stores index of current block within grid
- `blockDim.x` stores the number of blocks in the grid
- `gridDim.x` stores the number of blocks in the grid



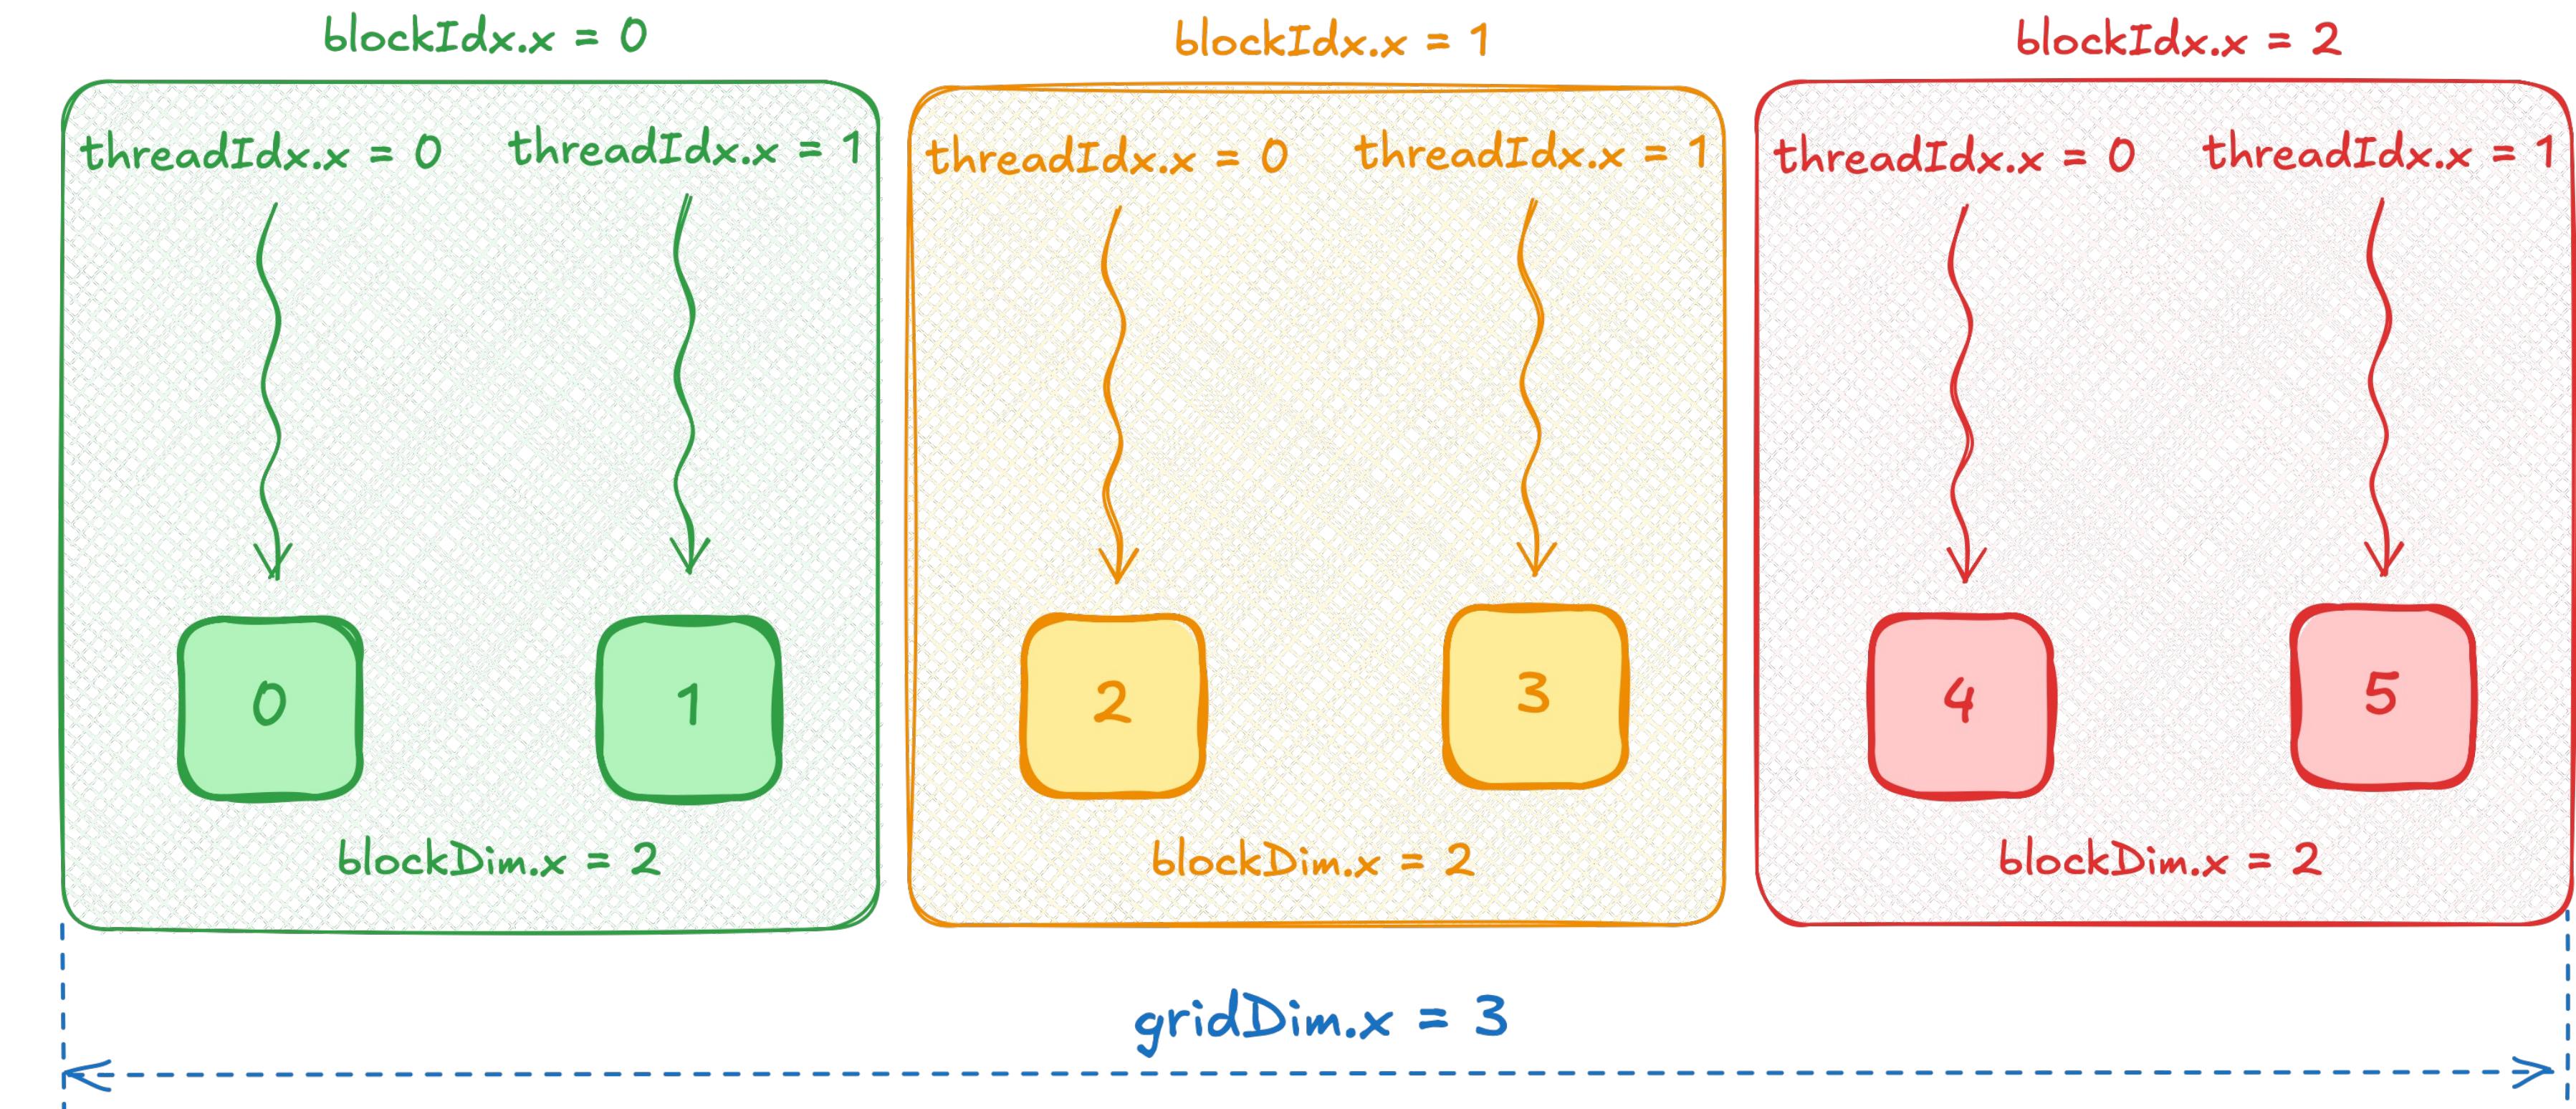
$$thread\_index = 2 * 0 + 1 = 1$$

$$thread\_index = 2 * 1 + 0 = 2$$

$$thread\_index = 2 * 2 + 1 = 5$$

# Which Block Size to Use?

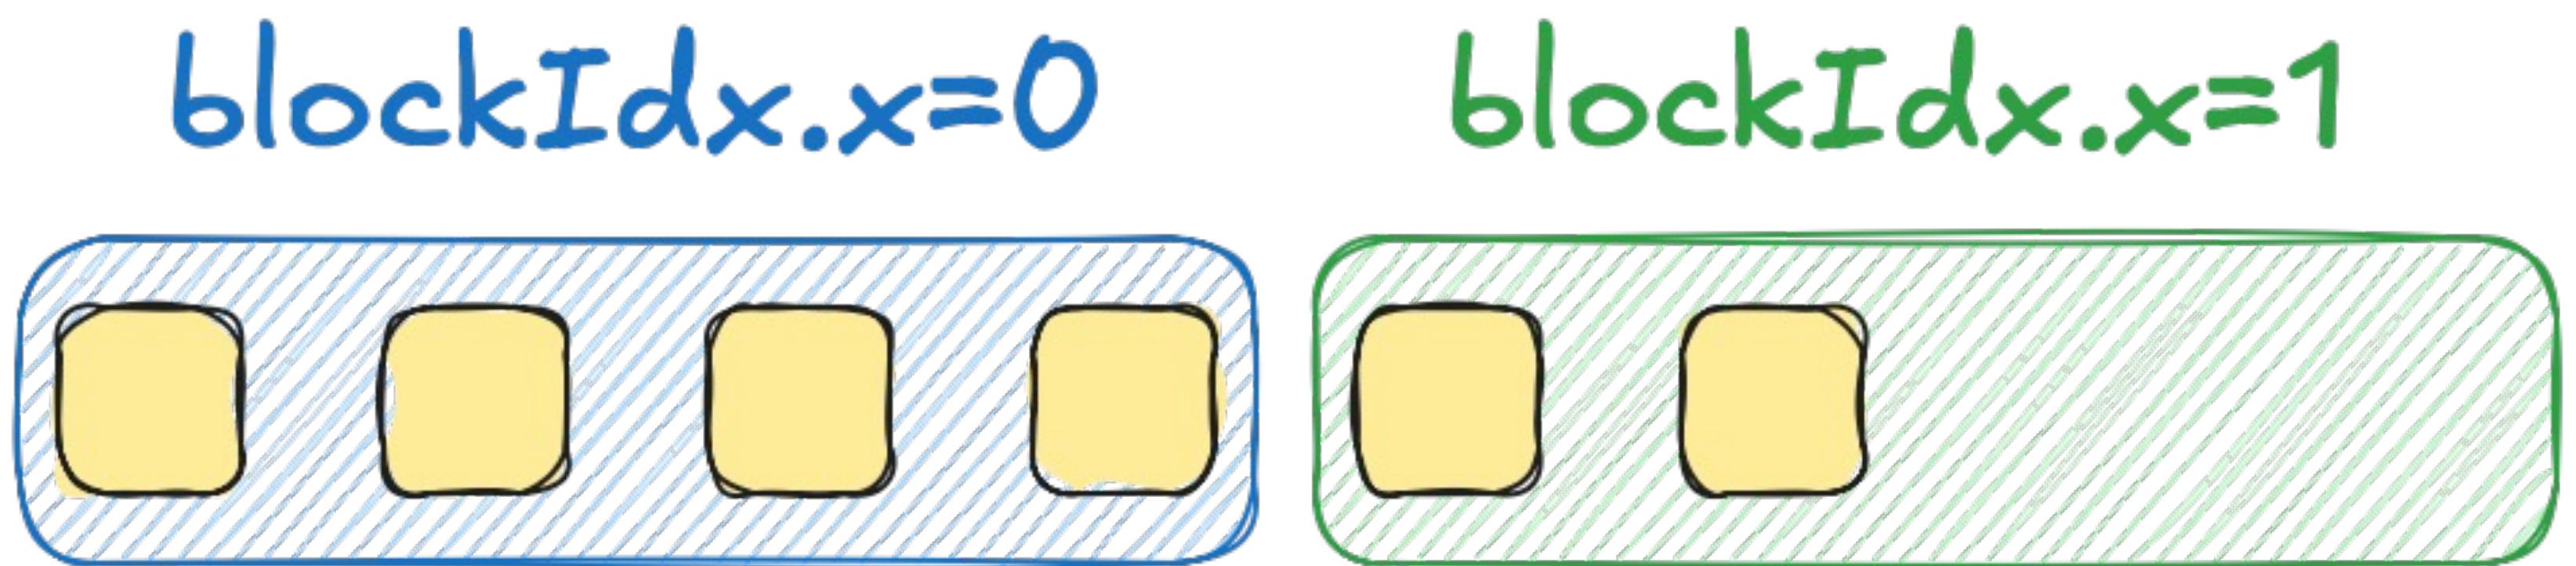
- Thread block size doesn't depend on problem size
- There's no block size that fits all the kernels
- As a rule of thumb:
  - Use block sizes that are multiple of **32**
  - Use **256** as a good default
  - Profile for further tuning



```
void simulate(ach::temperature_grid_f in,
              float *out,
              cudaStream_t stream)
{
    int block_size = 256;
    ...
    grid_kernel<<<grid_size, block_size, 0, stream>>>(in, out);
}
```

# Which Grid Size to Use?

- We need the number number of thread blocks (grid size)
- Grid size frequently depend on problem size
- As a rule of thumb, use `cuda::ceil_div` to compute grid size



```
void simulate(ach::temperature_grid_f in,
              float *out,
              cudaStream_t stream)
{
    int block_size = 256;
    int grid_size = cuda::ceil_div(in.size(), block_size);

    grid_kernel<<<grid_size, block_size, 0, stream>>>(in, out);
}
```

```
int ceil_div(int a, int b)
{
    return (a + b - 1) / b;
}
```

# CUDA Kernels

```
__global__
void grid_kernel(ach::temperature_grid_f in, float *out)
{
    int thread_index = blockDim.x * blockIdx.x + threadIdx.x;
    int number_of_threads = blockDim.x * gridDim.x;

    for (int id = thread_index; id < in.size(); id += number_of_threads)
    {
        out[id] = ach::compute(id, in);
    }
}

void simulate(ach::temperature_grid_f in,
              float *out,
              cudaStream_t stream)
{
    int block_size = 256;
    int grid_size = cuda::ceil_div(in.size(), block_size);

    grid_kernel<<<grid_size, block_size, 0, stream>>>(in, out);
}
```

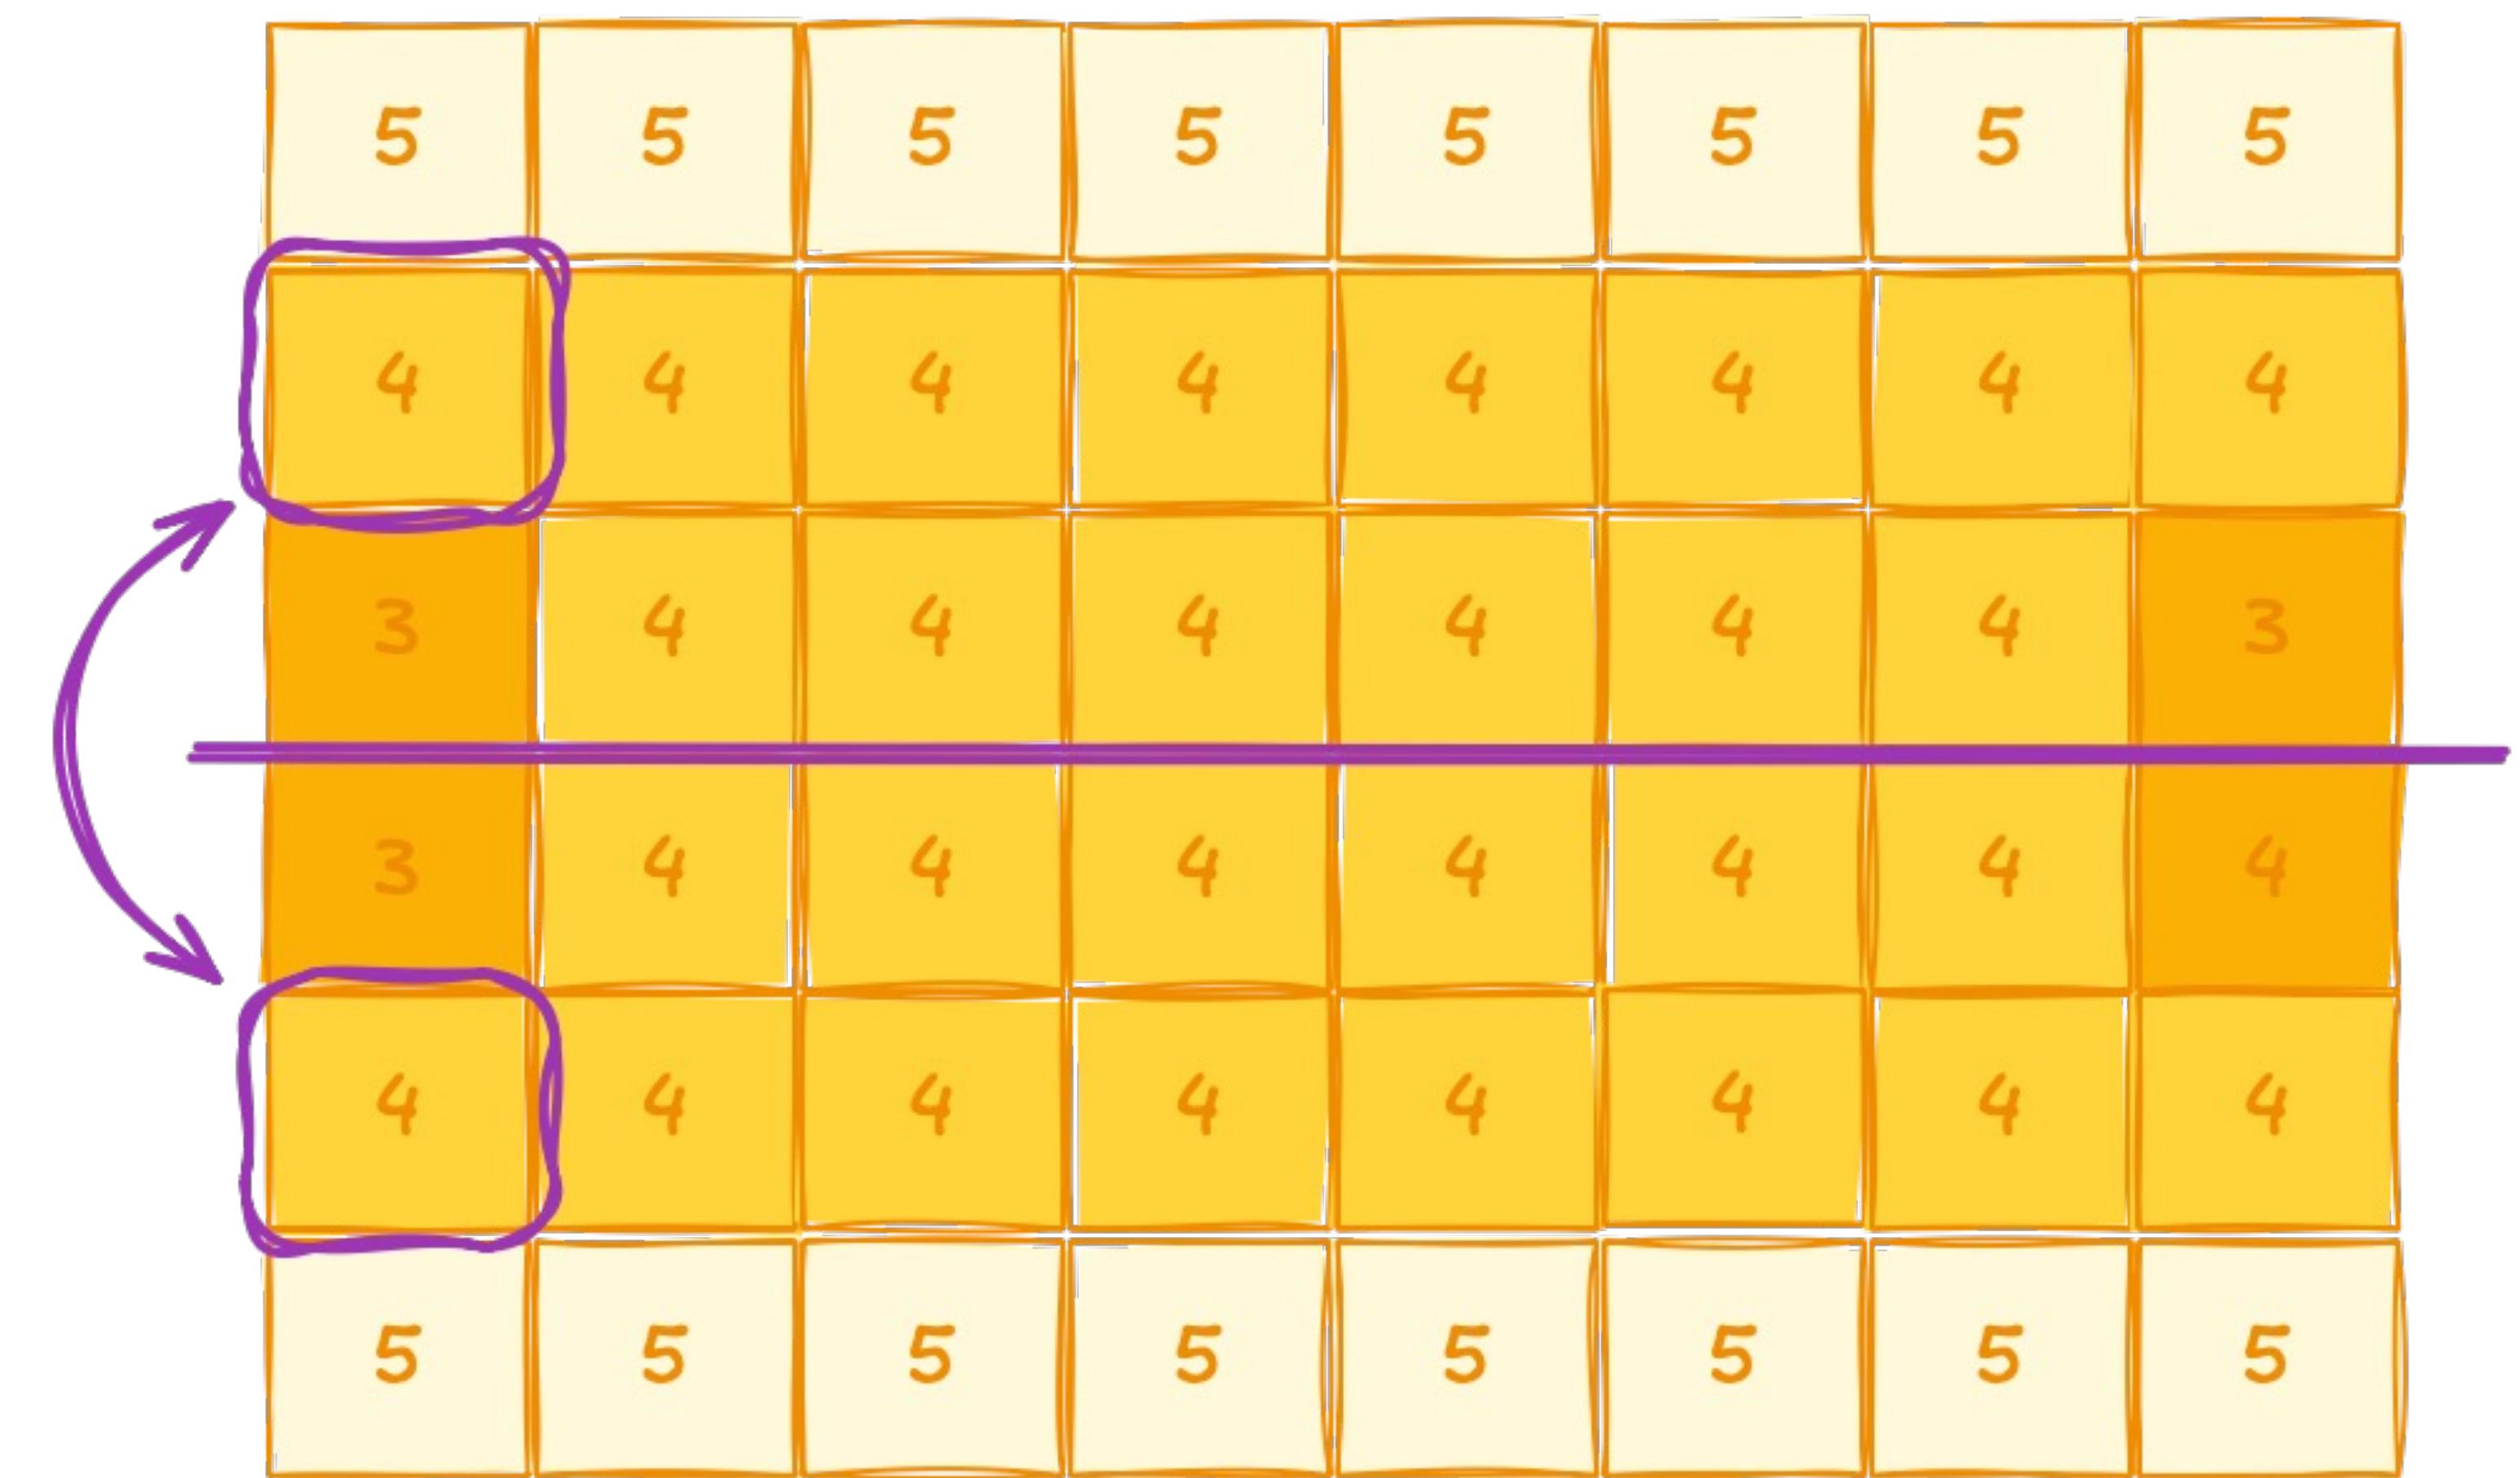
- 10 compute steps take:
  - cub: 0.000043
  - 1 thread: 4.15
  - 2 threads: 2.6
  - 256 threads: 0.037
  - 5'120'000 threads: 0.0003

# Exercise: Detect Asymmetry

- Boundary conditions are symmetric
- Remaining cells should be symmetric as well
- Rewrite the function below as a kernel to verify that

```
void symmetry_check(ach::temperature_grid_f temp, int row)
{
    int column = 0;

    float top = temp(row, column);
    float bottom = temp(temp.extent(0) - 1 - row, column);
    float diff = abs(top - bottom);
    if (diff > 0.1) {
        printf("Error: asymmetry in %d\n", column);
    }
}
```



# Exercise: Detect Asymmetry

## Solution

```
__global__
void symmetry_check_kernel(ach::temperature_grid_f temp, int row)
{
    int column = 0;

    float top = temp(row, column);
    float bottom = temp(temp.extent(0) - 1 - row, column);
    float diff = abs(top - bottom);
    if (diff > 0.1) {
        printf("Error: asymmetry in %d\n", column);
    }
}

void symmetry_check(ach::temperature_grid_f temp, cudaStream_t stream)
{
    int target_row = 0;
    symmetry_check_kernel<<<1, 1, 0, stream>>>(temp, target_row);
}
```

Annotate function with  
`__global__` specifier

# Exercise: Detect Asymmetry

## Solution

```
--global__
void symmetry_check_kernel(ach::temperature_grid_f temp, int row)
{
    int column = 0;

    float top = temp(row, column);
    float bottom = temp(temp.extent(0) - 1 - row, column);
    float diff = abs(top - bottom);
    if (diff > 0.1) {
        printf("Error: asymmetry in %d\n", column);
    }
}

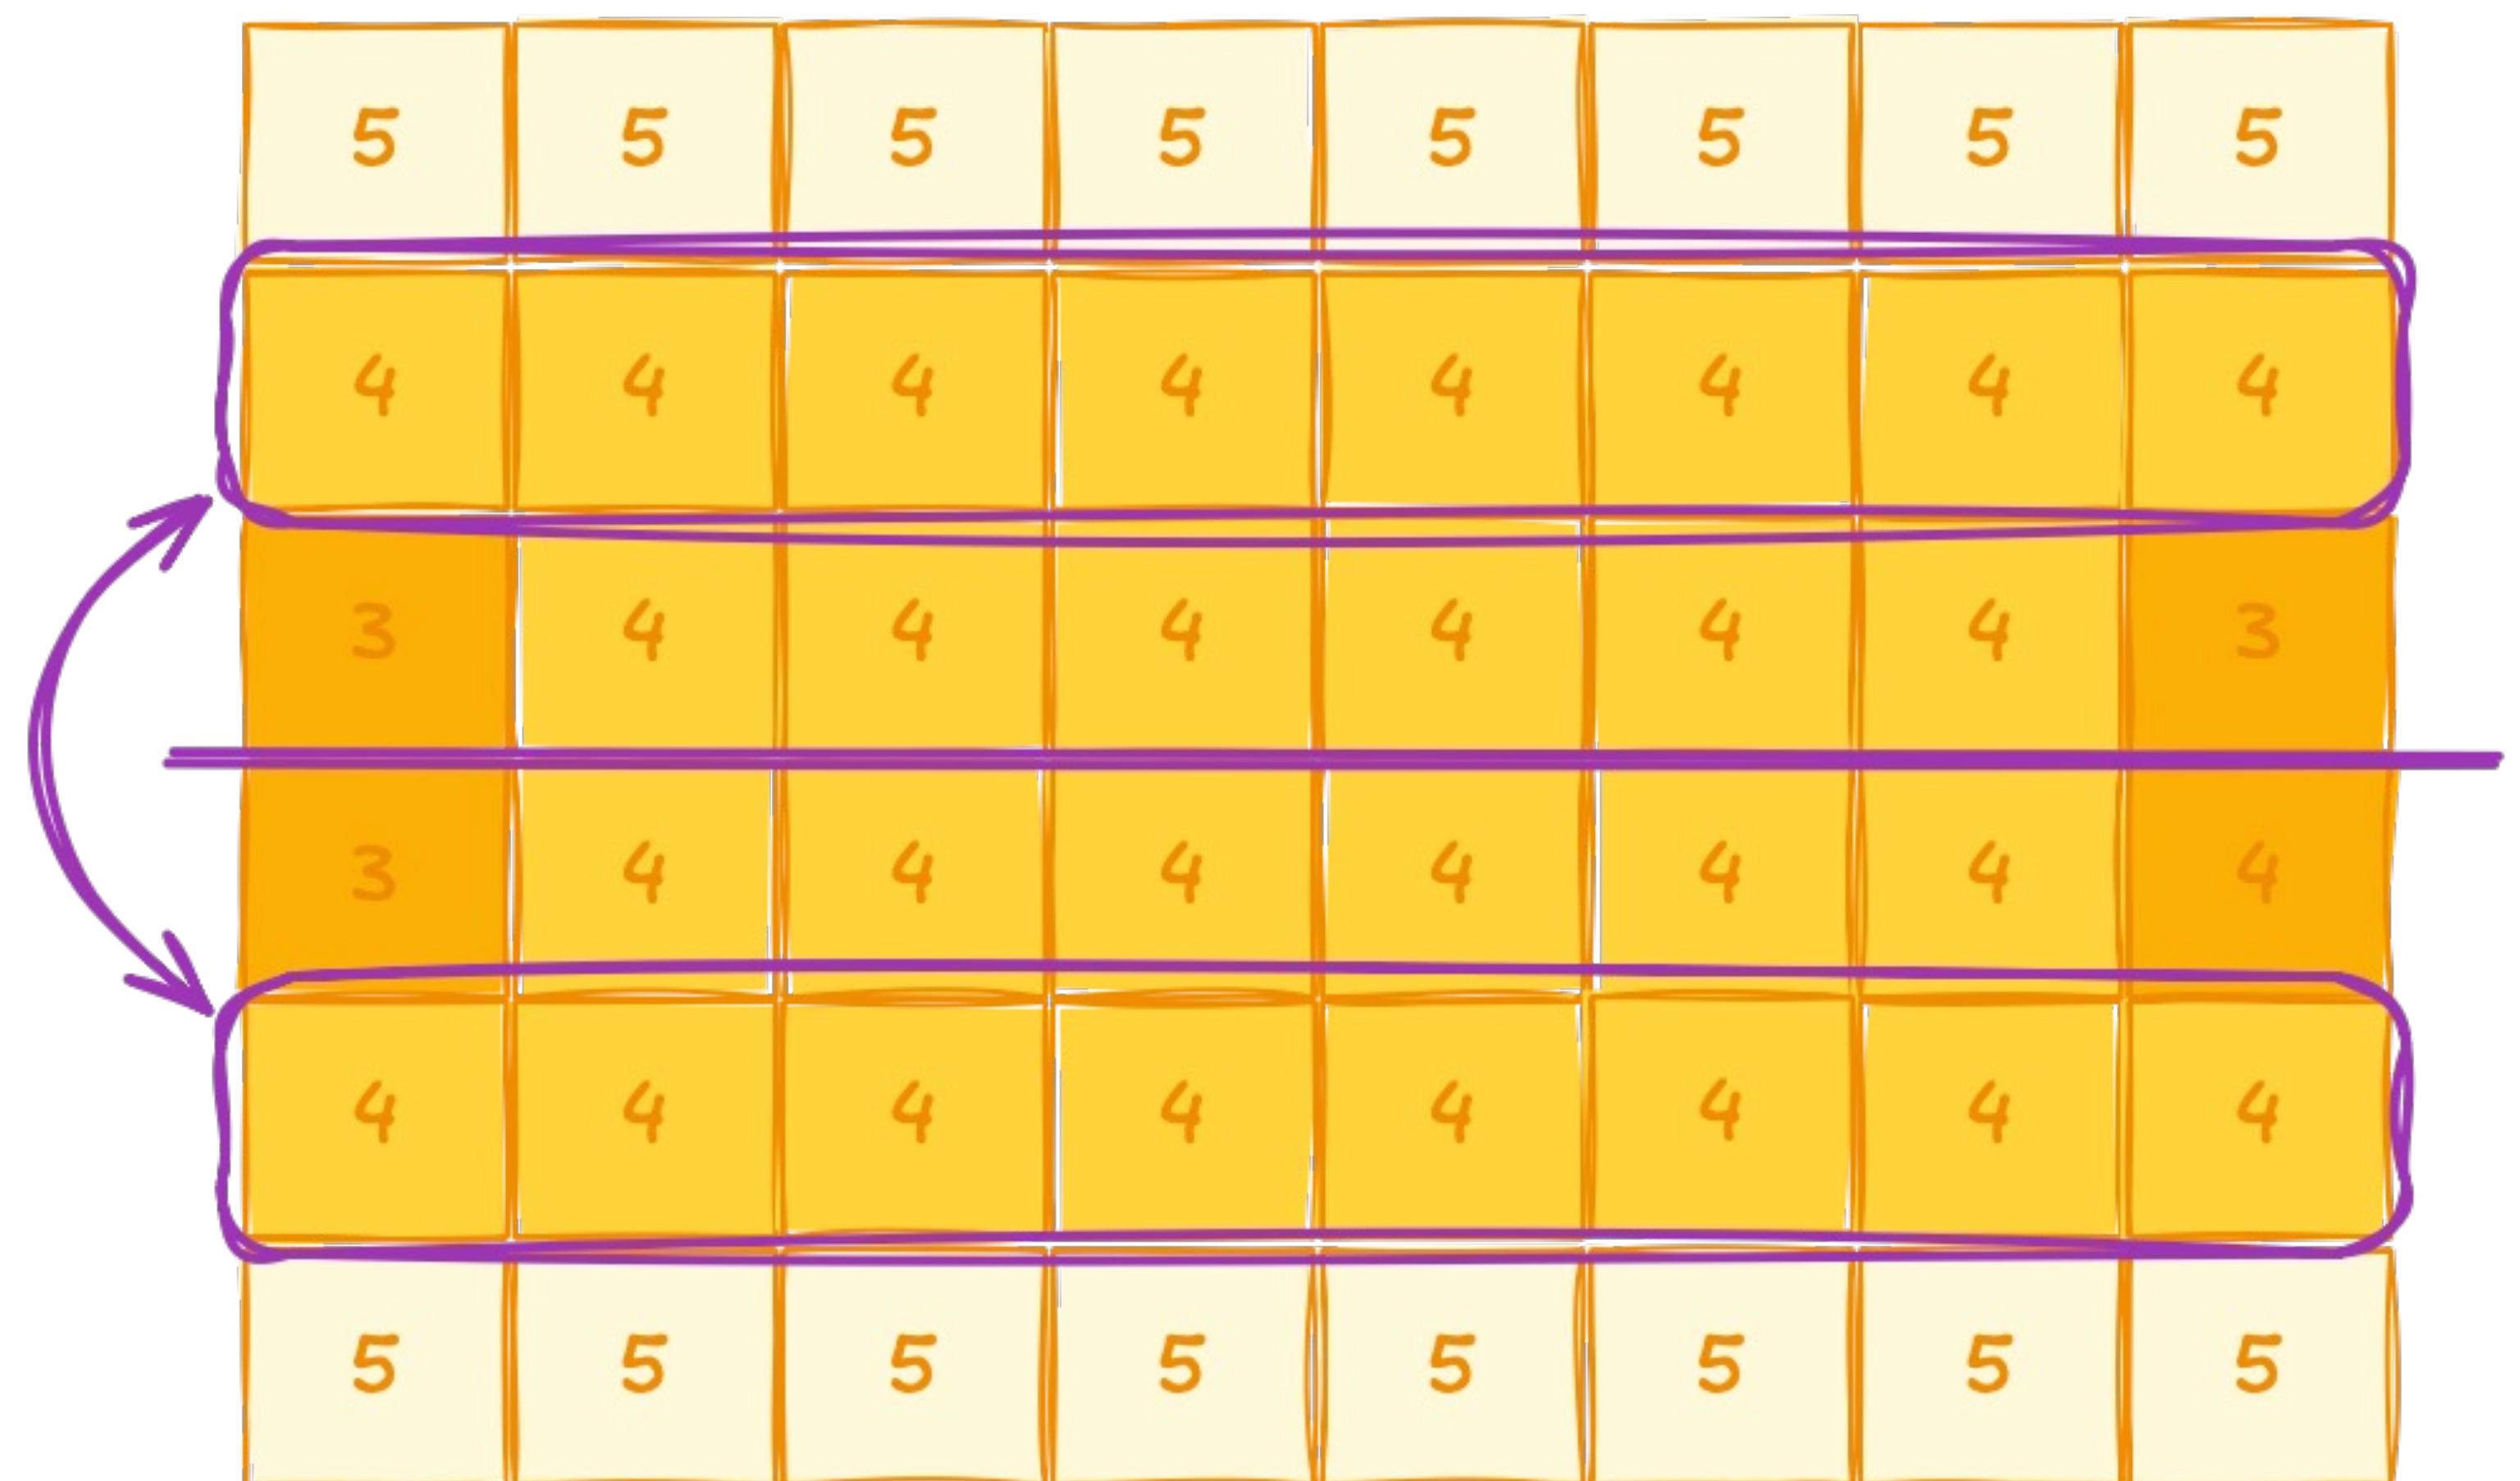
void symmetry_check(ach::temperature_grid_f temp, cudaStream_t stream)
{
    int target_row = 0;
    symmetry_check_kernel<<<1, 1, 0, stream>>>(temp, target_row); | Use triple chevron to launch
}                                            | the kernel
```

# Exercise: Detect Row Asymmetry

- Entire row should be symmetric, not just one cell
- Rewrite the kernel below to verify that it's the case
- Assign each thread to verify exactly one column

```
__global__
void symmetry_check(ach::temperature_grid_f temp, int row)
{
    int column = 0;

    float top = temp(row, column);
    float bottom = temp(temp.extent(0) - 1 - row, column);
    float diff = abs(top - bottom);
    if (diff > 0.1) {
        printf("Error: asymmetry in %d\n", column);
    }
}
```



# Exercise: Detect Row Asymmetry

## Solution

```
__global__ void row_symmetry_check_kernel(ach::temperature_grid_f temp, int row)
{
    int column = blockIdx.x * blockDim.x + threadIdx.x;

    if (abs(temp(row, column) - temp(temp.extent(0) - 1 - row, column)) > 0.1)
    {
        printf("Error: asymmetry in %d / %d\n", column, temp.extent(1));
    }
}

void symmetry_check(ach::temperature_grid_f temp, cudaStream_t stream)
{
    int width      = temp.extent(1);
    int block_size = 256;
    int grid_size  = cuda::ceil_div(width, block_size);

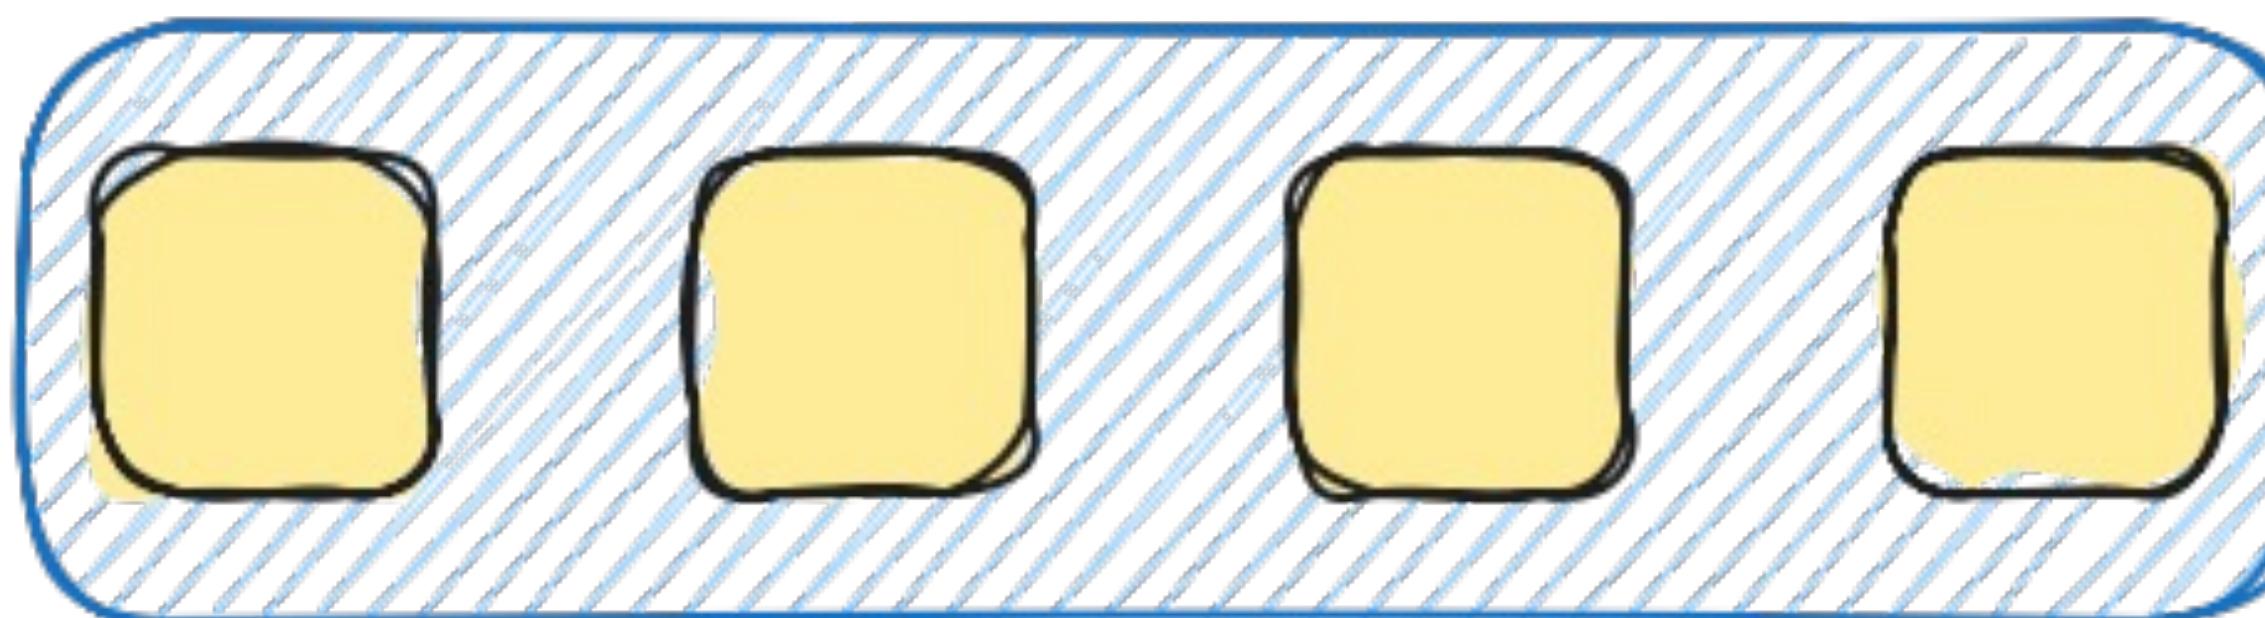
    int target_row = 0;
    row_symmetry_check_kernel<<<grid_size, block_size, 0, stream>>>(temp, target_row);
}
```

# Exercise: Detect Row Asymmetry

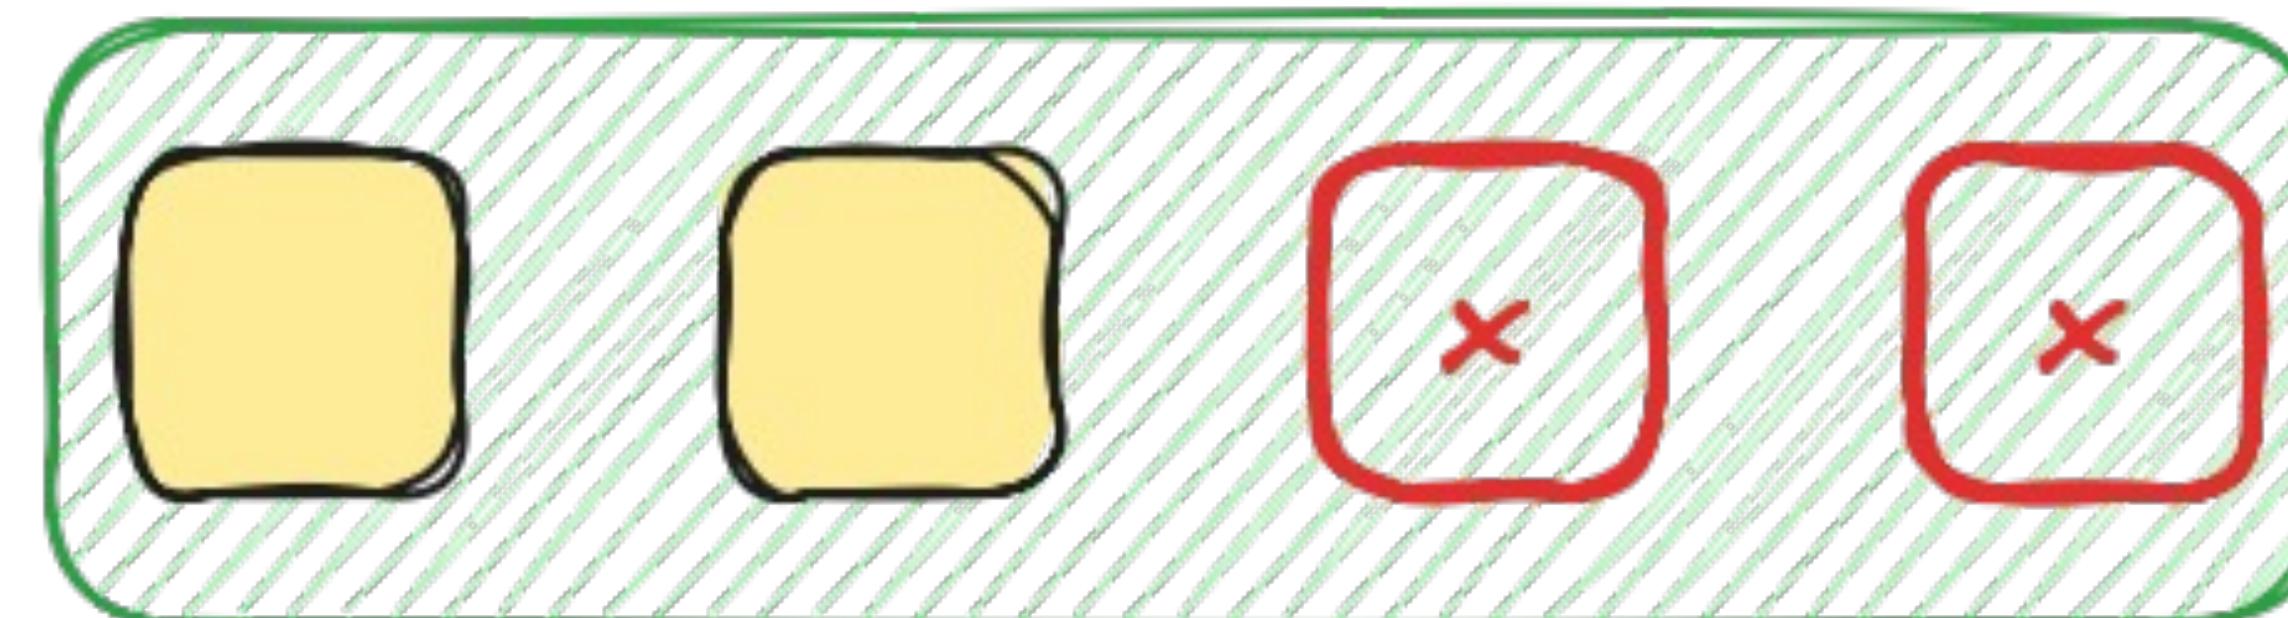
Bug in the solution

```
--global__ void symmetry_check_kernel(ach::temperature_grid_f temp, int row)
{
    int column = blockIdx.x * blockDim.x + threadIdx.x;
```

blockIdx.x=0



blockIdx.x=1

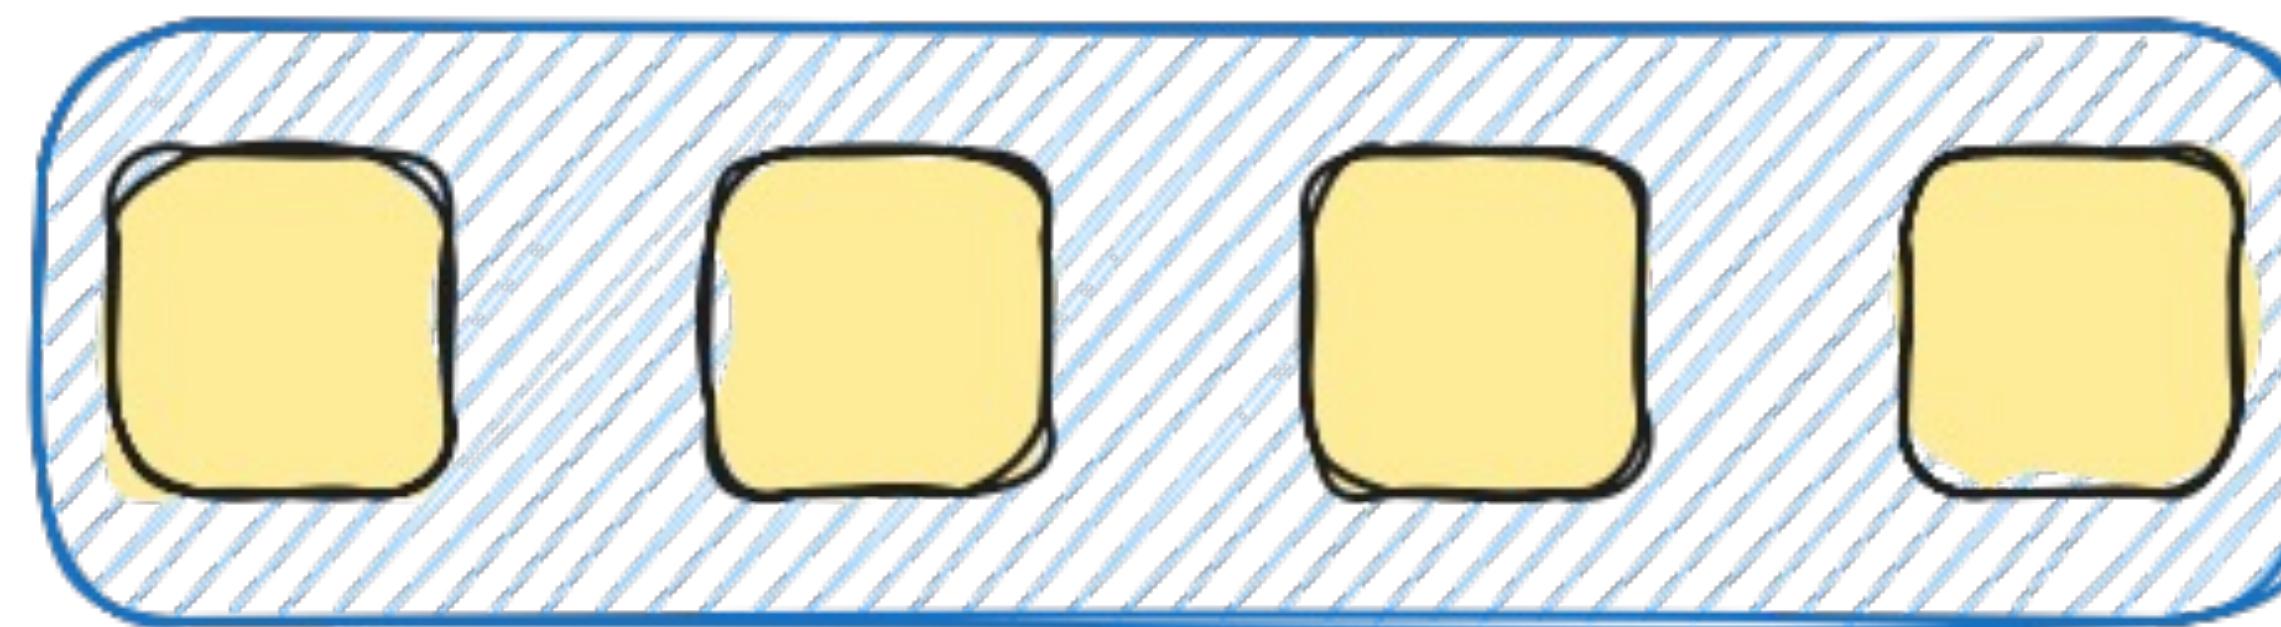


```
int width      = temp.extent(1);
int block_size = 1024;
int grid_size  = cuda::ceil_div(width, block_size);

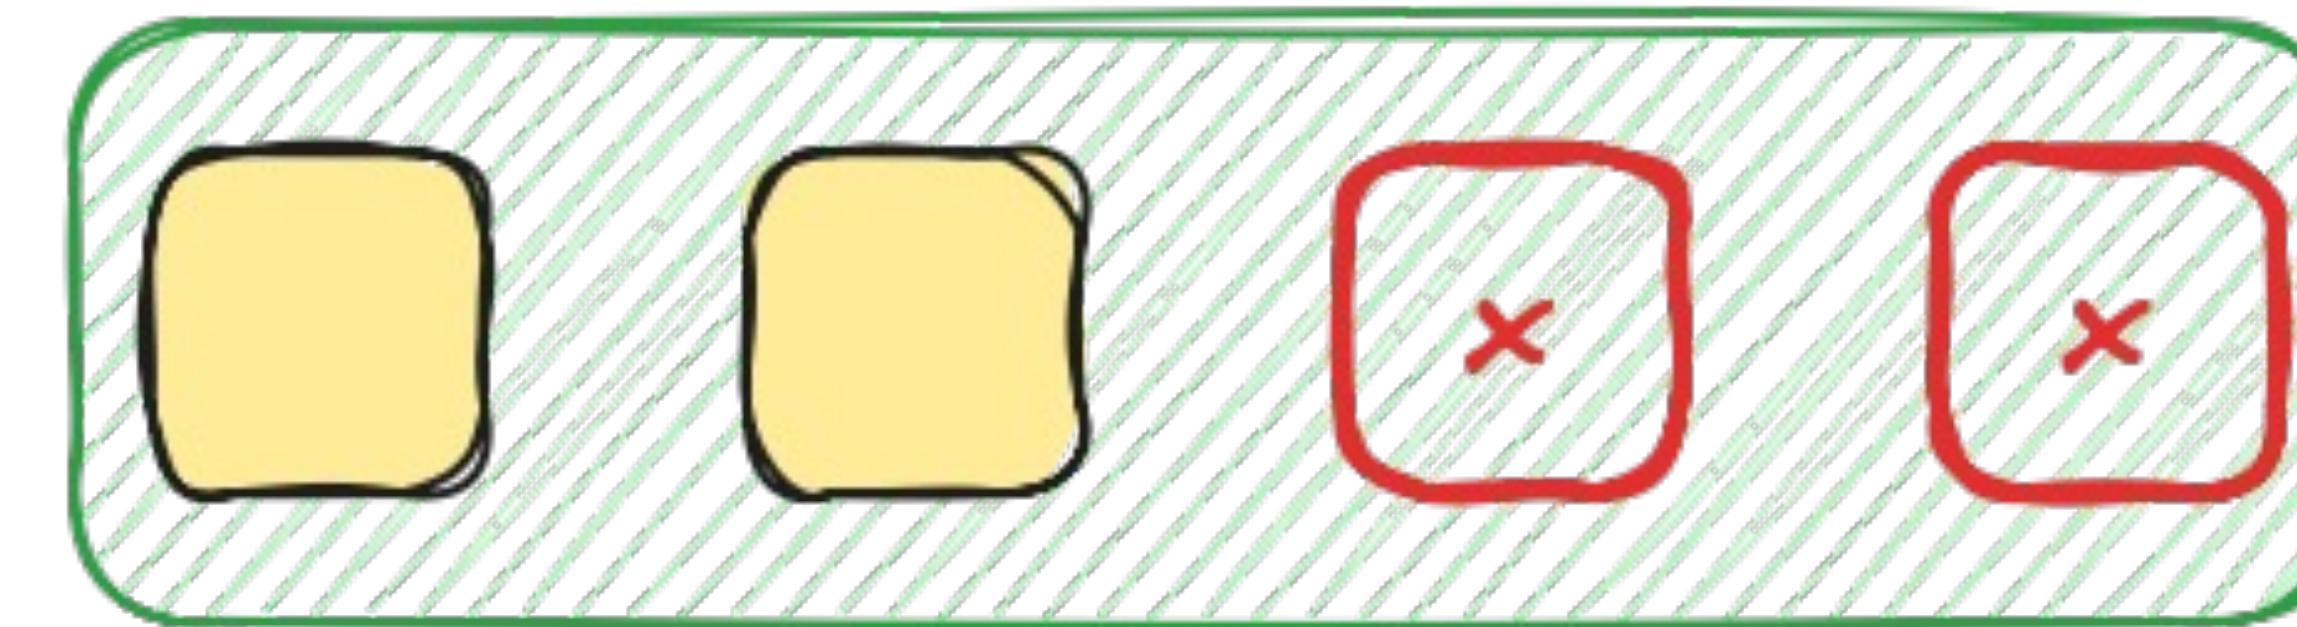
int target_row = 0;
symmetry_check_kernel<<<grid_size, block_size, 0, stream>>>(temp, target_row);
}
```

# Fixed Row Symmetry Check

blockIdx.x=0



blockIdx.x=1



```
--global__ void symmetry_check_kernel(ach::temperature_grid_f temp, int row)
{
    int column = blockIdx.x * blockDim.x + threadIdx.x;

    if (column < temp.extent(1)) | Check if a given thread is within bounds of problem size
    {
        if (abs(temp(row, column) - temp(temp.extent(0) - 1 - row, column)) > 0.1)
        {
            printf("Error: asymmetry in %d / %d\n", column, temp.extent(1));
        }
    }
}
```

# Compute Sanitizer

```
nvcc --extended-lambda -g -G -o /tmp/a.out ./src/row_symmetry_check.cu
```

```
compute-sanitizer /tmp/a.out
```

```
===== Invalid __global__ read of size 4 bytes
===== at symmetry_check_kernel(cuda::std::mdspan<>, int) in row_symmetry_check.cu:7
===== by thread (905,0,0) in block (4,0,0)
===== Address 0x7c5287588004 is out of bounds
===== and is 5 bytes after the nearest allocation at 0x7c5286200000 of size 20,480,000 bytes
===== Saved host backtrace up to driver entry point at kernel launch time
```

```
__global__ void symmetry_check_kernel(ach::temperature_grid_f temp, int row)
{
    int column = blockIdx.x * blockDim.x + threadIdx.x;

    if (abs(temp(row, column) - temp(temp.extent(0) - 1 - row, column)) > 0.1)
    {
        printf("Error: asymmetry in %d / %d\n", column, temp.extent(1));
    }
}
```

- One way to detect the bug would be compute sanitizer

# Assertions

```
#define CCCL_ENABLE_ASSERTIONS
#include <cuda/std/mdspan>

__global__ void row_symmetry_check_kernel(ach::temperature_grid_f temp, int row)
{
    int column = blockIdx.x * blockDim.x + threadIdx.x;

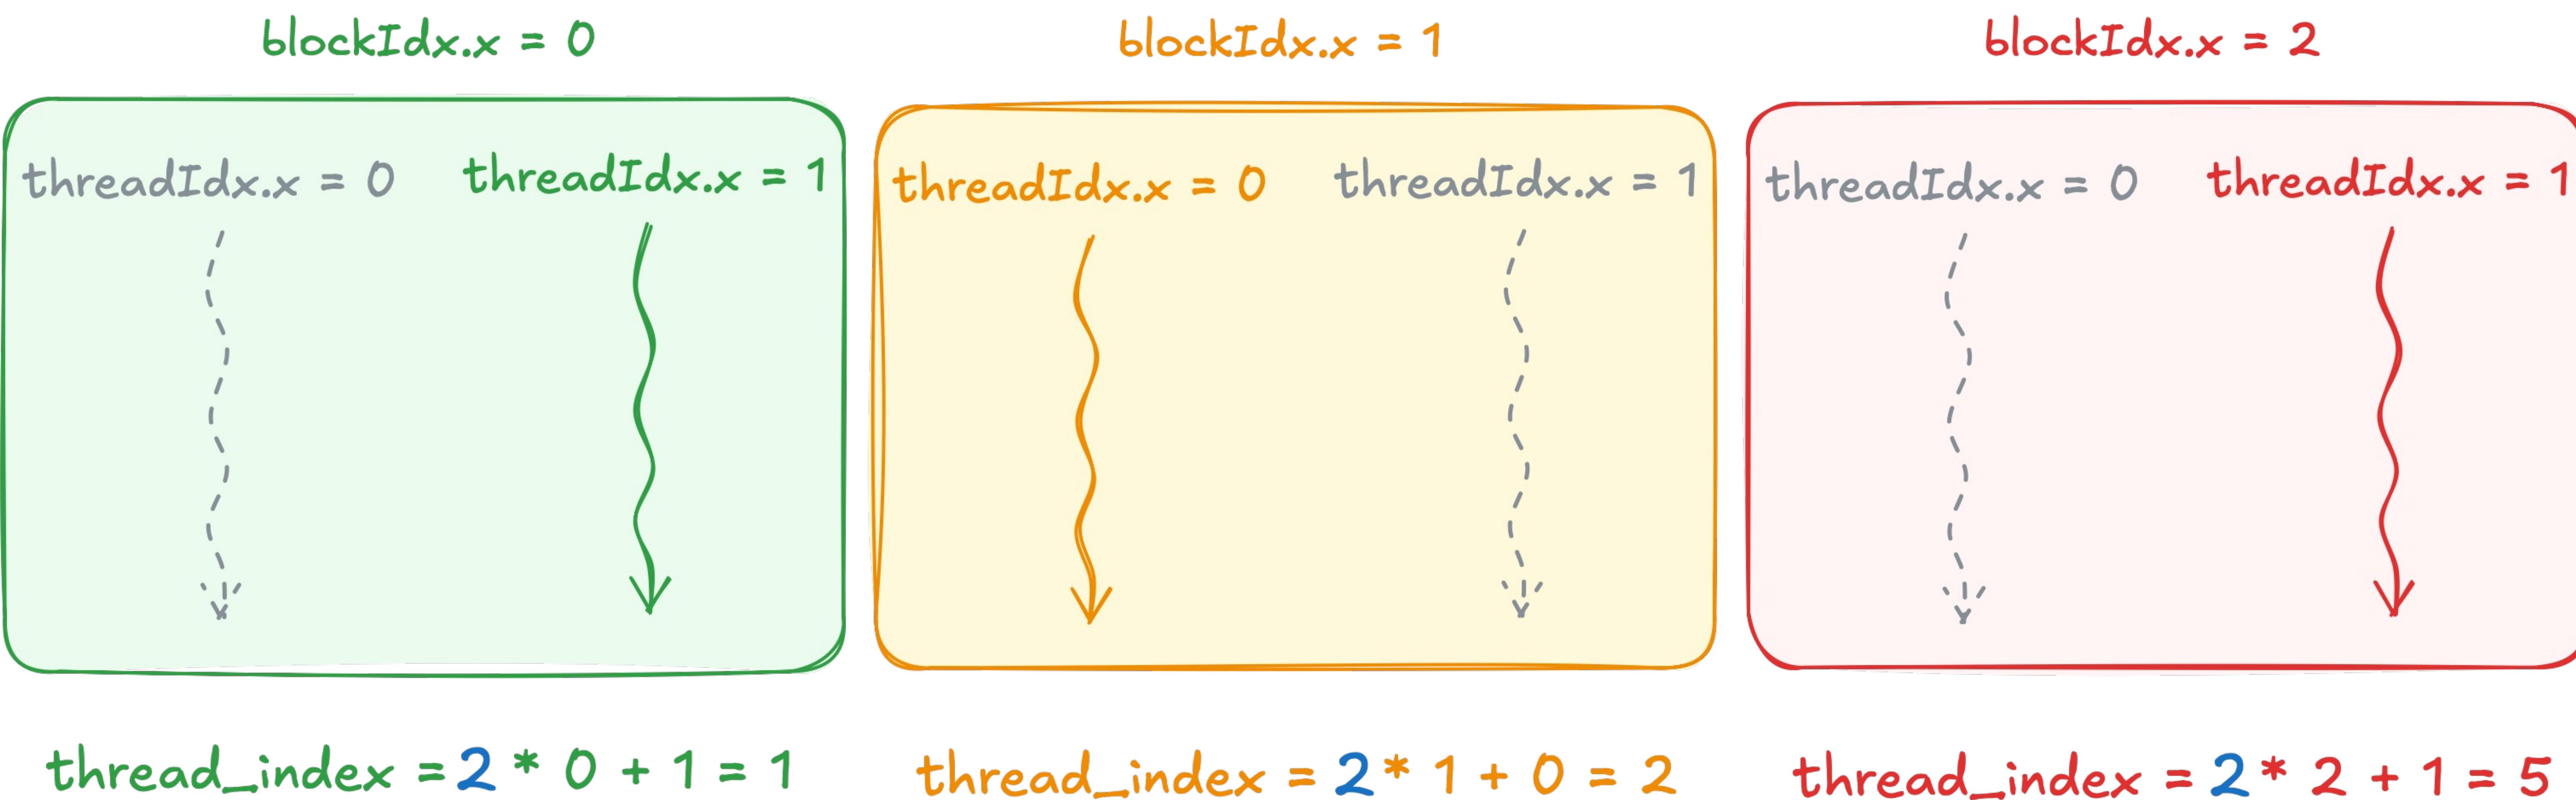
    if (abs(temp(row, column) - temp(temp.extent(0) - 1 - row, column)) > 0.1)
    {
        printf("Error: asymmetry in %d / %d\n", column, temp.extent(1));
    }
}
```

- Another way to detect the bug would be enabling assertions
- Runtime checks affect performance, so OOB access checks are disable by default
- You could enable them by defining **CCCL\_ENABLE\_ASSERTIONS** before including any header in debug

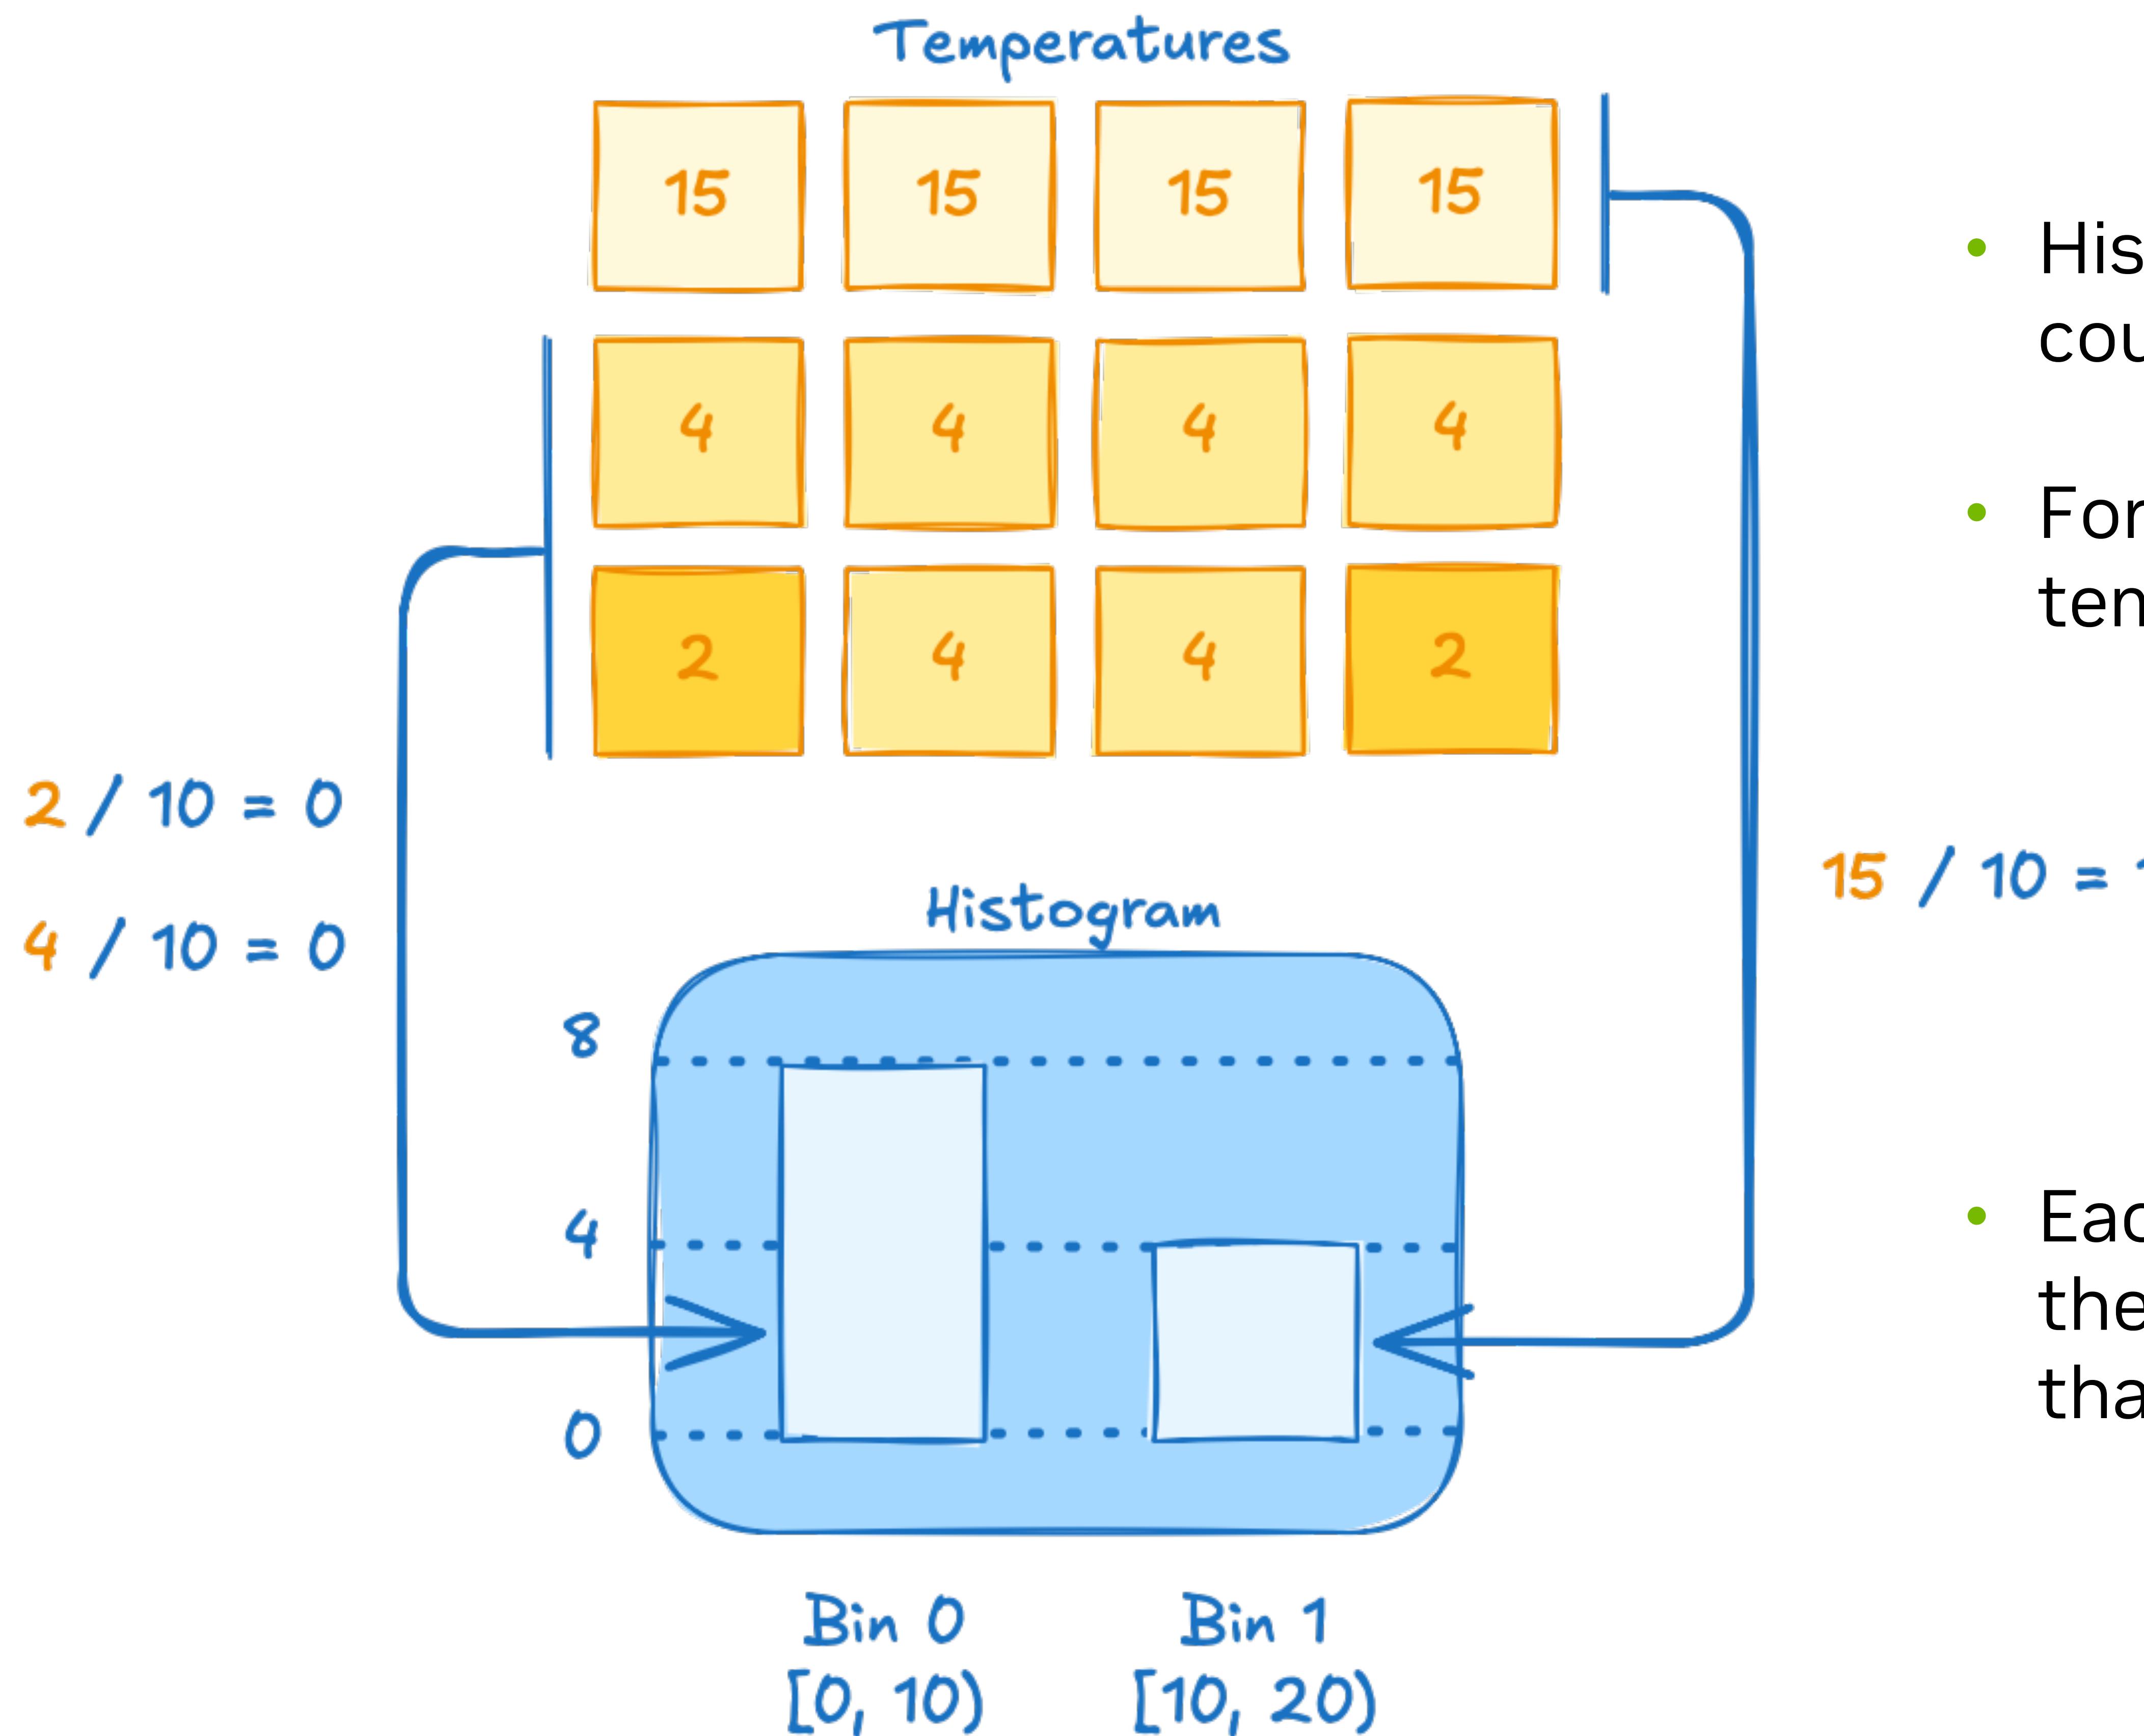
```
mdspan.h:134: __index: block: [0,0,0], thread: [3,0,0]
Assertion `cuda::std::mdspan subscript out of range!' failed.
```

# Why Thread Hierarchy?

- We just fixed a bug caused by splitting grid into blocks
- It's reasonable to ask a question: why adding all this complexity with thread hierarchy?



# Histogram



- Histogram groups data into bins by counting how many values fall into each bin
- For our case, it shows overall distribution of temperatures
- Each bin covers a temperature range, and the bar height shows how many cells lie in that range

# span

- `cuda::std::span` is like one-dimensional `cuda::std::mdspan`

```
__global__ void histogram_kernel(  
    cuda::std::span<float> temperatures,  
    cuda::std::span<int>    histogram)  
{  
    ...  
}
```

- Includes size information
- Safer alternative to raw pointers
- Reducing common pointer errors
- Lightweight view over a sequence
- Does not own the data, just references it

# Interface

```
cuda::std::array<int, 4> a{1, 2, 3, 4}; | constructor takes pointer and size
cuda::std::span<int> span(a.data(), 3);

std::printf("span[0] = %d\n", span[0]); // 1
std::printf("span[1] = %d\n", span[1]); // 2

std::printf("span.size() = %zu\n", span.size()); // 3
```

# Interface Recap

```
cuda::std::array<int, 4> a{1, 2, 3, 4};  
cuda::std::span<int> span(a.data(), 3);
```

```
std::printf("span[0] = %d\n", span[0]); // 1  
std::printf("span[1] = %d\n", span[1]); // 2
```

provides subscript operator to access elements

```
std::printf("span.size() = %zu\n", span.size()); // 3
```

# Interface

```
cuda::std::array<int, 4> a{1, 2, 3, 4};  
cuda::std::span<int> span(a.data(), 3);  
  
std::printf("span[0] = %d\n", span[0]); // 1  
std::printf("span[1] = %d\n", span[1]); // 2  
  
std::printf("span.size() = %zu\n", span.size()); // 3
```

returns the number of elements in the span

# First Attempt at Implementing Histogram Kernel

```
--global__ void histogram_kernel(  
    cuda::std::span<float> temperatures,  
    cuda::std::span<int>    histogram)  
{  
    int cell = blockIdx.x * blockDim.x + threadIdx.x; | each thread bins exactly one cell
```

# First Attempt at Implementing Histogram Kernel

```
constexpr float bin_width = 10;

__global__ void histogram_kernel(
    cuda::std::span<float> temperatures,
    cuda::std::span<int> histogram)
{
    int cell = blockIdx.x * blockDim.x + threadIdx.x;
    int bin = static_cast<int>(temperatures[cell] / bin_width); | divide temperature by bin width (10
                                                               degrees) to compute the bin index
```



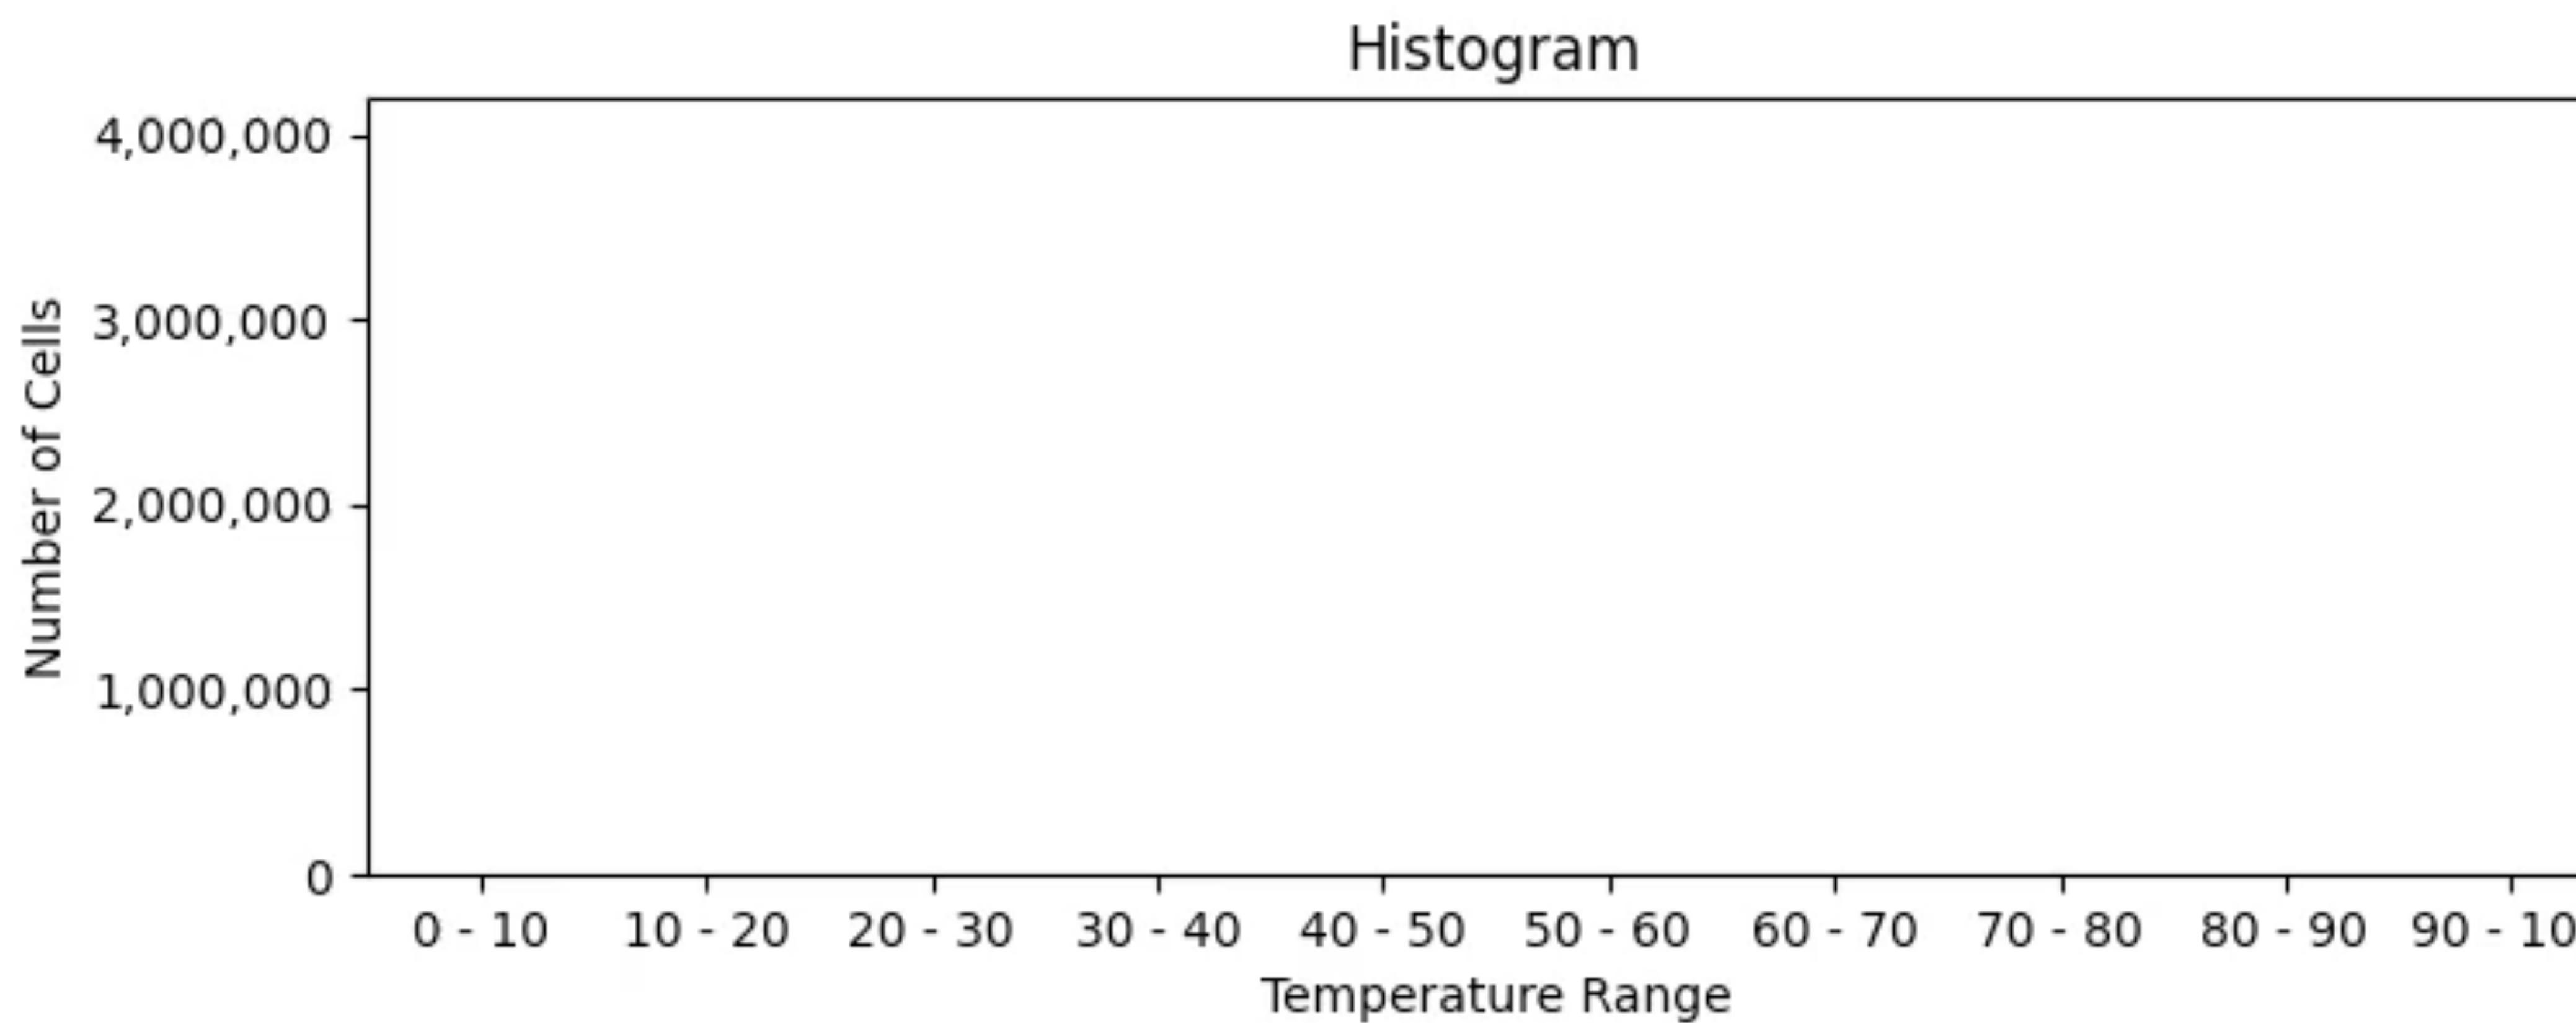
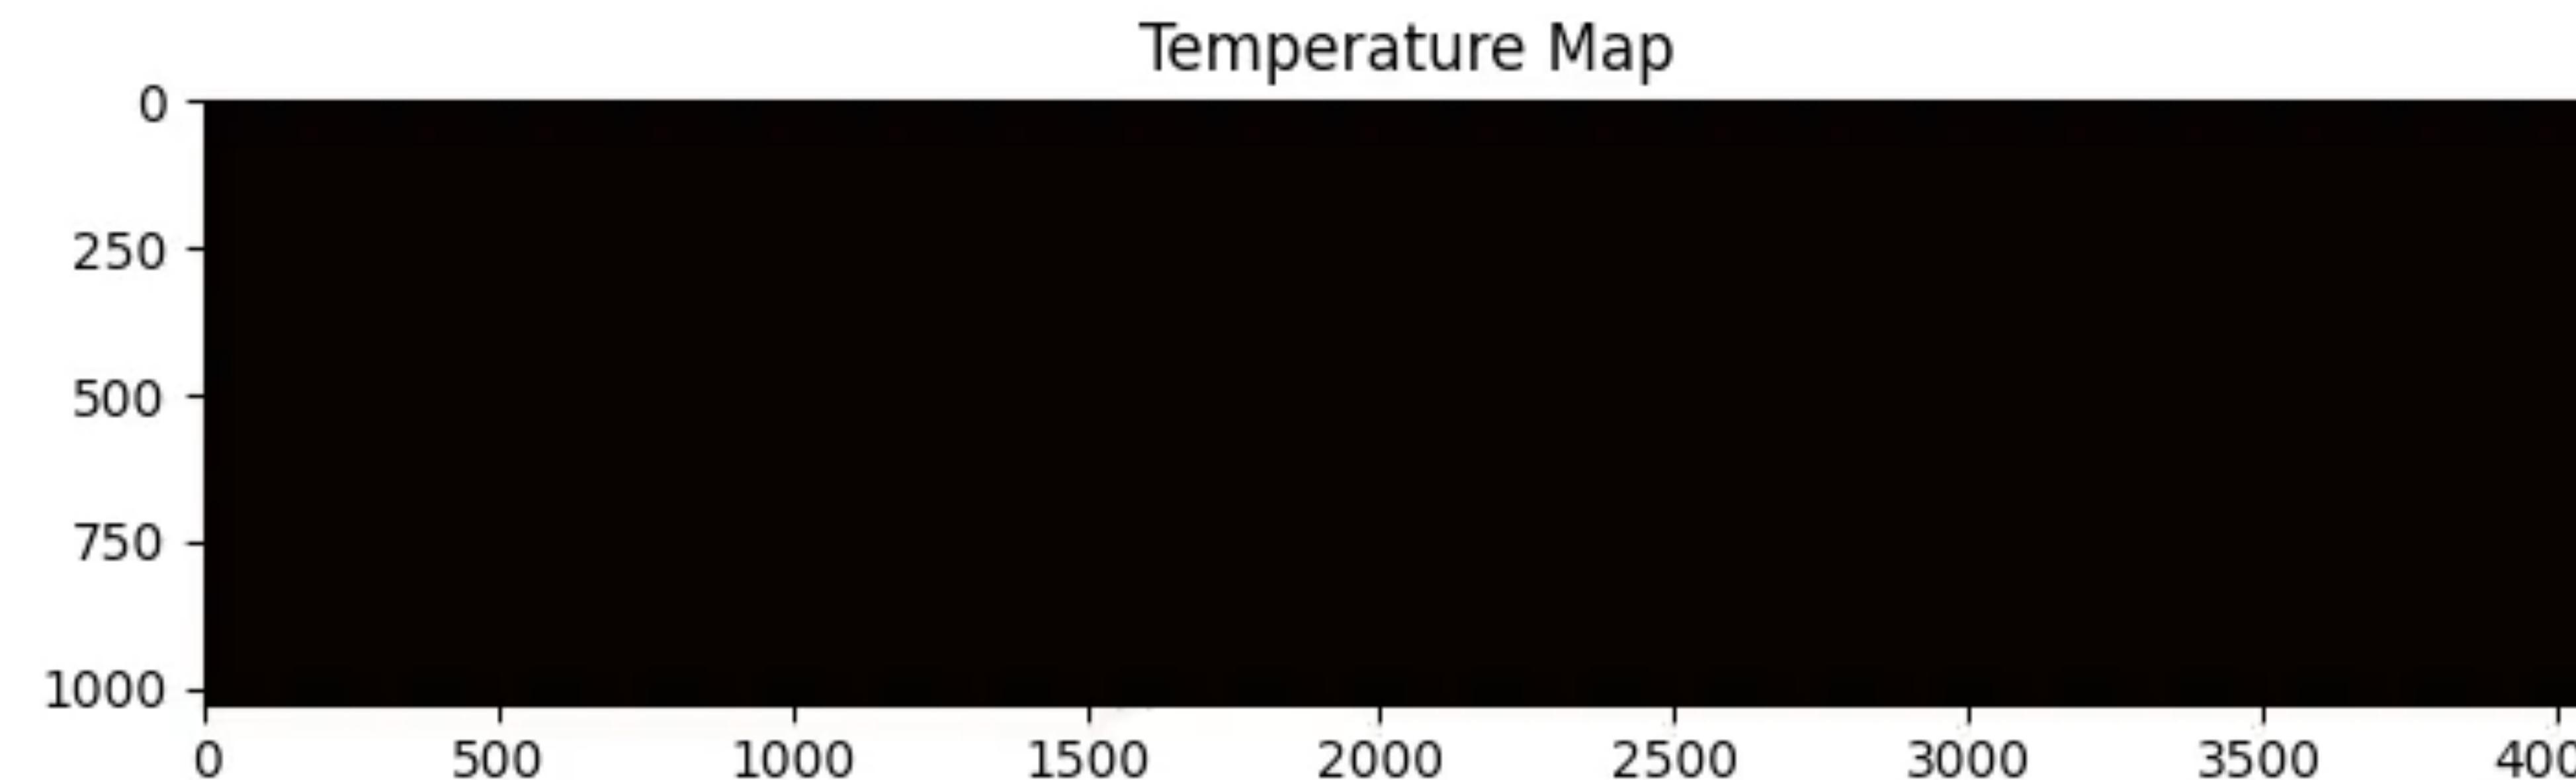
- In production code, you should check if cell index is within bounds
- We'll omit this check for simplicity and make sure input is multiple of block size

# First Attempt at Implementing Histogram Kernel

```
--global__ void histogram_kernel(  
    cuda::std::span<float> temperatures,  
    cuda::std::span<int>    histogram)  
{  
    int cell = blockIdx.x * blockDim.x + threadIdx.x;  
    int bin = static_cast<int>(temperatures[cell] / bin_width);  
  
    int old_count = histogram[bin];  
    int new_count = old_count + 1;  
    histogram[bin] = new_count;  
}
```

- Read old value of selected bin
- Modify the value we just read (increment)
- Write modified value back into memory

# Something Went Wrong



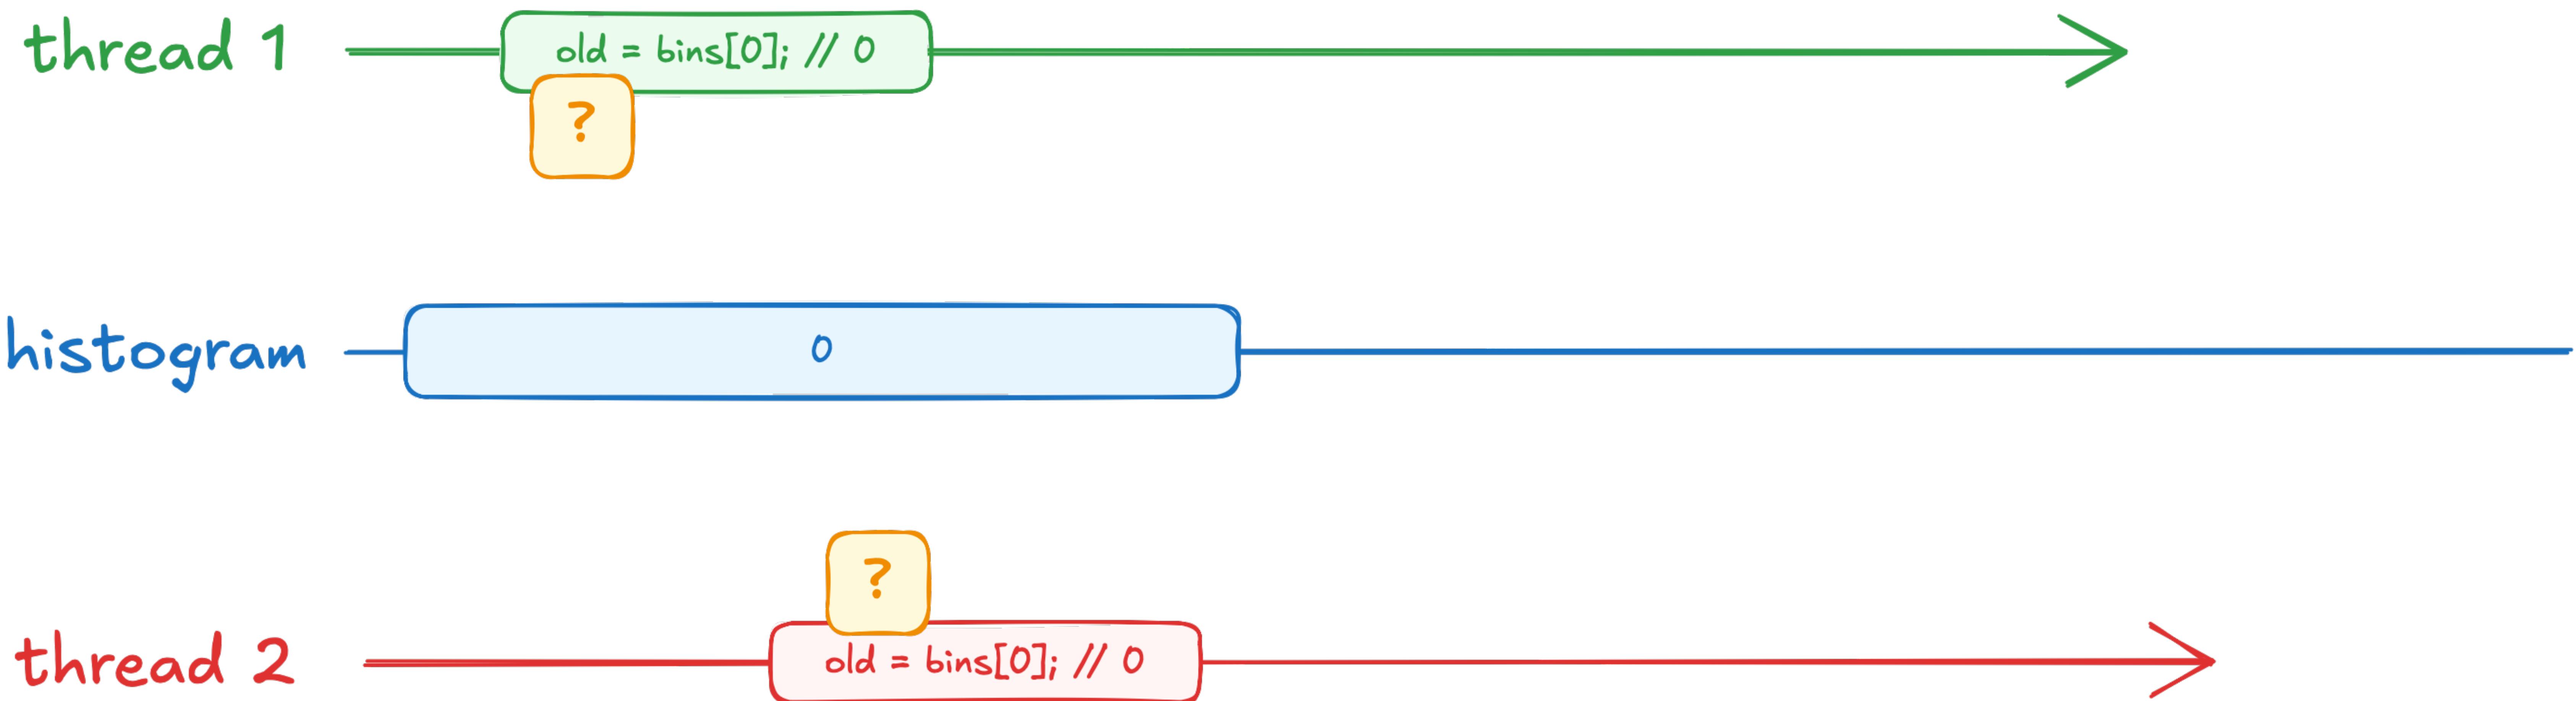
- We have about 4M cells
- We should see bars adding up to 4M
- But resulting histogram looks empty
- What happened?

# Source of the Data Race

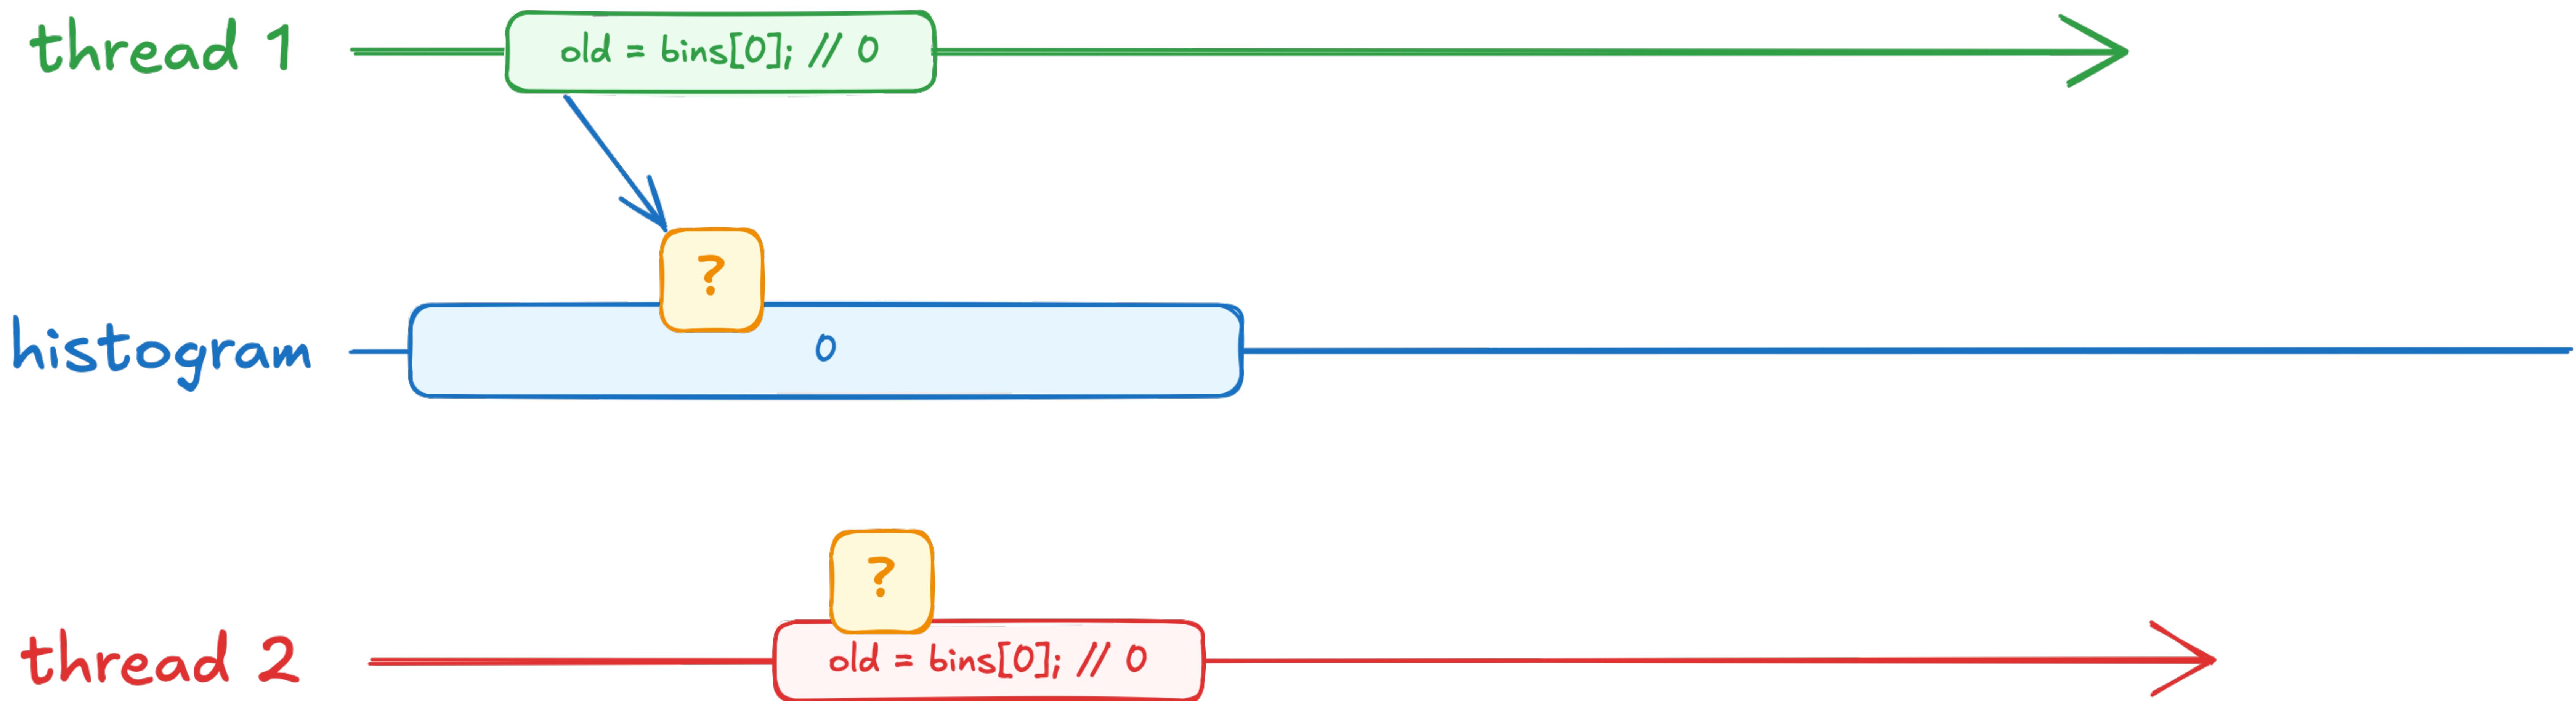
```
--global__ void histogram_kernel(  
    cuda::std::span<float> temperatures,  
    cuda::std::span<int>    histogram)  
{  
    int cell = blockIdx.x * blockDim.x + threadIdx.x;  
    int bin = static_cast<int>(temperatures[cell] / bin_width);  
    int old_count = histogram[bin];  
    int new_count = old_count + 1;  
    histogram[bin] = new_count;  
}
```

- Our kernel has a data race in it
- Millions of threads read and write the same memory location!

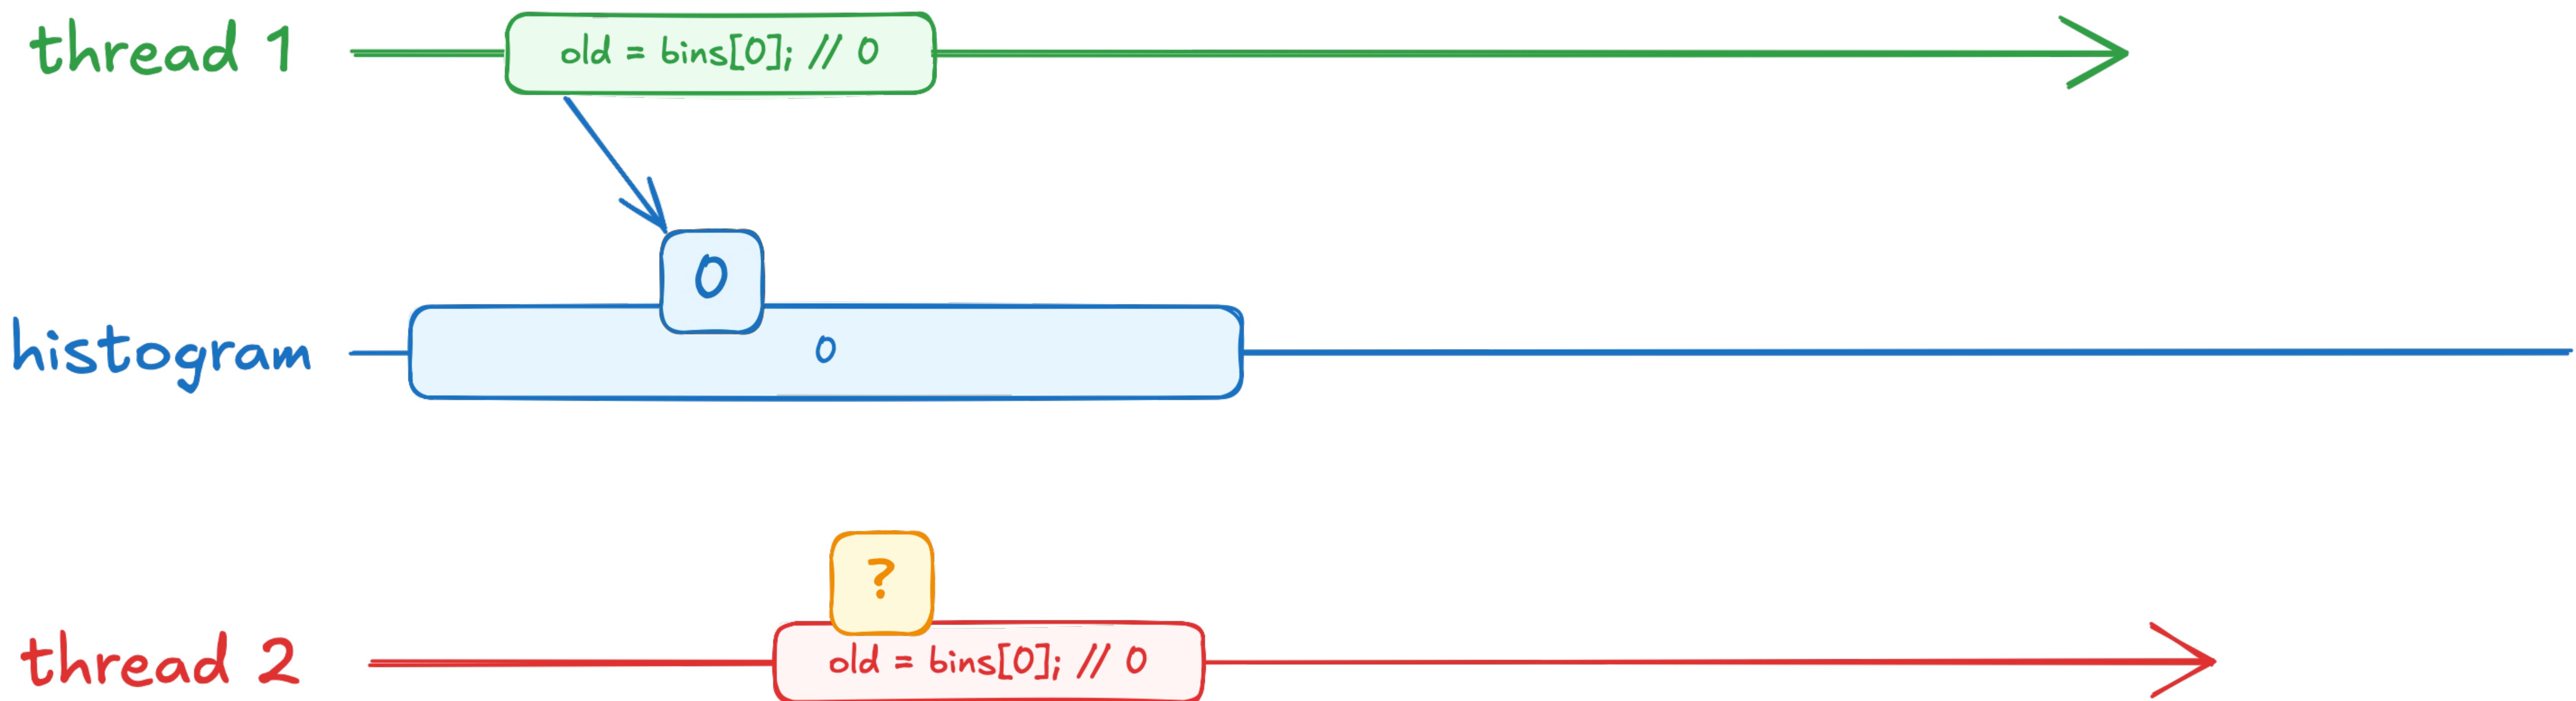
# Source of the Data Race



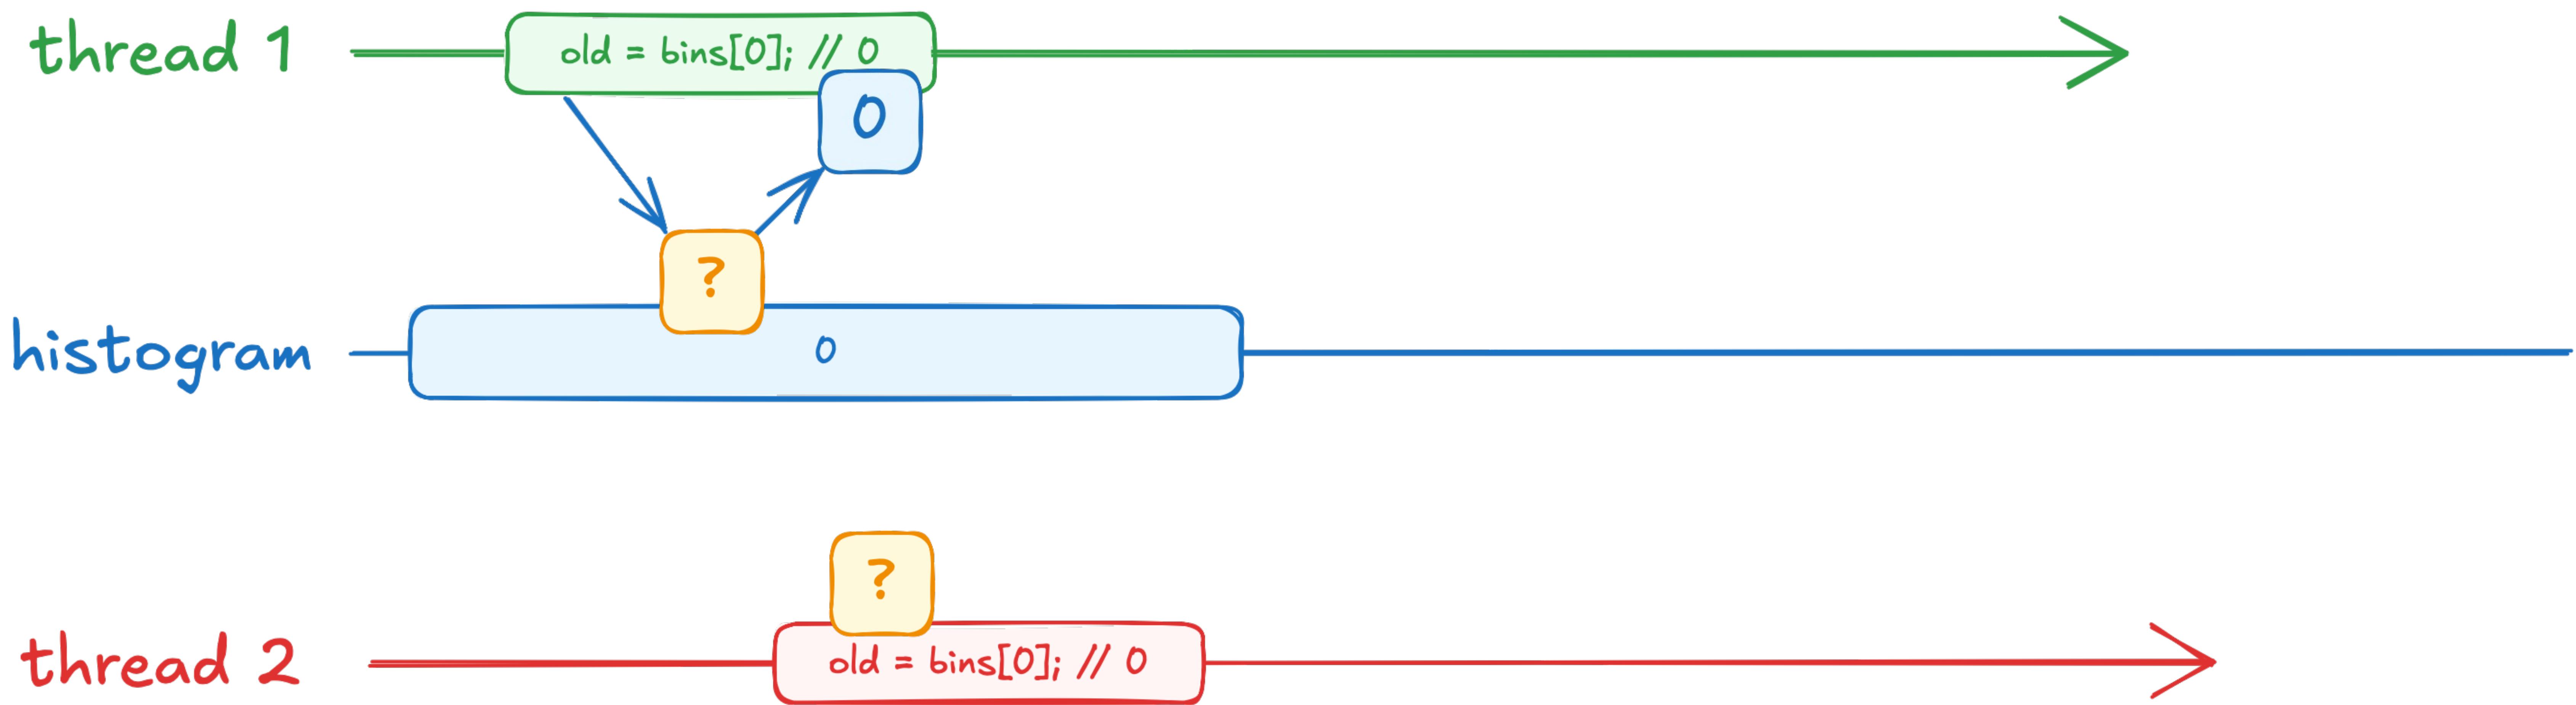
# Source of the Data Race



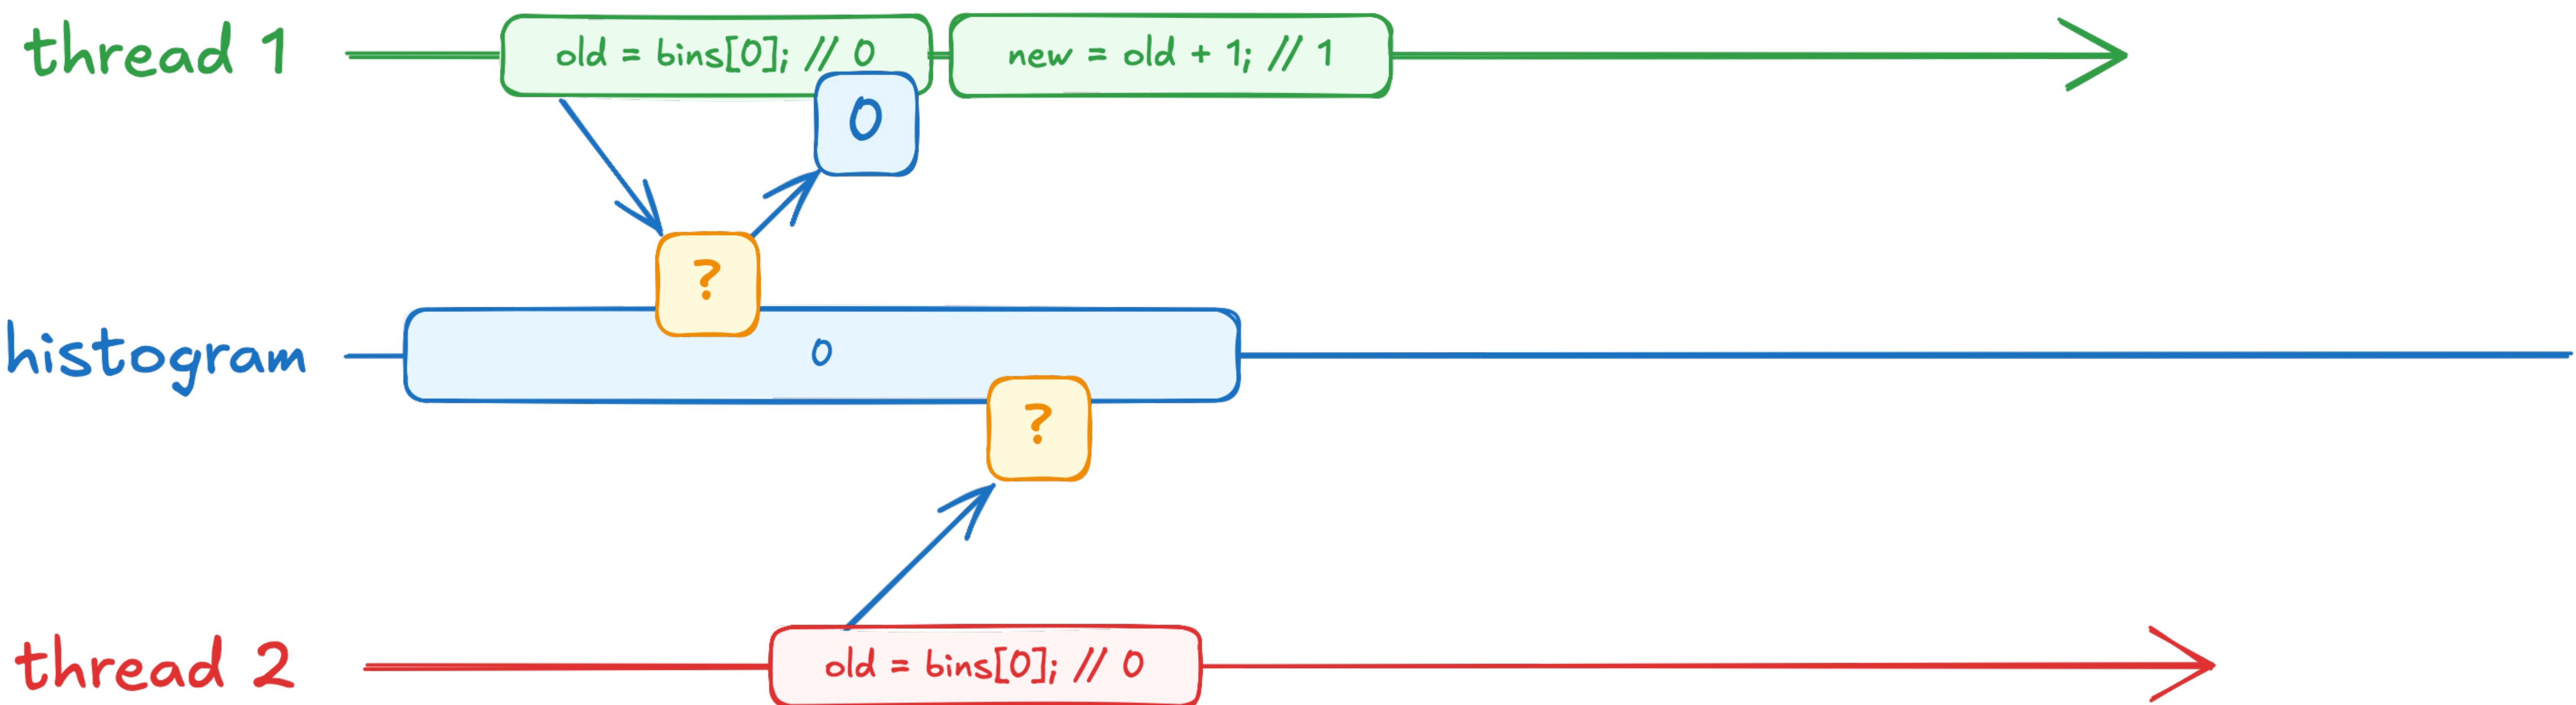
# Source of the Data Race



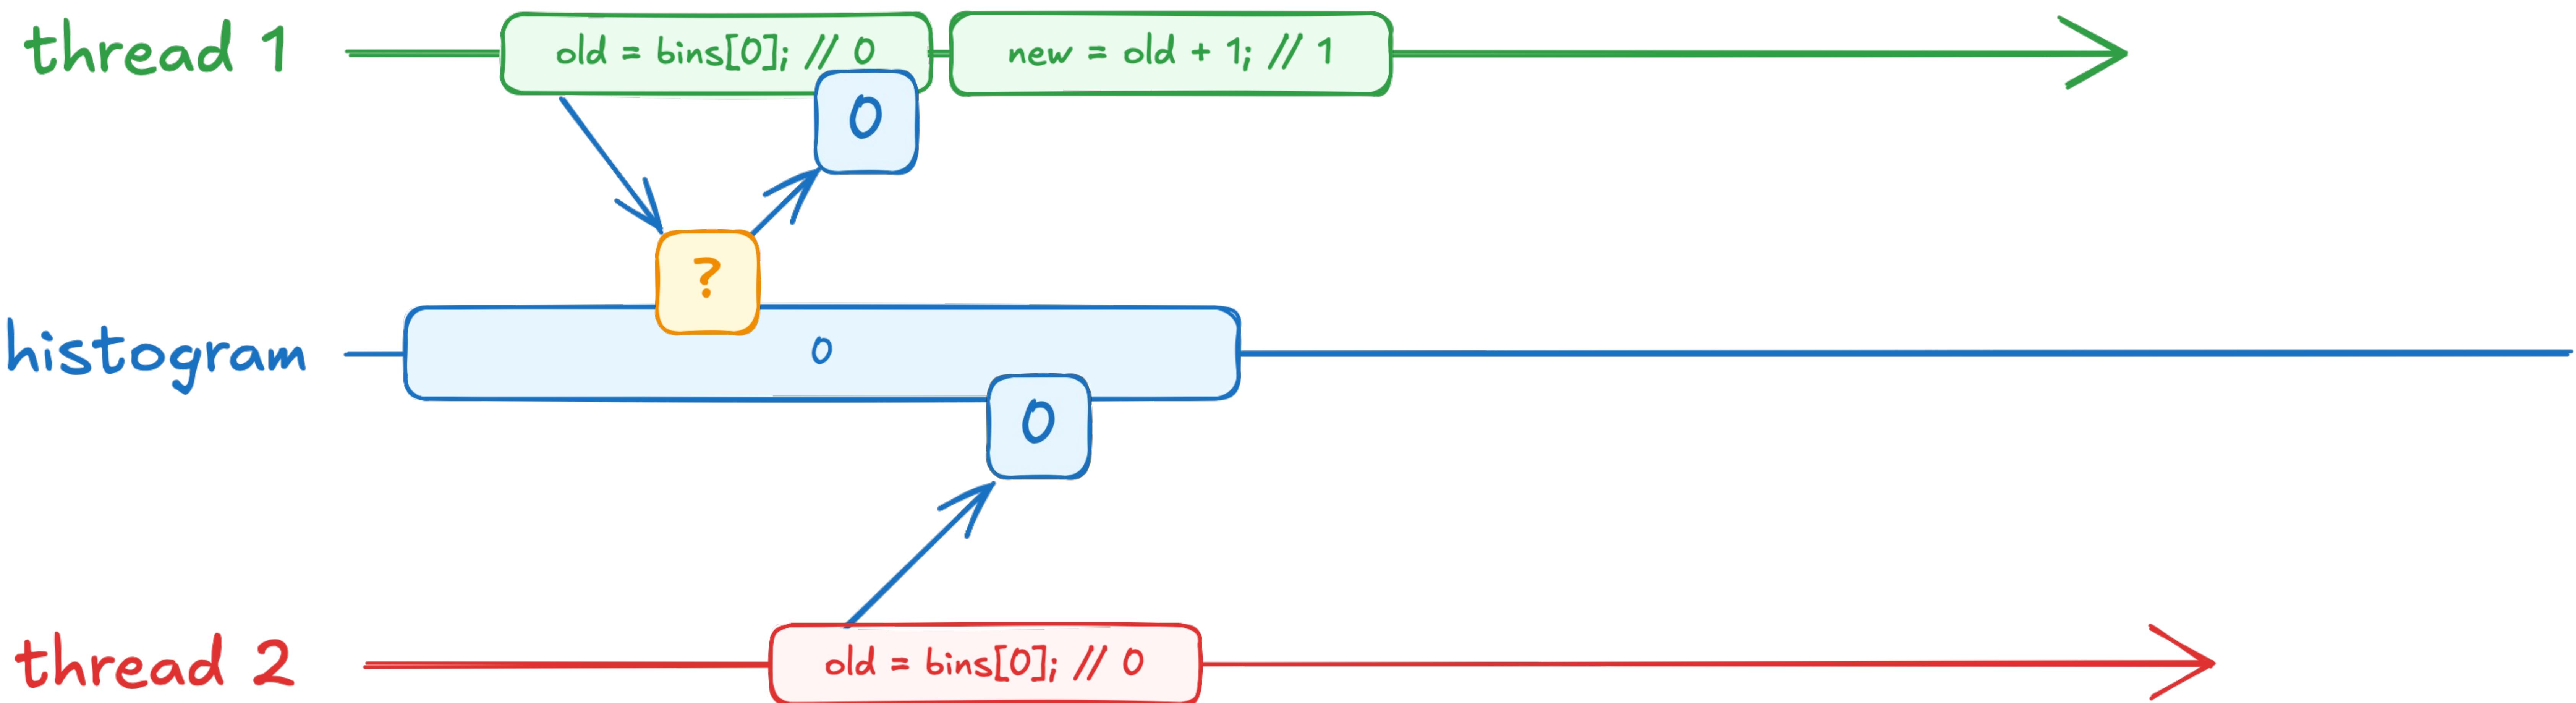
# Source of the Data Race



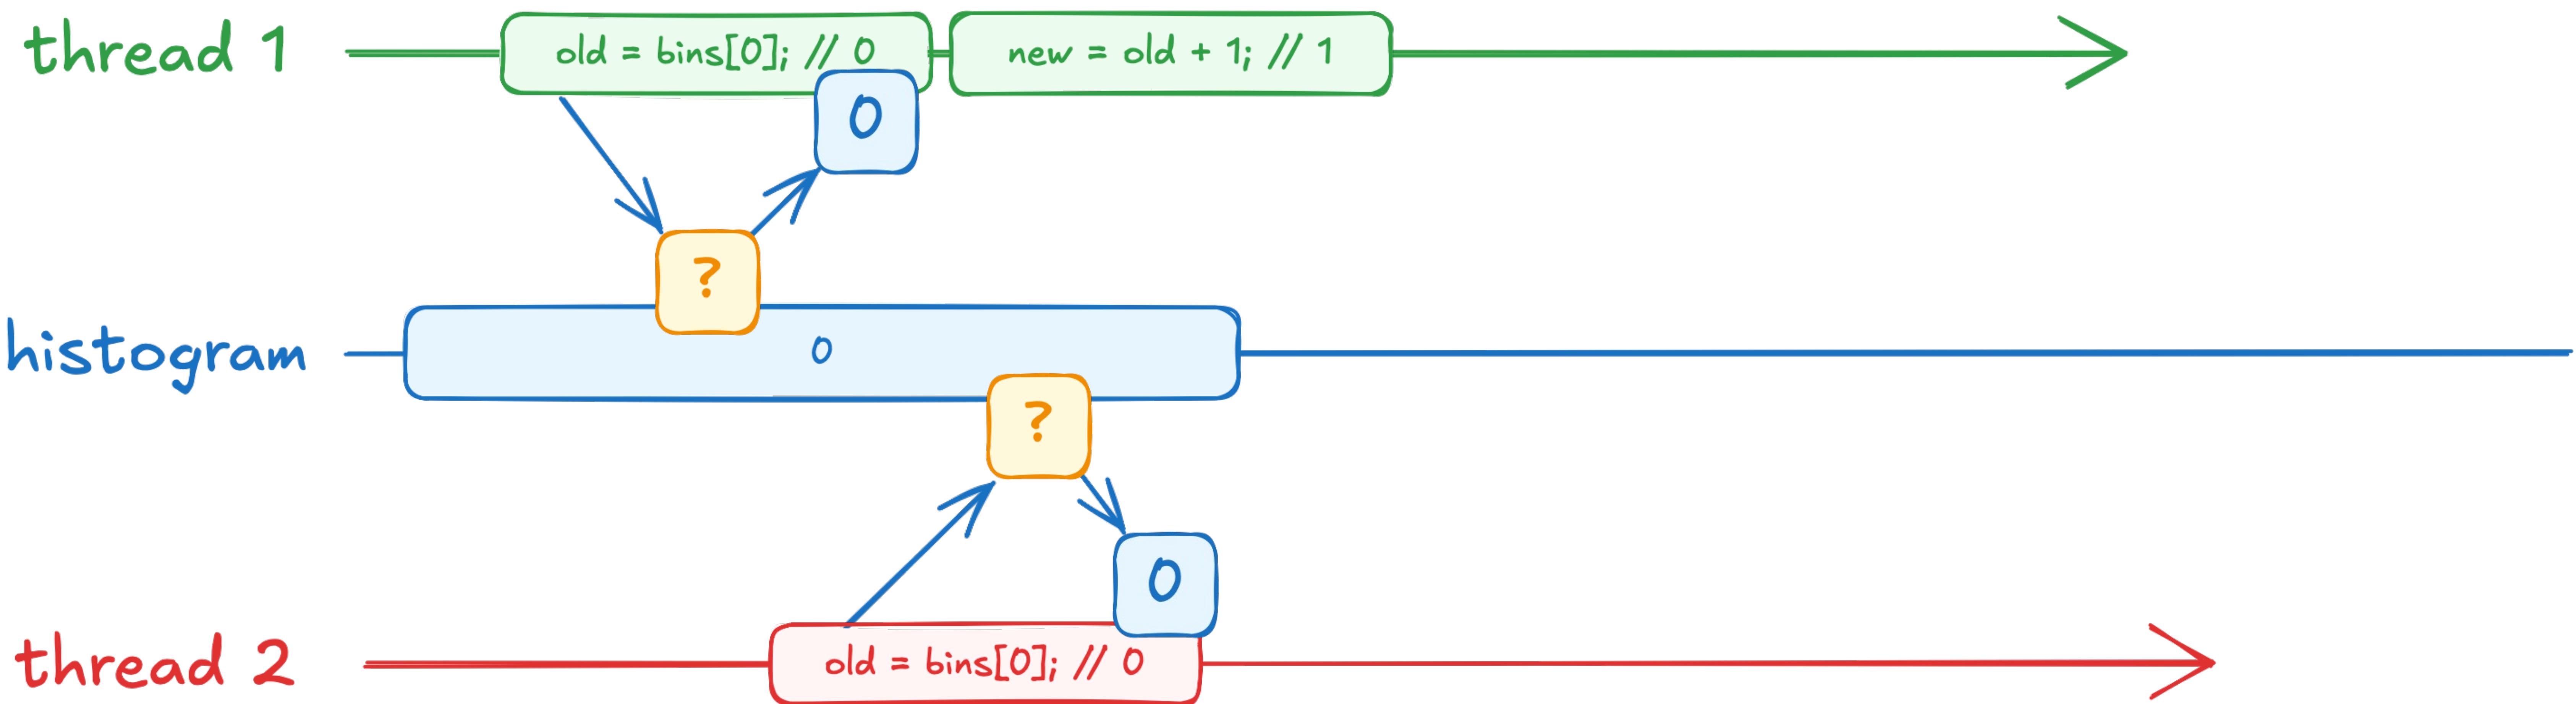
# Source of the Data Race



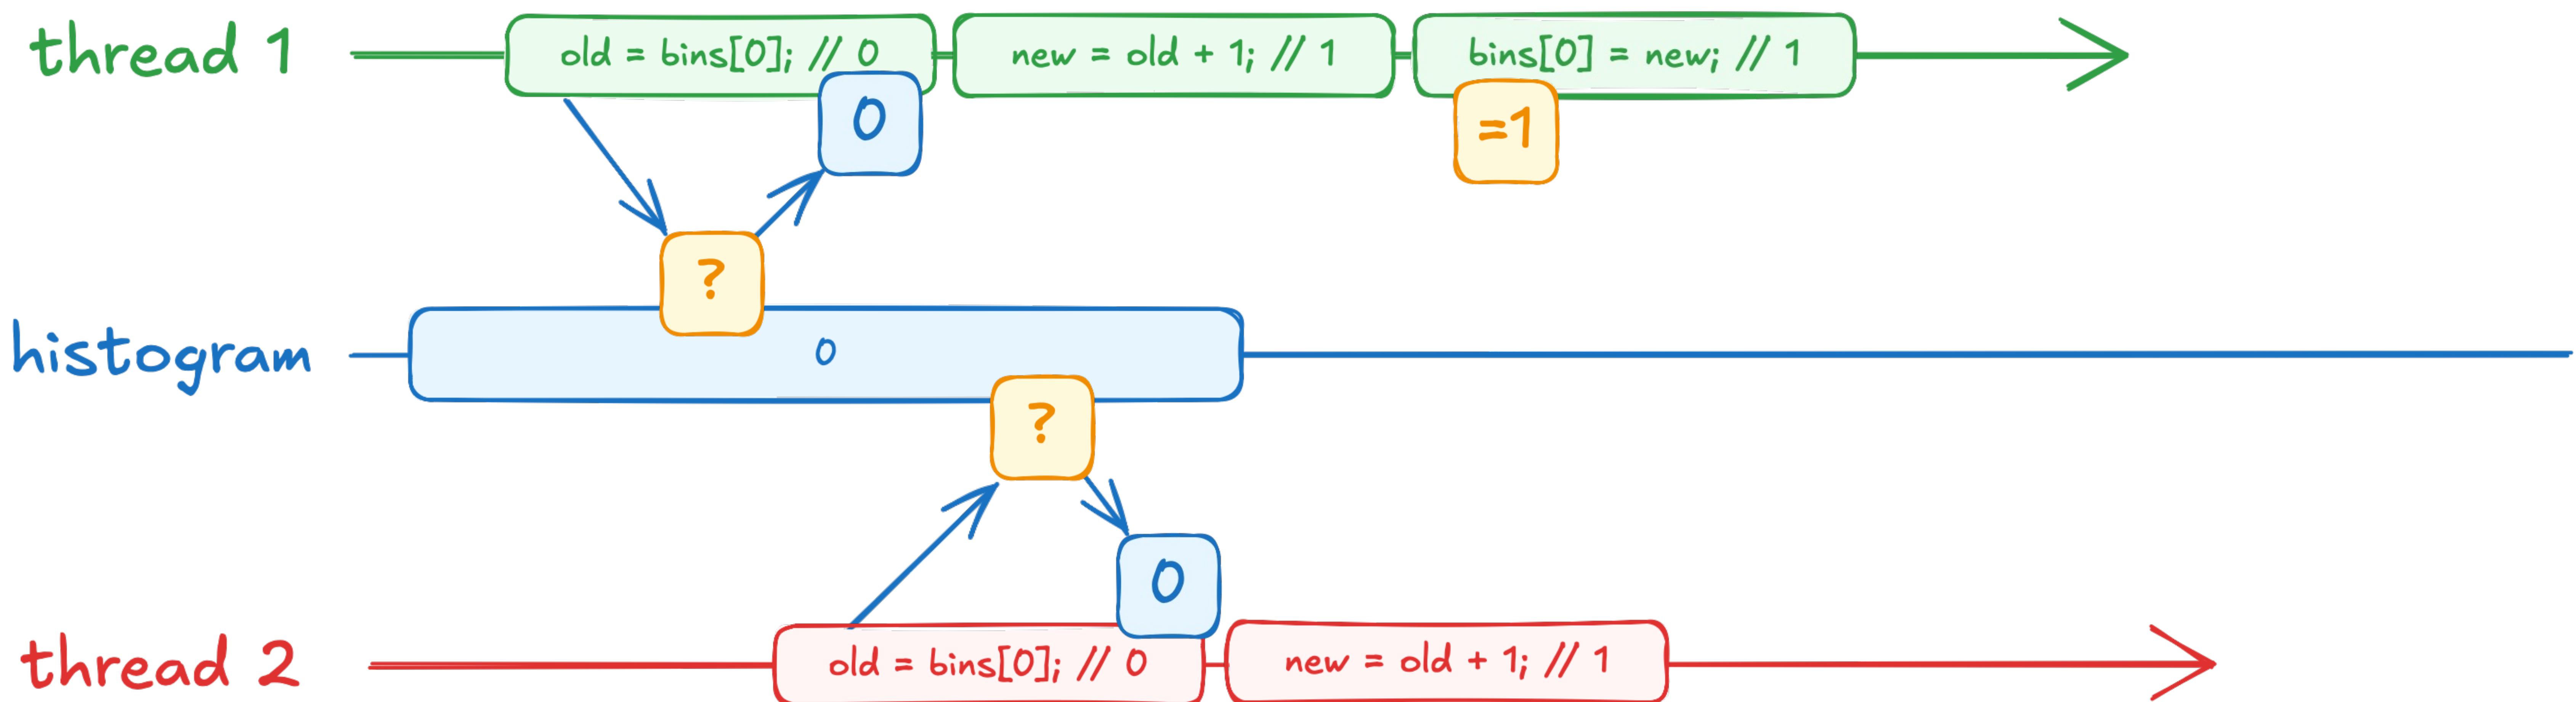
# Source of the Data Race



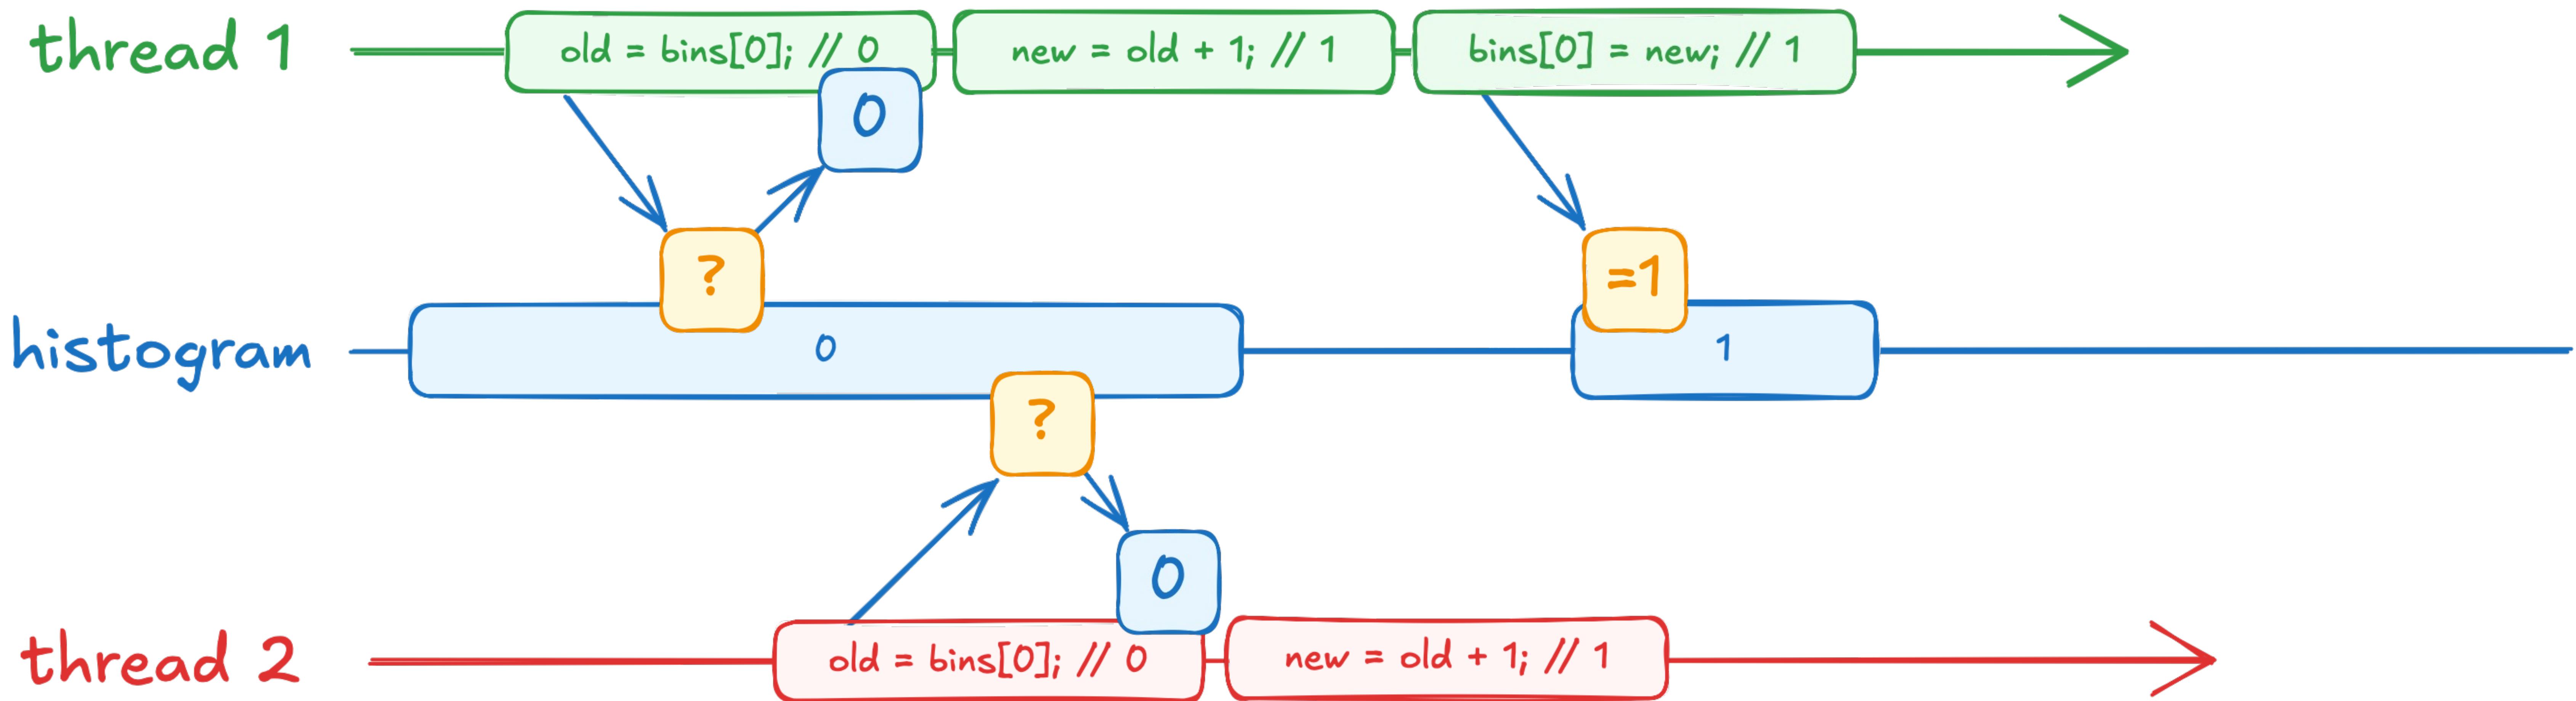
# Source of the Data Race



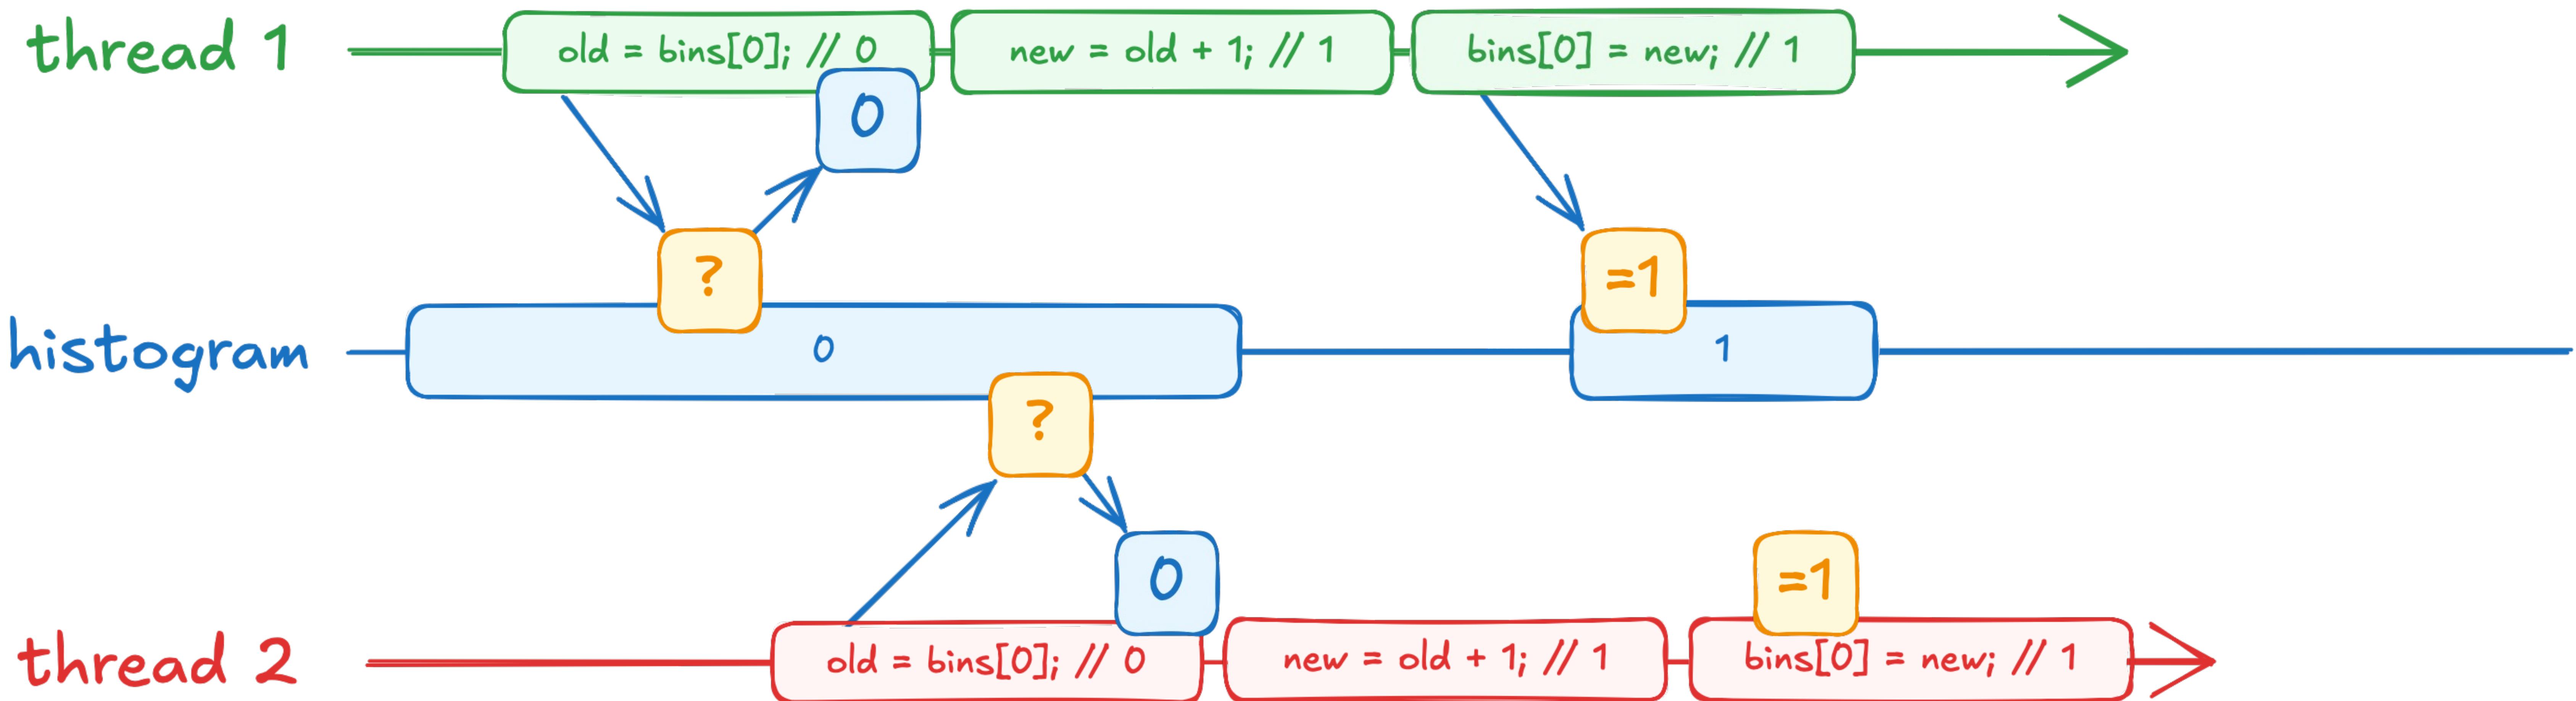
# Source of the Data Race



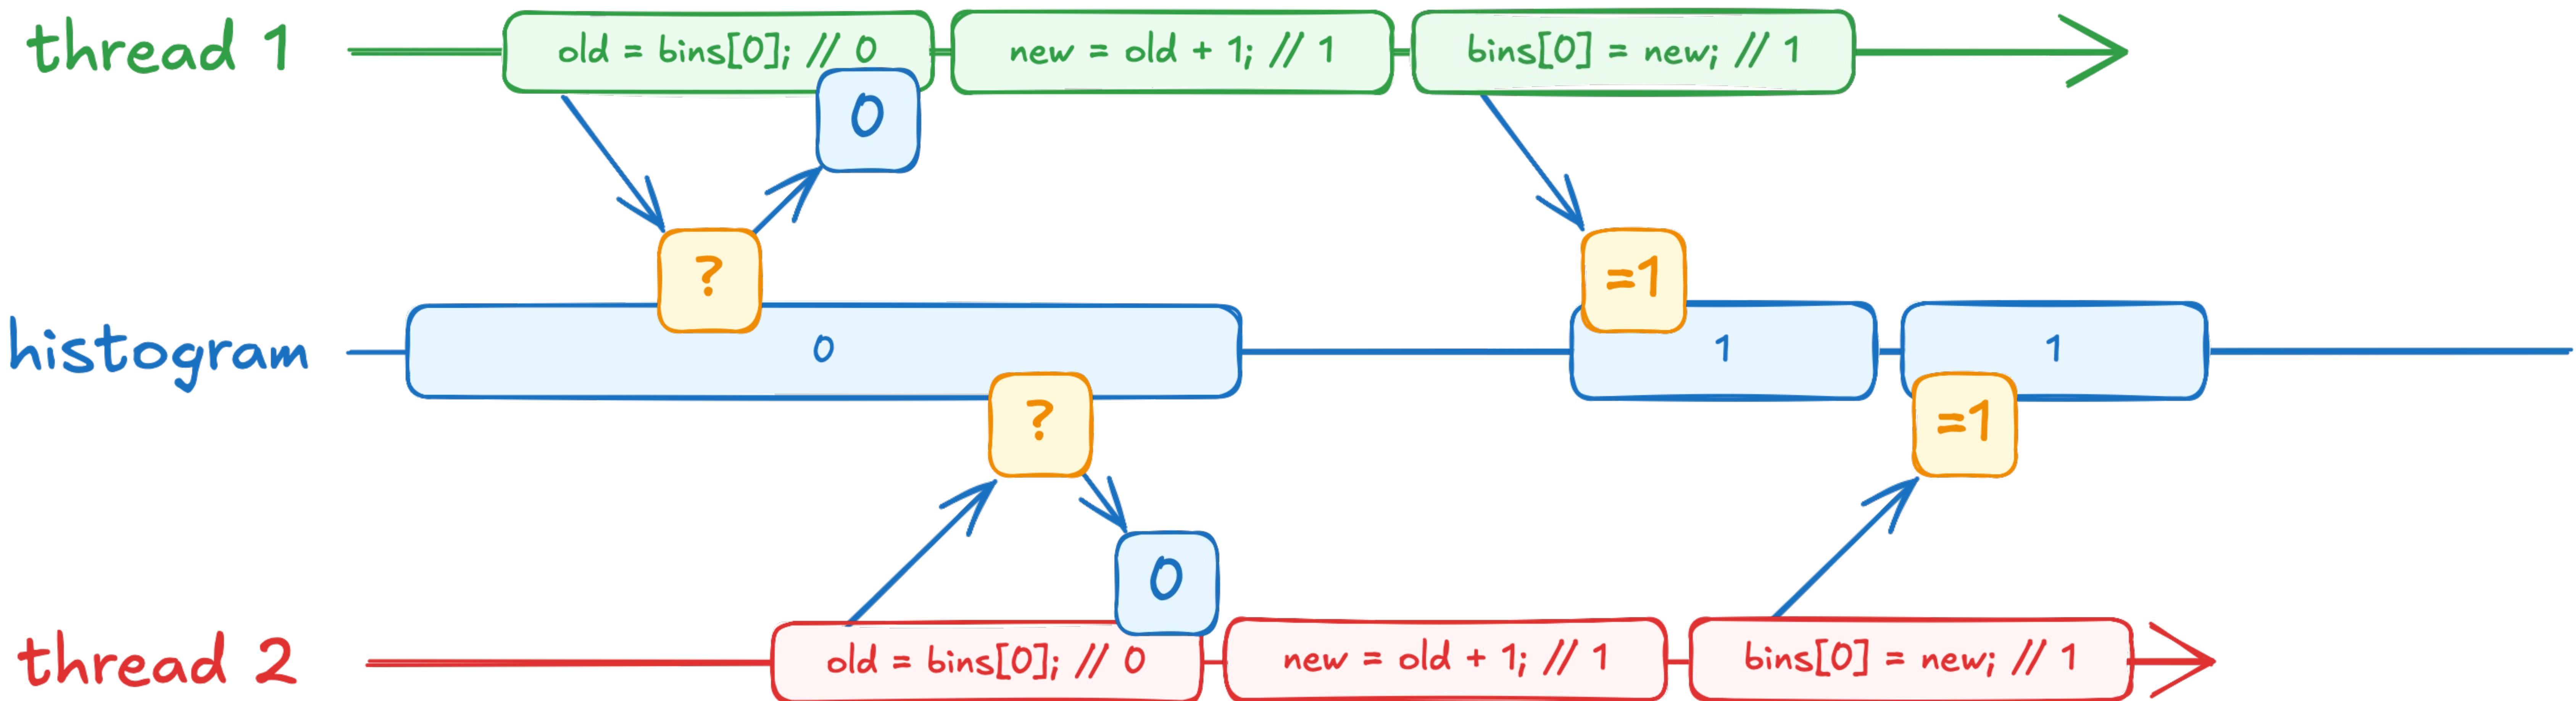
# Source of the Data Race



# Source of the Data Race



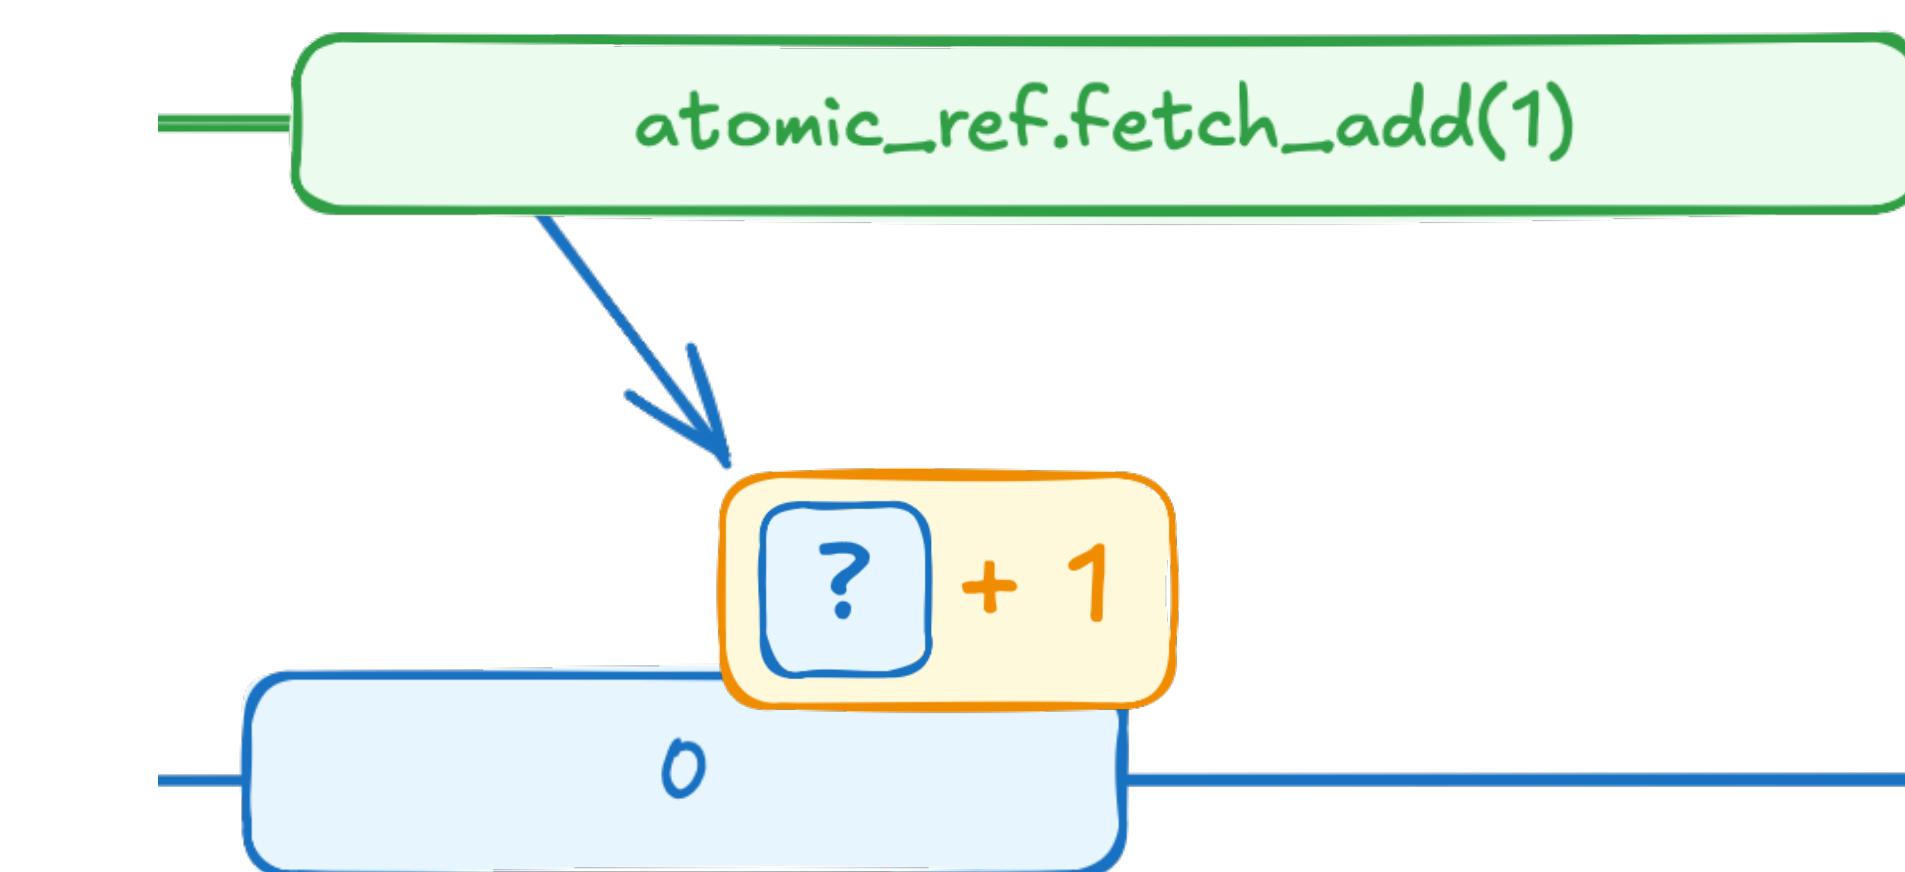
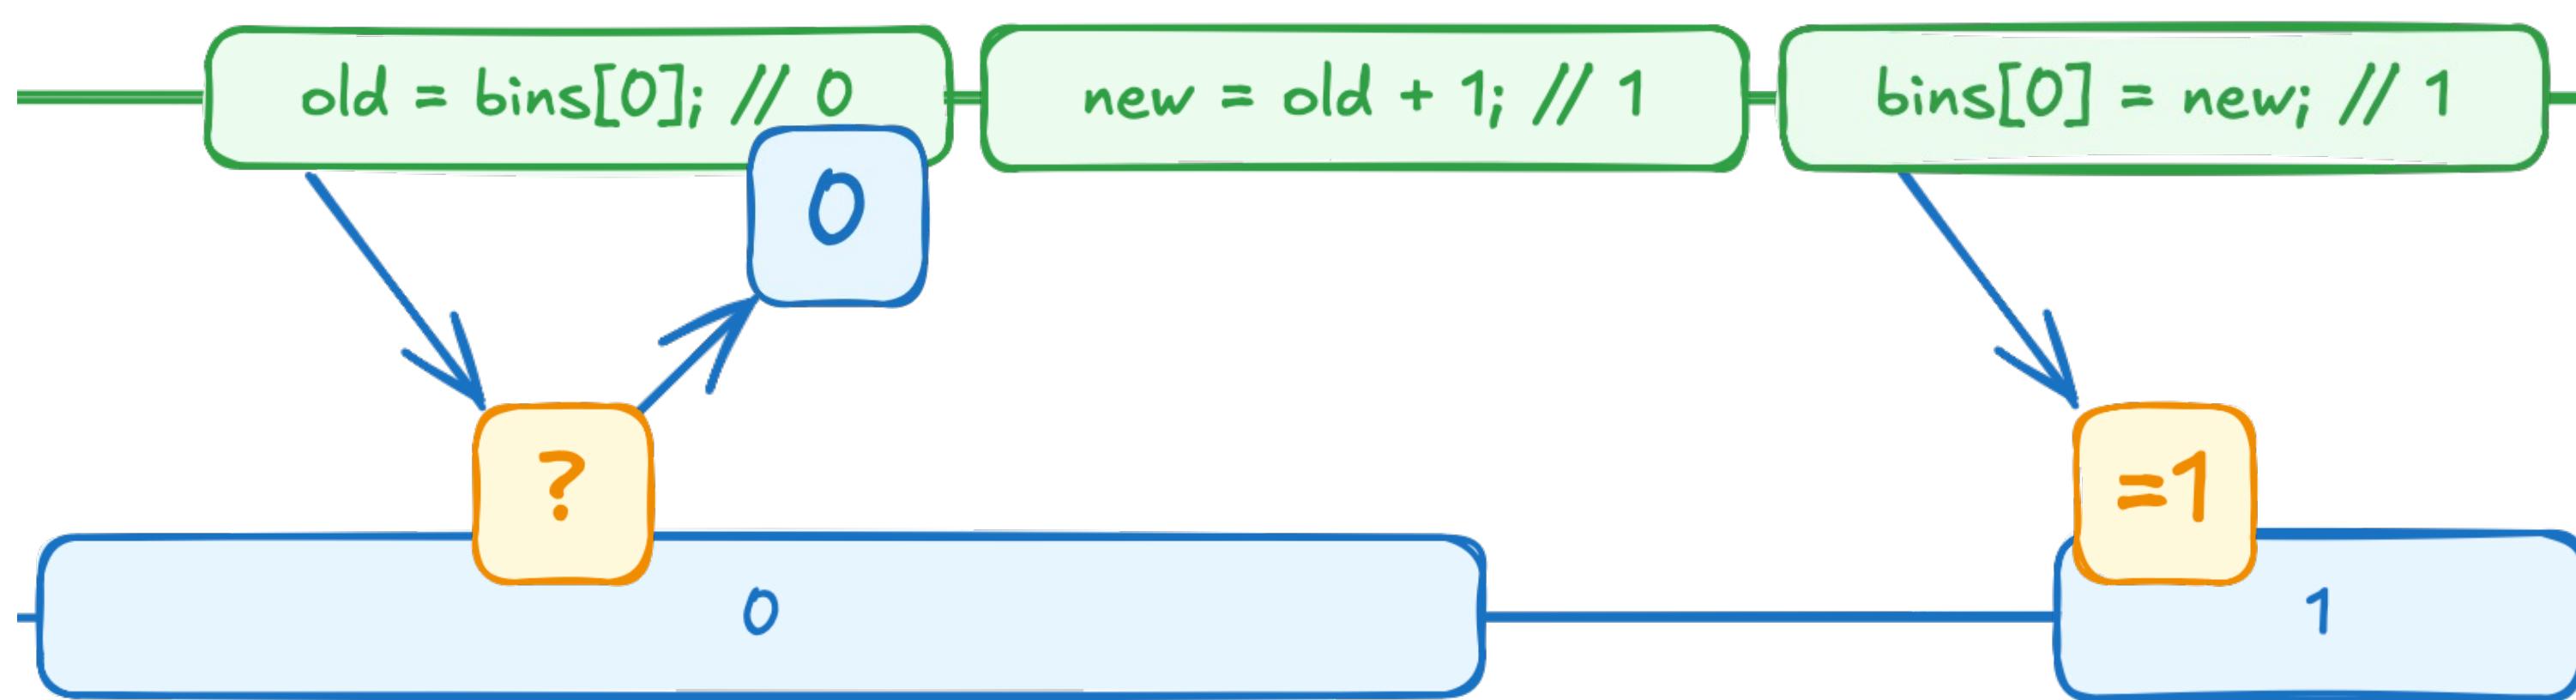
# Source of the Data Race



# Atomic Memory Operations

To fix data race, we have to make 1. read 2. modify 3. write steps appear as a single **indivisible** read-modify-write operation

```
int old_count = histogram[bin];  
  
int new_count = old_count + 1;           histogram[bin].fetch_add(1);  
  
histogram[bin] = new_count;
```

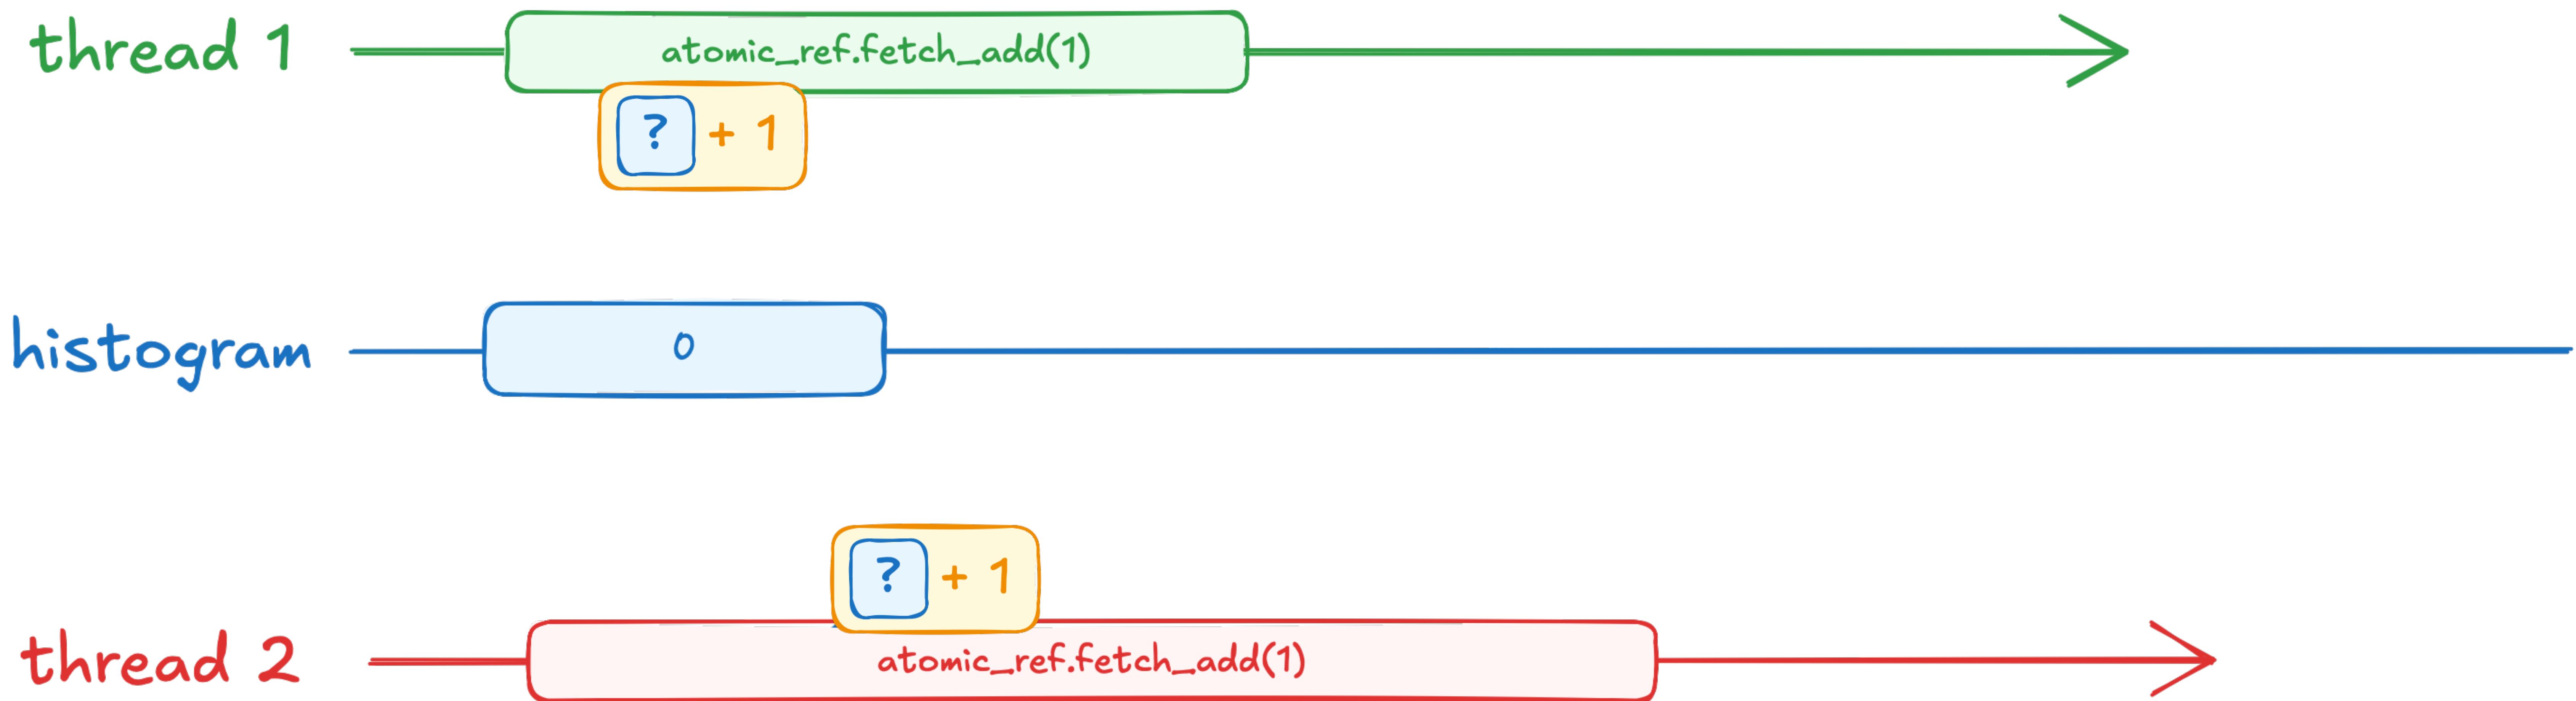


# Atomic Memory Operations

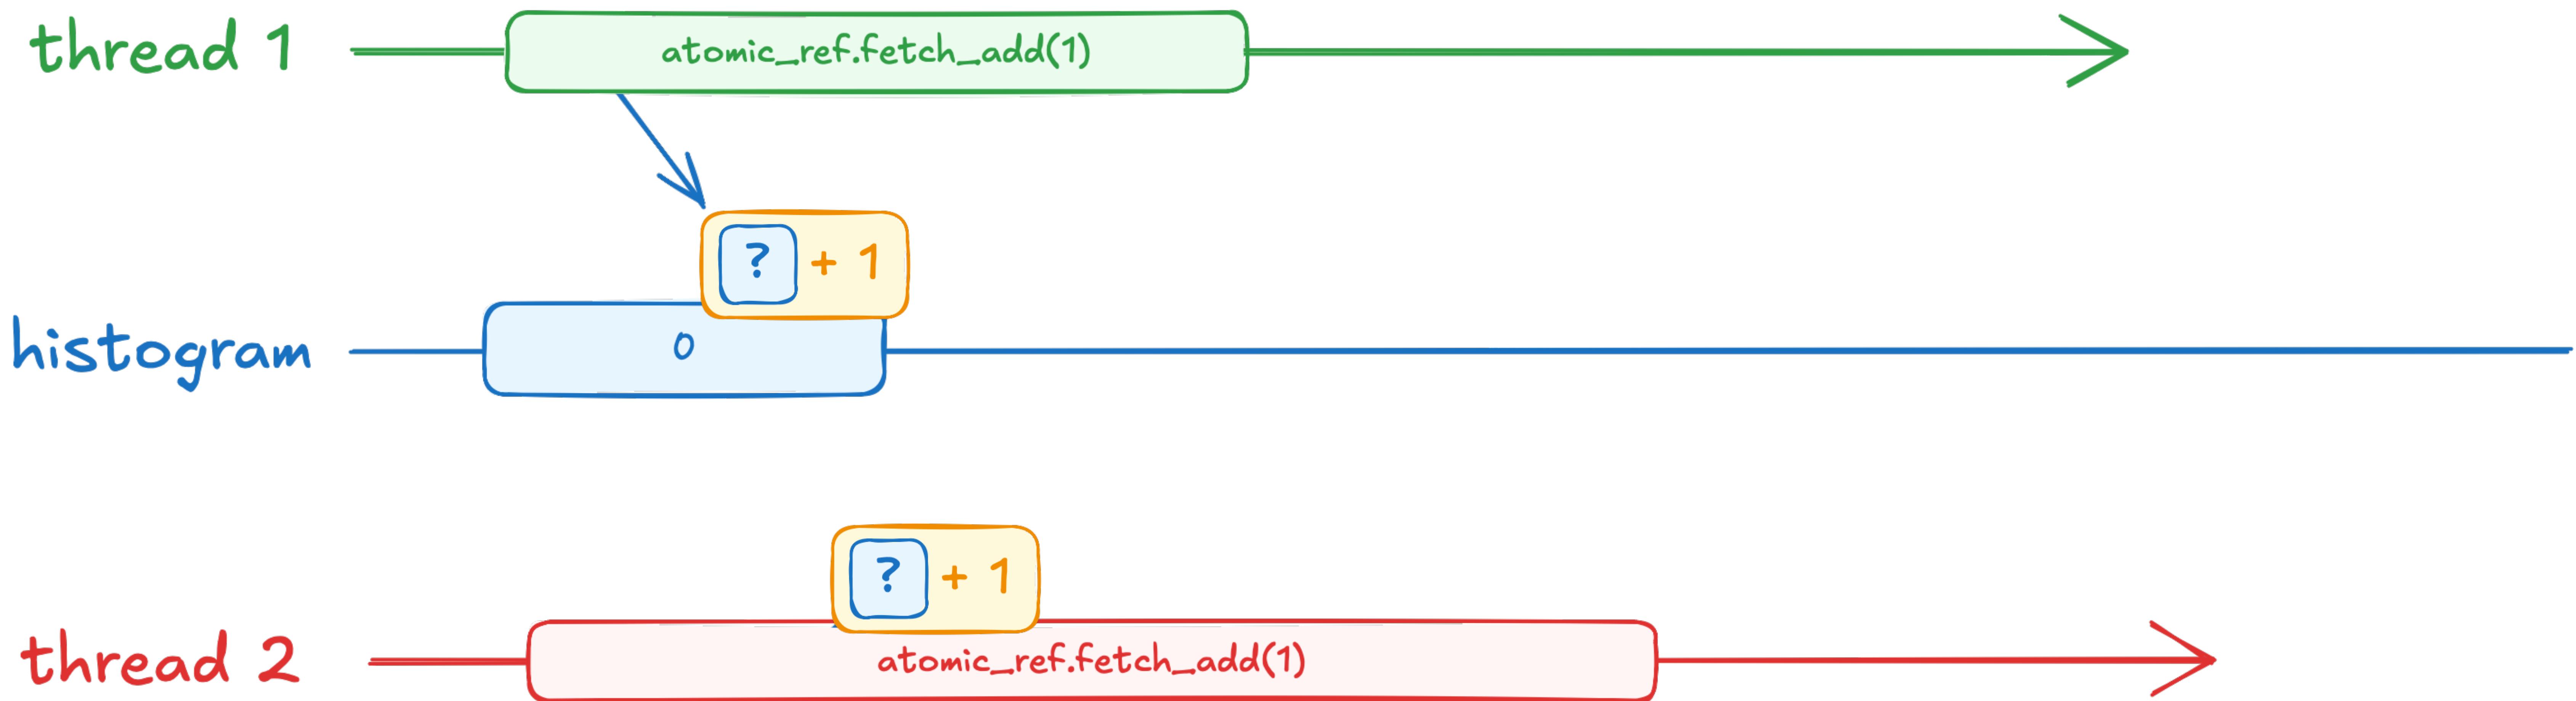
```
cuda::std::atomic_ref<int> ref(count[0]);  
  
ref.fetch_add(3);  
  
ref.fetch_sub(2);  
  
ref.fetch_and(1);  
  
...
```

- Atomics memory operations are indivisible
- `atomic_ref` applies atomic operations to the object it references
- For our example, atomic increment appears as indivisible read-modify-write operation

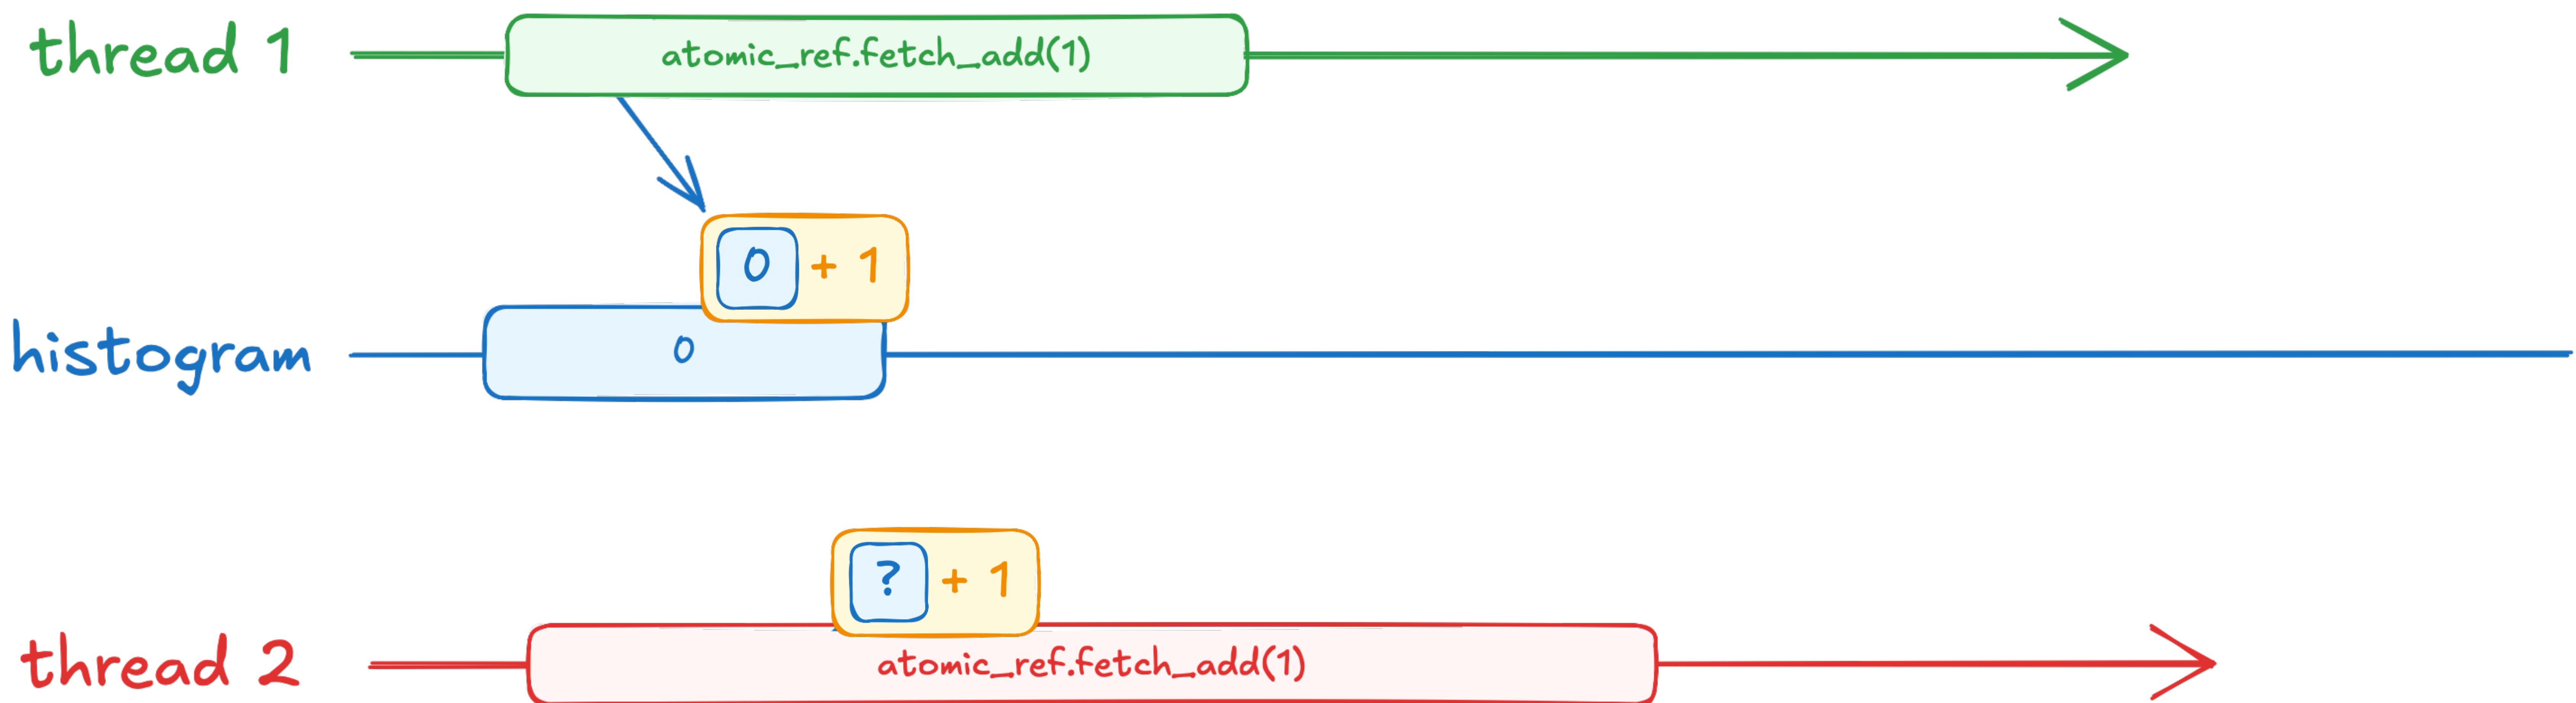
# Atomic Memory Operations



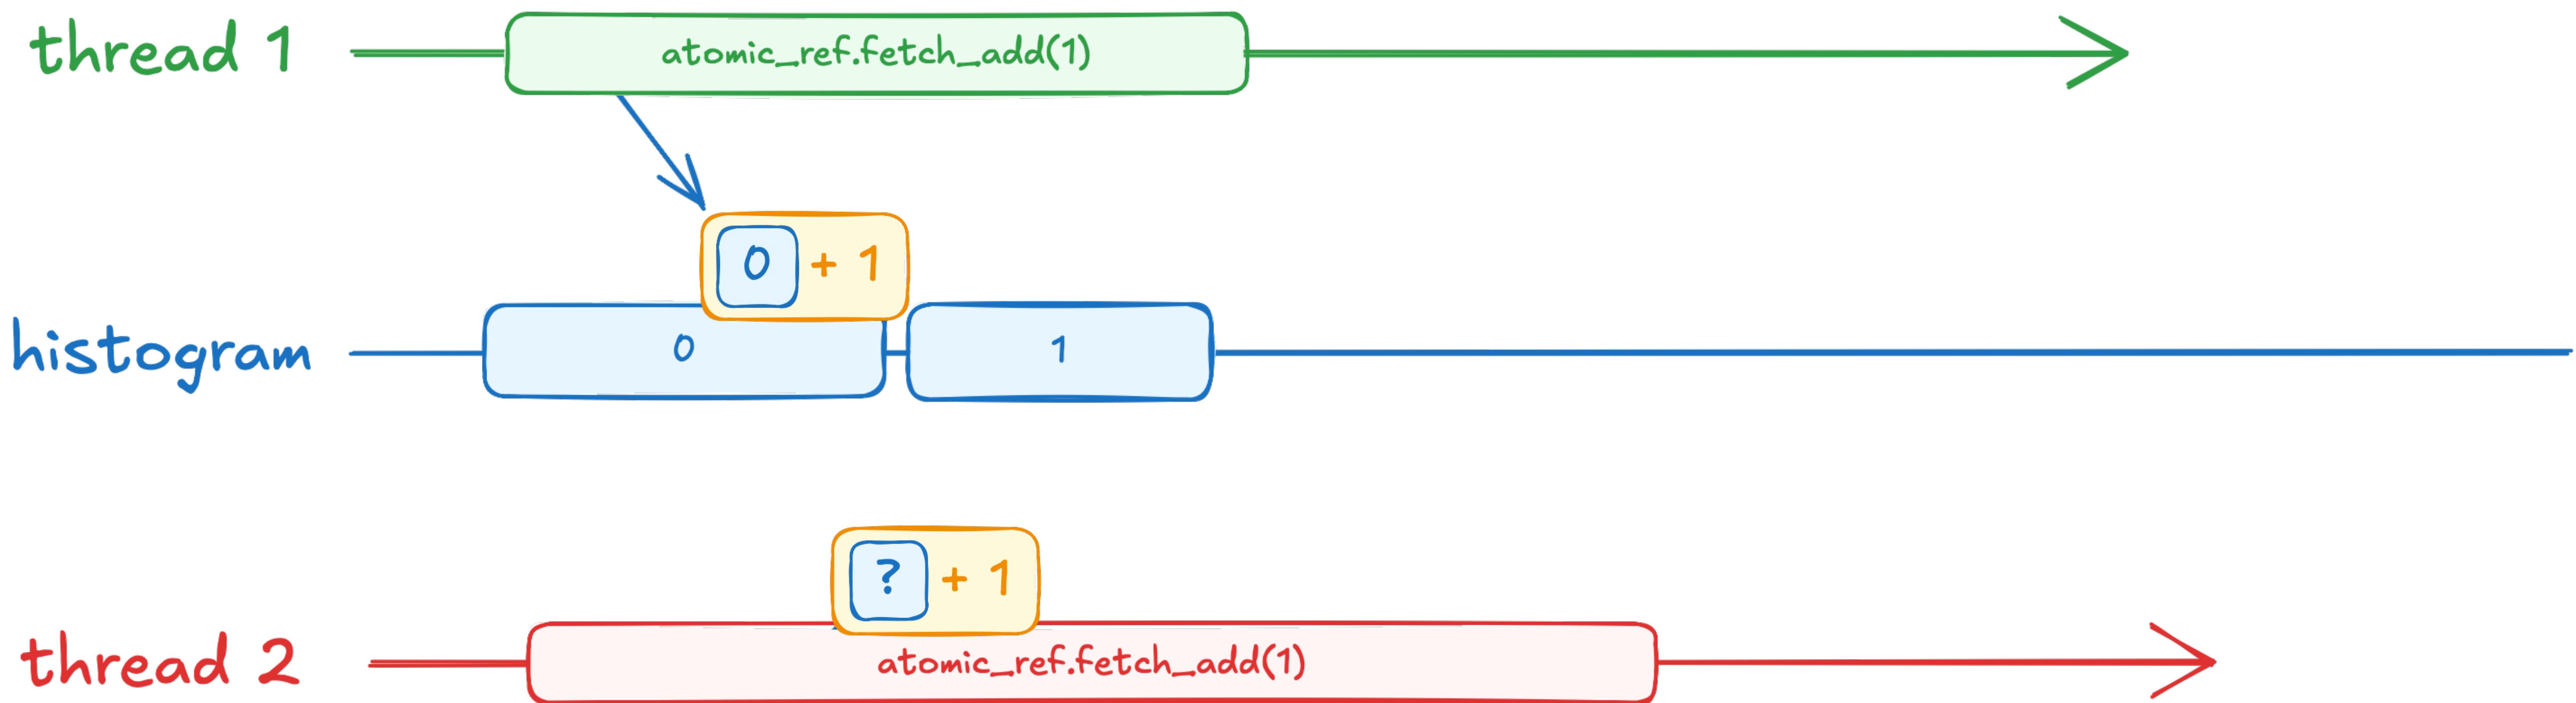
# Atomic Memory Operations



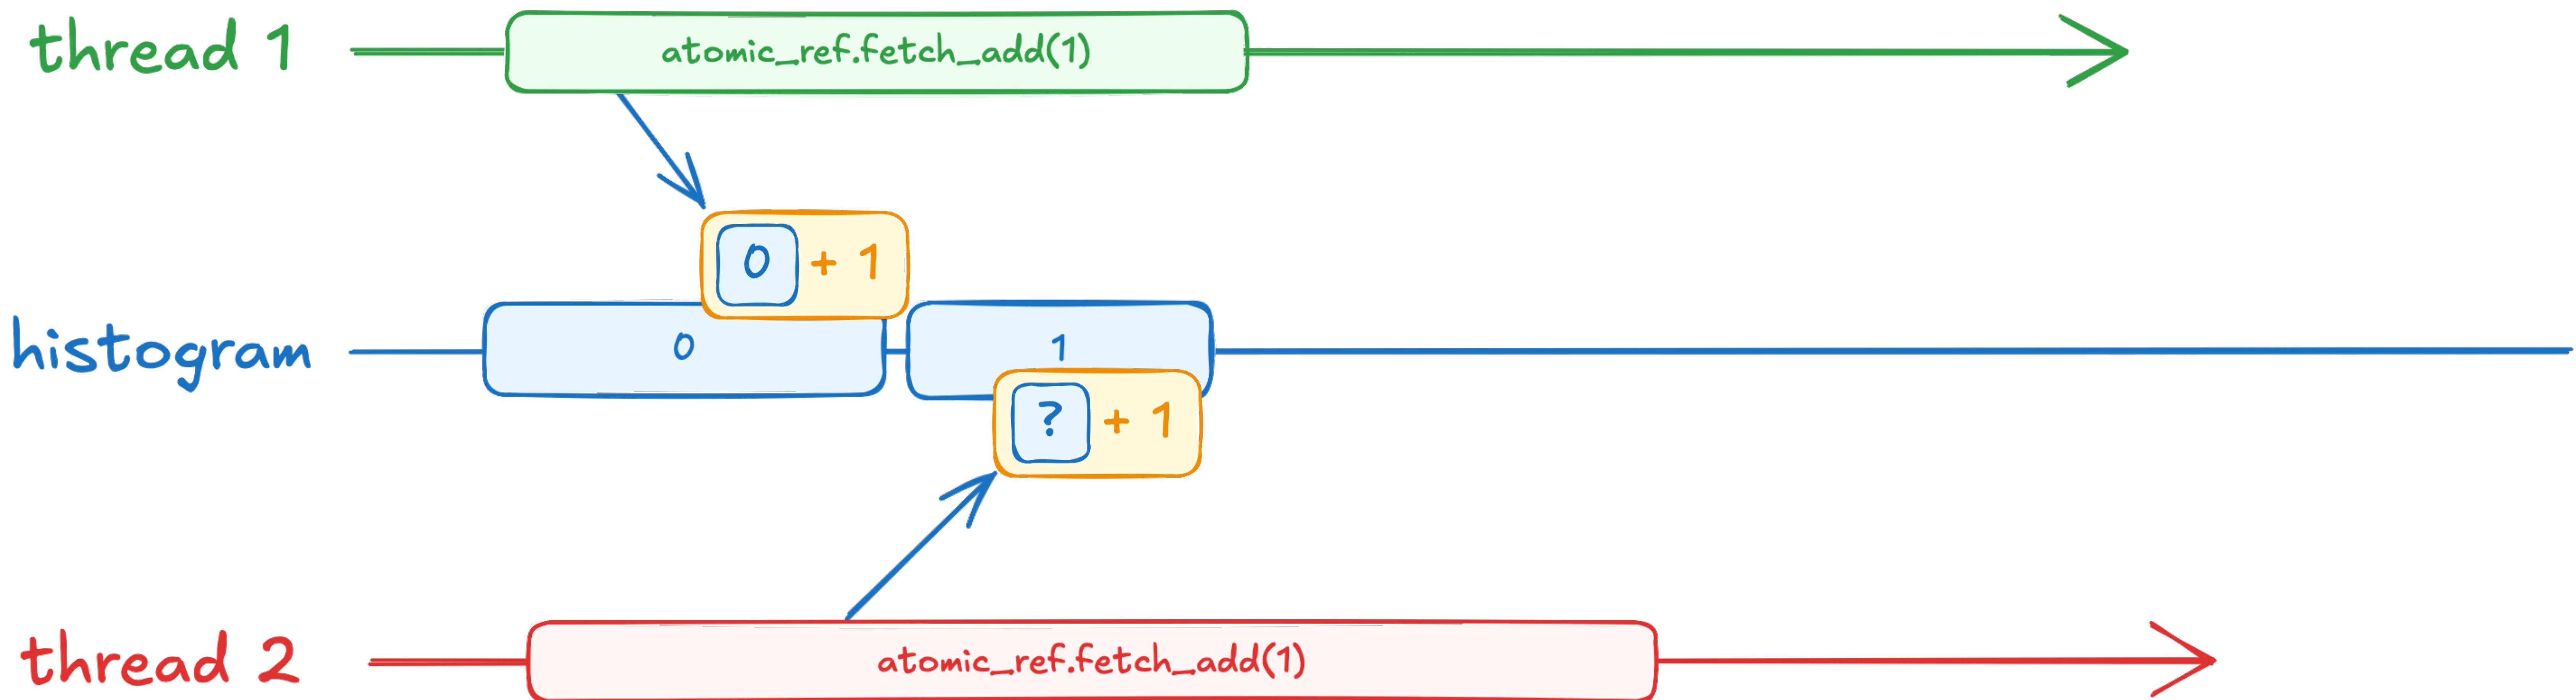
# Atomic Memory Operations



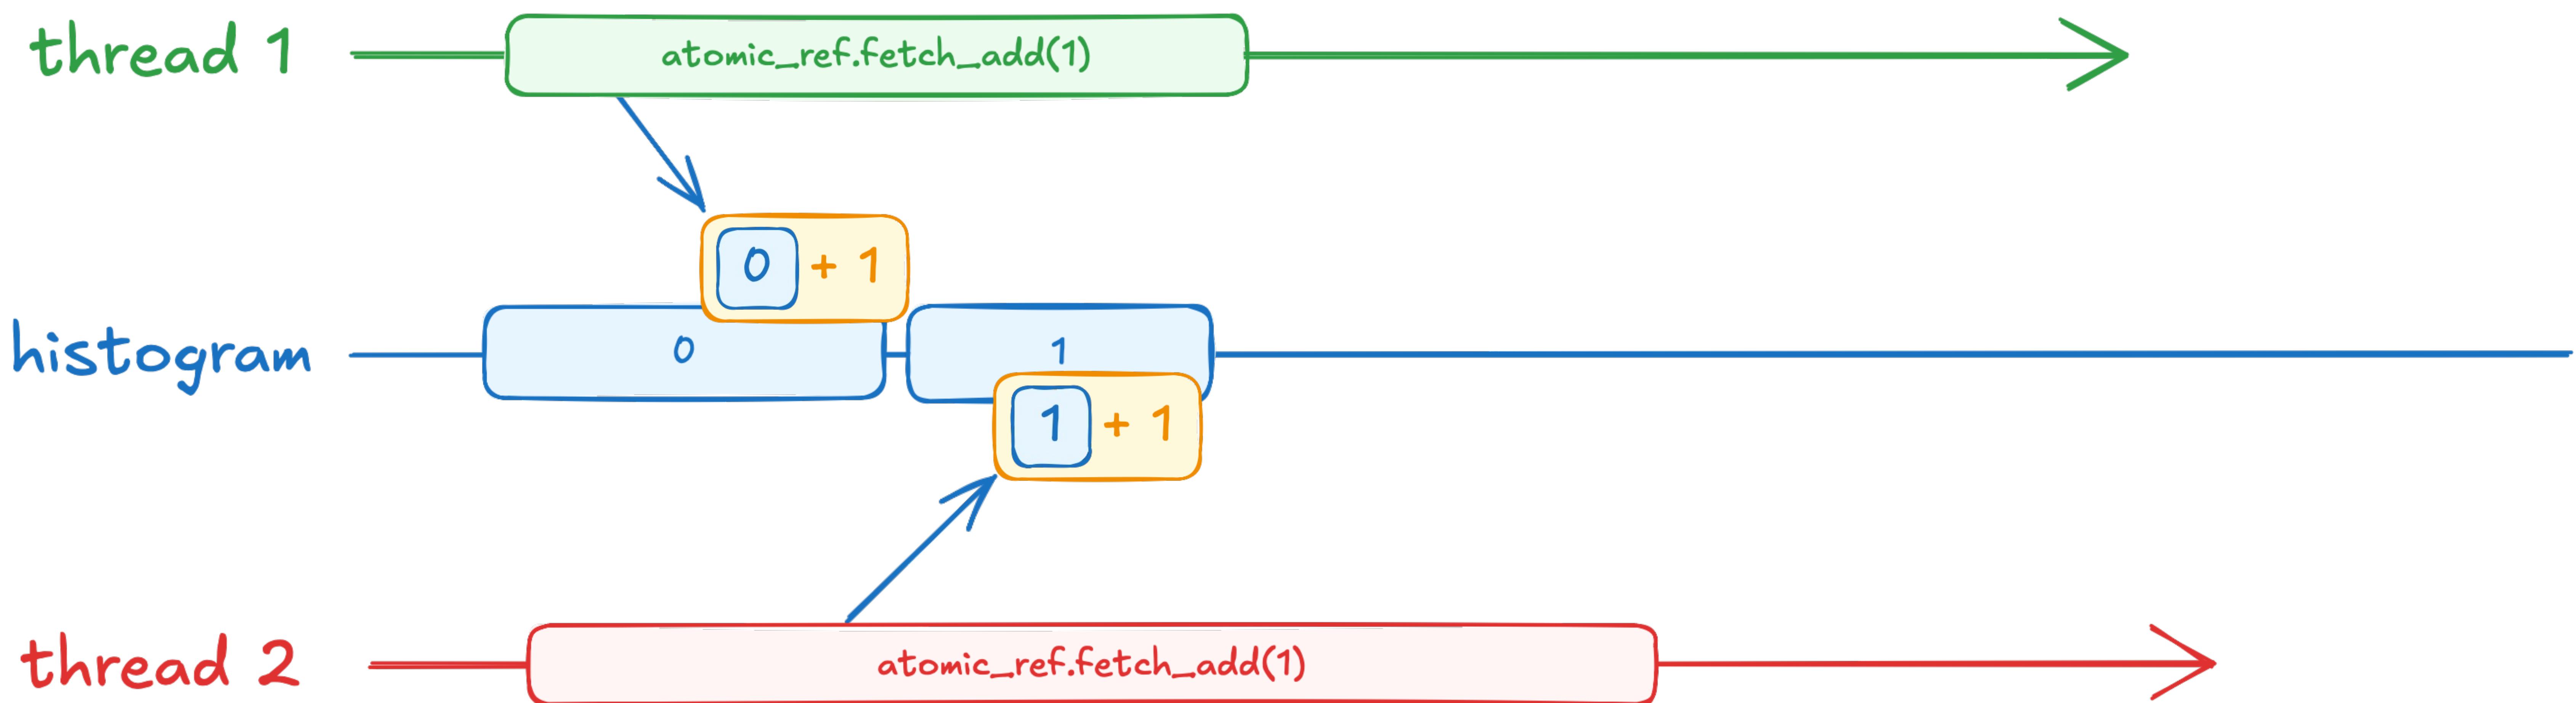
# Atomic Memory Operations



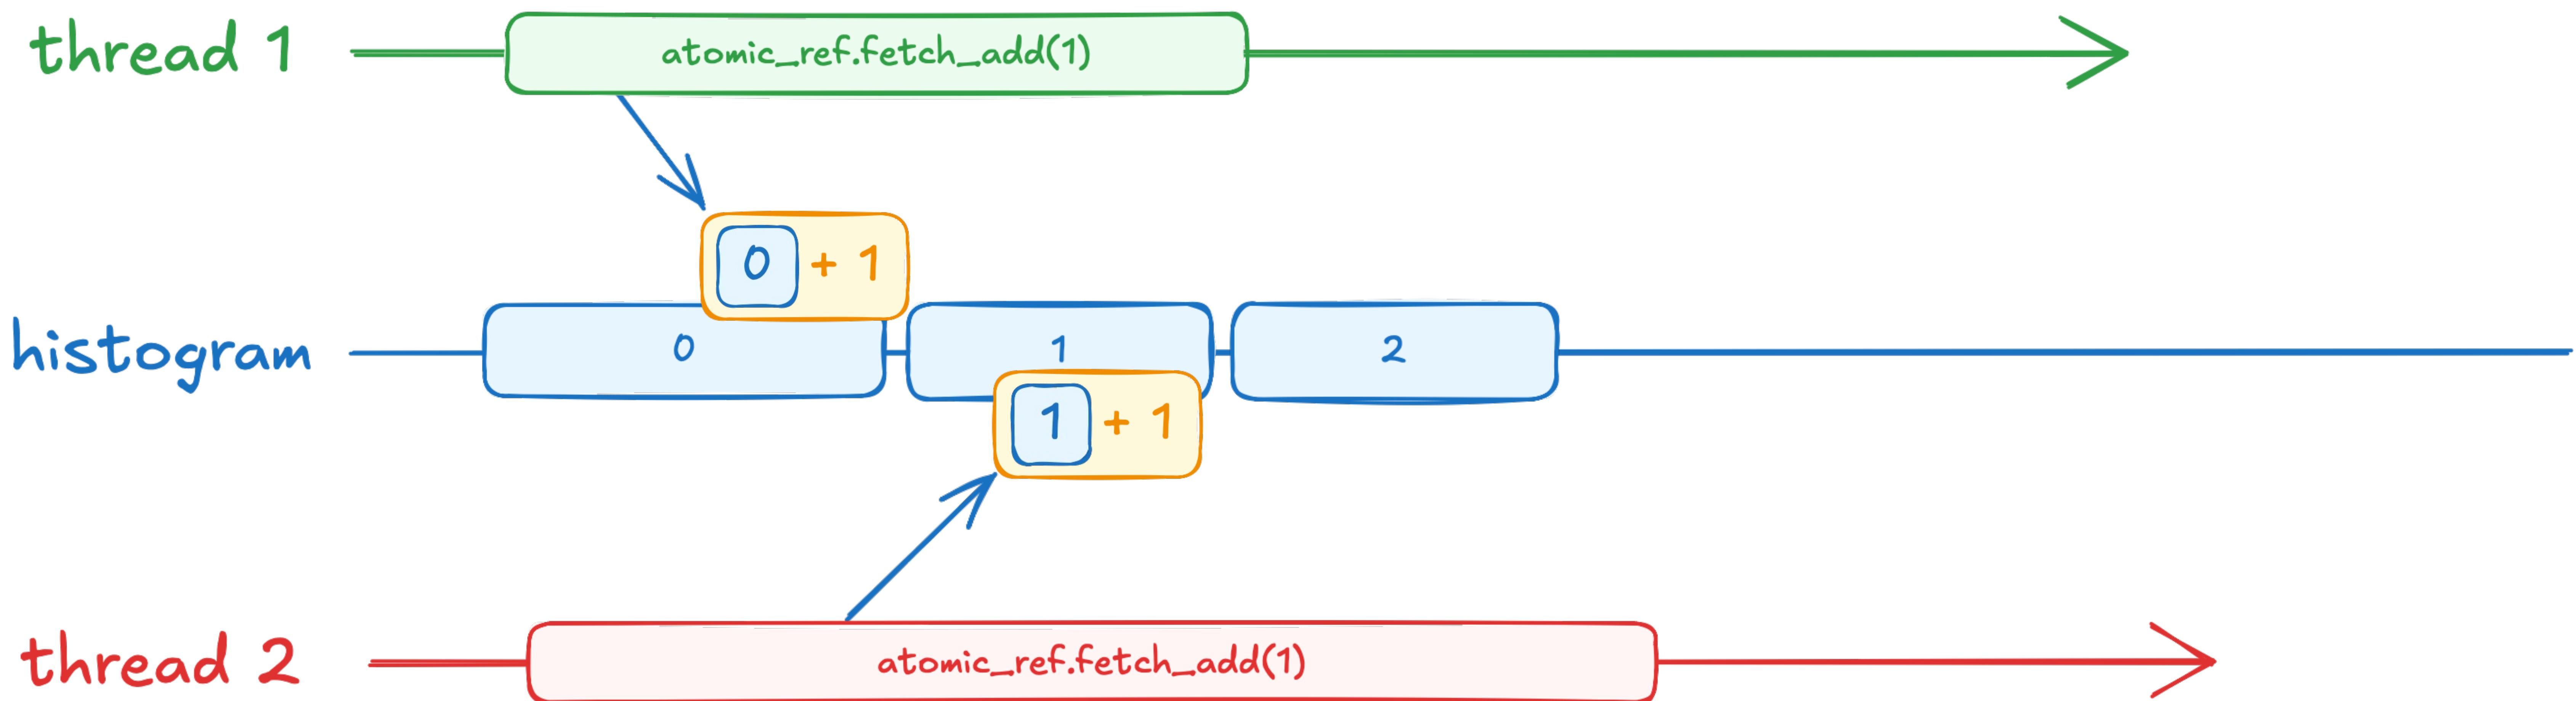
# Atomic Memory Operations



# Atomic Memory Operations



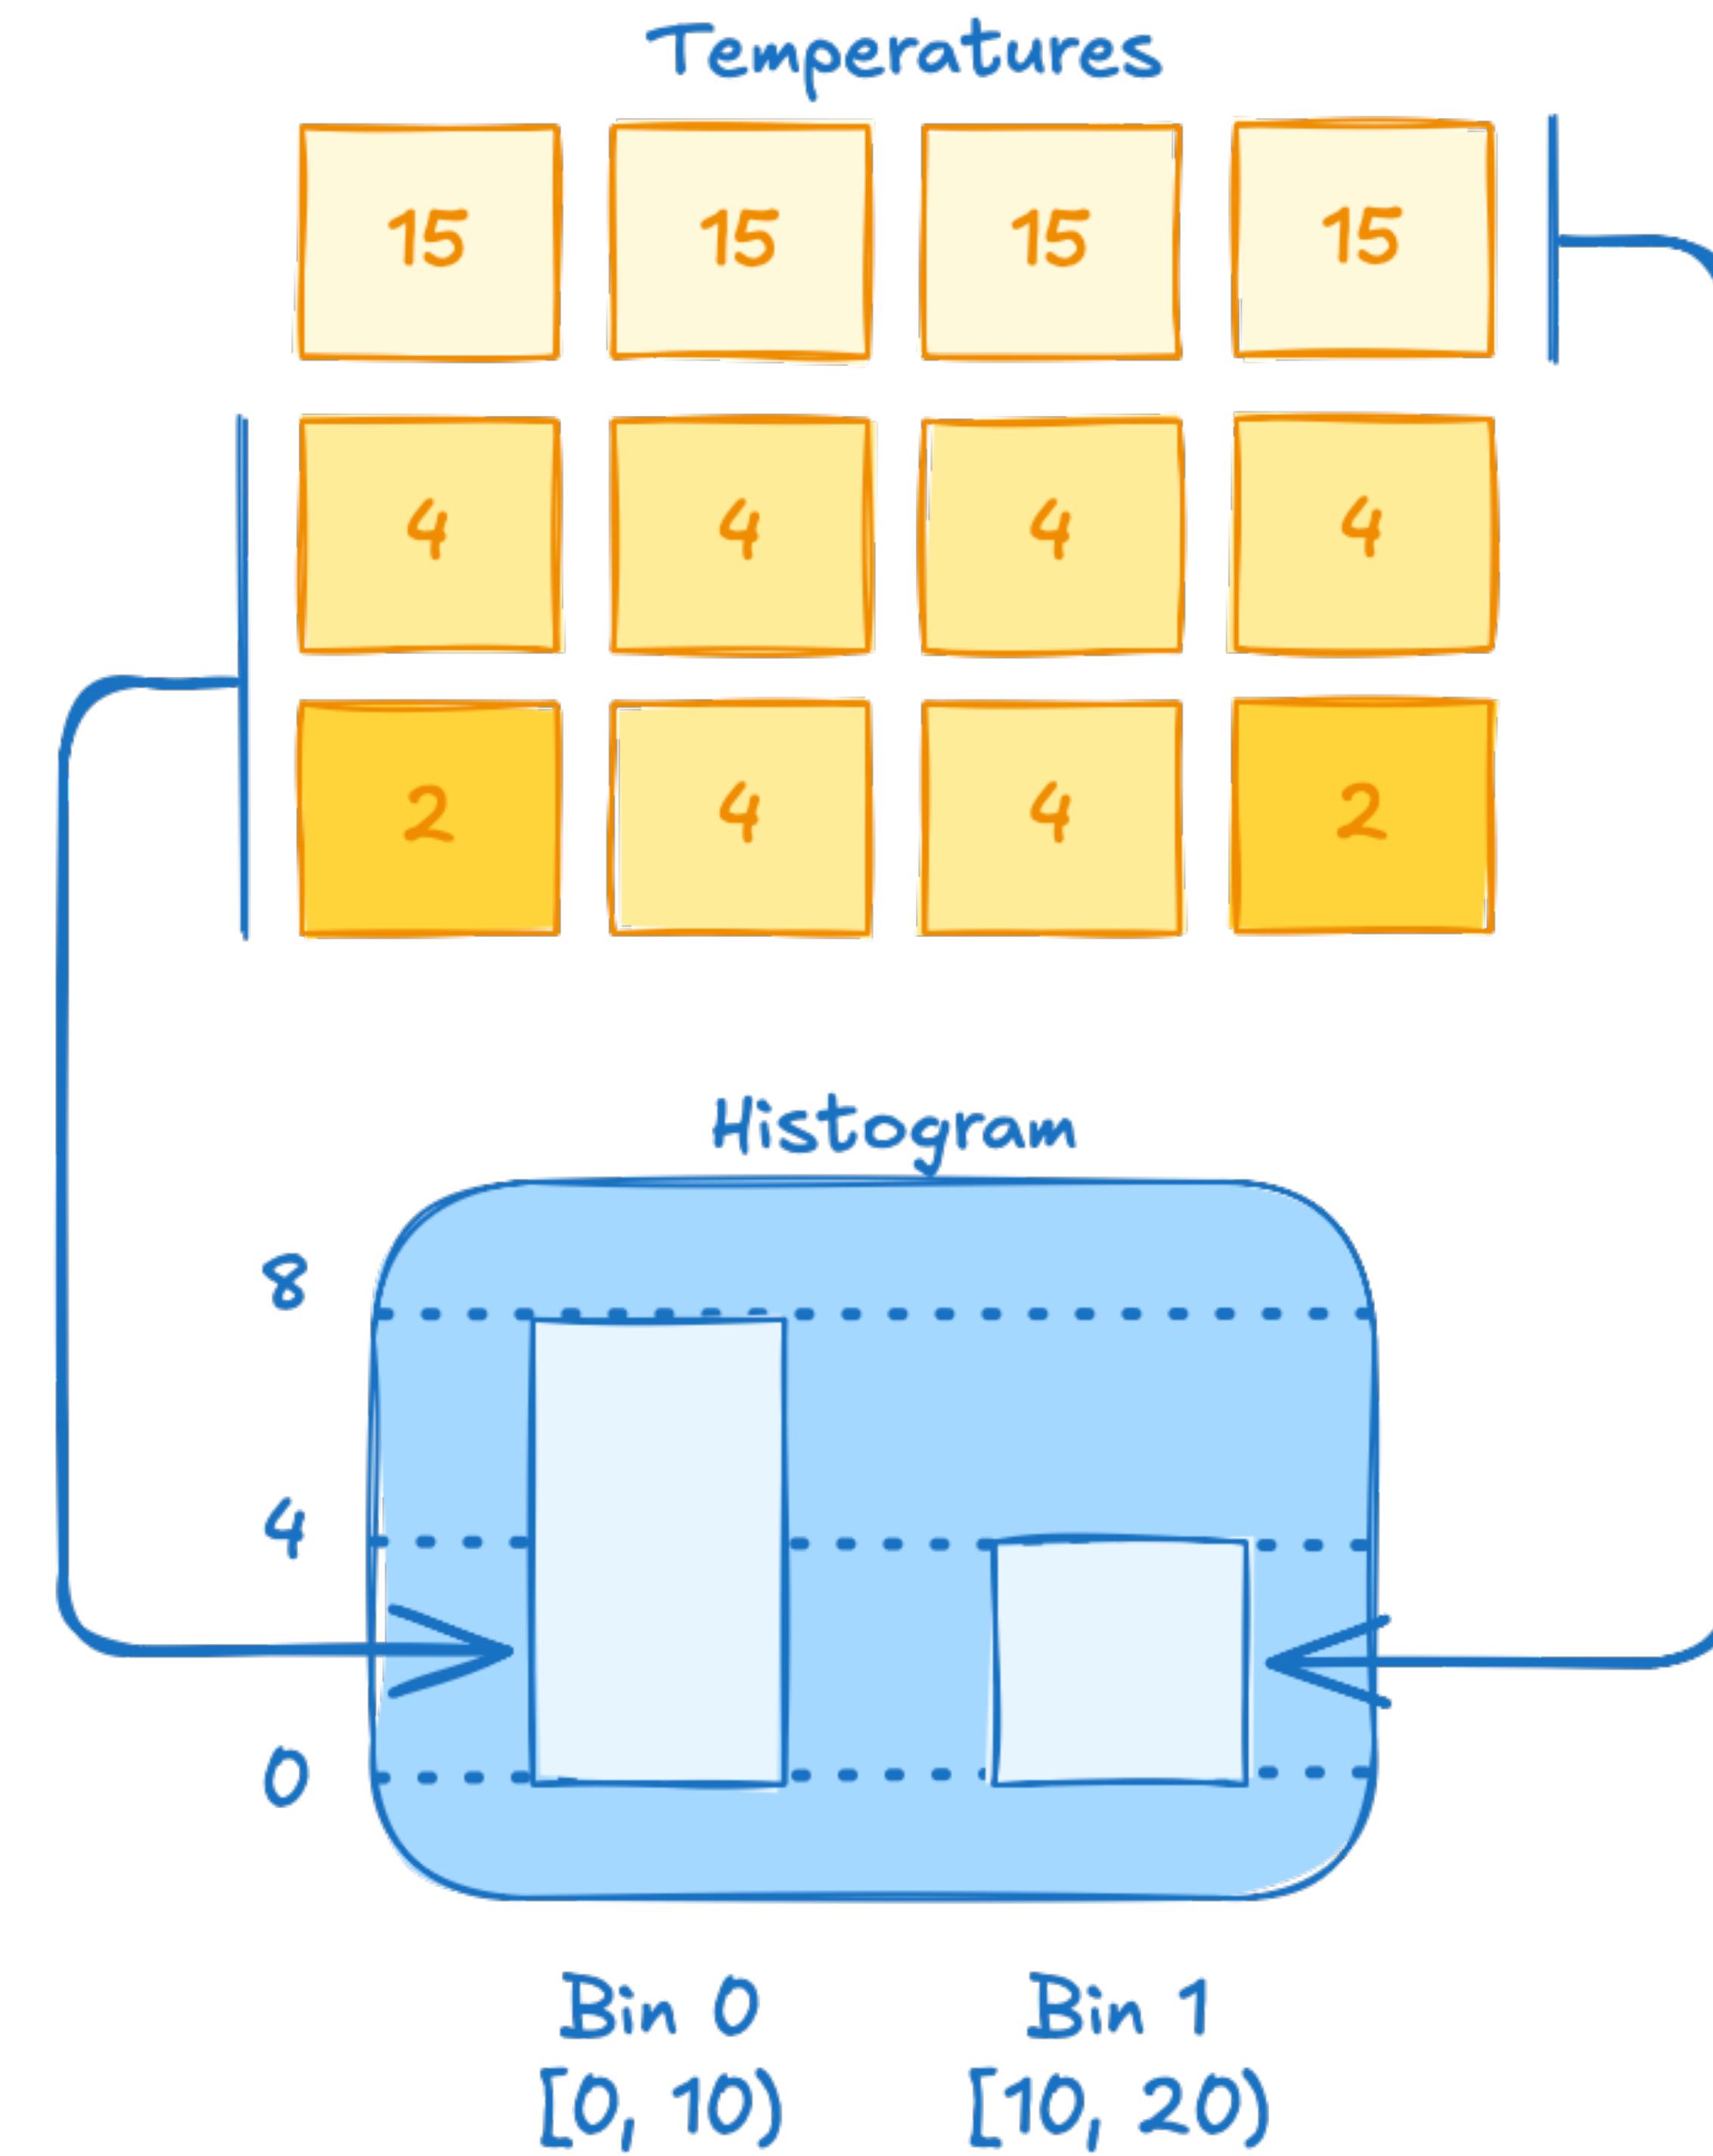
# Atomic Memory Operations



# Exercise: Fix Histogram

- Use atomics to fix data race in histogram

```
__global__ void histogram_kernel(
    cuda::std::span<float> temperatures,
    cuda::std::span<int> histogram)
{
    int cell = blockIdx.x * blockDim.x + threadIdx.x;
    int bin = static_cast<int>(temperatures[cell] / bin_width);
    int old_count = histogram[bin];
    int new_count = old_count + 1;
    histogram[bin] = new_count;
}
```



# Exercise: Fix Histogram

## Solution

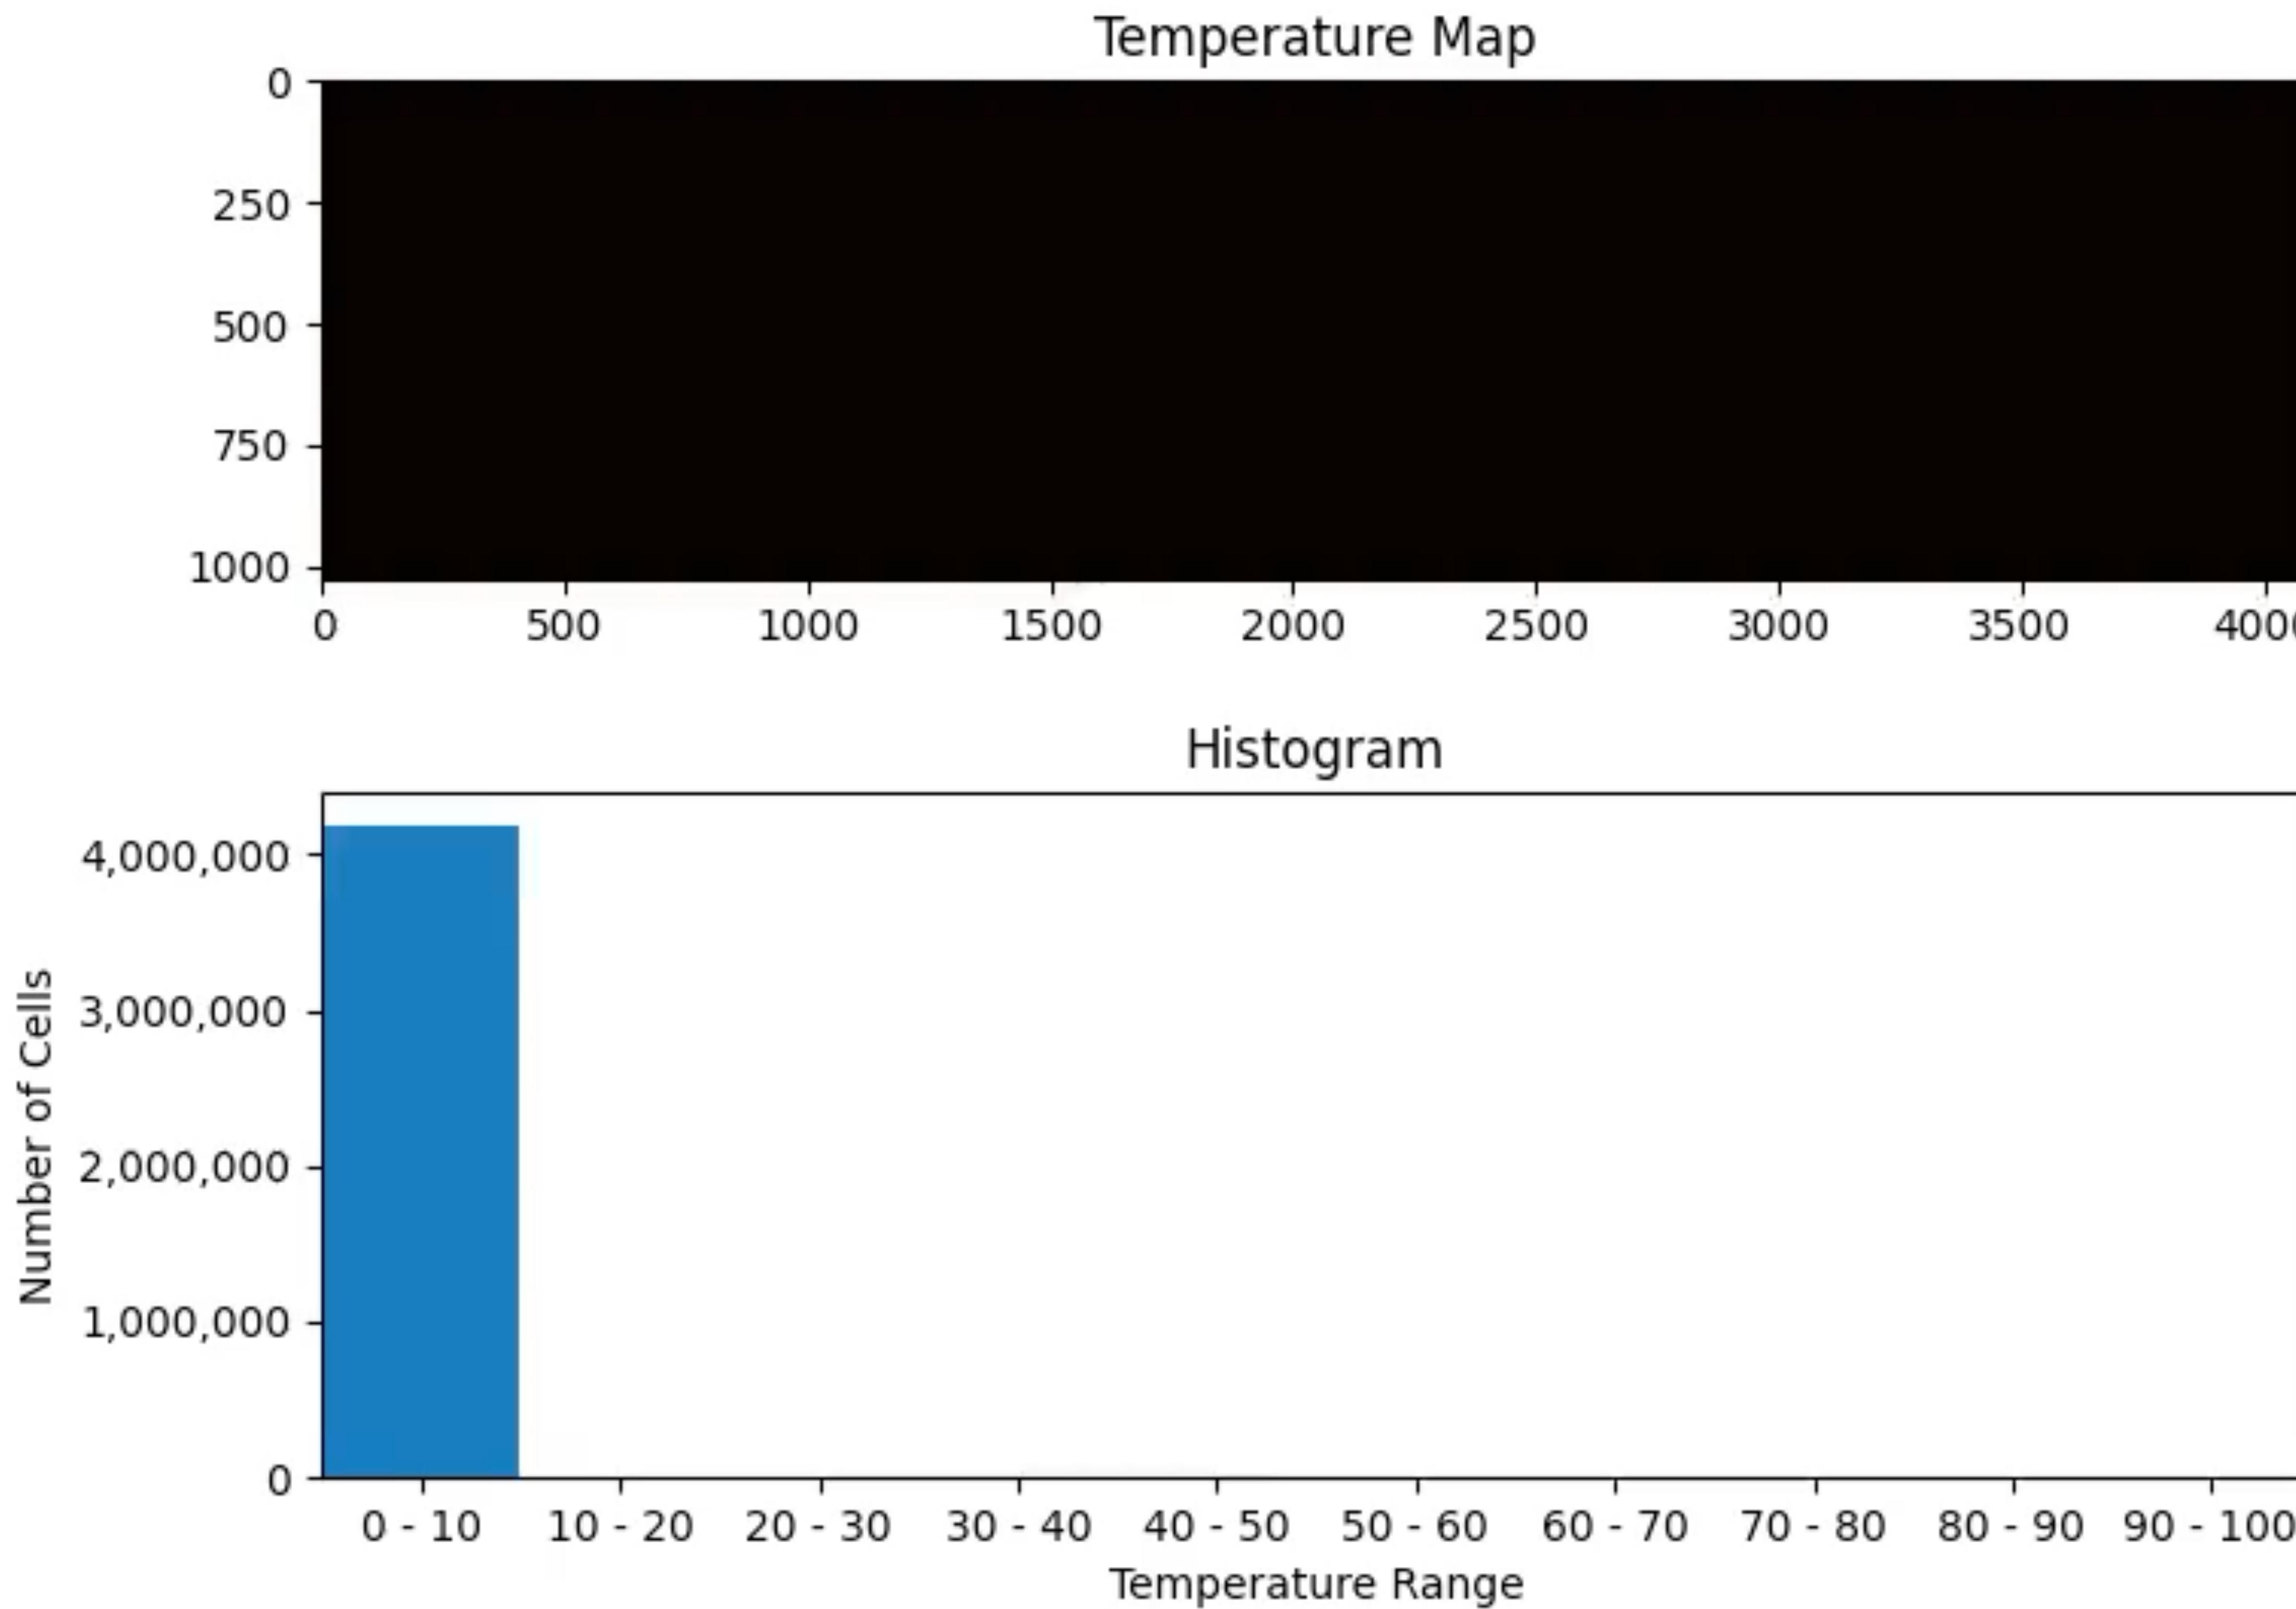
- Use atomics to fix data race in histogram

```
--global__ void histogram_kernel(  
    cuda::std::span<float> temperatures,  
    cuda::std::span<int> histogram)  
{  
    int cell = blockIdx.x * blockDim.x + threadIdx.x;  
    int bin = static_cast<int>(temperatures[cell] / bin_width);  
  
    cuda::std::atomic_ref<int> ref(histogram[bin]);  
    ref.fetch_add(1);  
}
```

- wrap reference to bin into atomic
- fetch add to bin using atomic ref

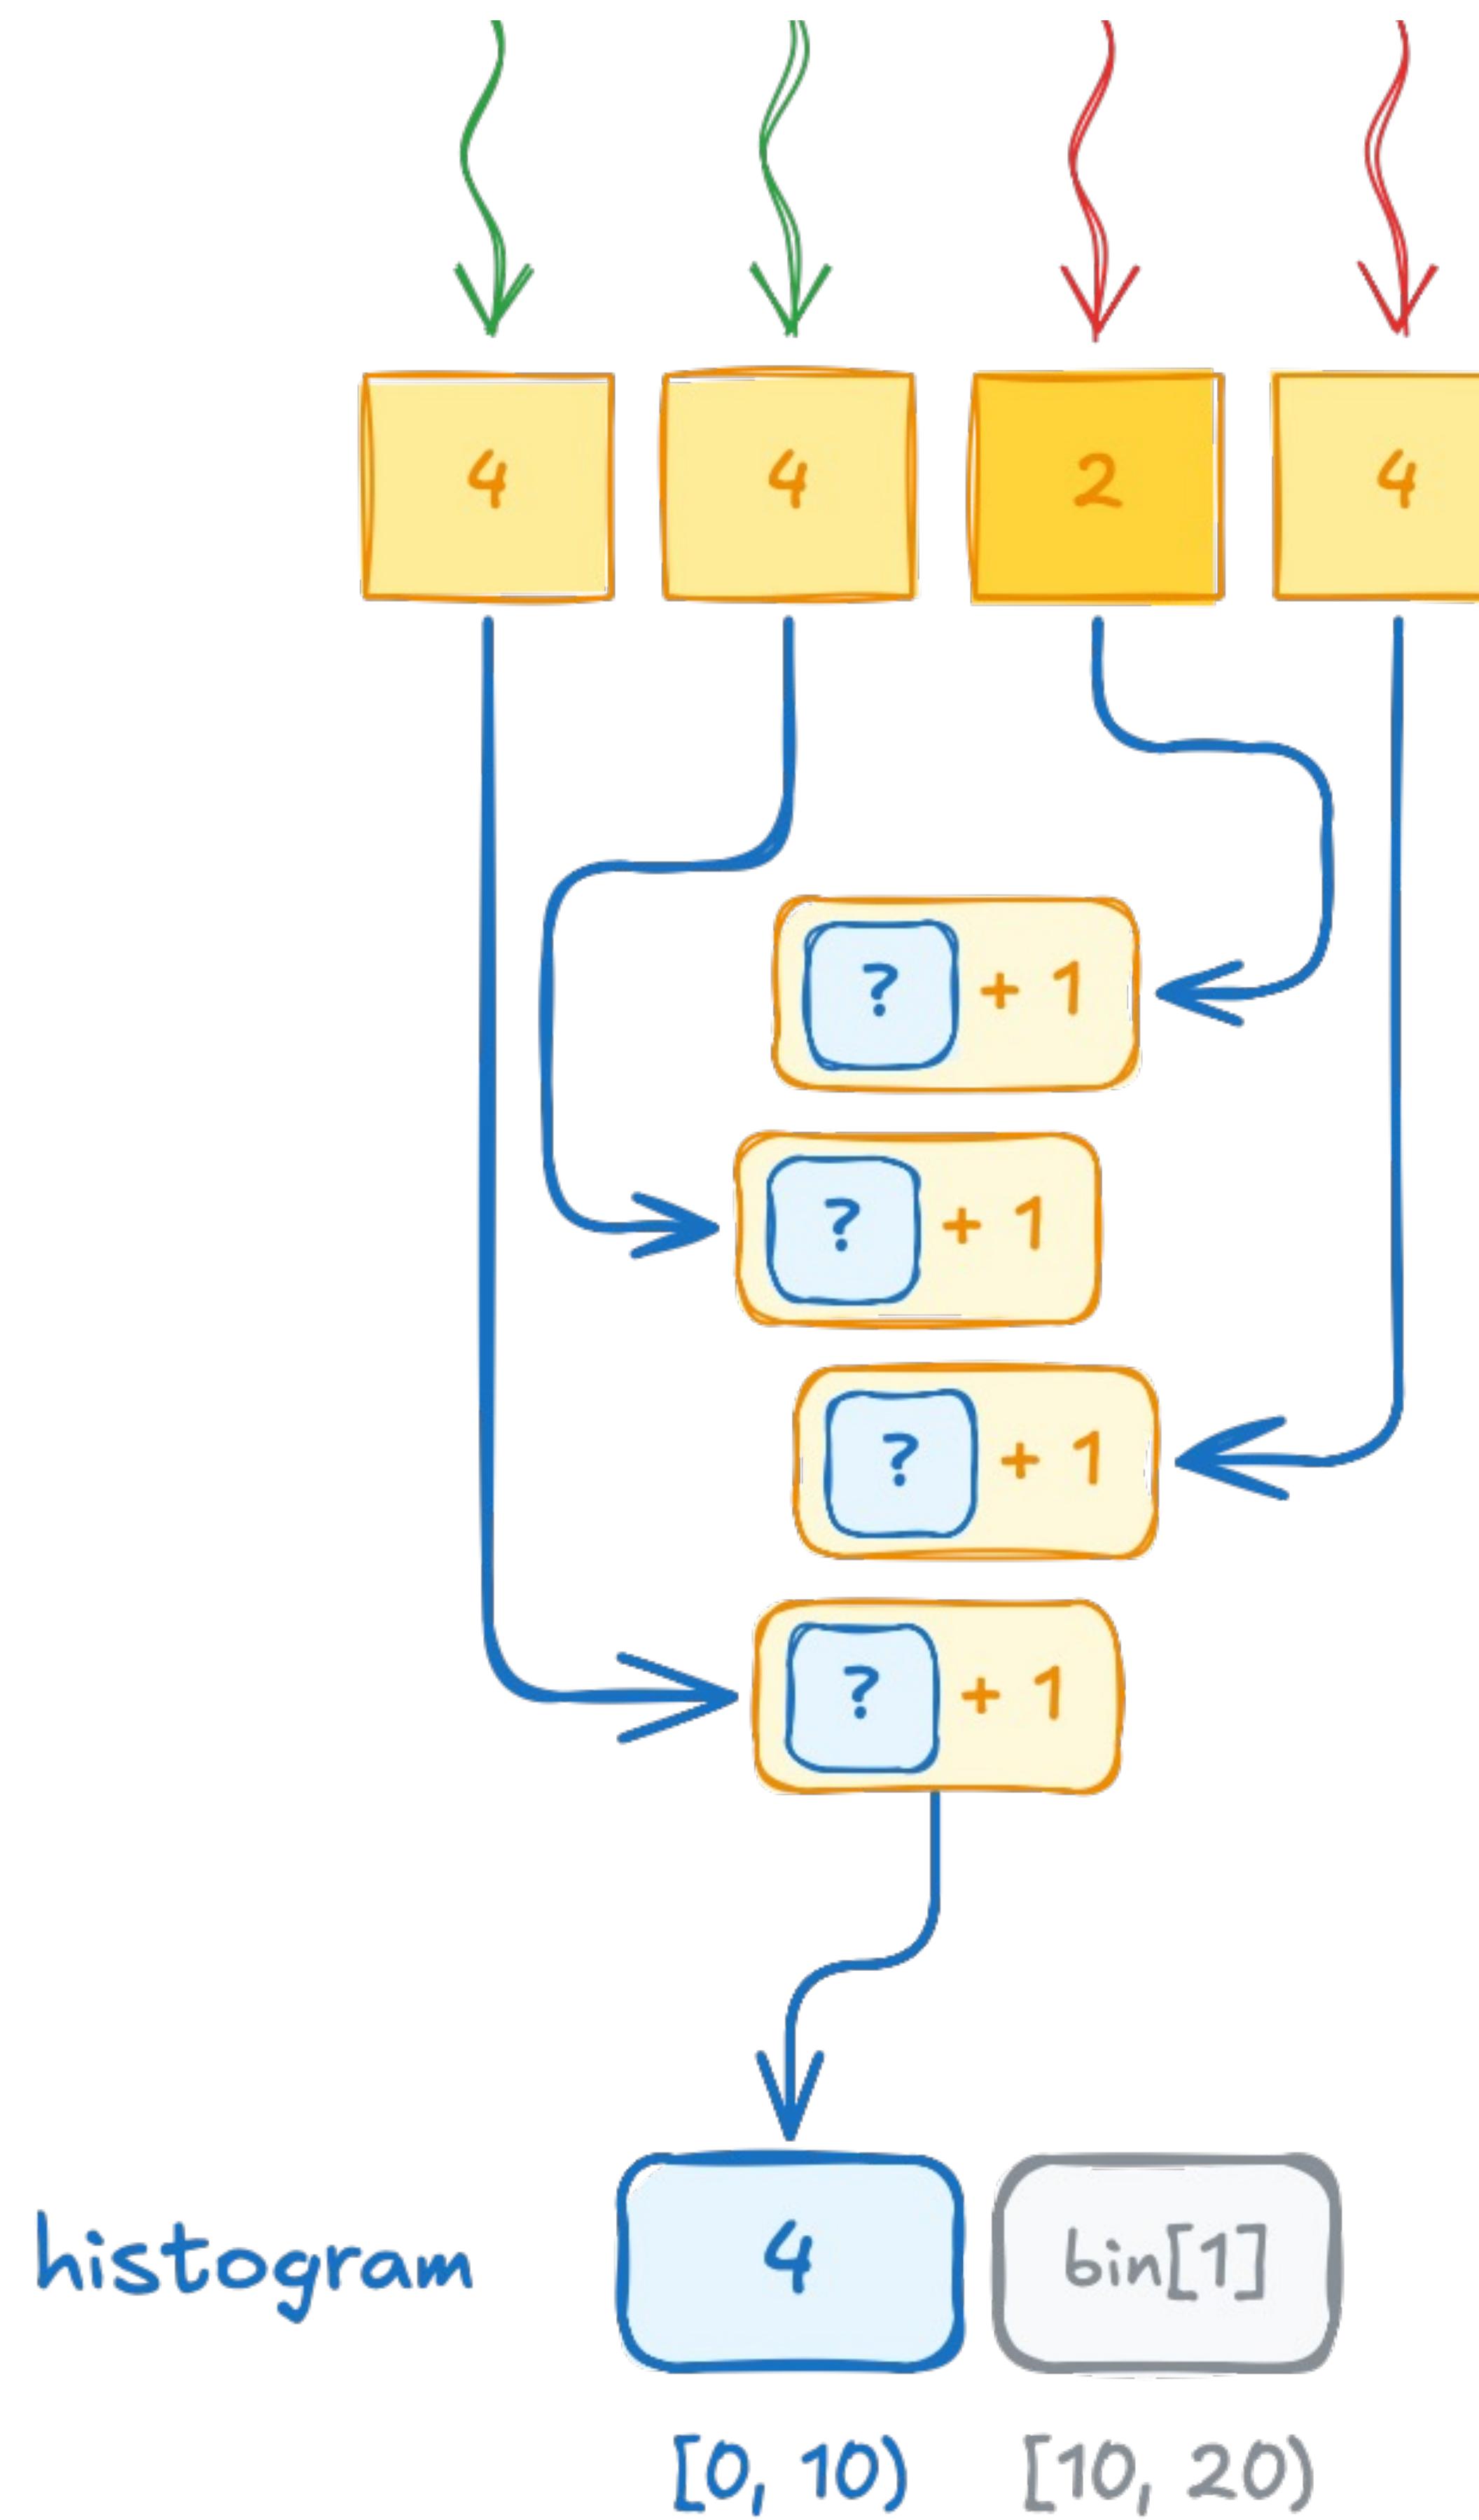
# Exercise: Fix Histogram

Solution



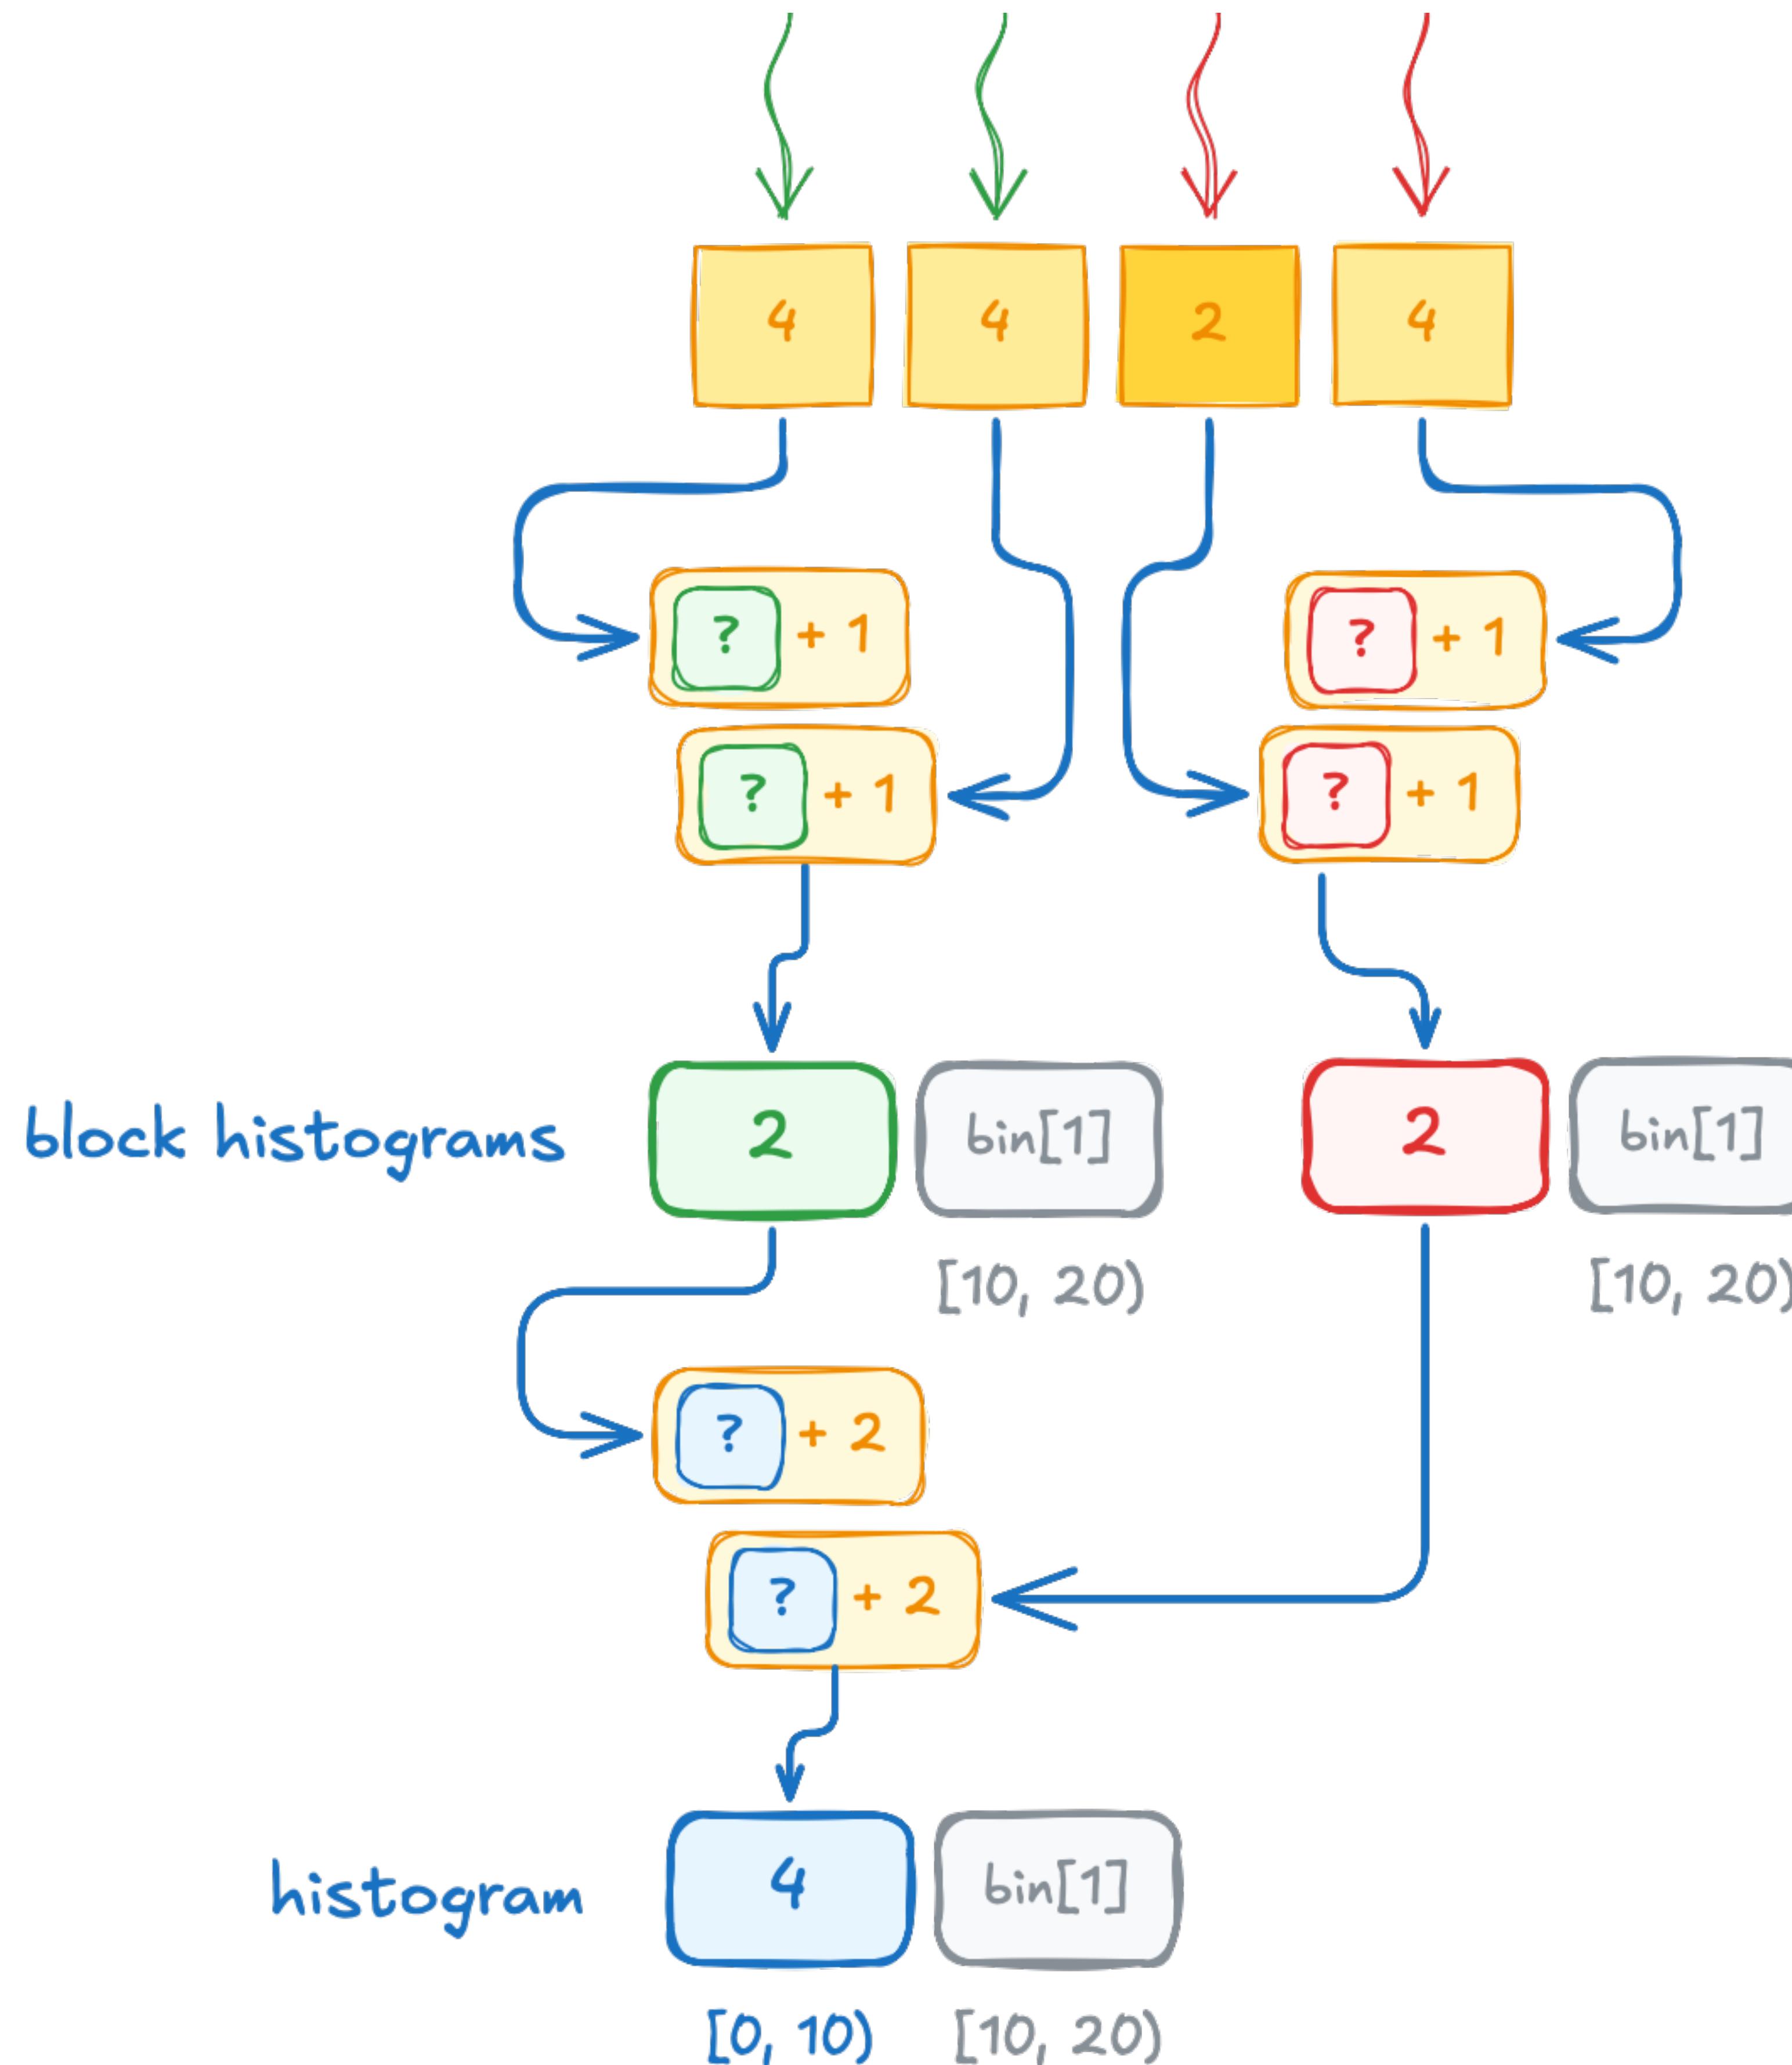
- We are getting correct results now!
- But the performance can't be worse
- Currently getting **~6 GB/s**

# New Source of Serialization



- Atomics let us regain functional correctness
- But this comes at a cost of serialization of memory accesses
- As you know, serialization should be avoided at all costs:
  - Consider thread block size = **256** and grid size = **16384**
  - If all cells mapping into the same bin,
  - we'll get a queue of **4194304** atomics
- So, what are our options?

# Privatized Histograms



- We could add privatized histograms, one per thread block:
  - Consider thread block size = **256** and grid size = **16384**
  - If all cells mapping into the same bin,
  - we'll get a queue of **256** atomics on private histograms followed by **16384** atomics on resulting histogram
- In limit, this reduces contention by ~ block size

# Implementing Privatized Histograms

```
__global__ void histogram_kernel(
    cuda::std::span<float> temperatures,
    cuda::std::span<int> block_histograms,
    cuda::std::span<int> histogram)
{
    cuda::std::span<int> block_histogram =
        block_histograms.subspan(
            blockIdx.x * histogram.size(),
            histogram.size());

    int cell = blockIdx.x * blockDim.x + threadIdx.x;
    int bin = static_cast<int>(temperatures[cell] / bin_width);

    cuda::std::atomic_ref<int> block_ref(block_histogram[bin]);
    block_ref.fetch_add(1);

    if (threadIdx.x < histogram.size())
    {
        cuda::std::atomic_ref<int> ref(histogram[threadIdx.x]);
        ref.fetch_add(block_histogram[threadIdx.x]);
    }
}
```

- Histogram kernel now takes another span with per-thread-block histograms

# Implementing Privatized Histograms

```
--global__ void histogram_kernel(
    cuda::std::span<float> temperatures,
    cuda::std::span<int> block_histograms,
    cuda::std::span<int> histogram)
{
    cuda::std::span<int> block_histogram = block_histograms.subspan(
        blockIdx.x * histogram.size(),
        histogram.size());

    int cell = blockIdx.x * blockDim.x + threadIdx.x;
    int bin = static_cast<int>(temperatures[cell] / bin_width);

    cuda::std::atomic_ref<int> block_ref(block_histogram[bin]);
    block_ref.fetch_add(1);

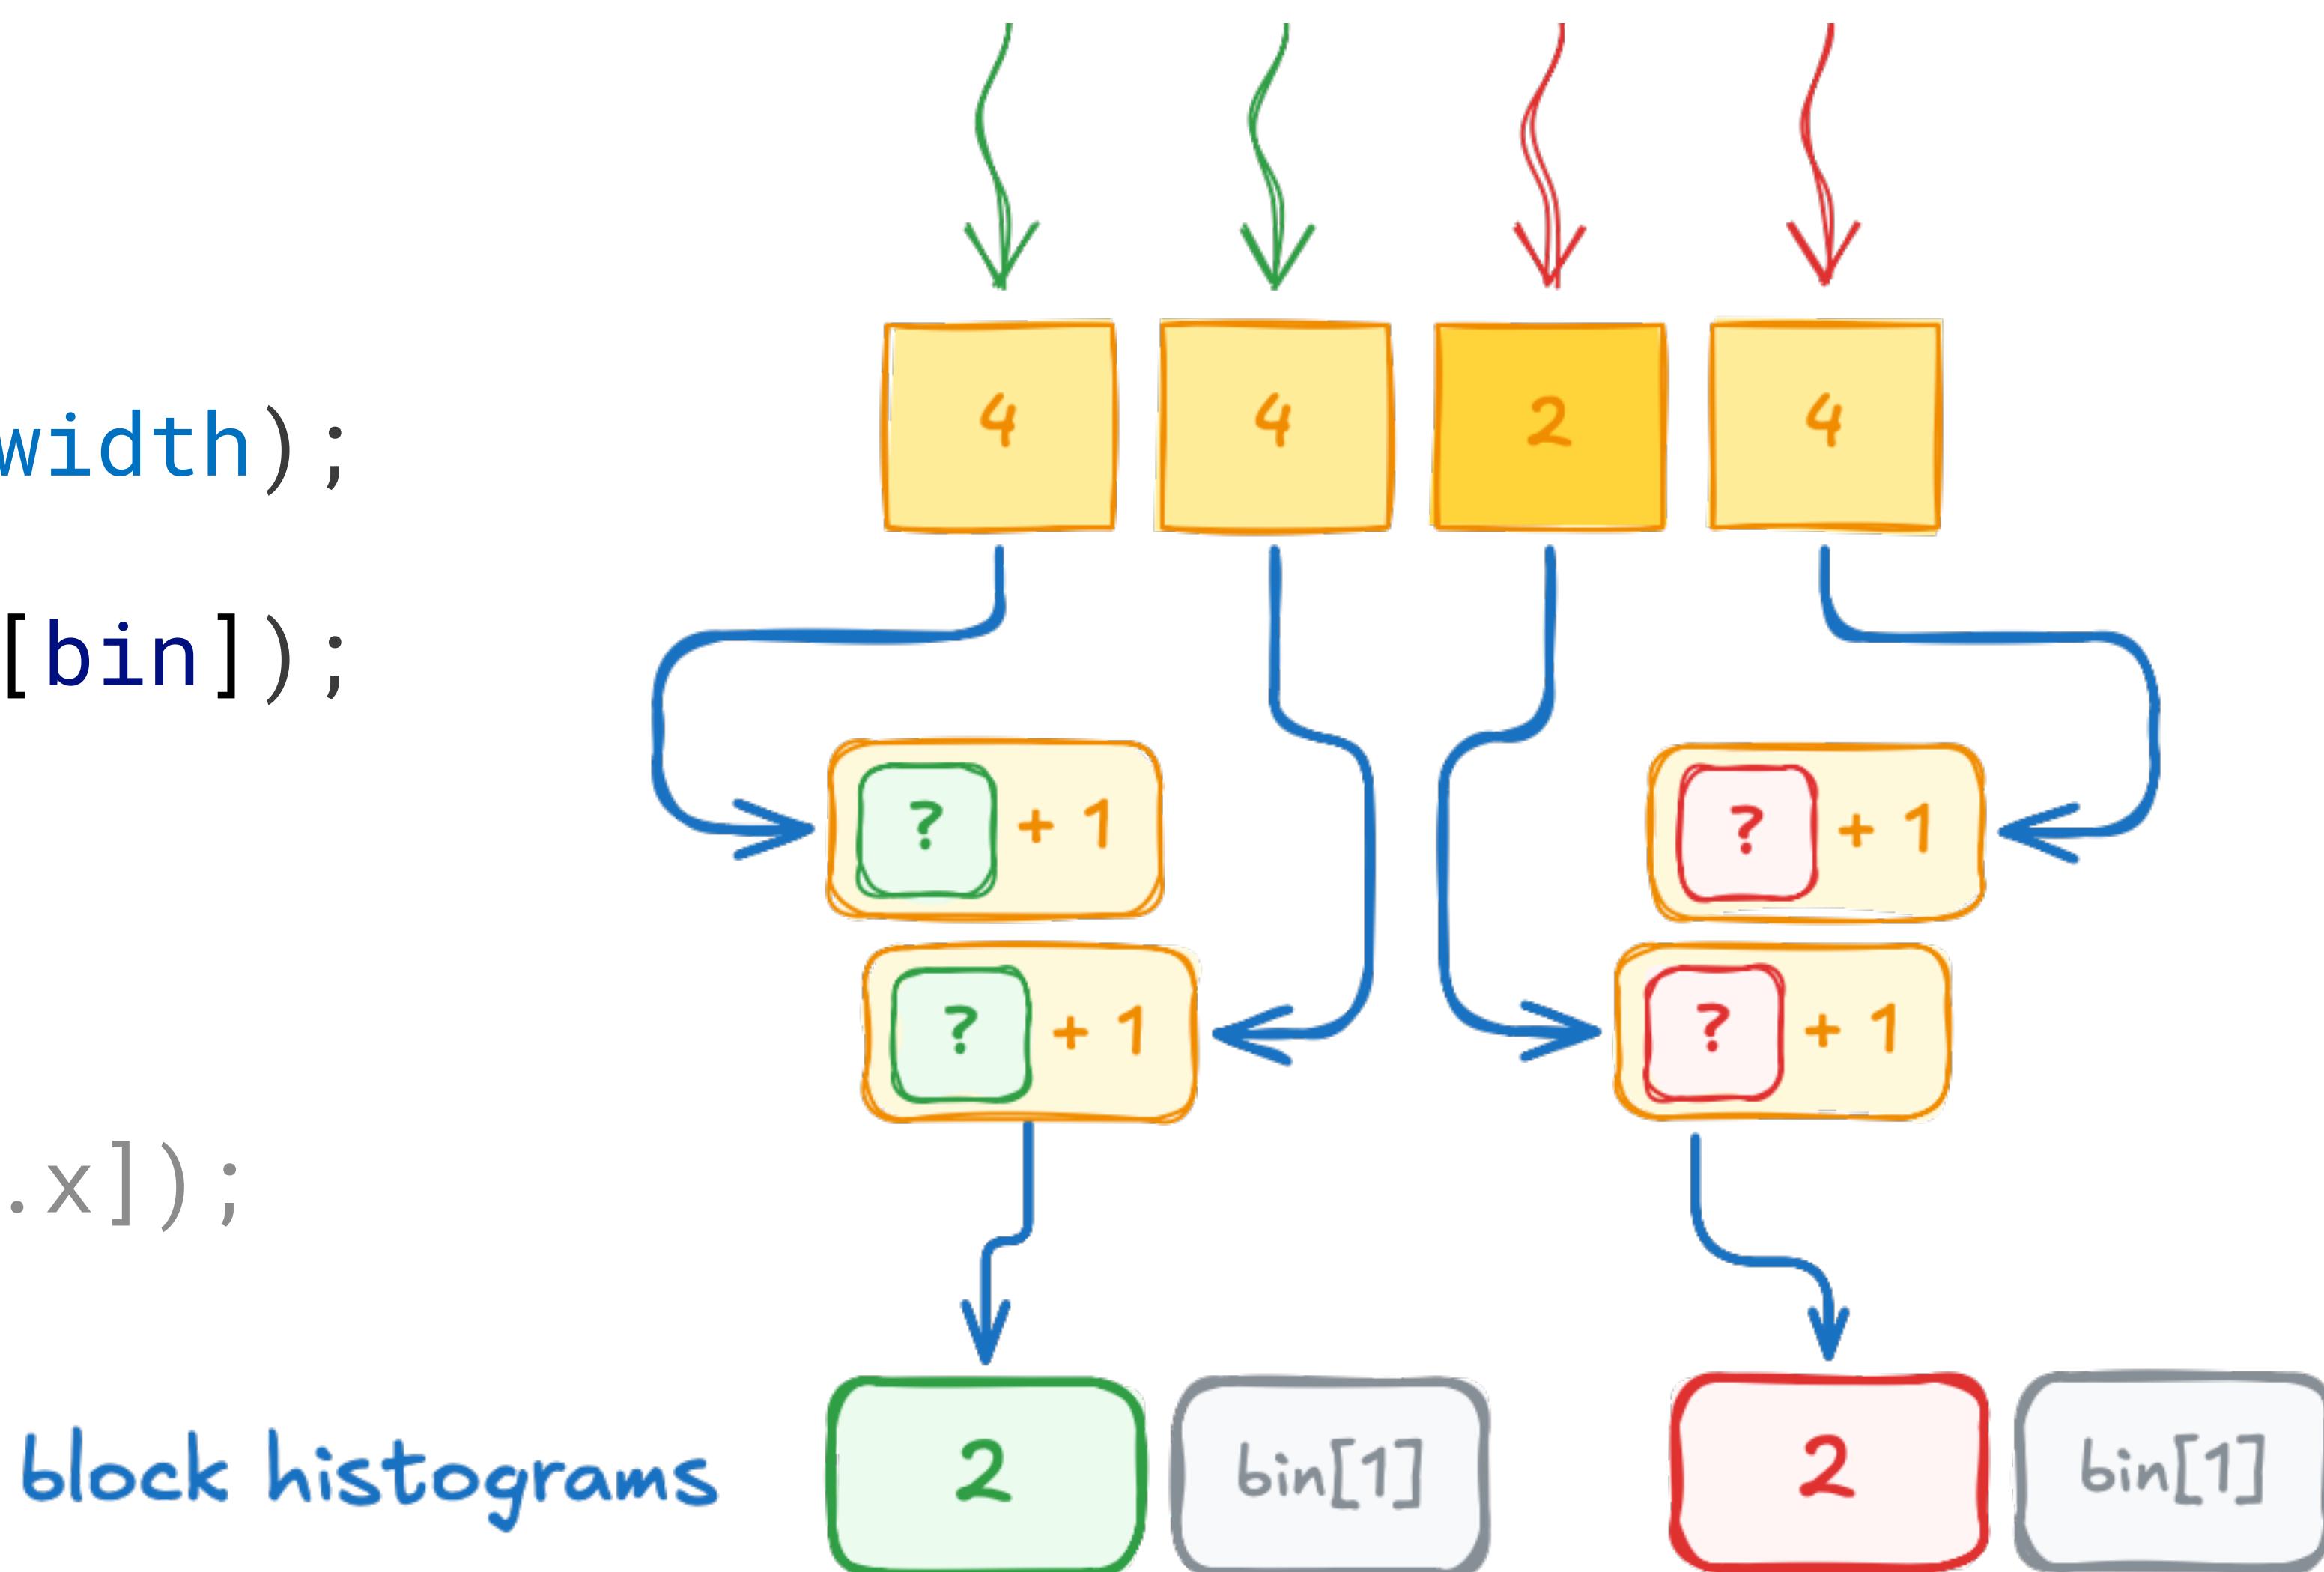
    if (threadIdx.x < histogram.size())
    {
        cuda::std::atomic_ref<int> ref(histogram[threadIdx.x]);
        ref.fetch_add(block_histogram[threadIdx.x]);
    }
}
```

- obtain a view over thread-block private histogram
- subspan offset to private histogram is block index multiplied by histogram size

# Implementing Privatized Histograms

```
--global__ void histogram_kernel(  
    cuda::span<float> temperatures,  
    cuda::span<int> block_histograms,  
    cuda::span<int> histogram)  
{  
    cuda::span<int> block_histogram =  
        block_histograms.subspan(  
            blockIdx.x * histogram.size(),  
            histogram.size());  
  
    int cell = blockIdx.x * blockDim.x + threadIdx.x;  
    int bin = static_cast<int>(temperatures[cell] / bin_width);  
  
    cuda::atomic_ref<int> block_ref(block_histogram[bin]);  
    block_ref.fetch_add(1);  
  
    if (threadIdx.x < histogram.size())  
    {  
        cuda::atomic_ref<int> ref(histogram[threadIdx.x]);  
        ref.fetch_add(block_histogram[threadIdx.x]);  
    }  
}
```

- Atomically increment thread-block histograms



# Implementing Privatized Histograms

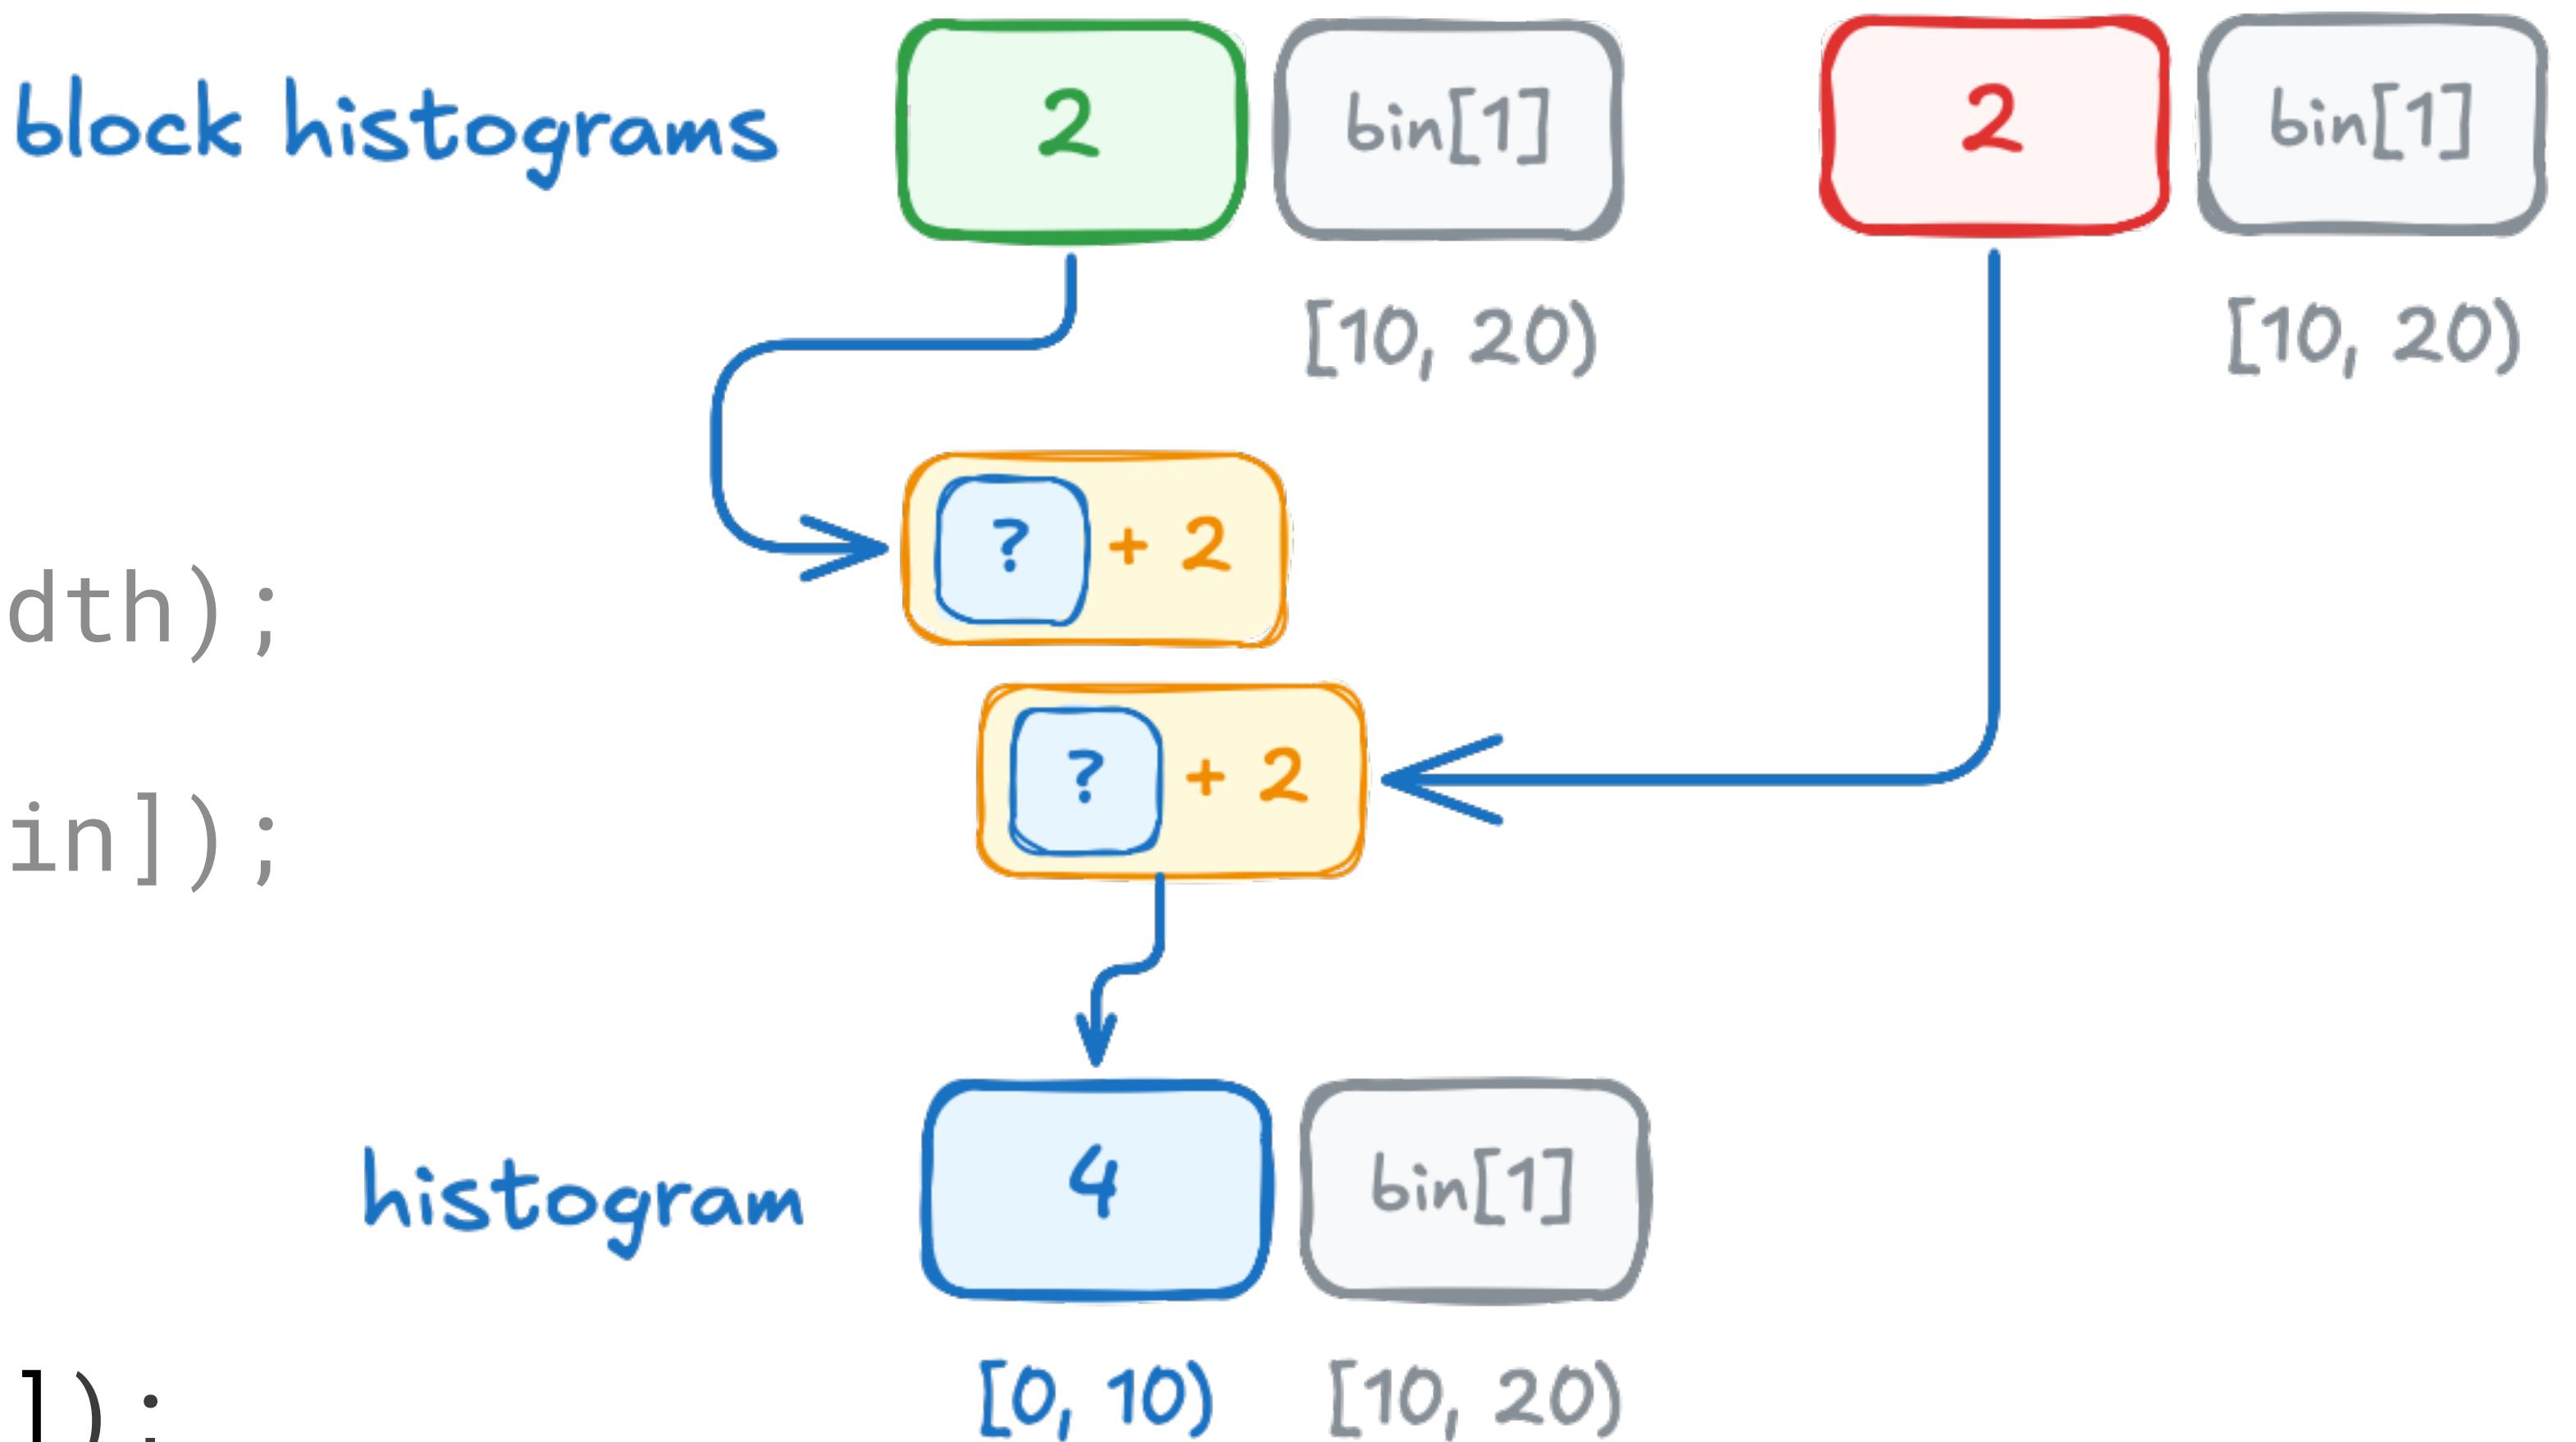
```
--global__ void histogram_kernel(
    cuda::span<float> temperatures,
    cuda::span<int> block_histograms,
    cuda::span<int> histogram)
{
    cuda::span<int> block_histogram =
        block_histograms.subspan(
            blockIdx.x * histogram.size(),
            histogram.size());

    int cell = blockIdx.x * blockDim.x + threadIdx.x;
    int bin = static_cast<int>(temperatures[cell] / bin_width);

    cuda::atomic_ref<int> block_ref(block_histogram[bin]);
    block_ref.fetch_add(1);

    if (threadIdx.x < histogram.size())
    {
        cuda::atomic_ref<int> ref(histogram[threadIdx.x]);
        ref.fetch_add(block_histogram[threadIdx.x]);
    }
}
```

- Contribute private histograms to resulting one



# Implementing Privatized Histograms

```
--global__ void histogram_kernel(
    cuda::std::span<float> temperatures,
    cuda::std::span<int> block_histograms,
    cuda::std::span<int> histogram)
{
    cuda::std::span<int> block_histogram =
        block_histograms.subspan(
            blockIdx.x * histogram.size(),
            histogram.size());

    int cell = blockIdx.x * blockDim.x + threadIdx.x;
    int bin = static_cast<int>(temperatures[cell] / bin_width);

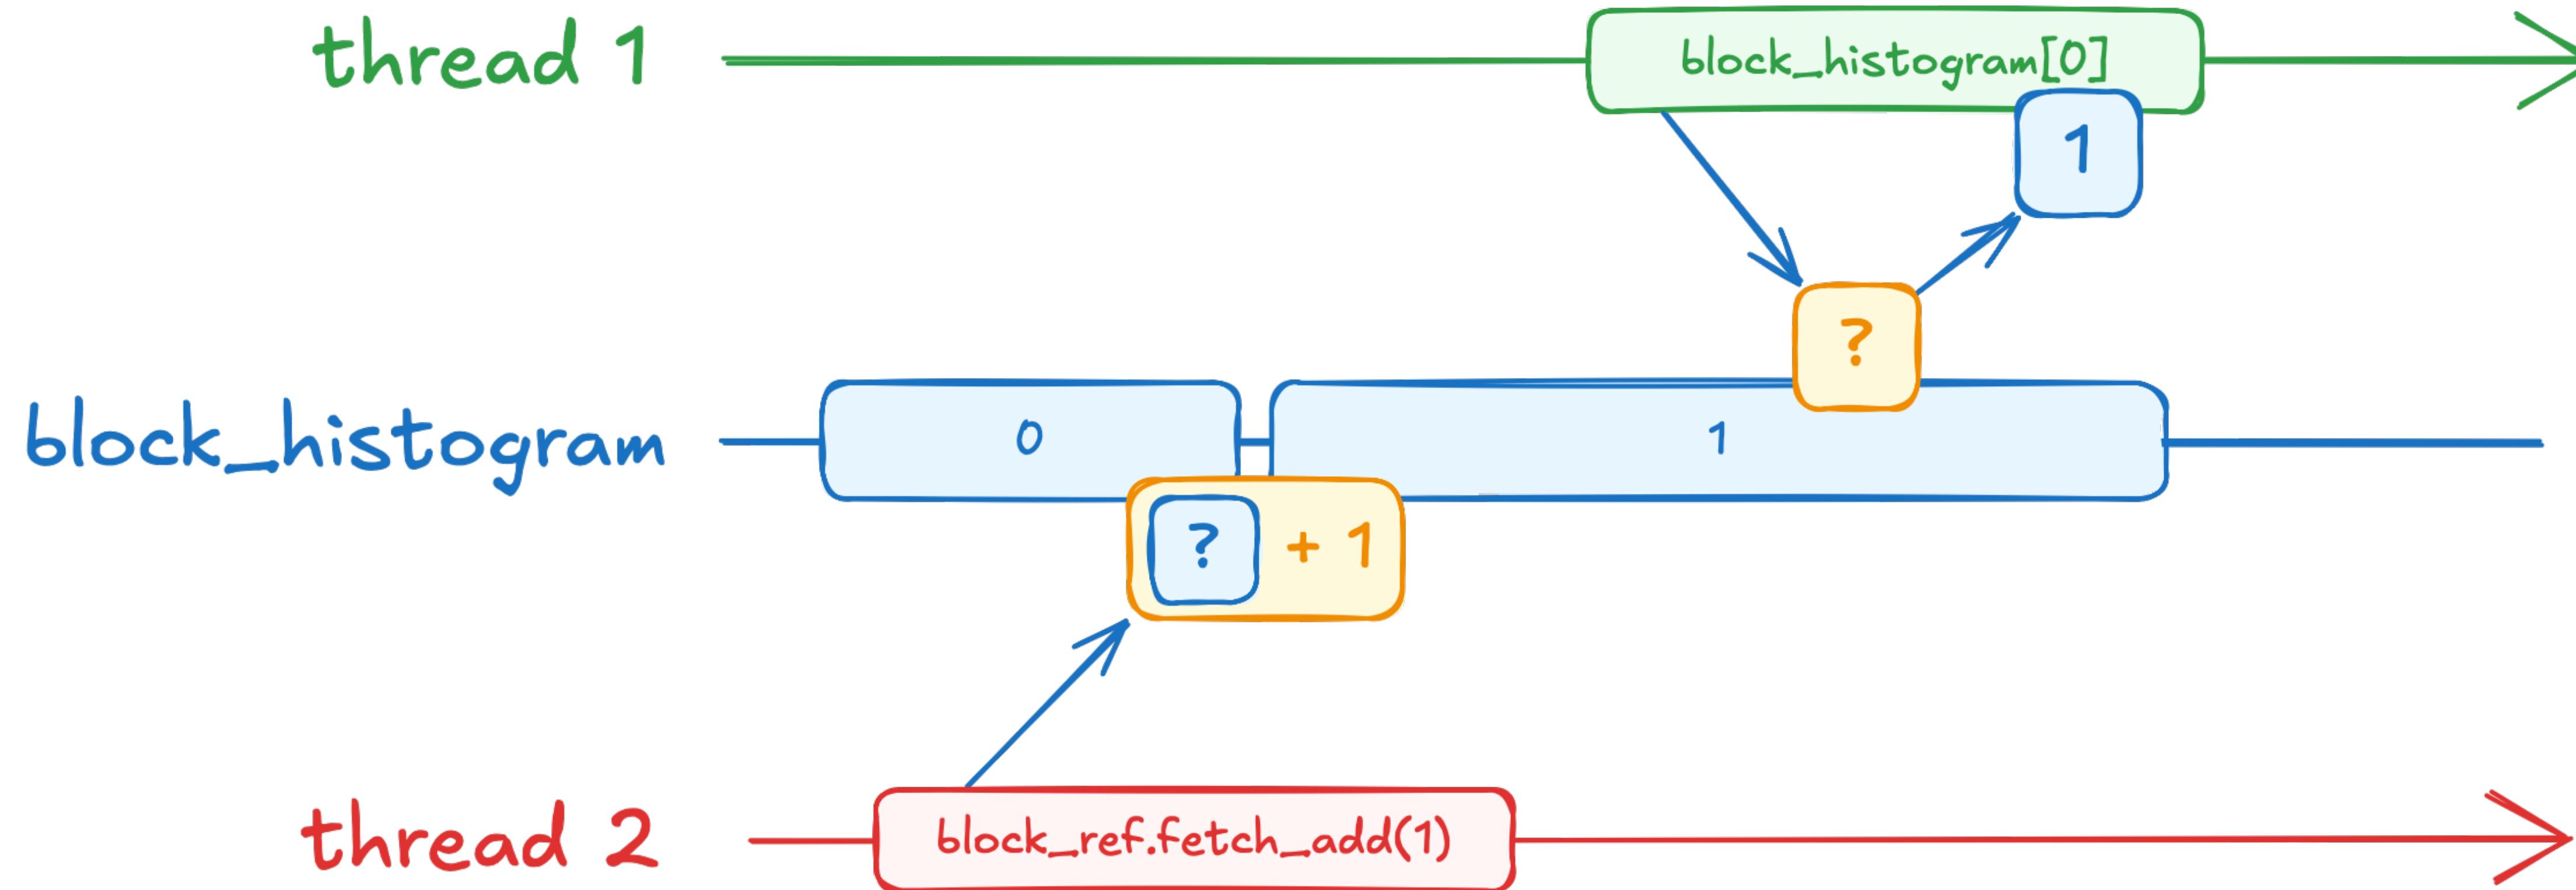
    cuda::std::atomic_ref<int> block_ref(block_histogram[bin]);
    block_ref.fetch_add(1);

    if (threadIdx.x < histogram.size())
    {
        cuda::std::atomic_ref<int> ref(histogram[threadIdx.x]);
        ref.fetch_add(block_histogram[threadIdx.x]);
    }
}
```



- Unfortunately, highlighted code has a bug in it
- We mix non-atomic read with atomic increment, which is an undefined behavior

# What We Expected

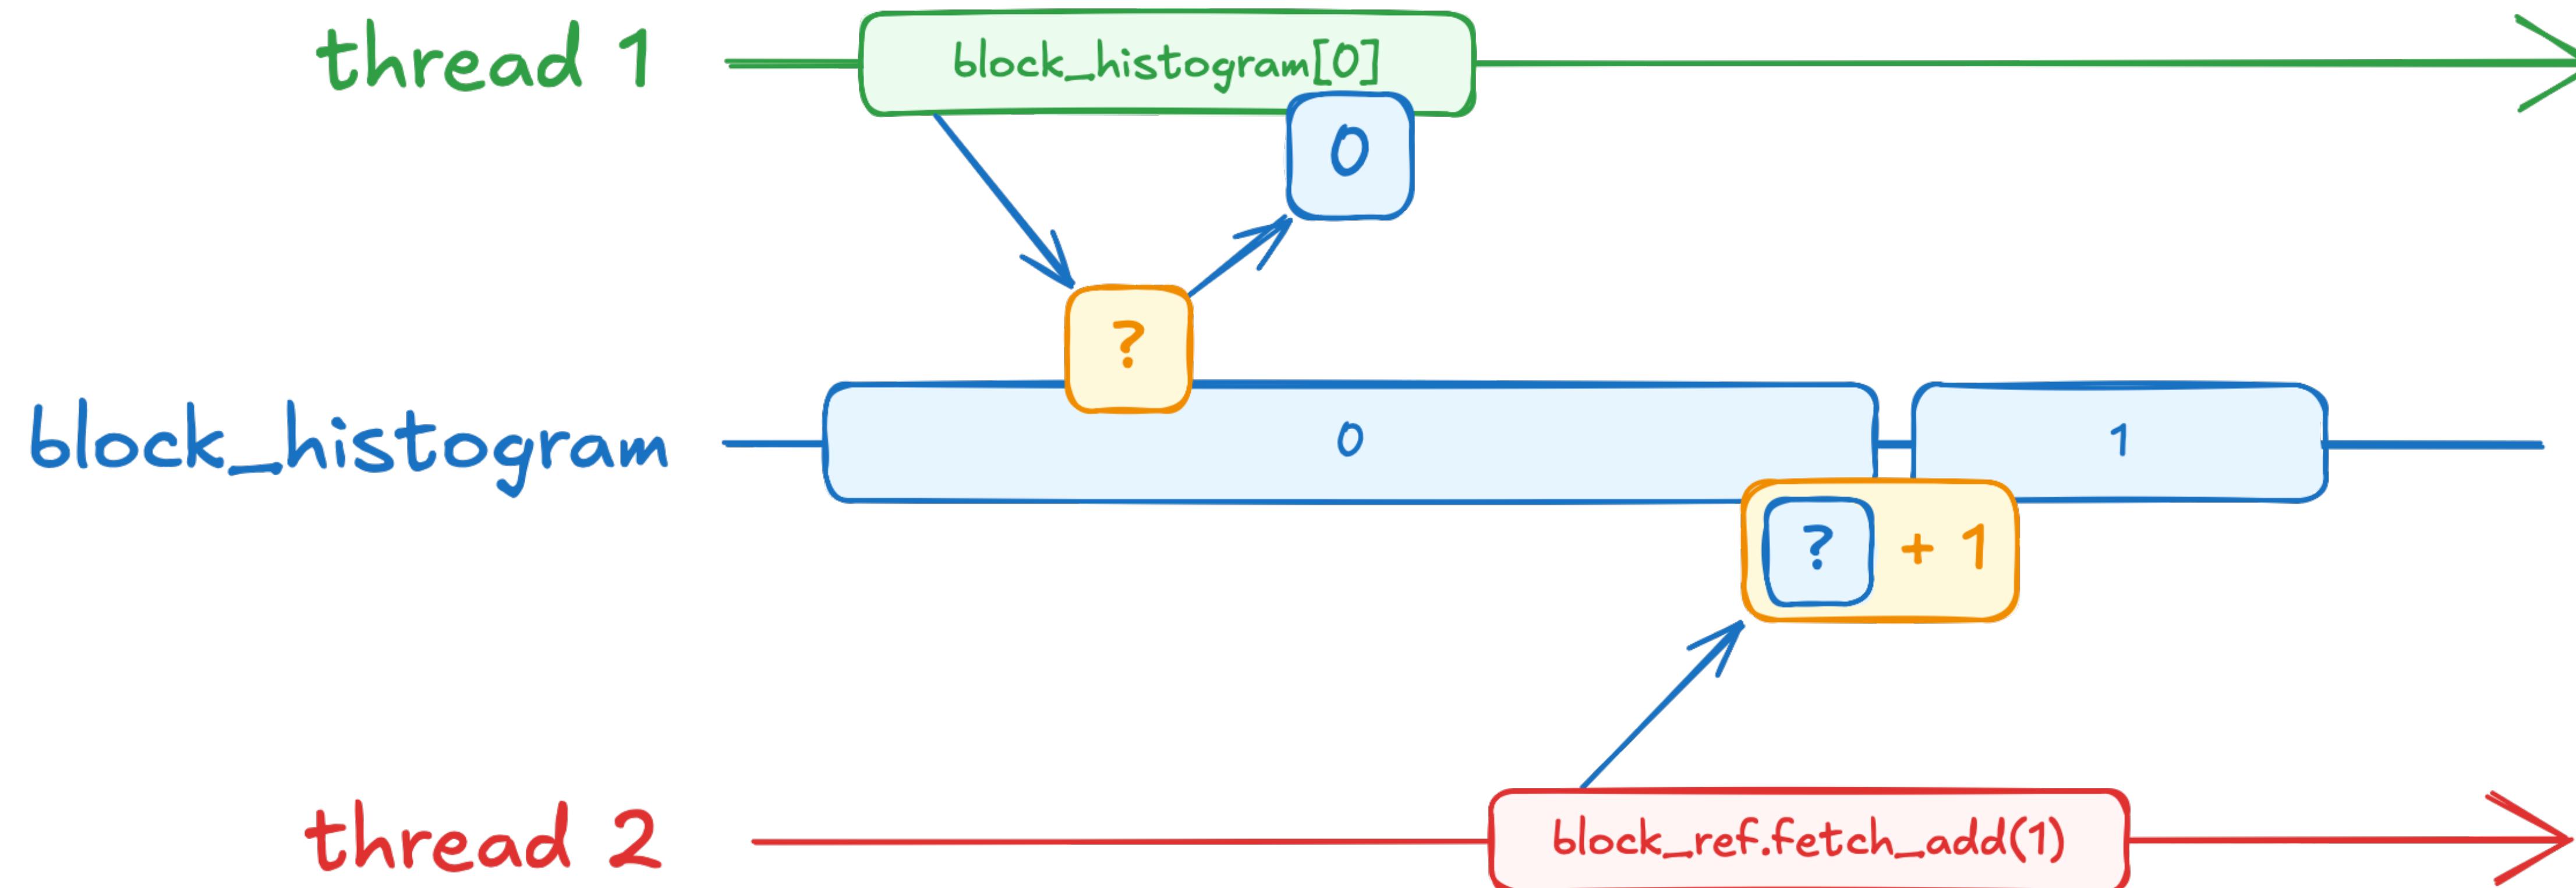


```
cuda::std::atomic_ref<int> block_ref(block_histogram[bin]);  
block_ref.fetch_add(1);
```

```
if (threadIdx.x < histogram.size())  
{  
    cuda::std::atomic_ref<int> ref(histogram[threadIdx.x]);  
    ref.fetch_add(block_histogram[threadIdx.x]);  
}
```

- We expected all block histogram increments to happen first
- And all block histogram reads to happen next

# What Can Happen

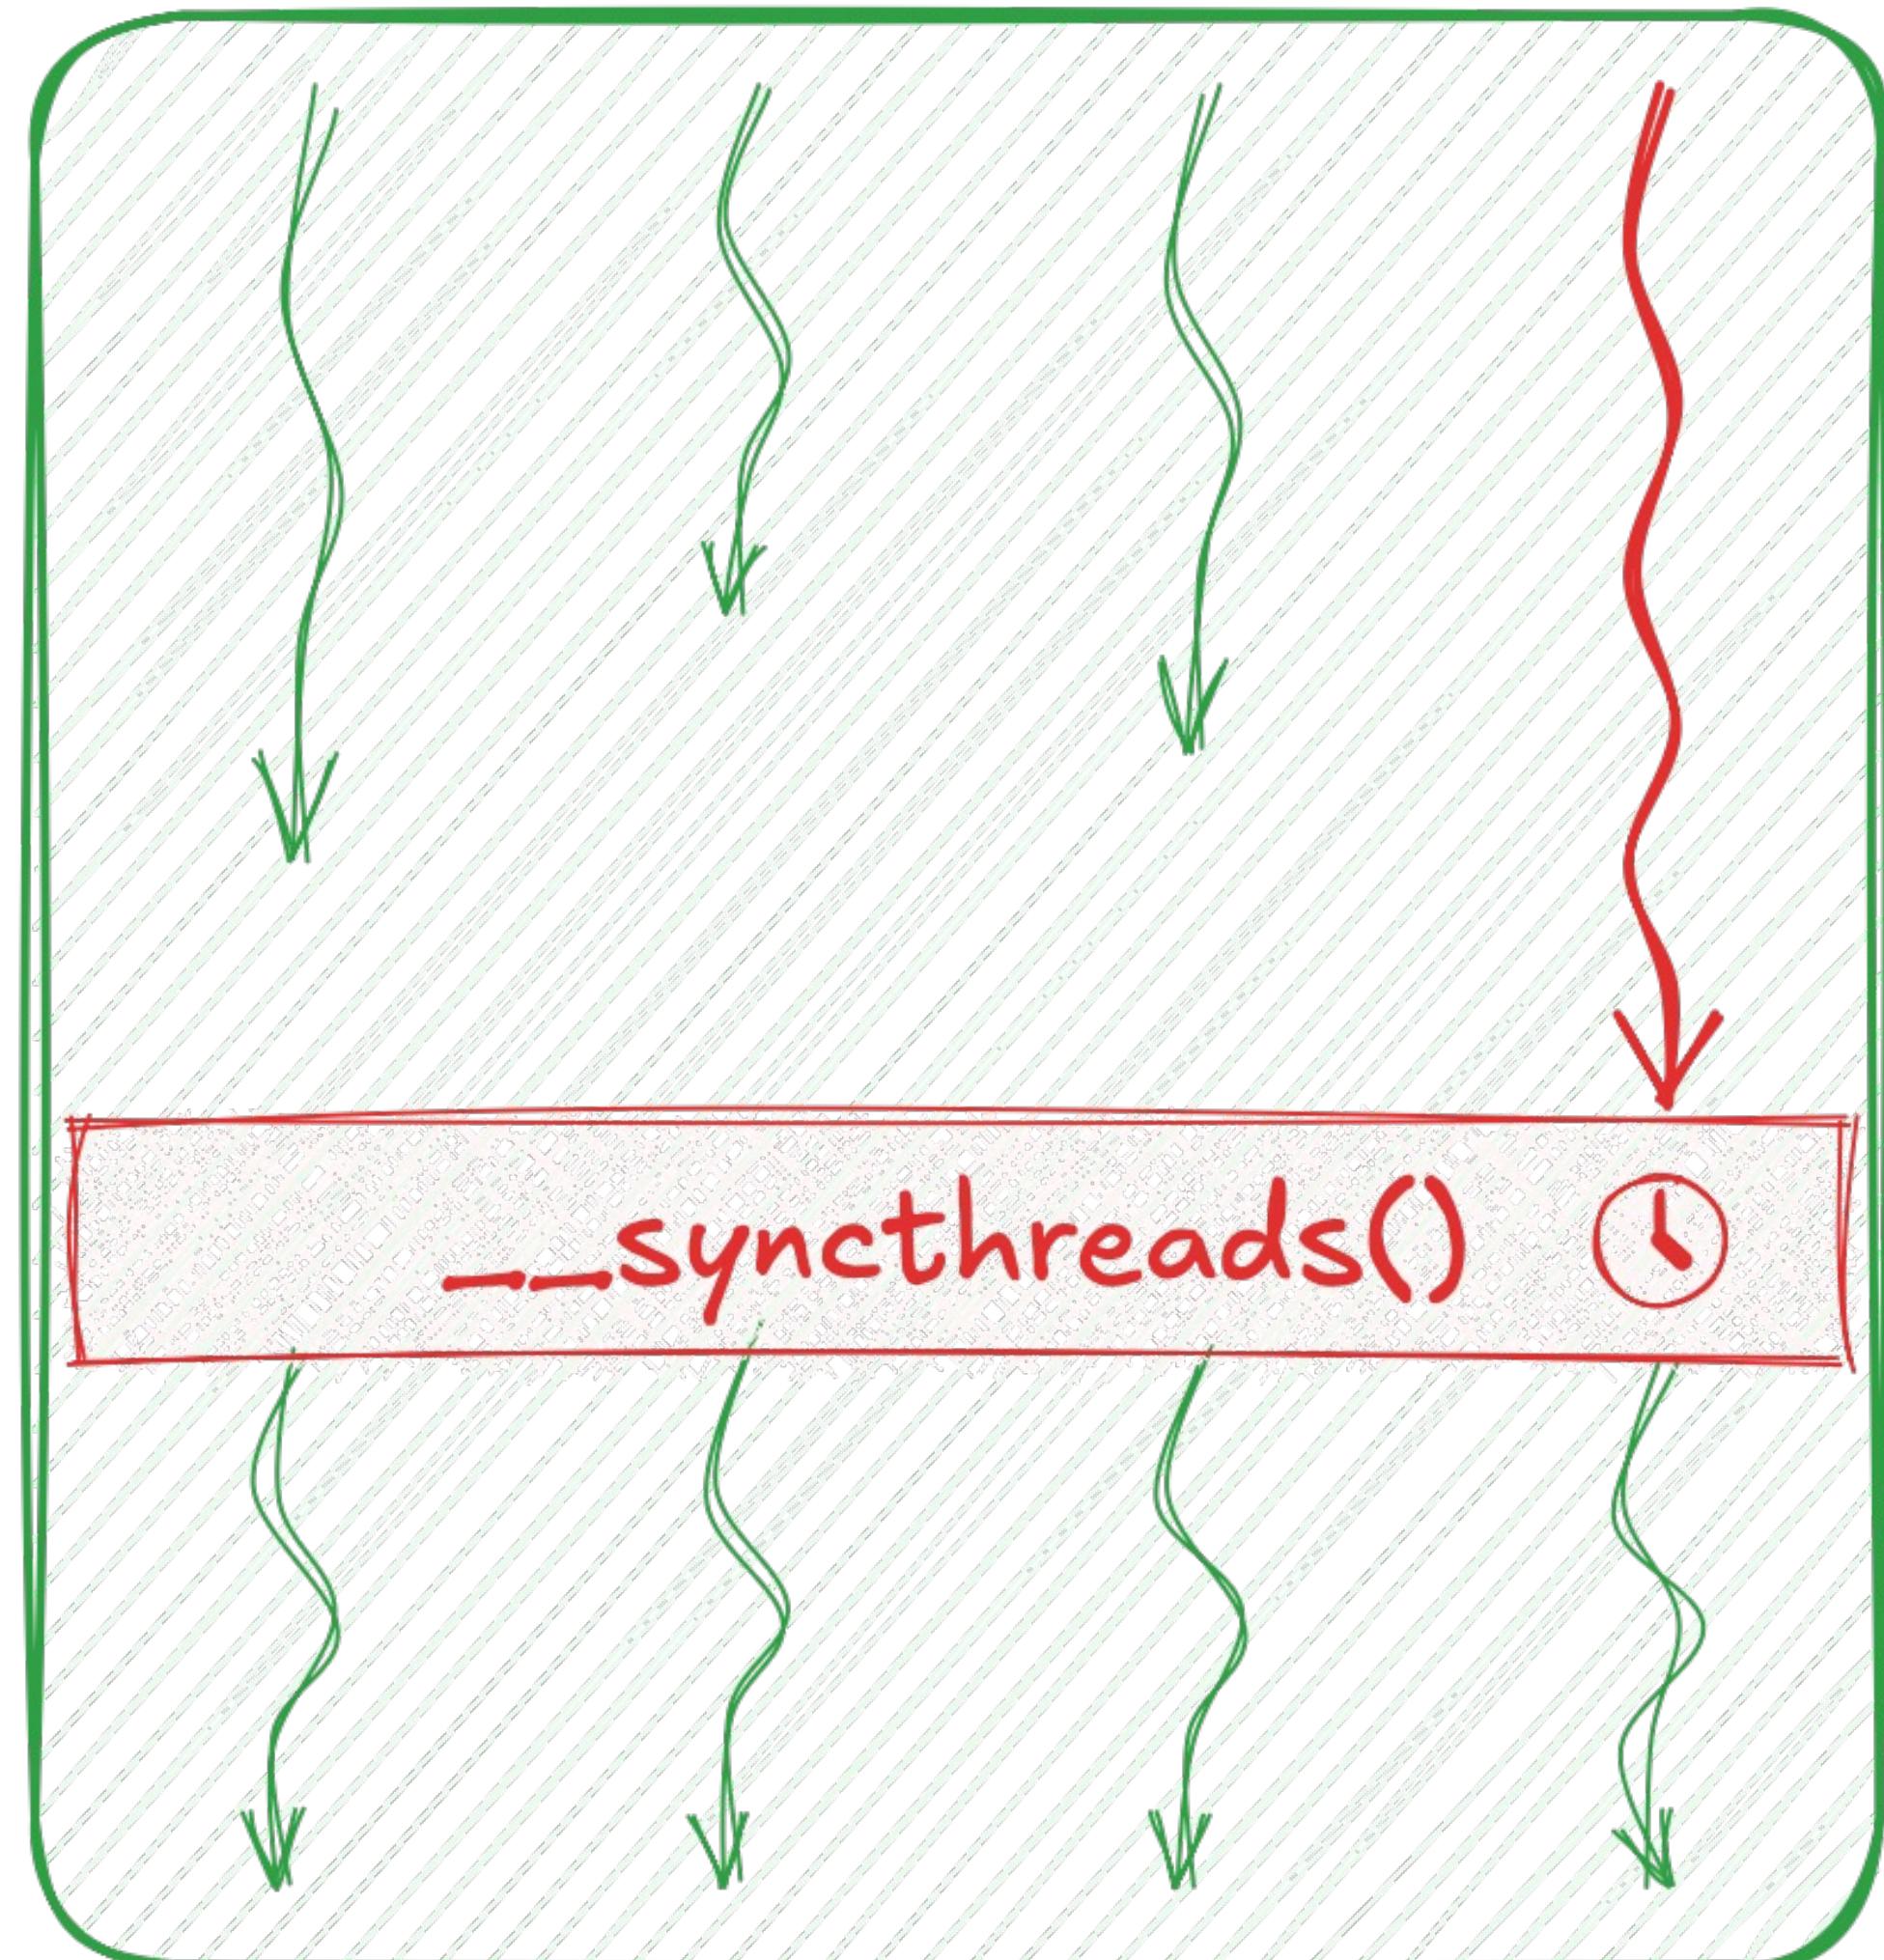


```
cuda::std::atomic_ref<int> block_ref(block_histogram[bin]);  
block_ref.fetch_add(1);  
  
if (threadIdx.x < histogram.size())  
{  
    cuda::std::atomic_ref<int> ref(histogram[threadIdx.x]);  
    ref.fetch_add(block_histogram[threadIdx.x]);  
}
```

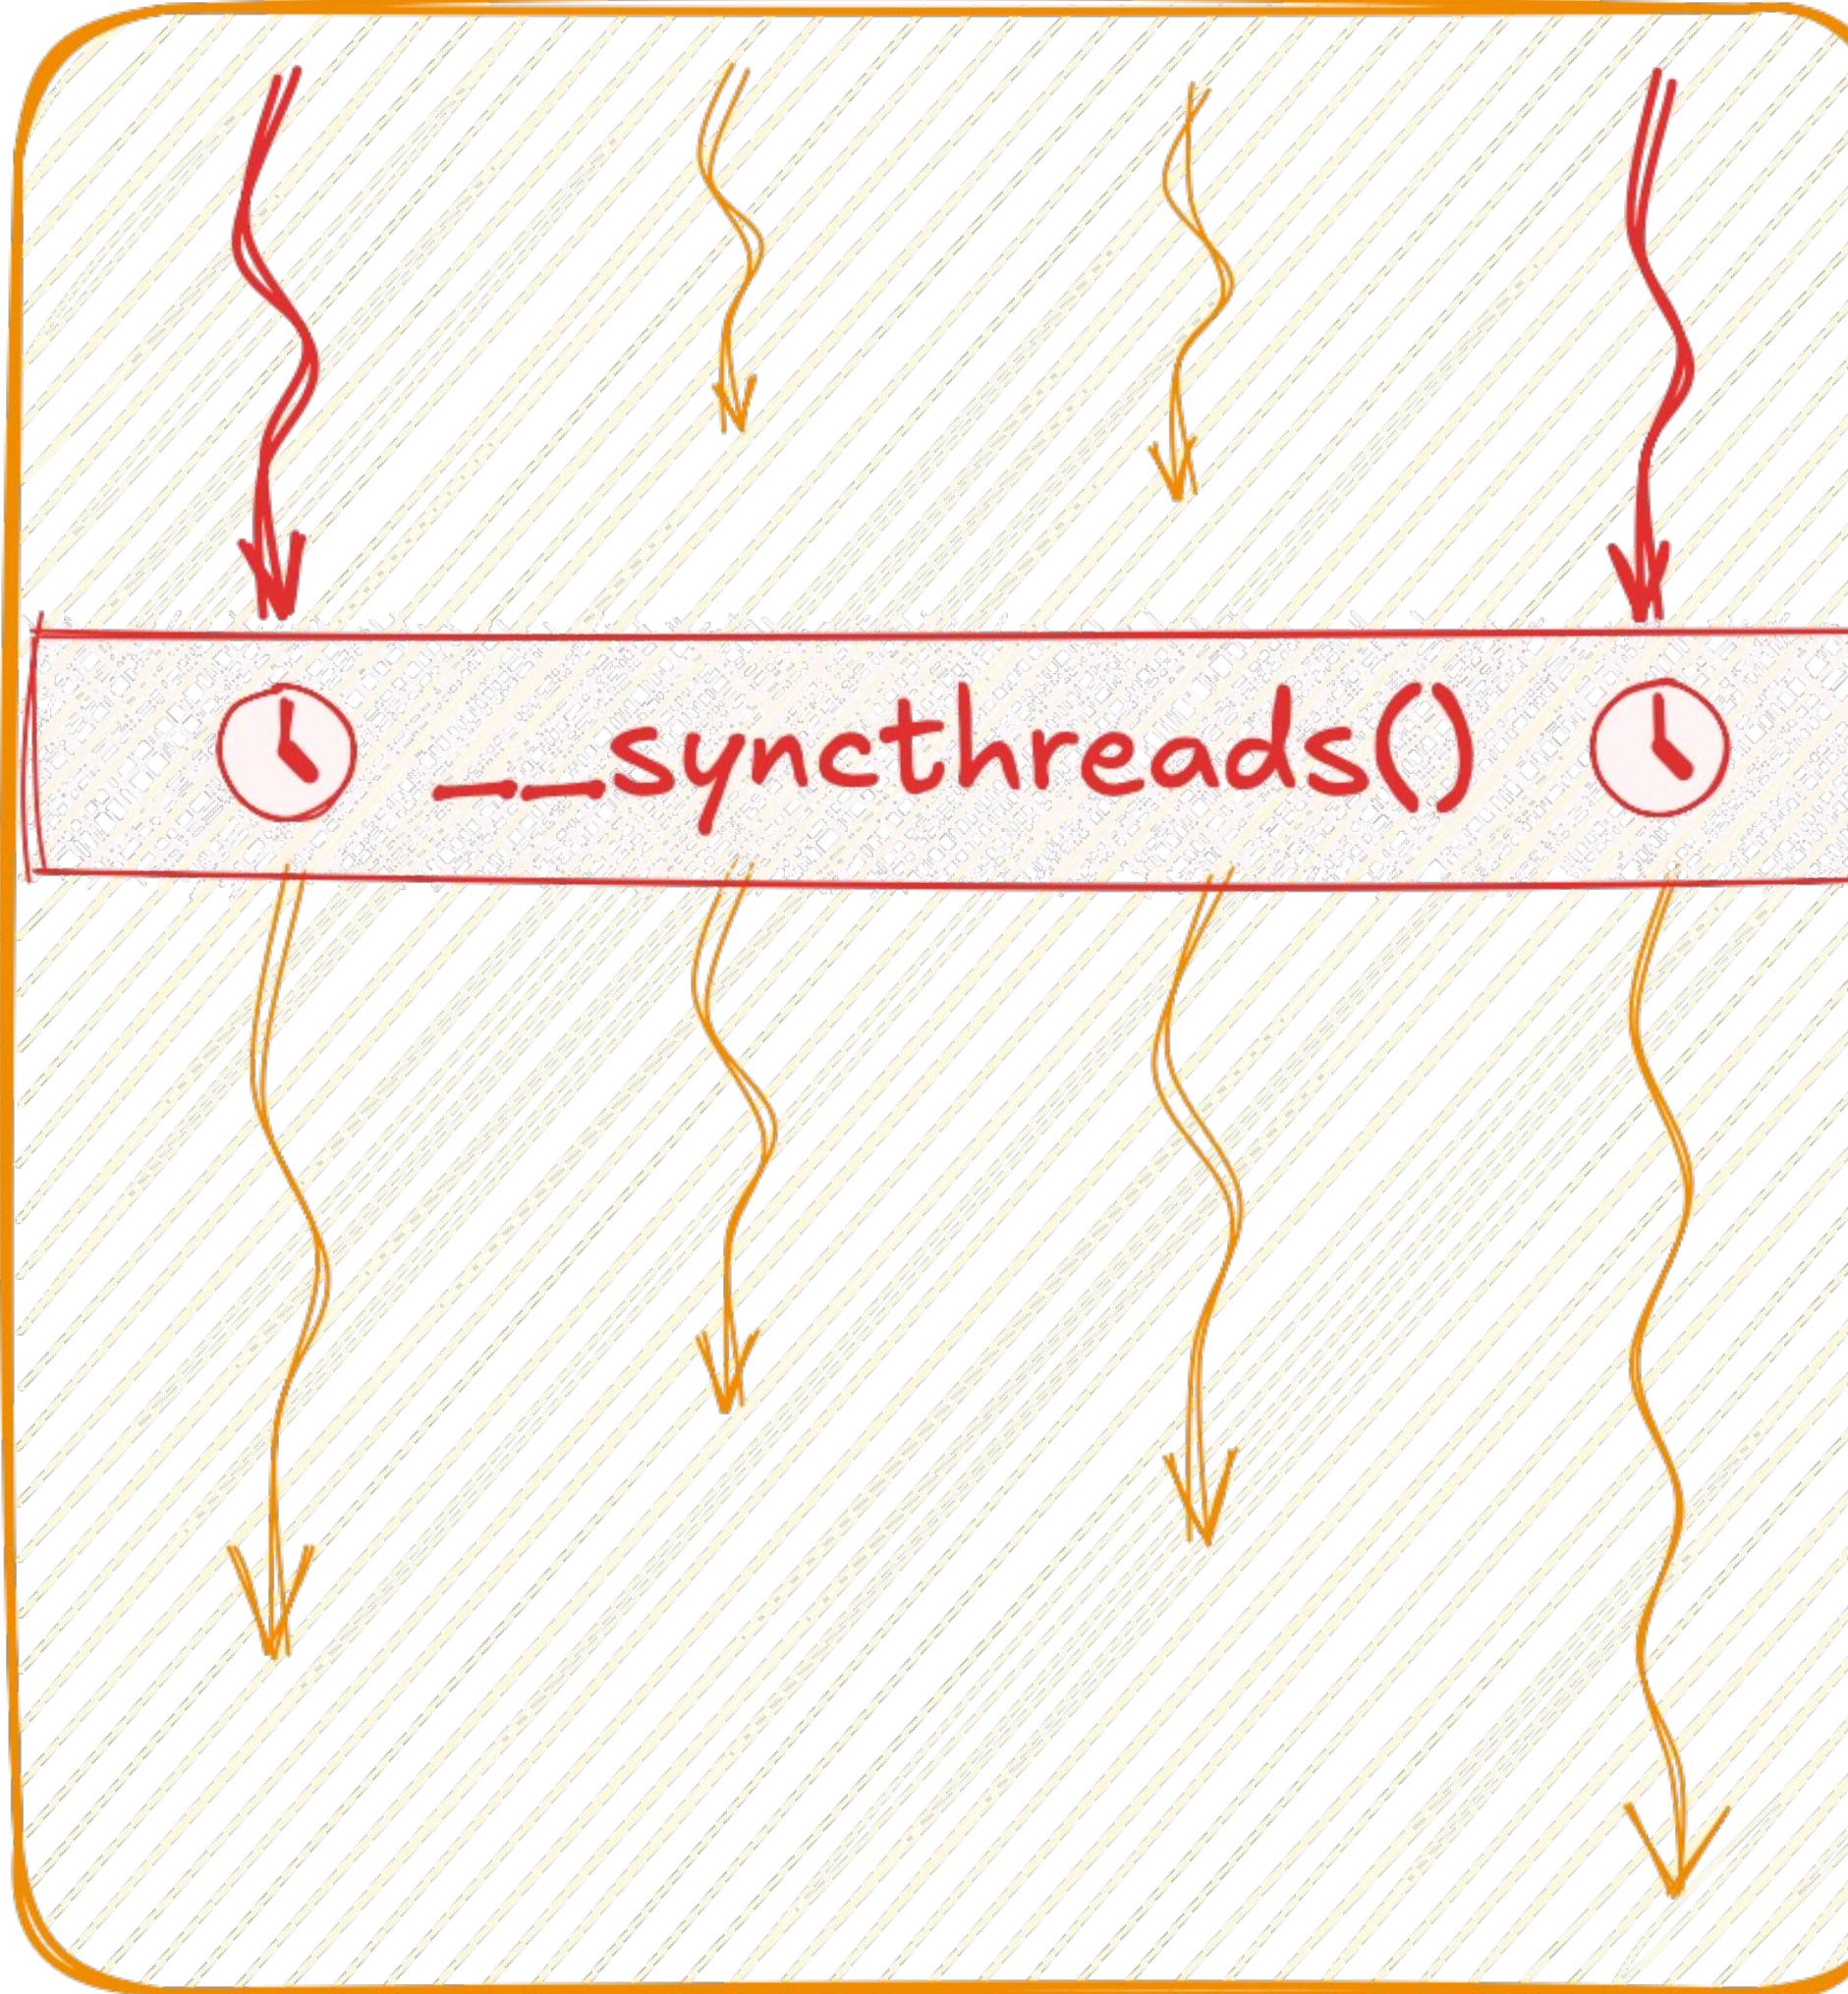
- In practice, we can't assume any order of operations performed by concurrent threads
- We need a mechanism to make sure all threads finished updating block histogram before reading it

# Synchronizing in a Thread Block

blockIdx.x = 0



blockIdx.x = 1



`--syncthreads()`

- invoked cooperatively by all threads in a thread block
- causes threads to wait until all other threads in the thread block arrive
- similar to `std::barrier`, but not allowed inside conditionals

# Vocabulary Types in **libcudacxx**

## Compound Types

- `cuda::std::pair`
- `cuda::std::tuple`
- ...

## Synchronization

- `cuda::std::atomic`
- `cuda::std::atomic_ref`
- `cuda::std::atomic_flag`
- ...

## Optional and Alternatives

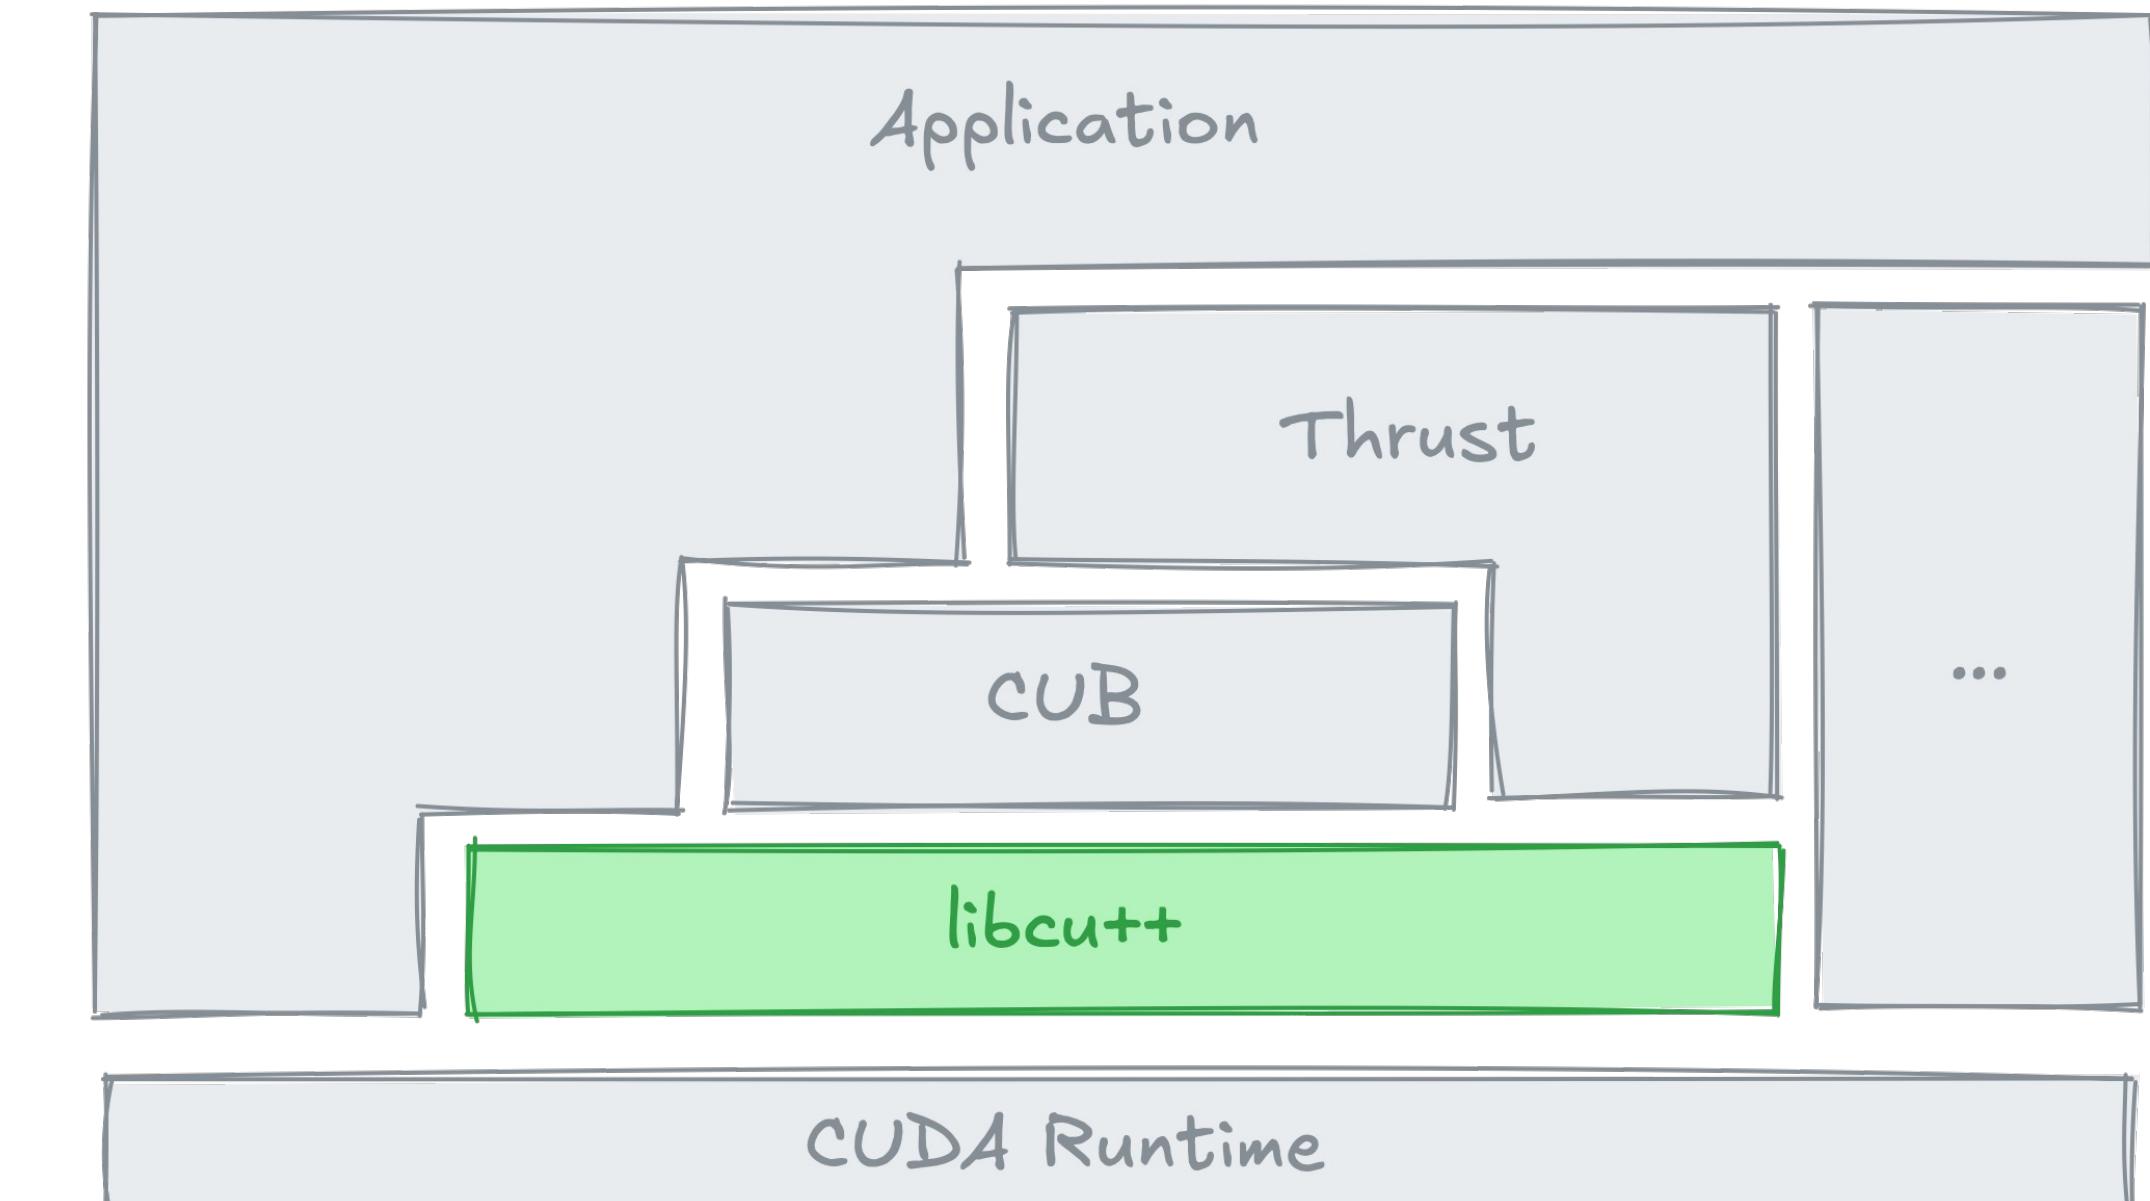
- `cuda::std::optional`
- `cuda::std::variant`
- ...

## Math

- `cuda::std::complex`
- ...

## CUDA Extensions

- `cuda::atomic`
- `cuda::atomic_ref`
- `cuda::atomic_flag`
- ...



<https://nvidia.github.io/cccl/libcudacxx>

# Thread Scopes

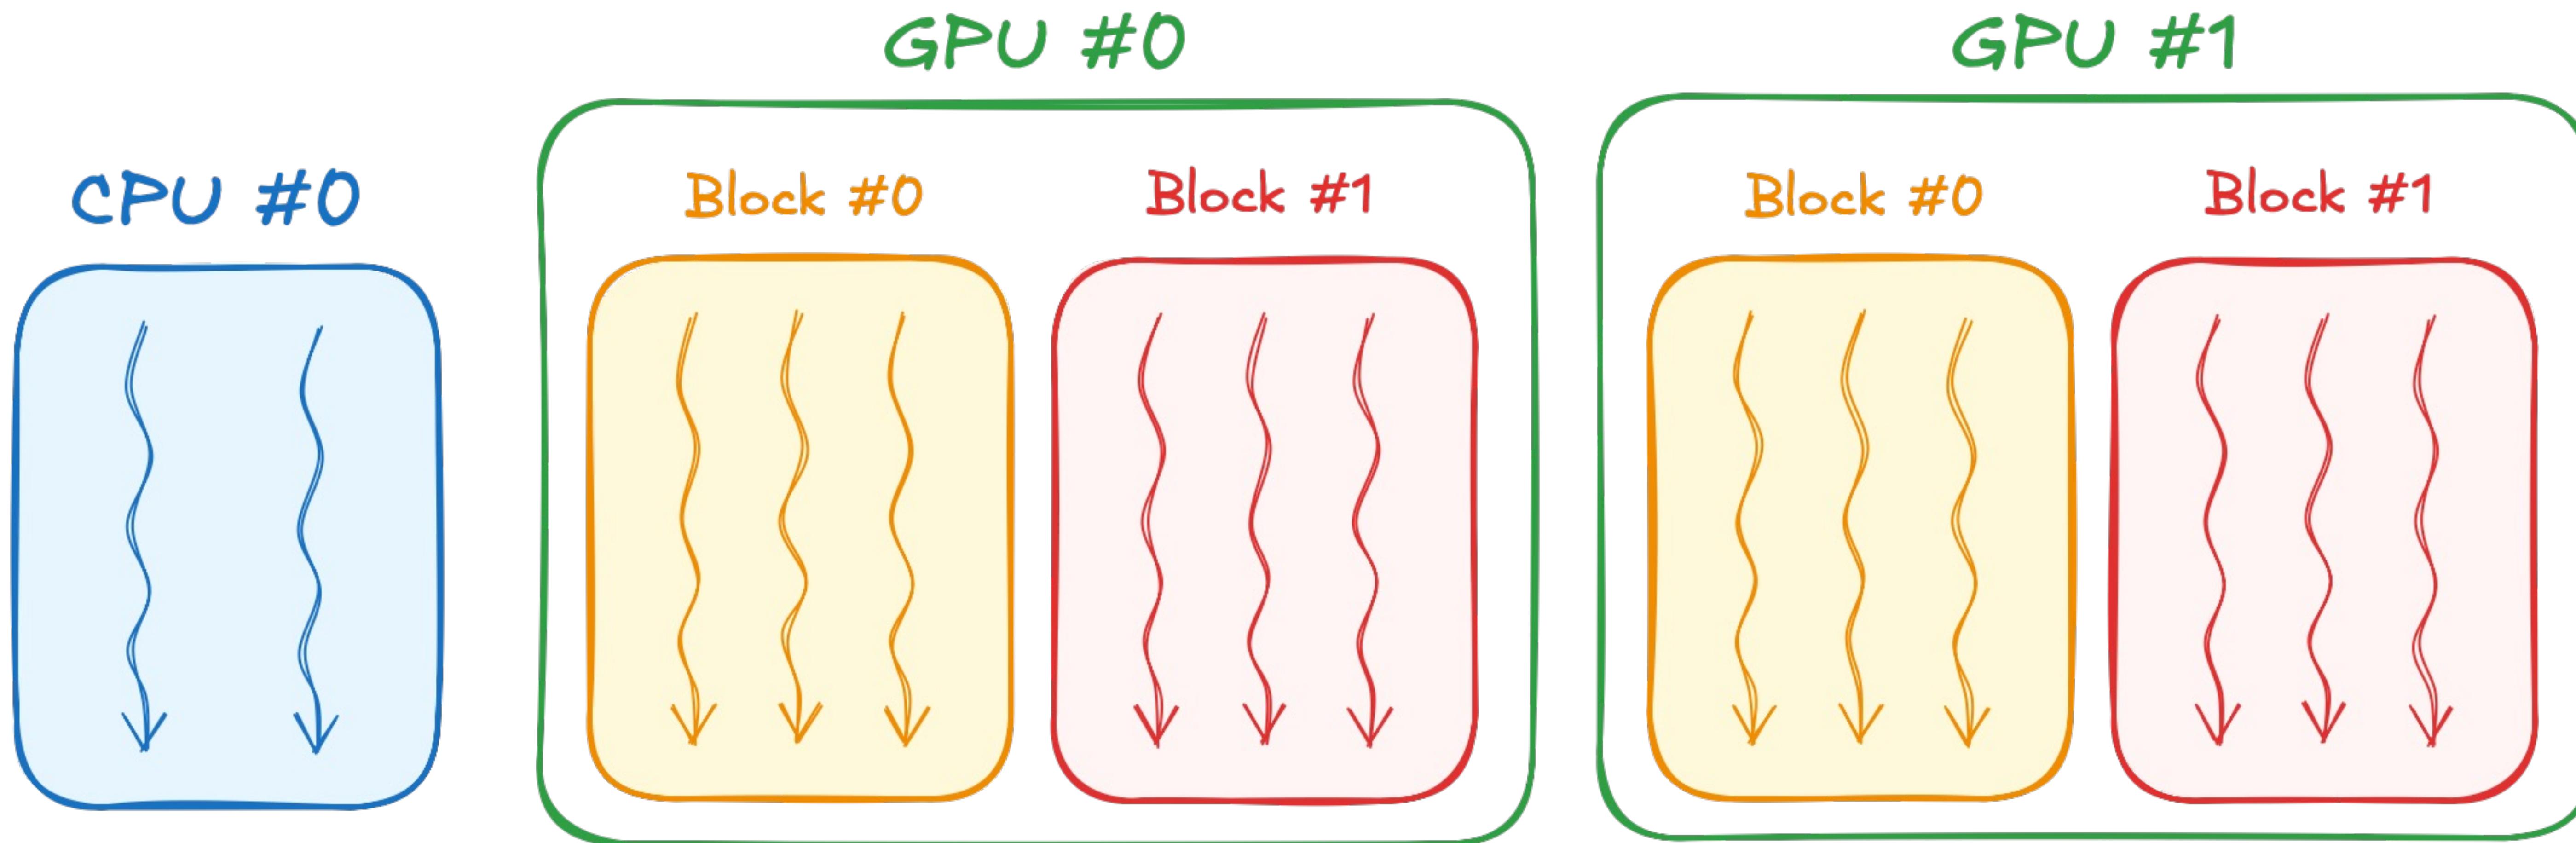
```
cuda::std::atomic_ref<int> ref(...);
```

- Both `cuda::std::atomic_ref` and `cuda::atomic_ref` have equivalent interface
- The only difference is that `cuda::atomic_ref` extends `cuda::std::atomic_ref` with additional `cuda::thread_scope` parameter
- Using appropriate thread scope can significantly affect performance

```
cuda::atomic_ref<int, thread_scope> ref(...);
```



# Thread Scope System

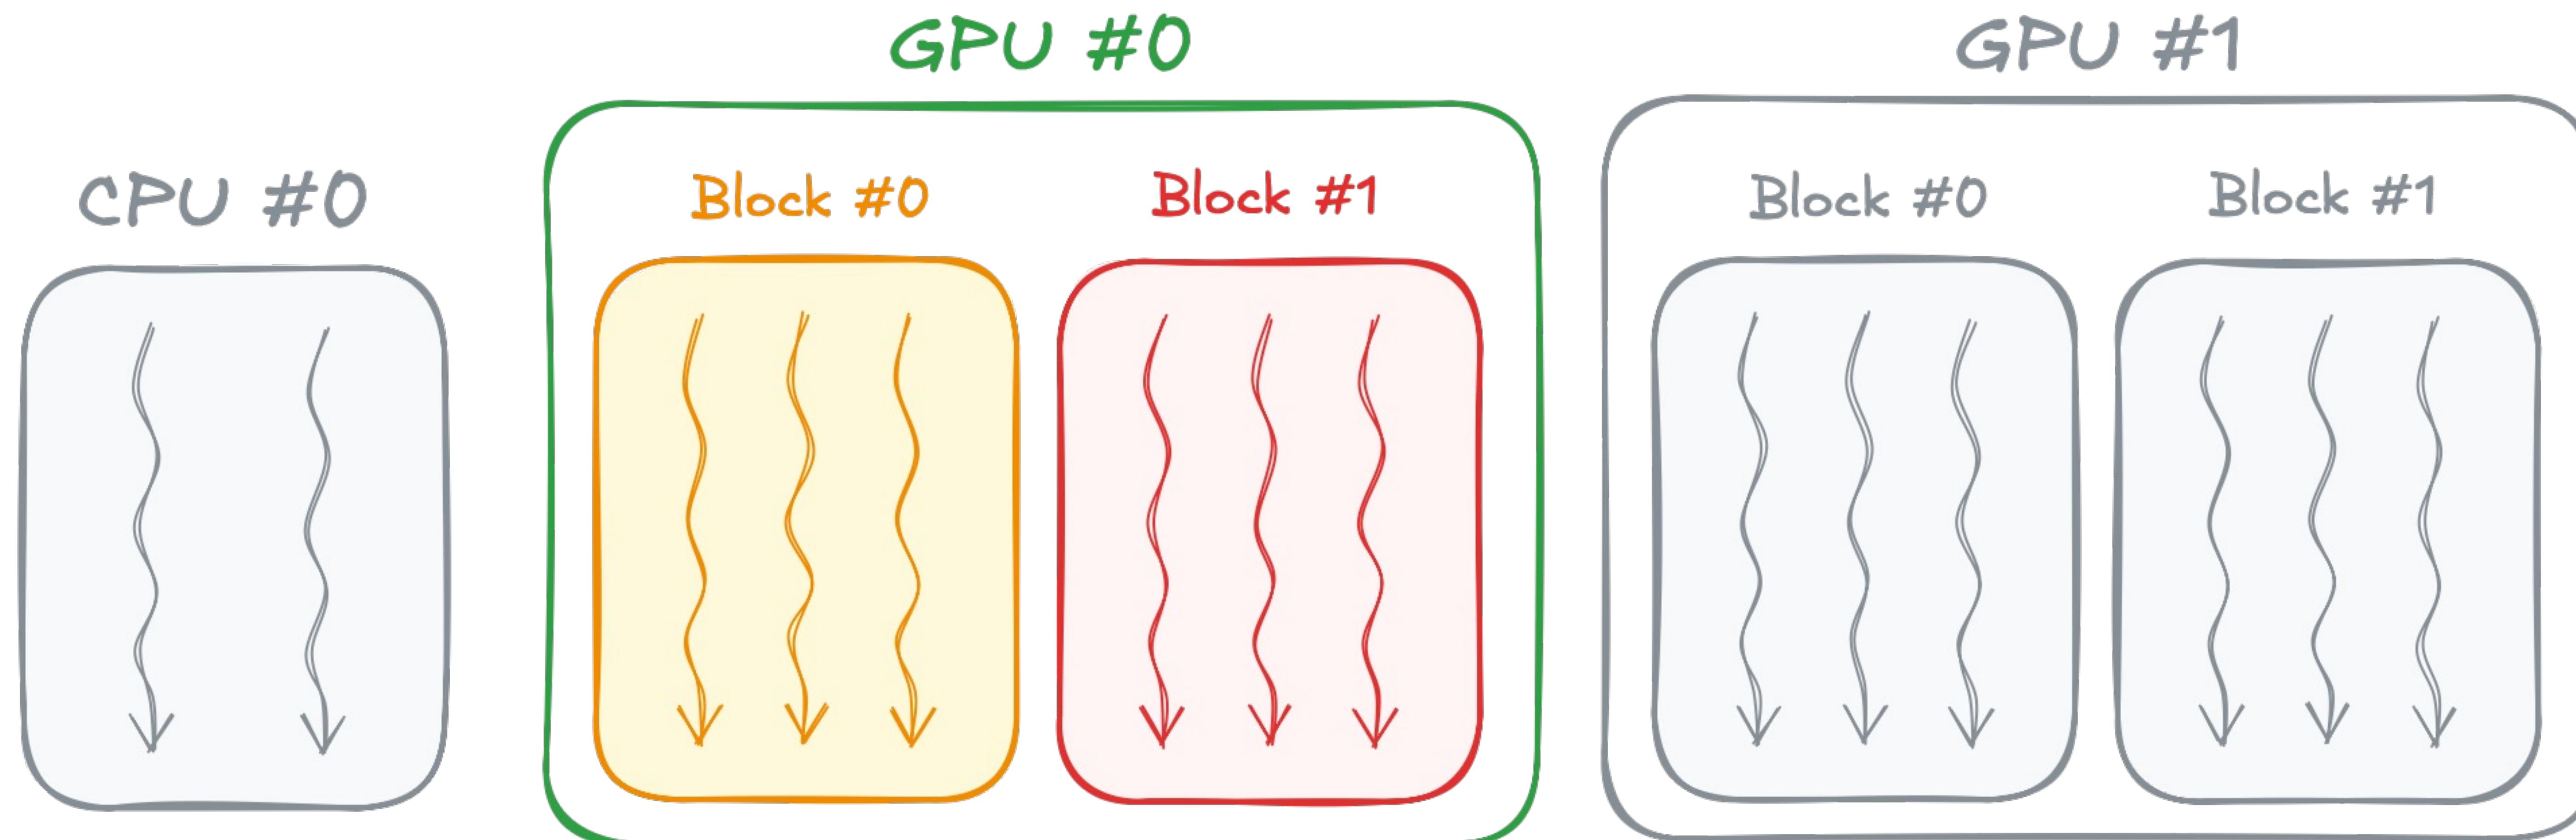


```
cuda::atomic_ref<int, thread_scope> ref(...);
```

- Equivalent to `cuda::std::atomic_ref`
- Each threads of a given system is related to each other thread by *system thread scope*

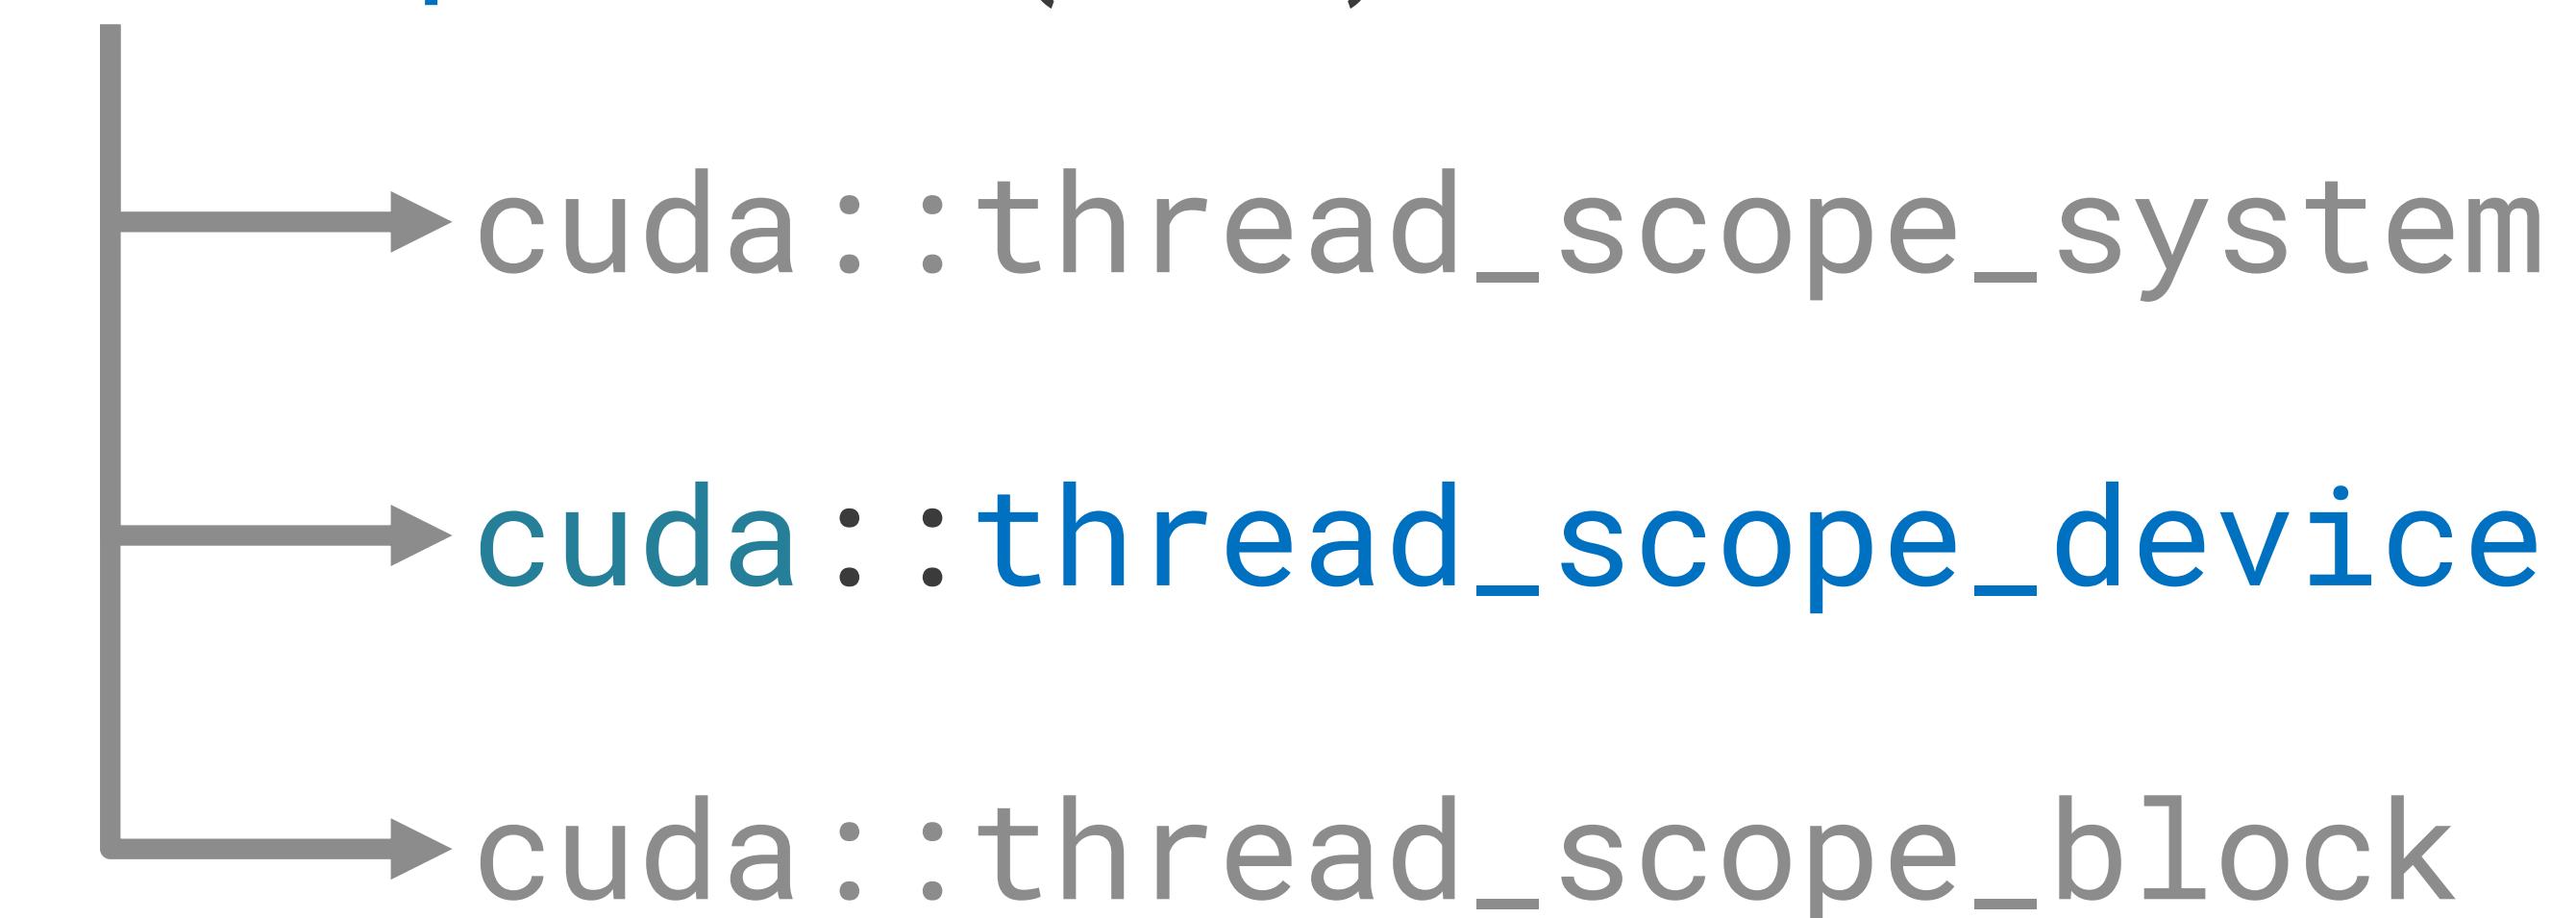


# Thread Scope Device

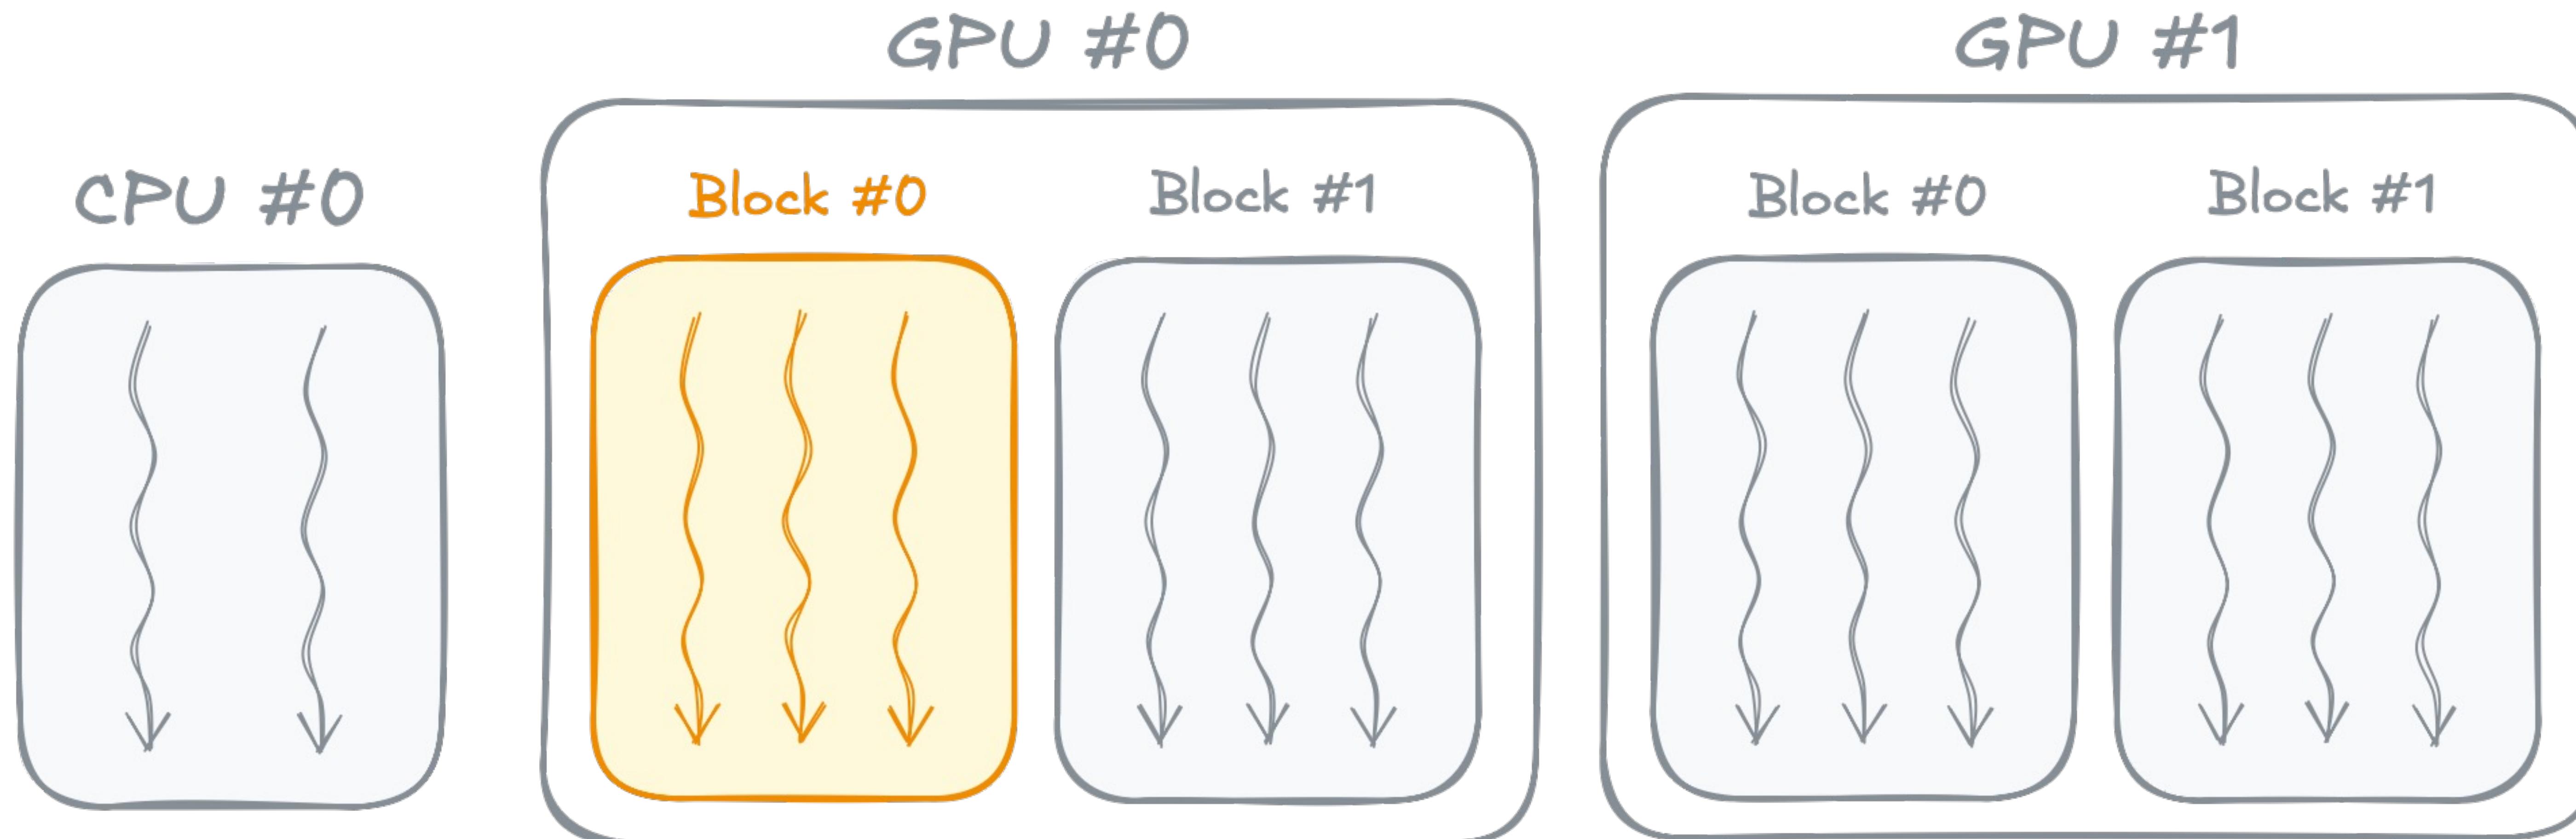


```
cuda::atomic_ref<int, thread_scope> ref(...);
```

- Each GPU thread is related to each other GPU thread on the same GPU by *device thread scope*

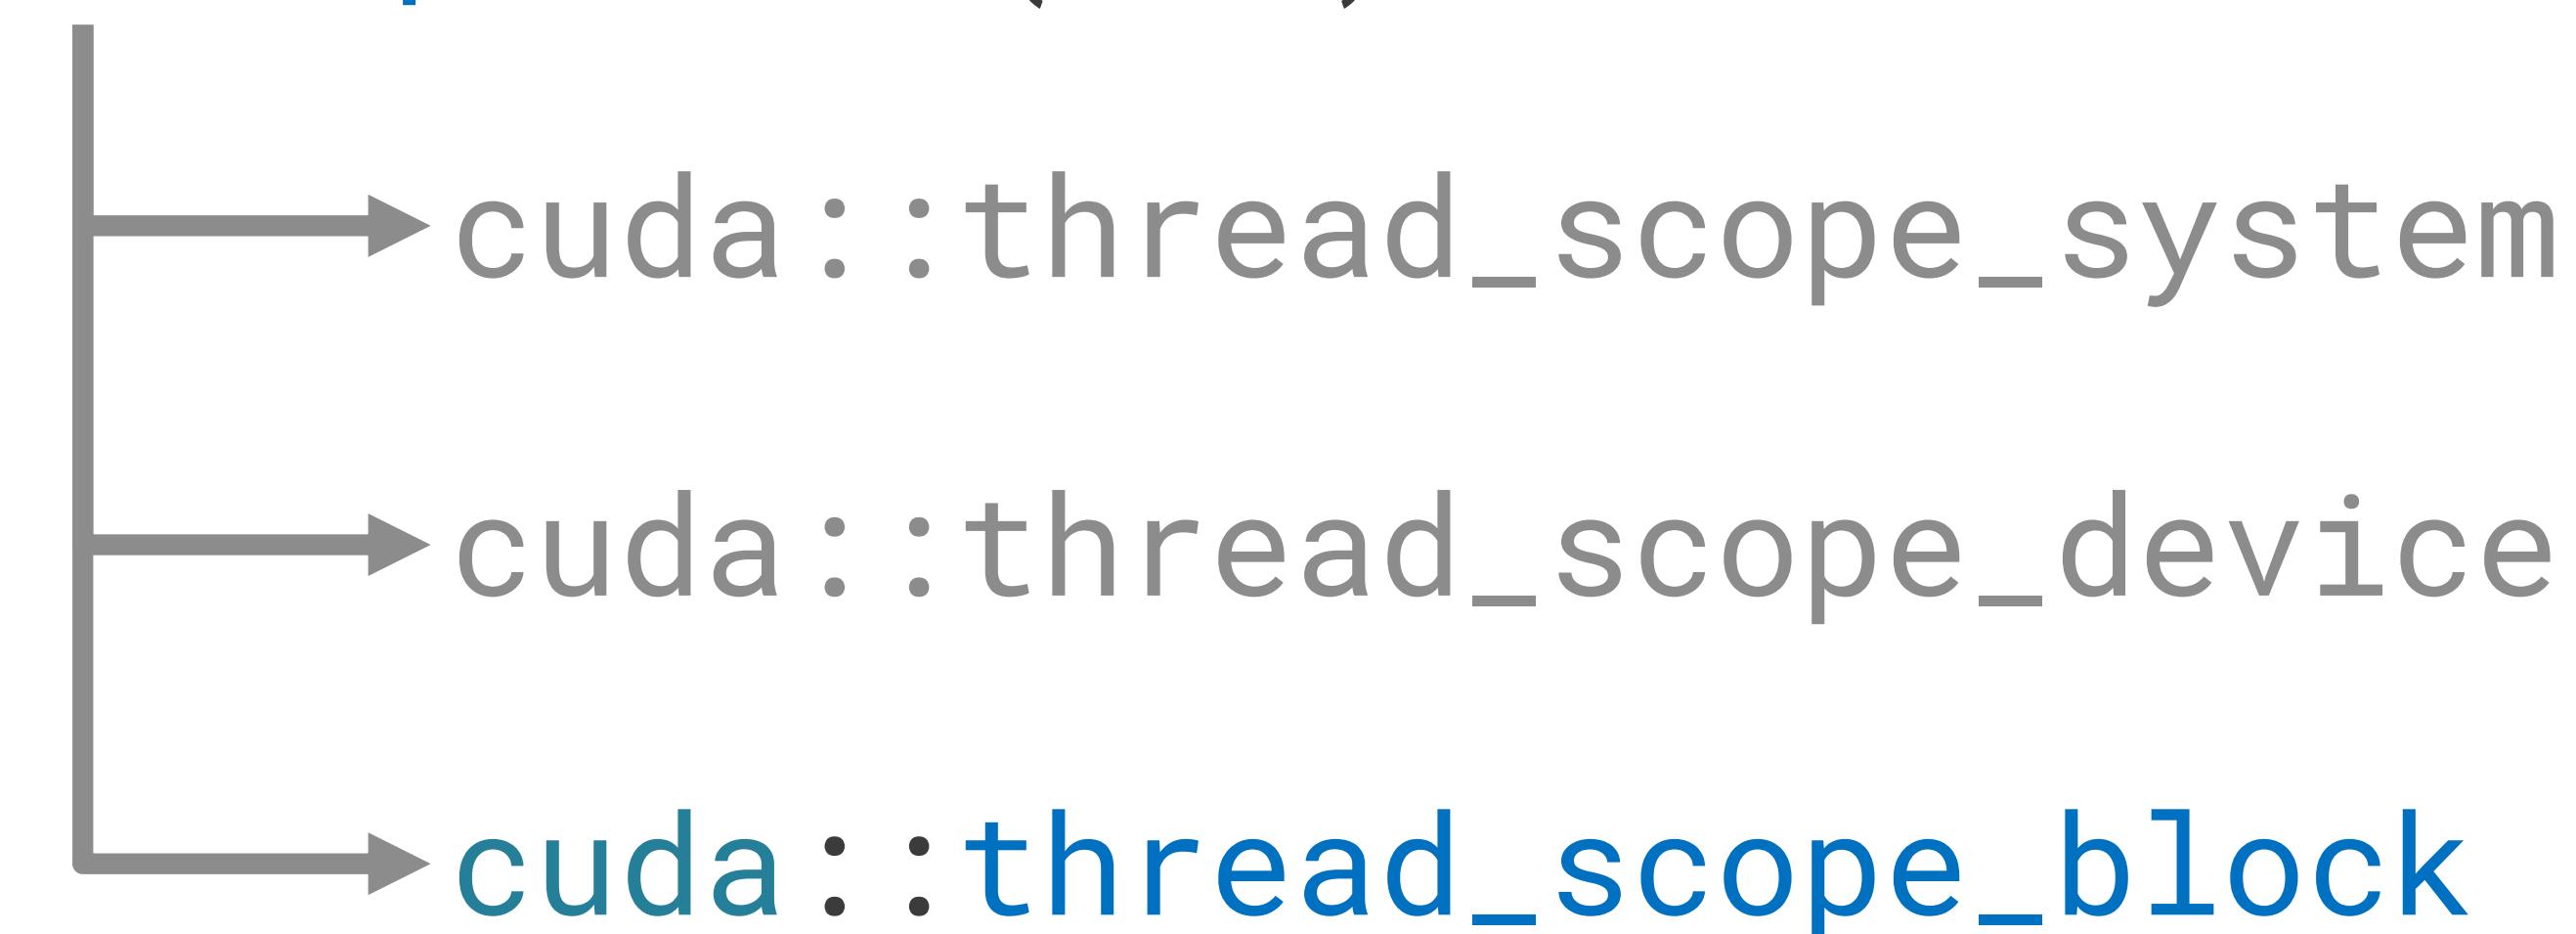


# Thread Scope Block



```
cuda::atomic_ref<int, thread_scope> ref(...);
```

- Each GPU thread is related to each other GPU thread in the same thread block by the *block thread scope*



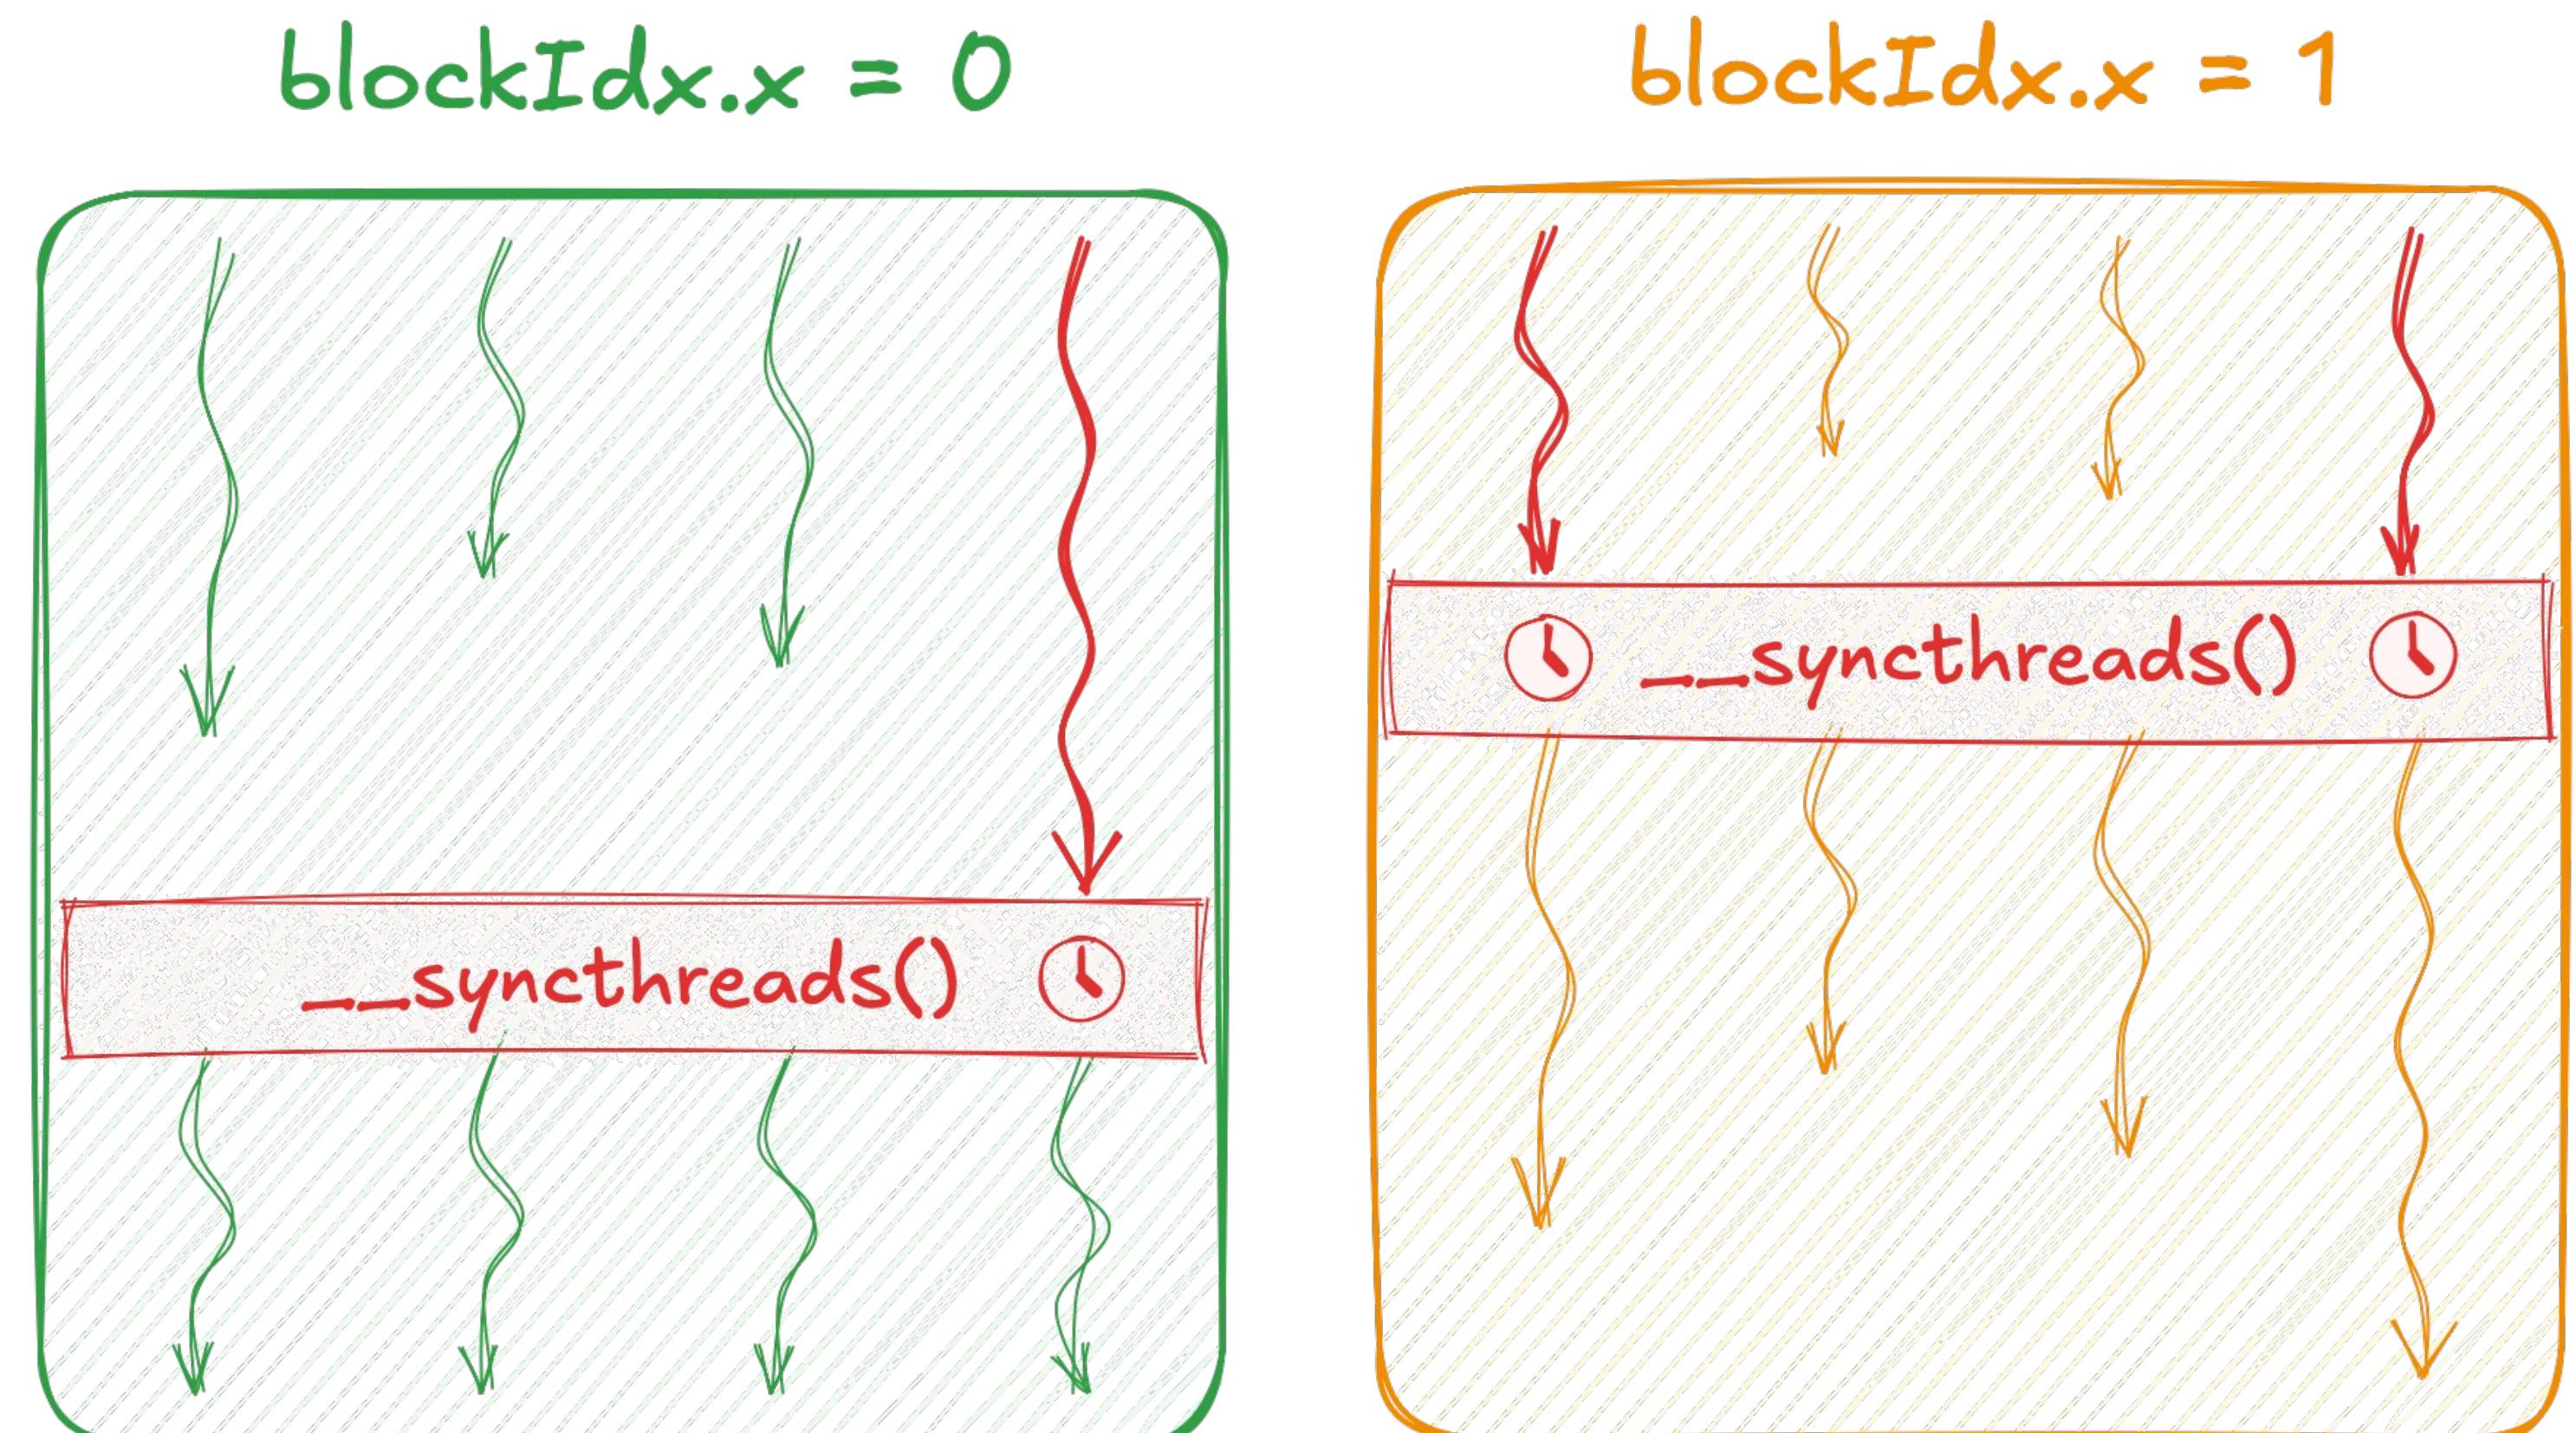
# Exercise: Fix Histogram

- Use `__syncthreads()` to make sure all threads updated block histogram before reading it
- Reduce thread scope of atomic operations for better performance (optional)

```
int cell = blockIdx.x * blockDim.x + threadIdx.x;
int bin = static_cast<int>(temperatures[cell] / bin_width);

cuda::std::atomic_ref<int> block_ref(block_histogram[bin]);
block_ref.fetch_add(1);

if (threadIdx.x < histogram.size())
{
    cuda::std::atomic_ref<int> ref(histogram[threadIdx.x]);
    ref.fetch_add(block_histogram[threadIdx.x]);
}
```



# Exercise: Fix Histogram

## Solution

```
int cell = blockIdx.x * blockDim.x + threadIdx.x;
int bin = static_cast<int>(temperatures[cell] / bin_width);

cuda::atomic_ref<int, cuda::thread_scope_block> block_ref(block_histogram[bin]);
block_ref.fetch_add(1);
__syncthreads();

if (threadIdx.x < histogram.size())
{
    cuda::atomic_ref<int, cuda::thread_scope_device> ref(histogram[threadIdx.x]);
    ref.fetch_add(block_histogram[threadIdx.x]);
}
```

- This change brings us from **6 GB/s** to **100 GB/s**
- This is much better, but we still have room for improvement

# Recap

```
cuda::std::span<int> block_histogram =
    block_histograms.subspan(
        blockIdx.x * histogram.size(),
        histogram.size());

int cell = blockIdx.x * blockDim.x + threadIdx.x;
int bin = static_cast<int>(temperatures[cell] / bin_width);

cuda::atomic_ref<int, cuda::thread_scope_block>
    block_ref(block_histogram[bin]);
block_ref.fetch_add(1);
__syncthreads();

if (threadIdx.x < num_bins)
{
    cuda::atomic_ref<int, cuda::thread_scope_device>
        ref(histogram[threadIdx.x]);
    ref.fetch_add(block_histogram[threadIdx.x]);
}
```

We've learned so many things!

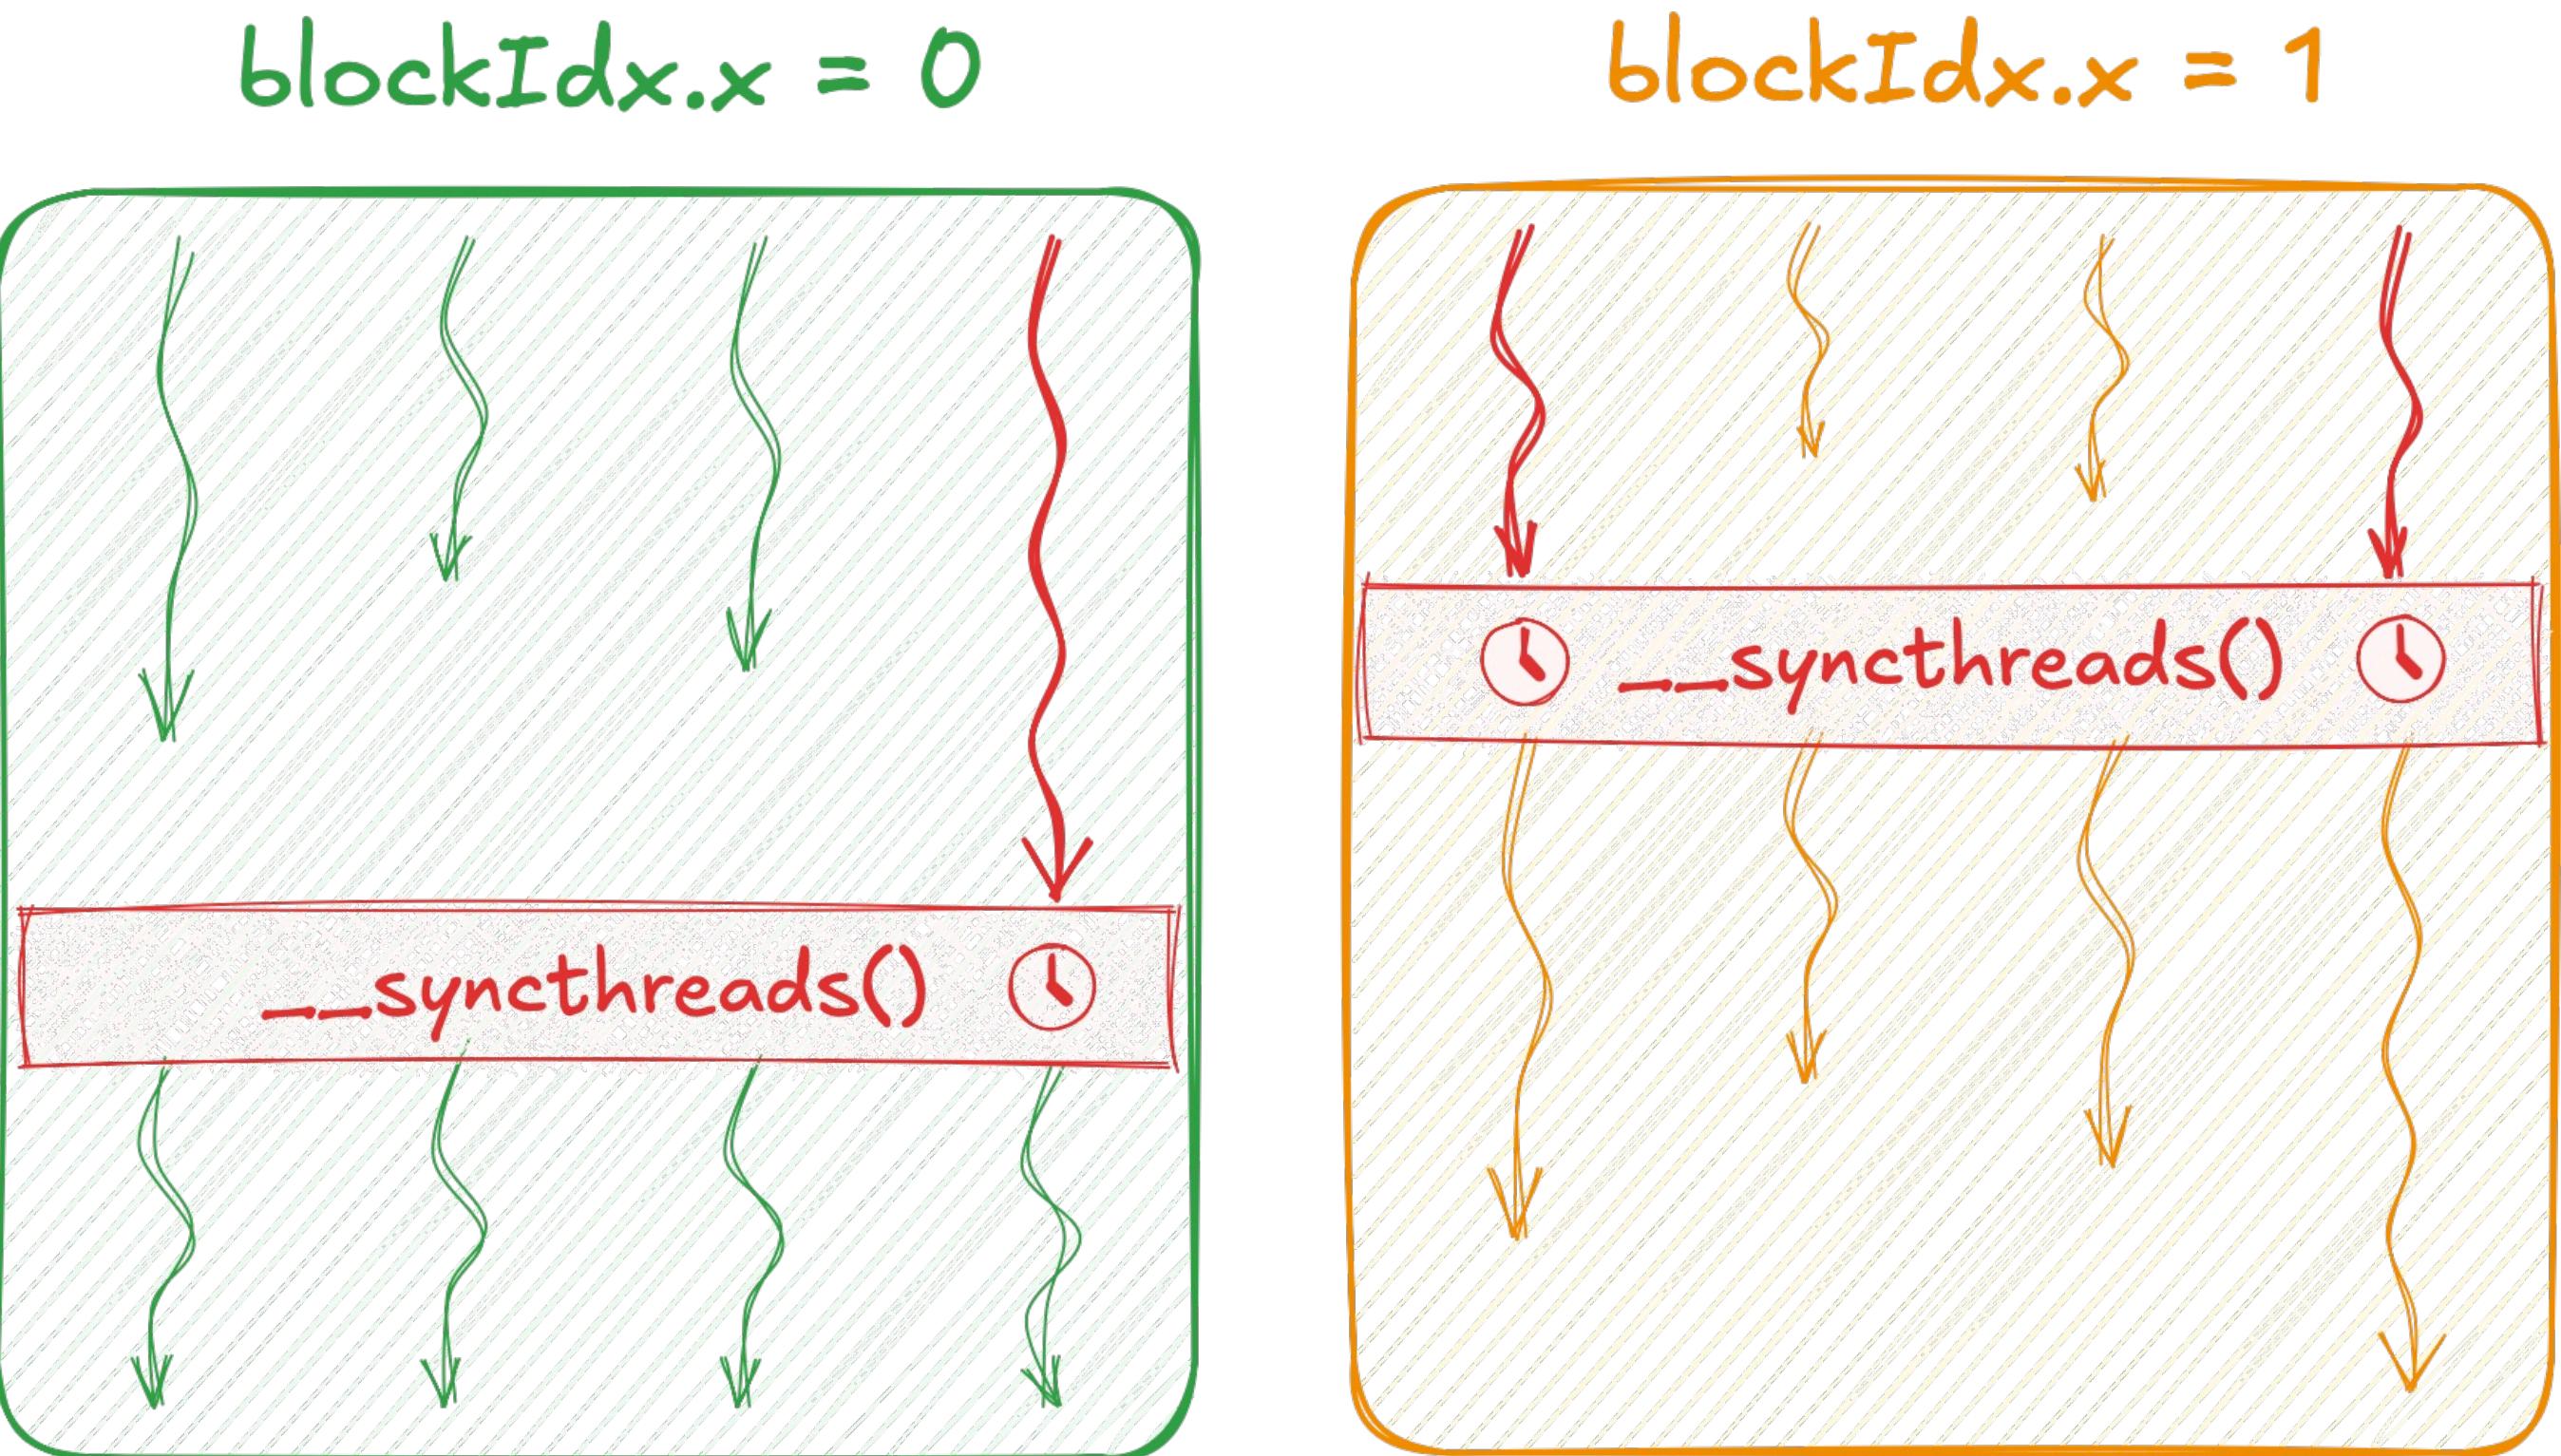
- Atomic operations
- Thread scopes
- Sub spans

But what about the initial question?

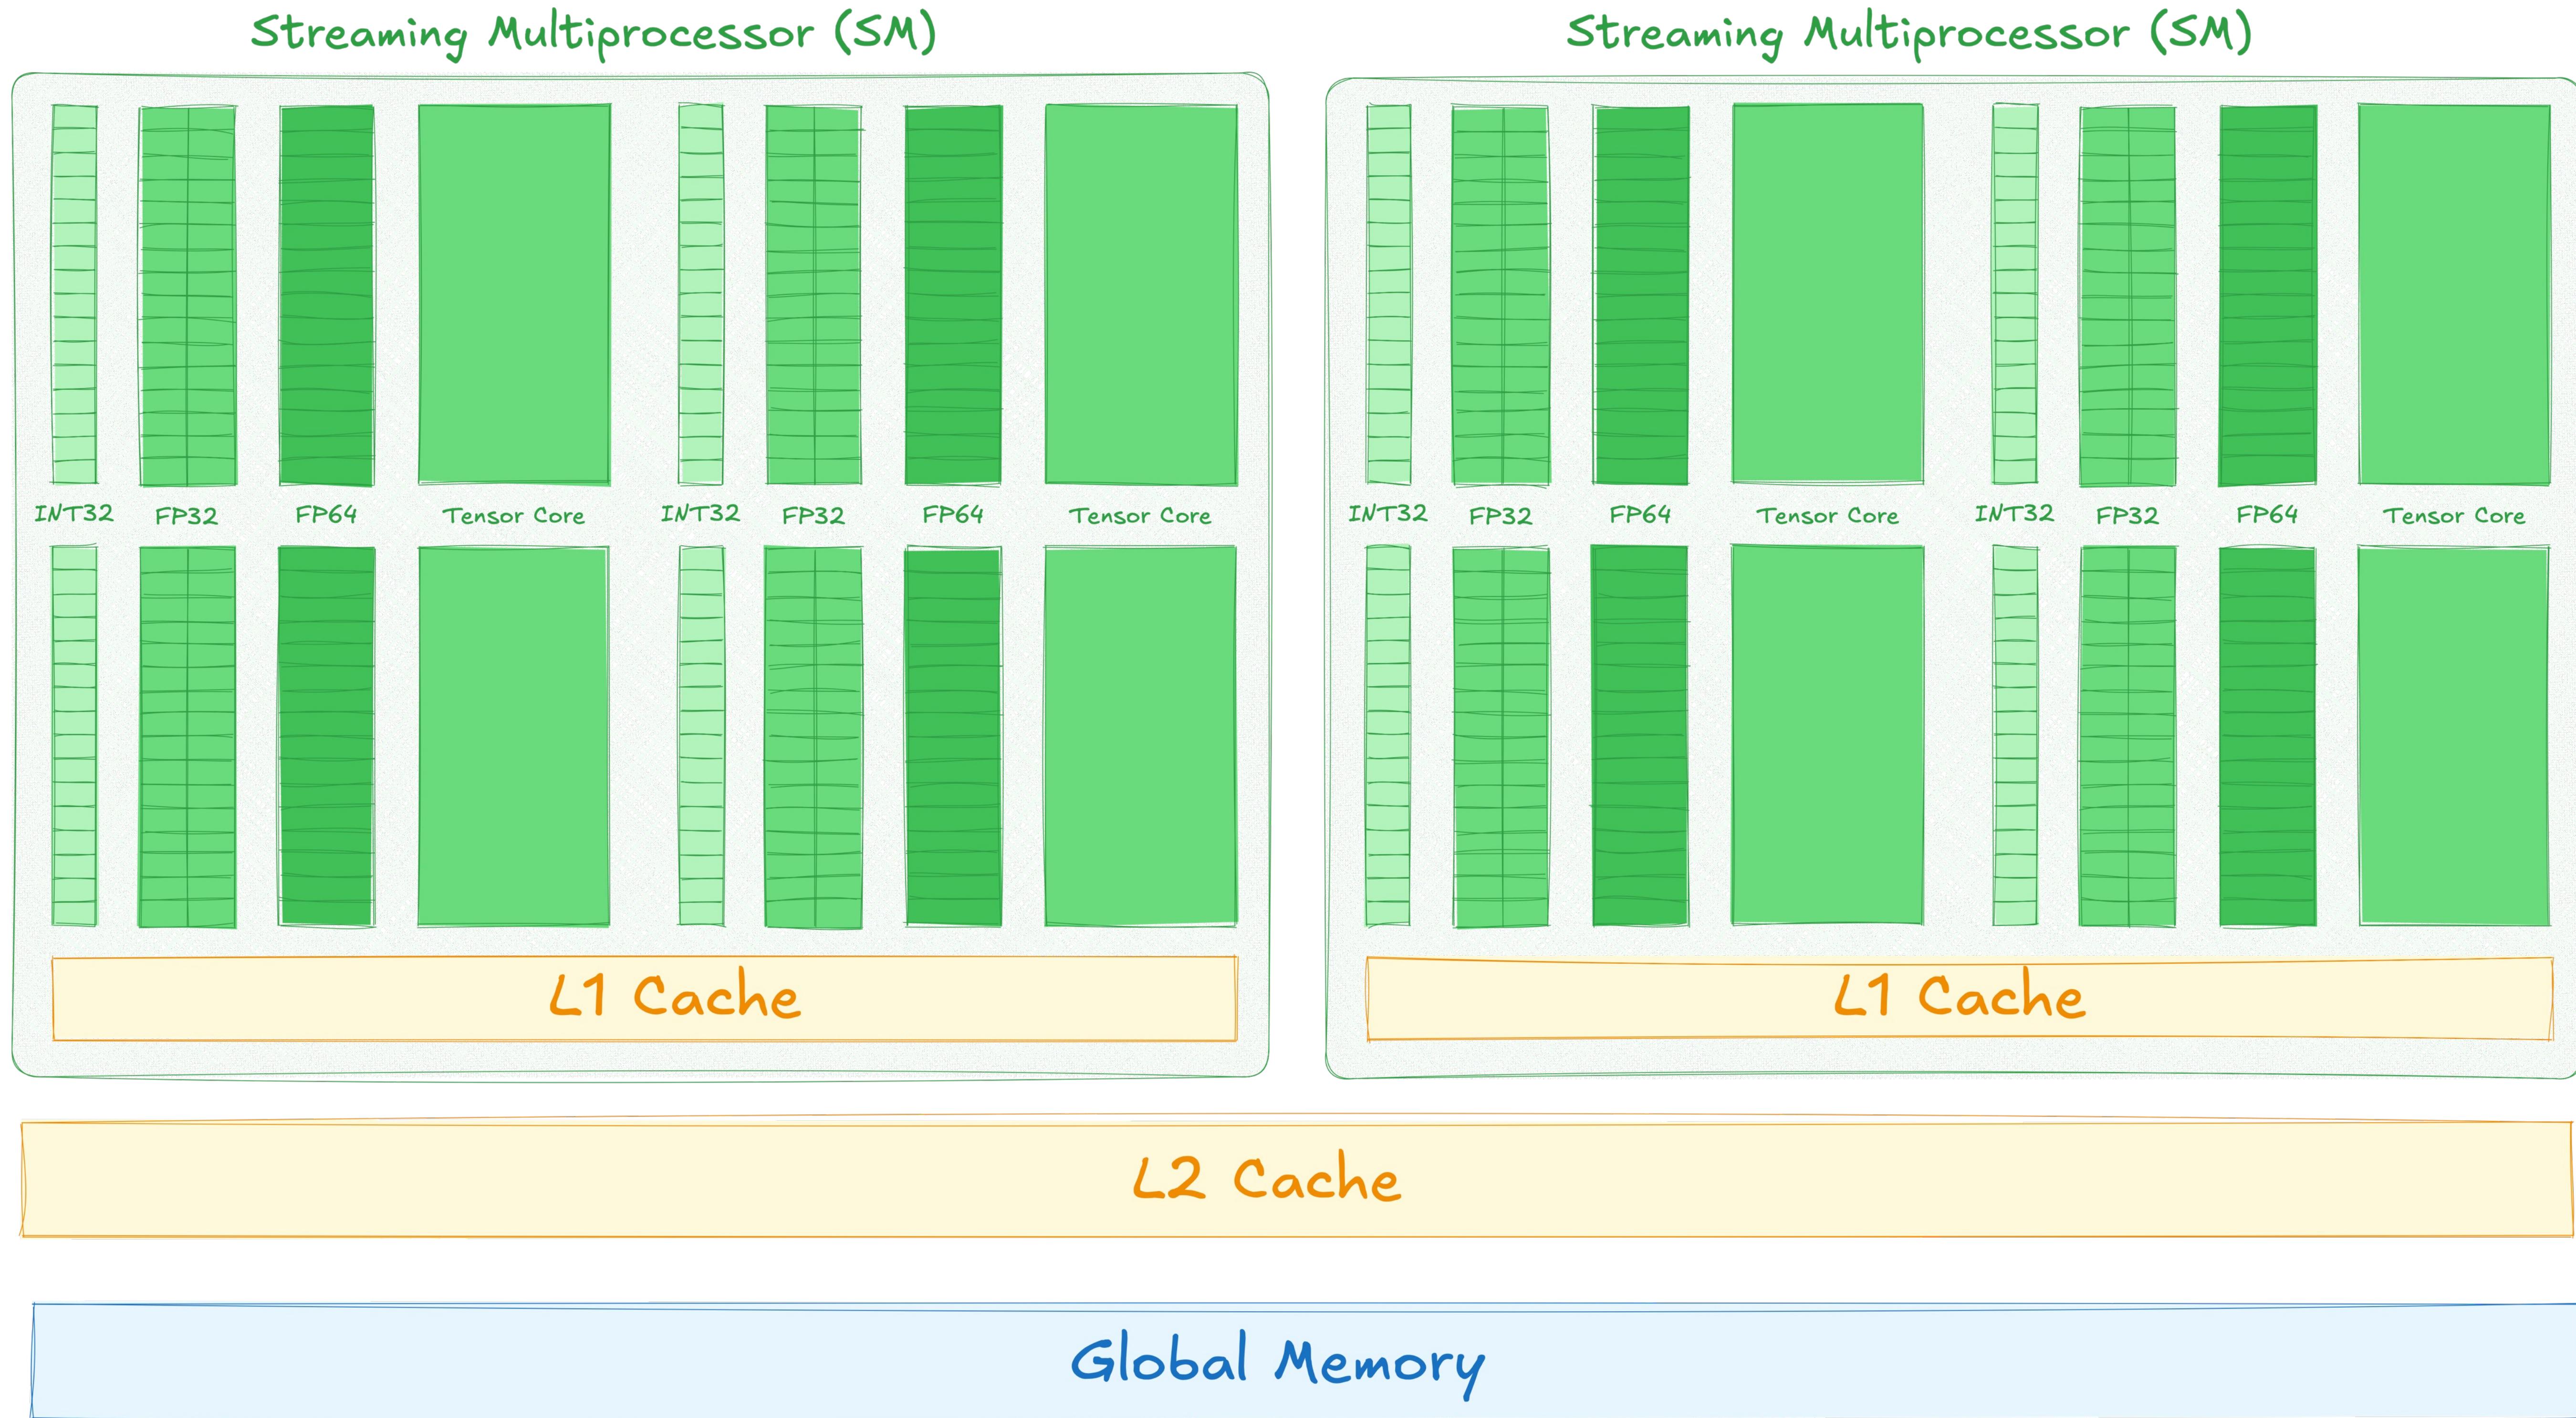
Why was thread hierarchy needed in the first place?

# Recap

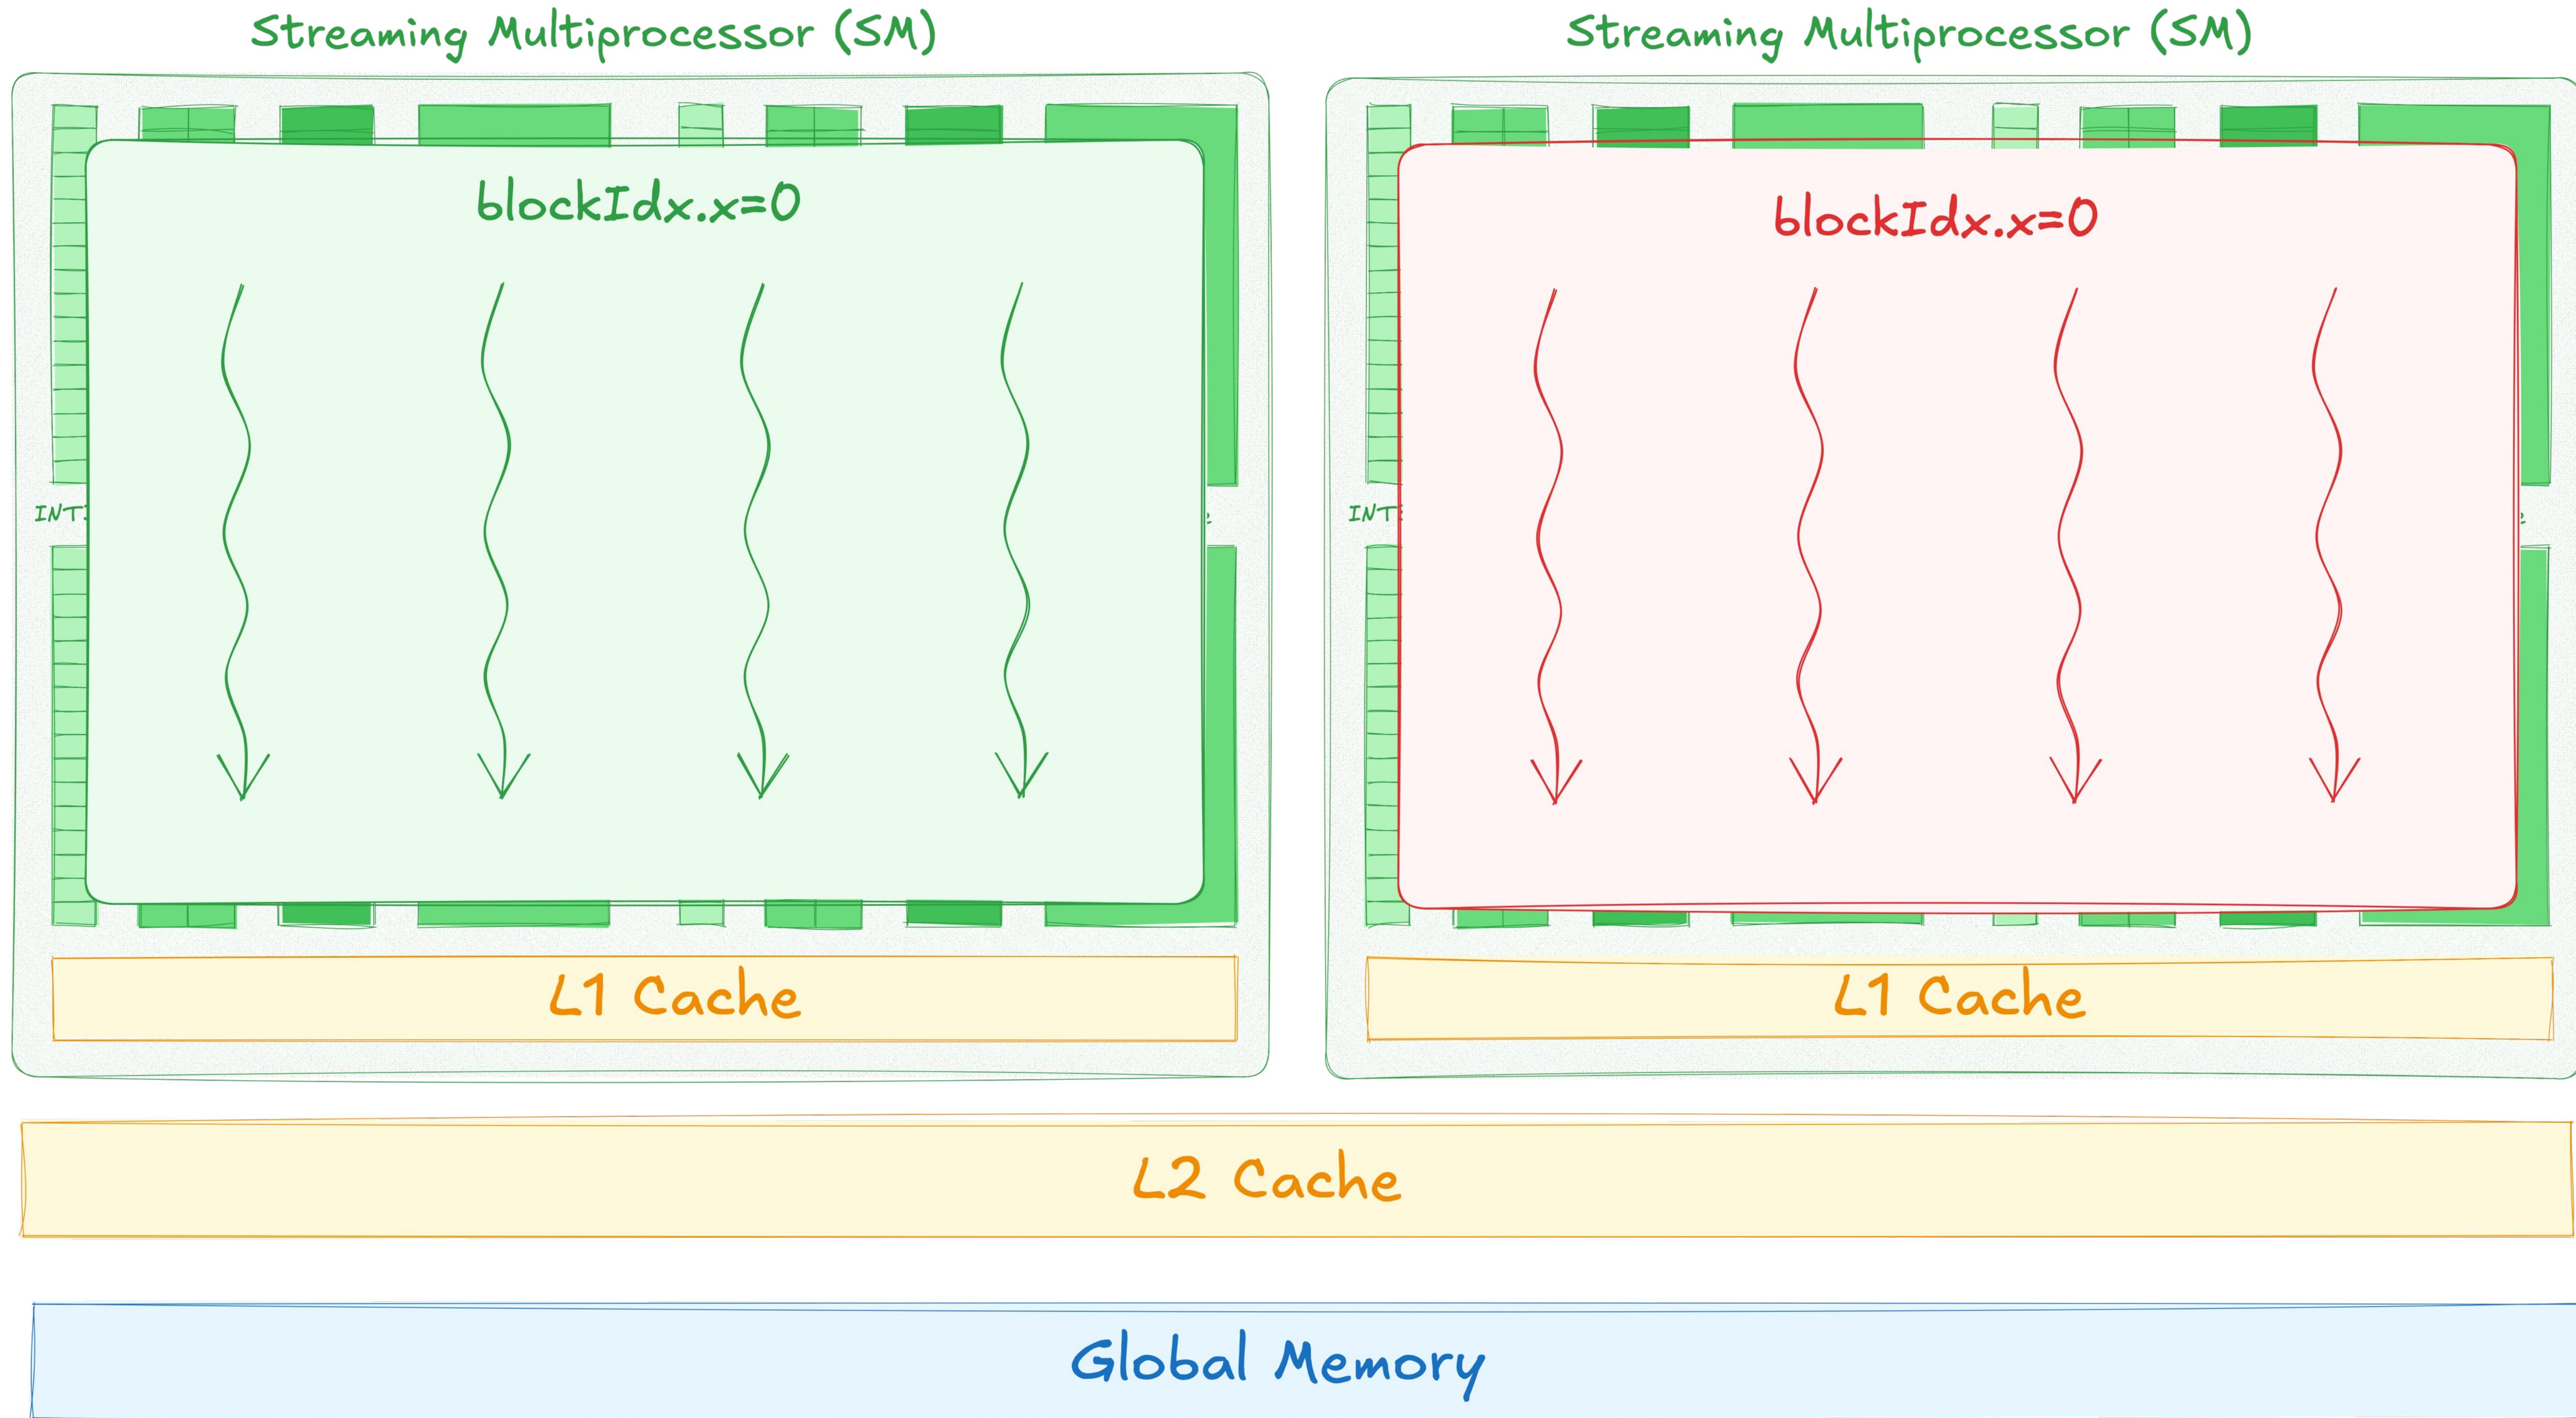
```
cuda::std::span<int> block_histogram =  
    block_histograms.subspan(  
        blockIdx.x * histogram.size(),  
        histogram.size());  
  
int cell = blockIdx.x * blockDim.x + threadIdx.x;  
int bin = static_cast<int>(temperatures[cell] / bin_width);  
  
cuda::atomic_ref<int, cuda::thread_scope_block>  
    block_ref(block_histogram[bin]);  
block_ref.fetch_add(1);  
__syncthreads();  
  
if (threadIdx.x < num_bins)  
{  
    cuda::atomic_ref<int, cuda::thread_scope_device>  
        ref(histogram[threadIdx.x]);  
    ref.fetch_add(block_histogram[threadIdx.x]);  
}
```



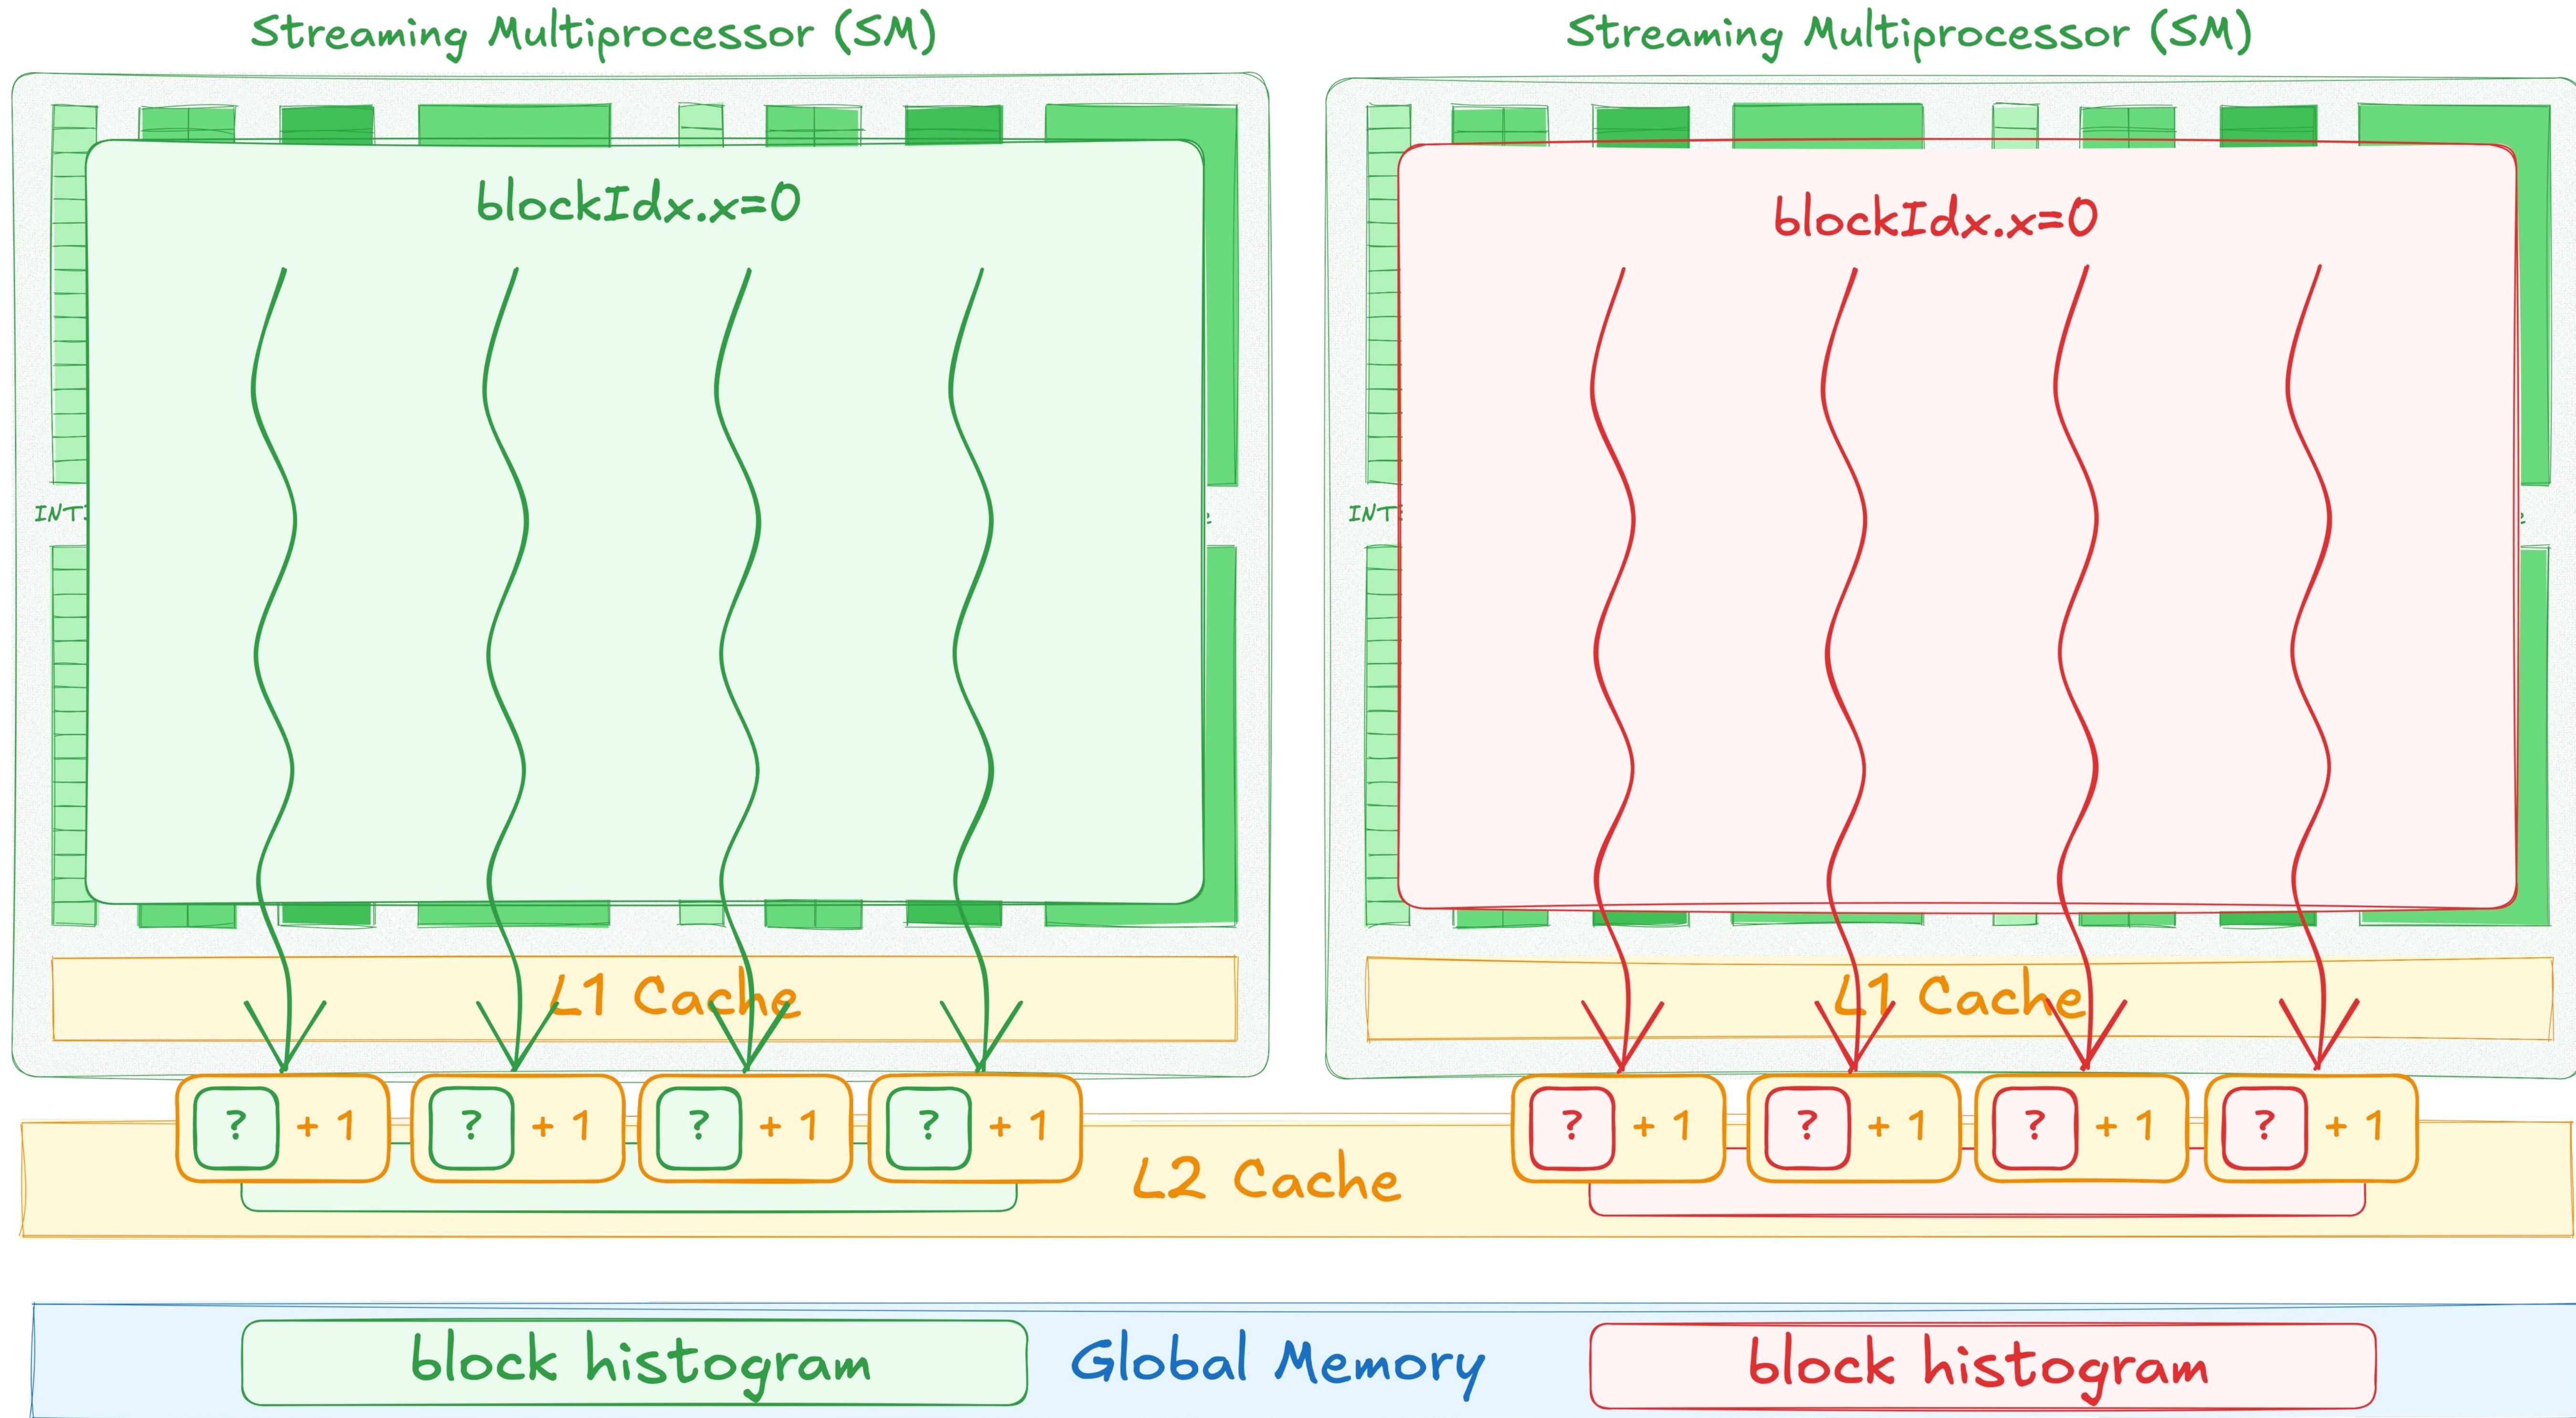
- Let's reconsider `__syncthreads`
- Synchronization is only available between threads of a given thread block
- This makes different groups of threads not equally capable, hence introducing non-uniformity in thread grouping
- Are there any other resources that are only accessible within certain groups of threads?



- GPUs consist of building blocks called **SMs**
- VRAM is called global memory (not local to any group of threads)

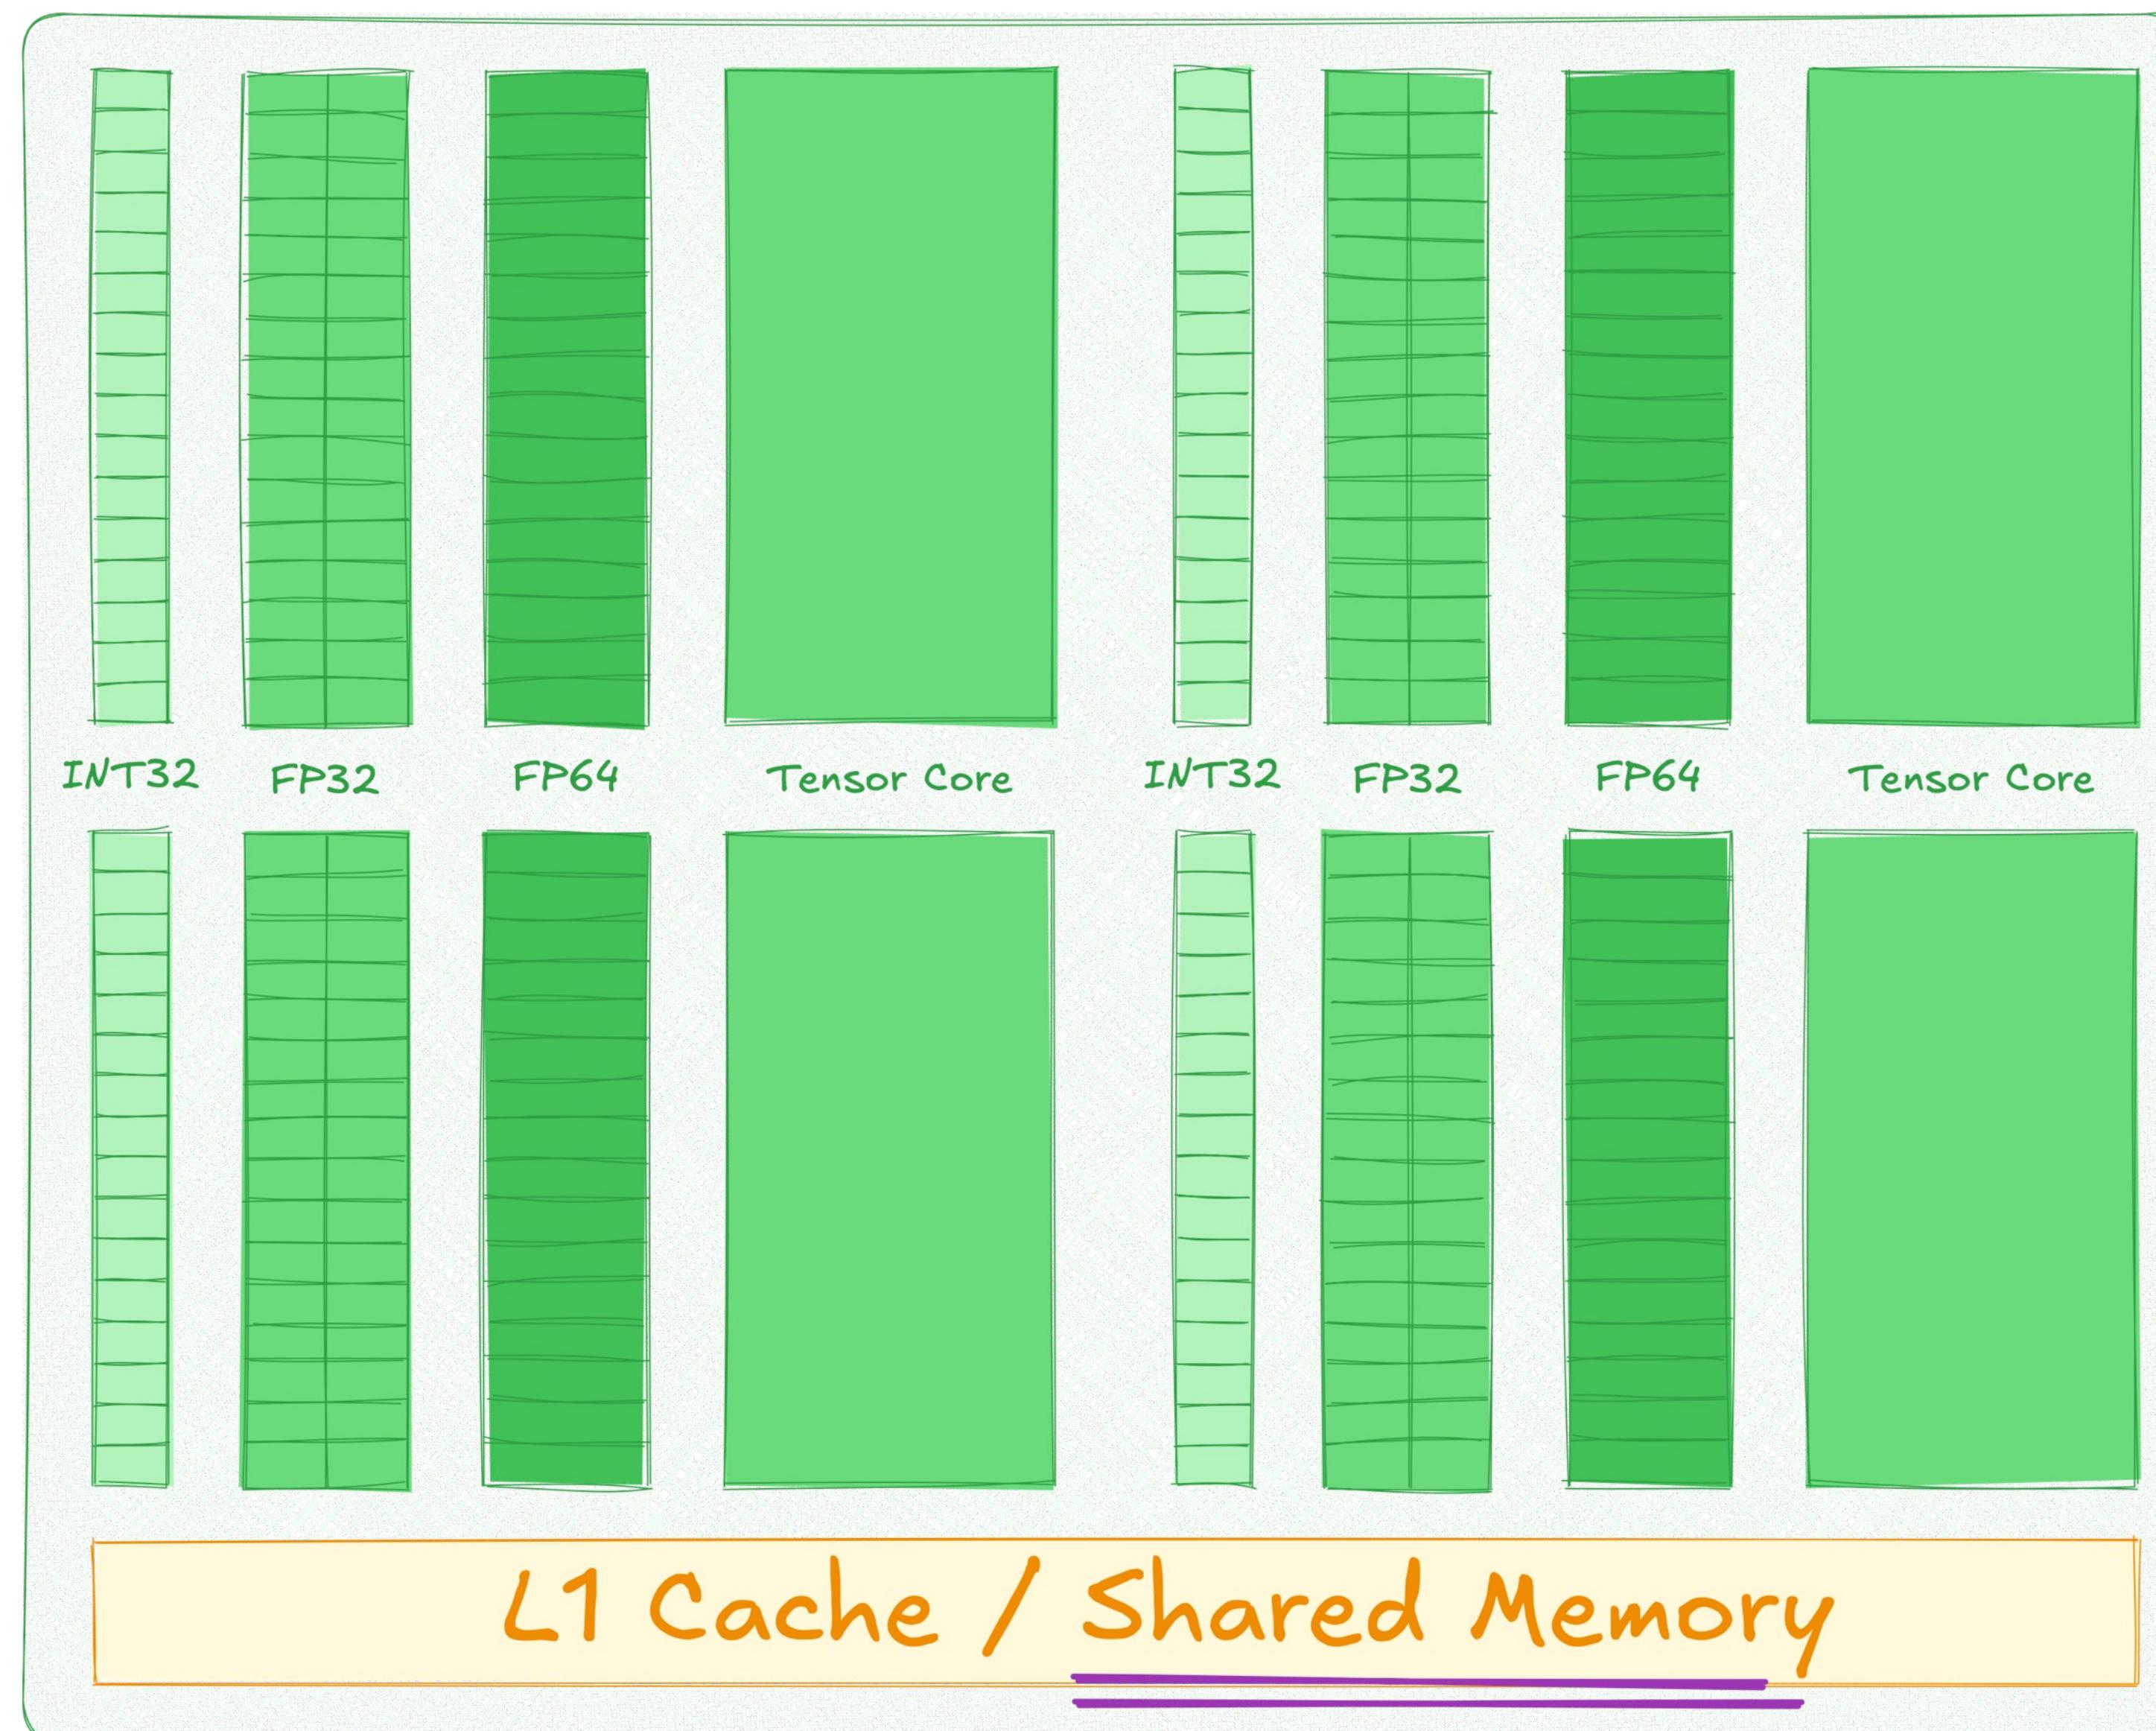


- Thread block never spans multiple **SMs**
- Some resources, like `__syncthreads` and **L1** are local to **SMs**

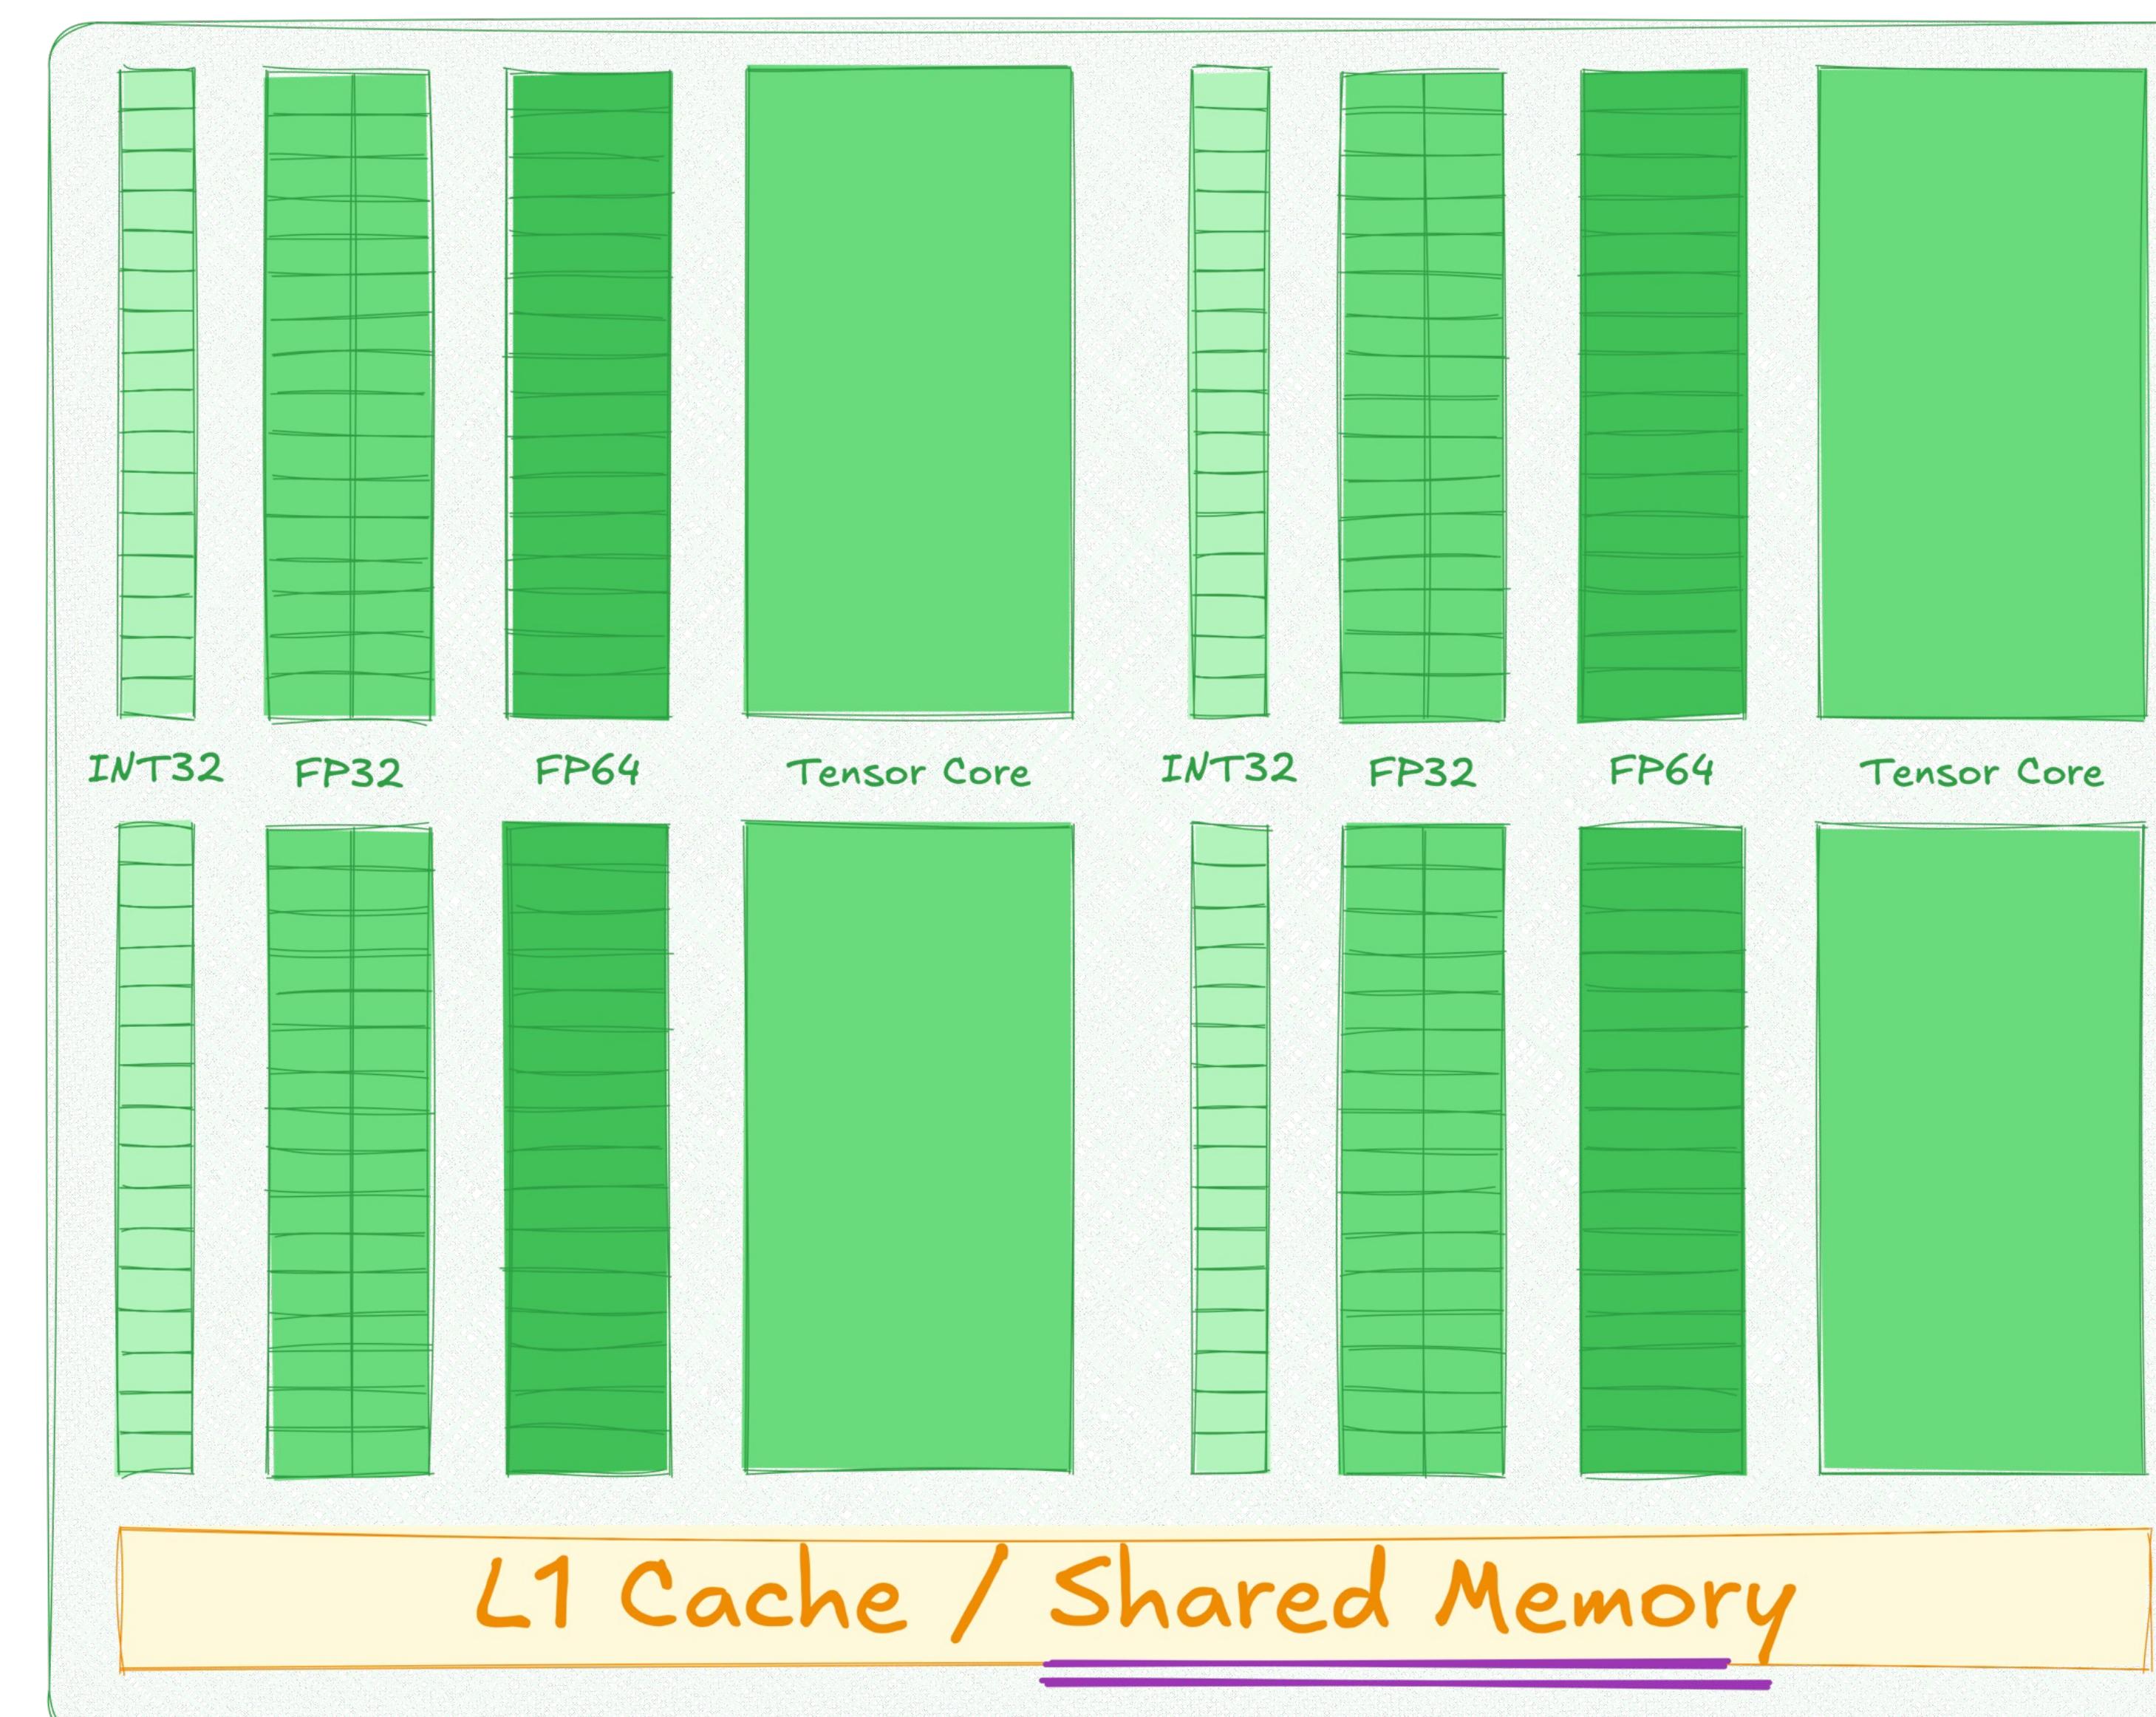


- Our kernel repeatedly issues atomic increments, which is likely going into L2
- But there's some inefficiency. We don't actually need block histograms to be written into global memory. We also waste quite a lot of space (one histogram per thread block)
- It'd be nice if we could allocate memory in L1

Streaming Multiprocessor (SM)



Streaming Multiprocessor (SM)



L2 Cache

- We actually can do that!
- CUDA provides a software-defined cache that's called shared memory
- Shared memory is only accessible within a given thread block and is collocated with the L1

# Shared Memory

```
__global__ void kernel()
{
    __shared__ int shared[4]; | • To allocate array in shared memory space, use
                            | __shared__ specifier
    shared[threadIdx.x] = threadIdx.x;

    __syncthreads();

    if (threadIdx.x == 0) {
        for (int i = 0; i < 4; i++) {
            std::printf("shared[%d] = %d\n", i, shared[i]);
        }
    }
}
```

# Shared Memory

```
__global__ void kernel()
{
    __shared__ int shared[4];

    shared[threadIdx.x] = threadIdx.x;
    __syncthreads();

    if (threadIdx.x == 0) {
        for (int i = 0; i < 4; i++) {
            std::printf("shared[%d] = %d\n", i, shared[i]);
        }
    }
}
```

- You can use it as if it was an ordinary array
- Just be mindful of other threads and avoid data races with `__syncthreads`

# Shared Memory

```
--global__ void kernel()
{
    __shared__ int shared[4];

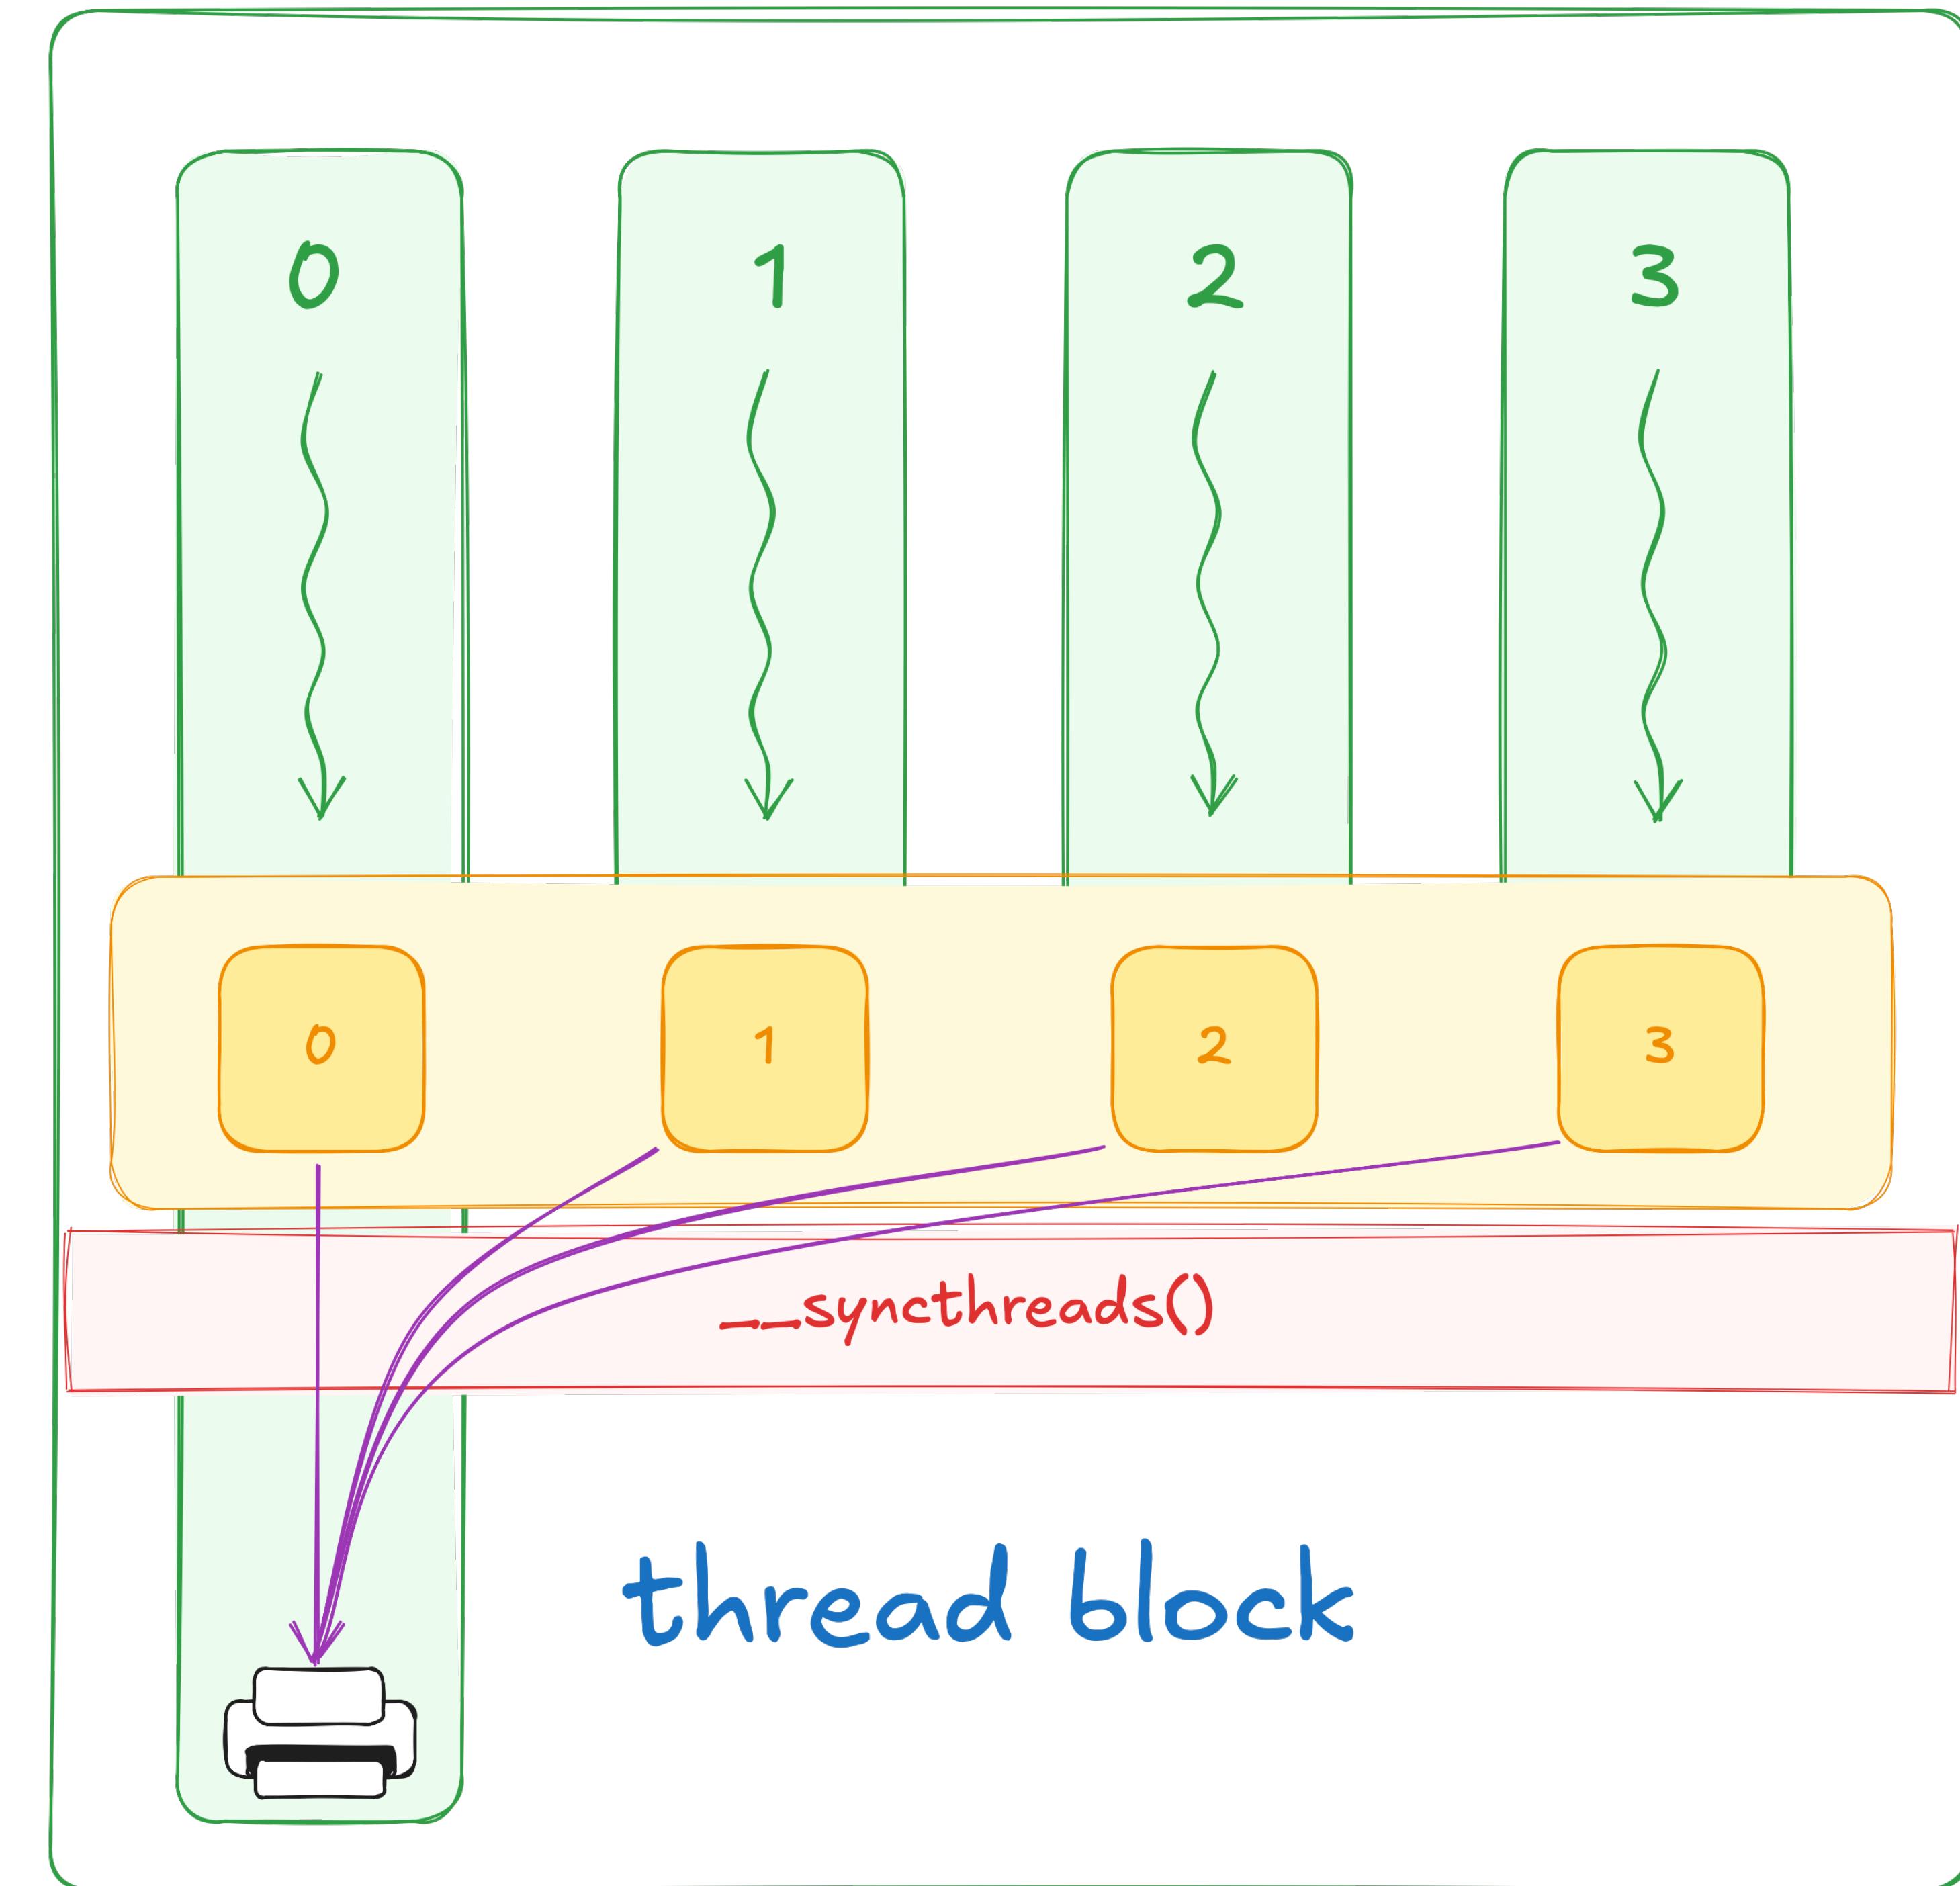
    shared[threadIdx.x] = threadIdx.x;
    __syncthreads();

    if (threadIdx.x == 0) {
        for (int i = 0; i < 4; i++) {
            std::printf("shared[%d] = %d\n", i, shared[i]);
        }
    }
}
```

thread local

shared

thread block



# Exercise: Optimize Histogram

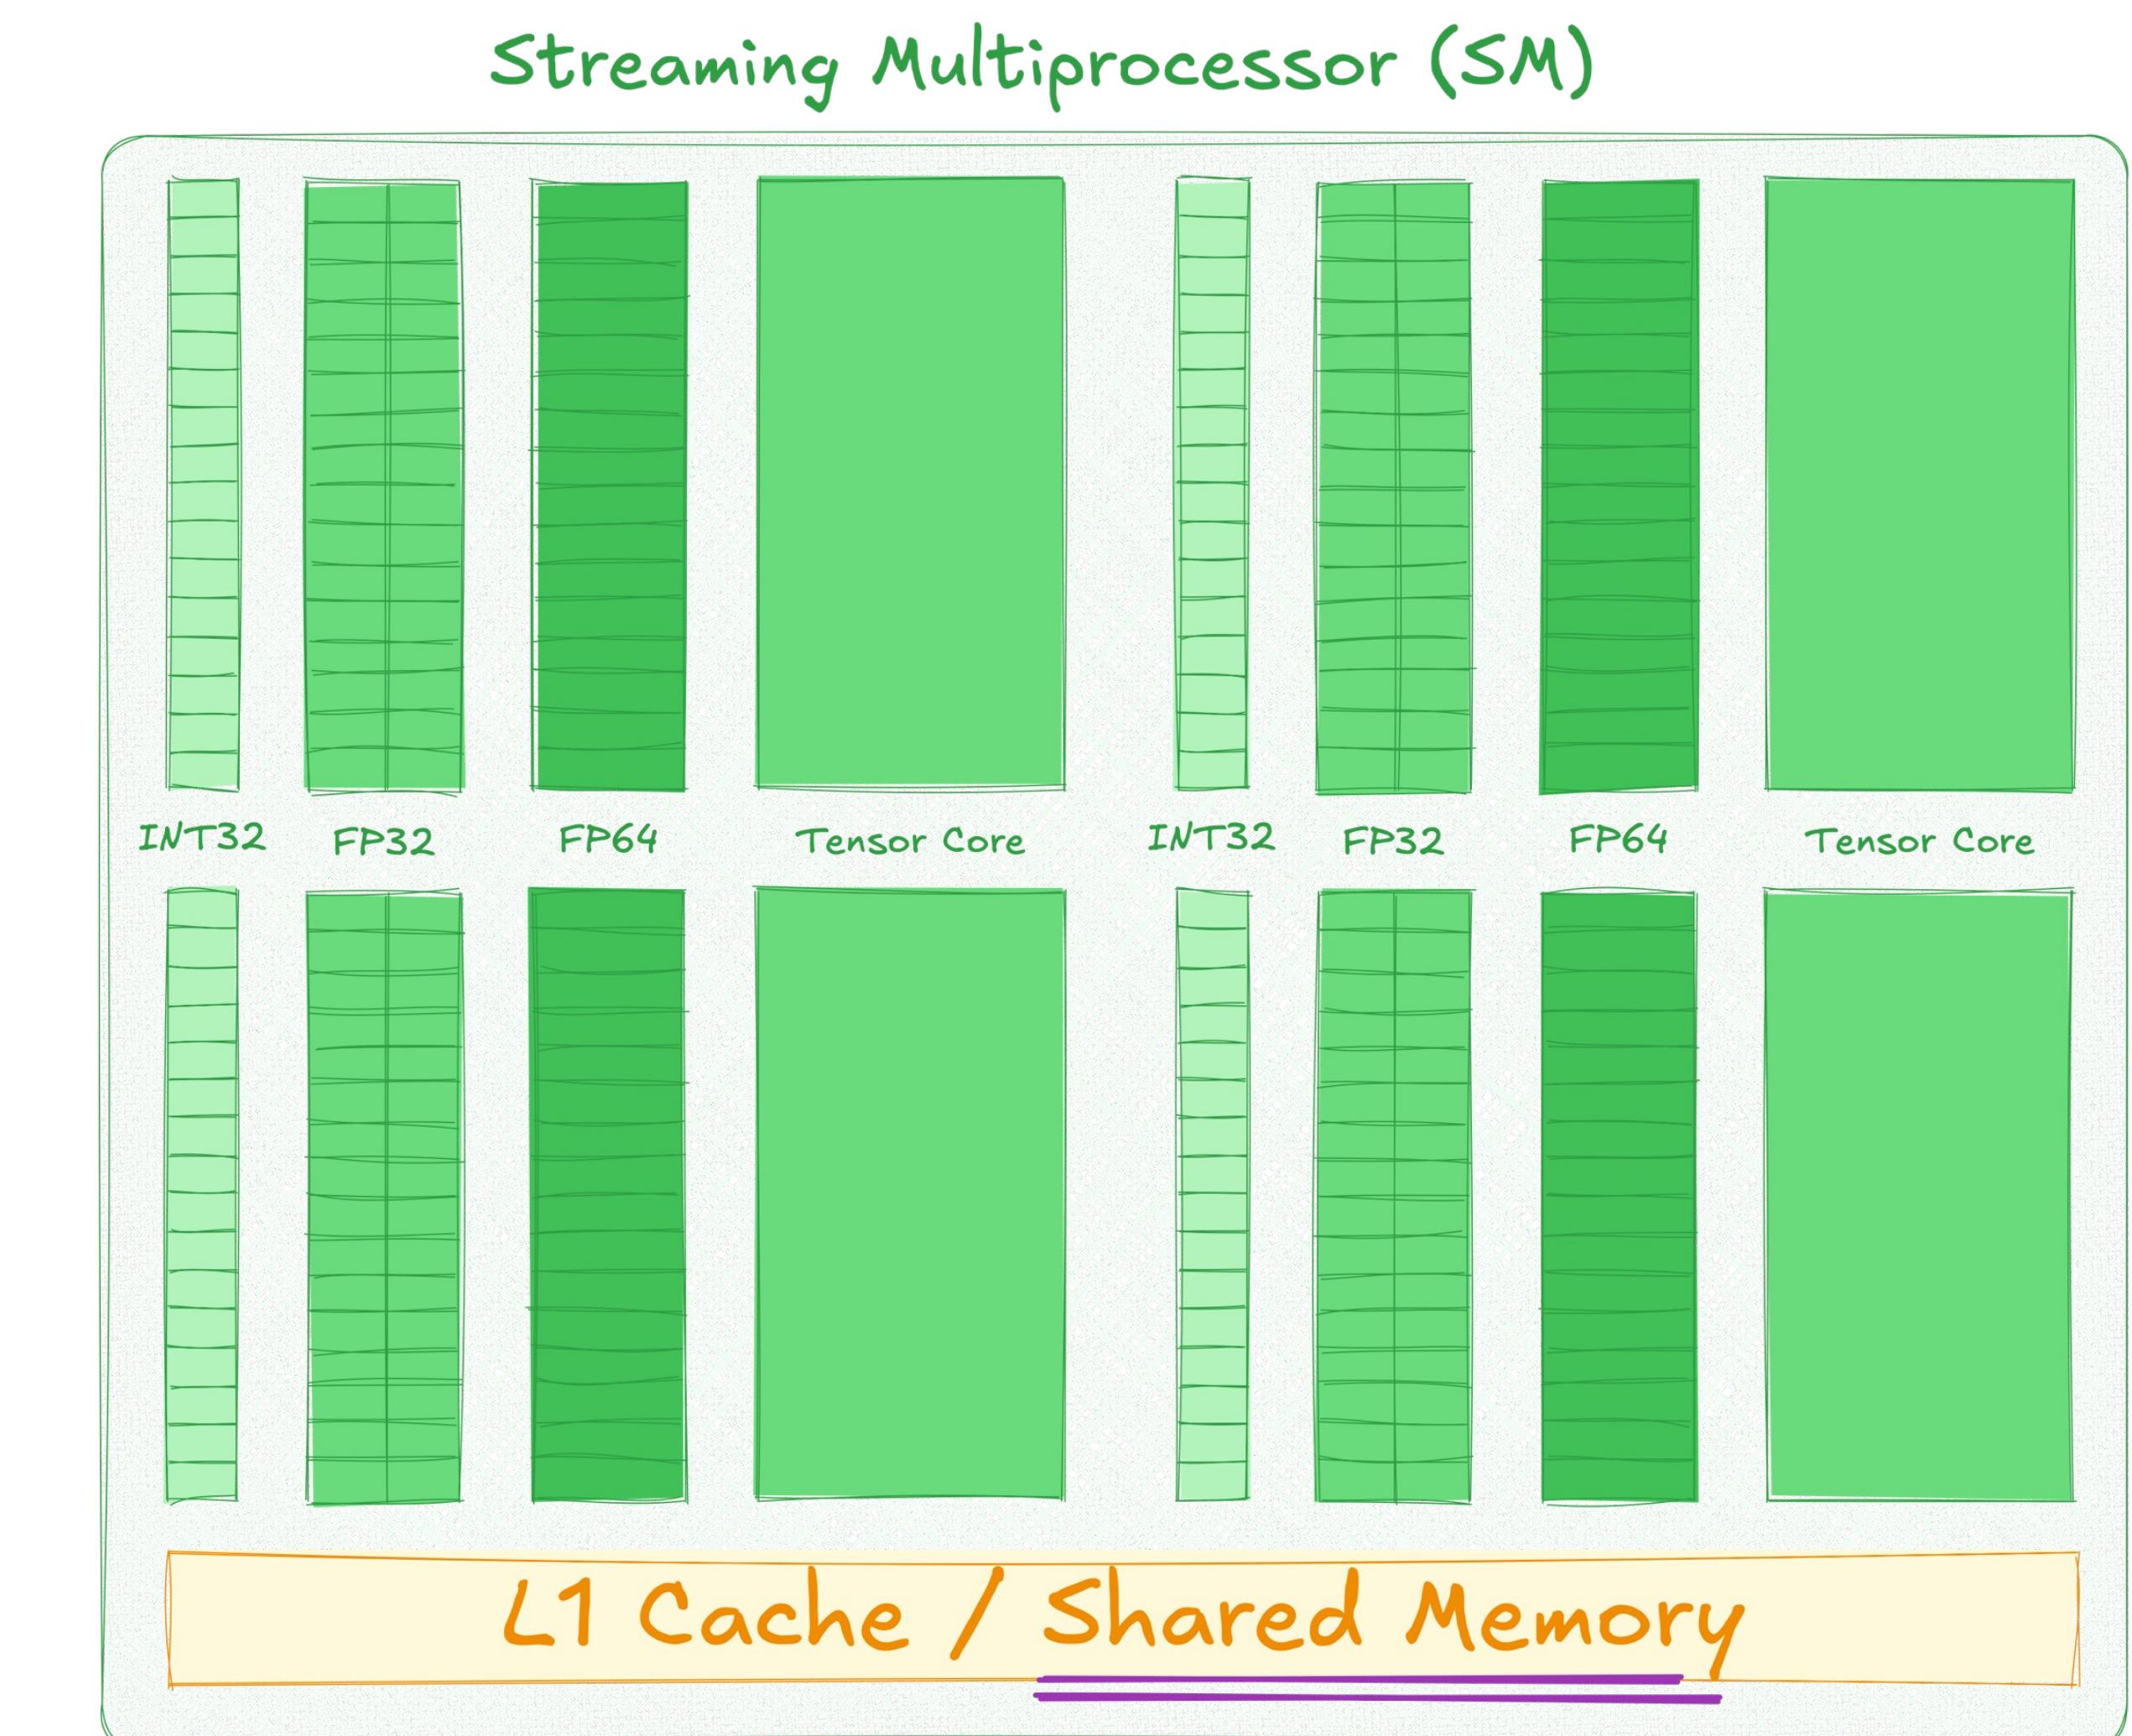
- Use shared instead of global memory to store block histogram

```
__global__ void histogram_kernel(
    cuda::std::span<float> temperatures,
    cuda::std::span<int> block_histograms,
    cuda::std::span<int> histogram)
{
    cuda::std::span<int> block_histogram =
        block_histograms.subspan(
            blockIdx.x * histogram.size(),
            histogram.size());

    int cell = blockIdx.x * blockDim.x + threadIdx.x;
    int bin = static_cast<int>(temperatures[cell] / bin_width);

    cuda::atomic_ref<int, cuda::thread_scope_block>
        block_ref(block_histogram[bin]);
    block_ref.fetch_add(1);
    __syncthreads();

    ...
}
```



# Exercise: Optimize Histogram

## Solution

```
__global__ void histogram_kernel(
    cuda::std::span<float> temperatures,
    cuda::std::span<int> histogram)
{
    __shared__ int block_histogram[num_bins];

    if (threadIdx.x < num_bins) {
        block_histogram[threadIdx.x] = 0;
    }
    __syncthreads();

    int cell = blockIdx.x * blockDim.x + threadIdx.x;
    int bin = static_cast<int>(temperatures[cell] / bin_width);

    cuda::atomic_ref<int, cuda::thread_scope_block>
        block_ref(block_histogram[bin]);
    block_ref.fetch_add(1);
    __syncthreads();
    ...
}
```

- Kernel doesn't need block histograms allocation now

# Exercise: Optimize Histogram

## Solution

```
--global__ void histogram_kernel(
    cuda::std::span<float> temperatures,
    cuda::std::span<int> histogram)
{
    __shared__ int block_histogram[num_bins];

    if (threadIdx.x < num_bins) {
        block_histogram[threadIdx.x] = 0;
    }
    __syncthreads();

    int cell = blockIdx.x * blockDim.x + threadIdx.x;
    int bin = static_cast<int>(temperatures[cell] / bin_width);

    cuda::atomic_ref<int, cuda::thread_scope_block>
        block_ref(block_histogram[bin]);
    block_ref.fetch_add(1);
    __syncthreads();
    ...
}
```

- We start by allocating block histogram in shared memory
- Then we initialize block histogram to zero by first threads
- And synchronize block to make sure other threads wait for block histogram to be properly initialized

# Exercise: Optimize Histogram

## Solution

```
--global__ void histogram_kernel(
    cuda::std::span<float> temperatures,
    cuda::std::span<int> histogram)
{
    __shared__ int block_histogram[num_bins];

    if (threadIdx.x < num_bins) {
        block_histogram[threadIdx.x] = 0;
    }
    __syncthreads();

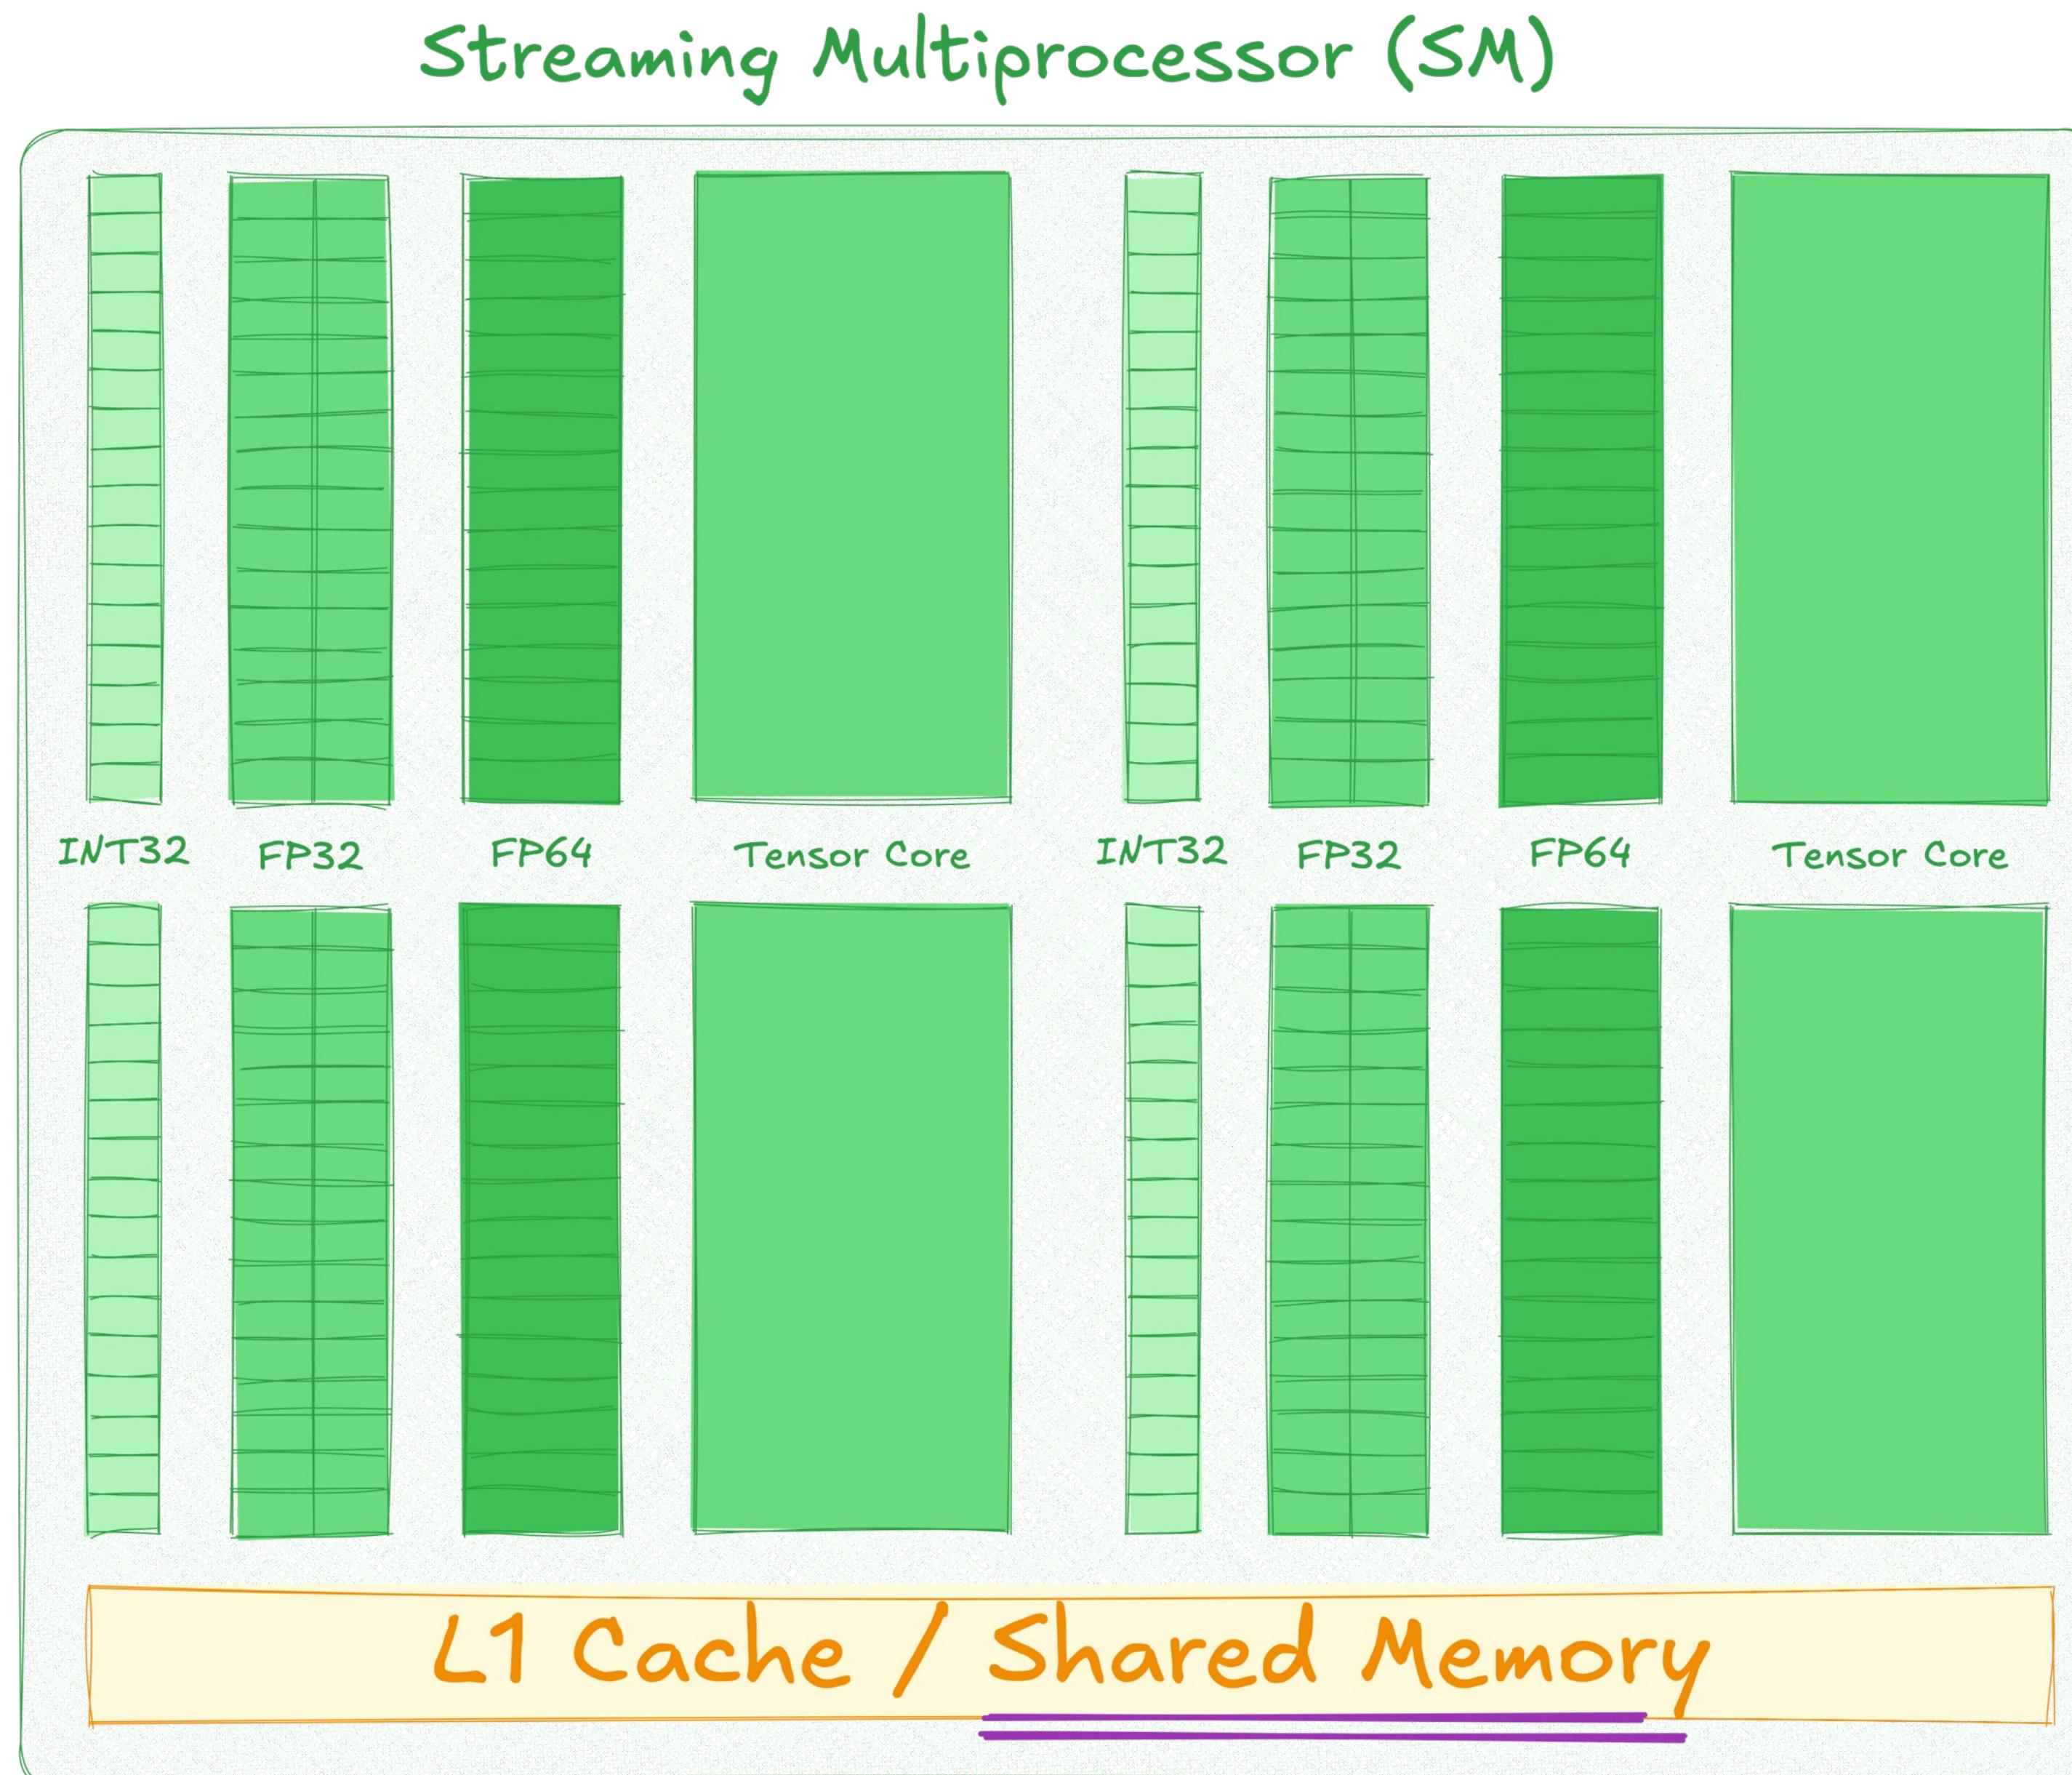
    int cell = blockIdx.x * blockDim.x + threadIdx.x;
    int bin = static_cast<int>(temperatures[cell] / bin_width);

    cuda::atomic_ref<int, cuda::thread_scope_block>
        block_ref(block_histogram[bin]);
    block_ref.fetch_add(1);
    __syncthreads();
    ...
}
```

- We then wrap shared memory reference to an atomic reference as usual

# Exercise: Optimize Histogram

## Solution



- Overall, this brings us from 100 GB/s to 400 GB/s

# Cooperative Libraries

## CUB

Cooperative general-purpose algorithms inside CUDA kernel

## cuBLASDx

Cooperative linear algebra functions inside CUDA kernel

## cuFFTDx

Cooperative fast Fourier transform inside CUDA kernel

...

- Authoring a CUDA kernel doesn't mean implementing every algorithm from scratch
- CUDA provides variety of cooperative libraries that make your kernels faster and development times shorter

# Cooperative Libraries

## CUB

Cooperative general-purpose algorithms inside CUDA kernel

## cuBLASDx

Cooperative linear algebra functions inside CUDA kernel

## cuFFTDx

Cooperative fast Fourier transform inside CUDA kernel

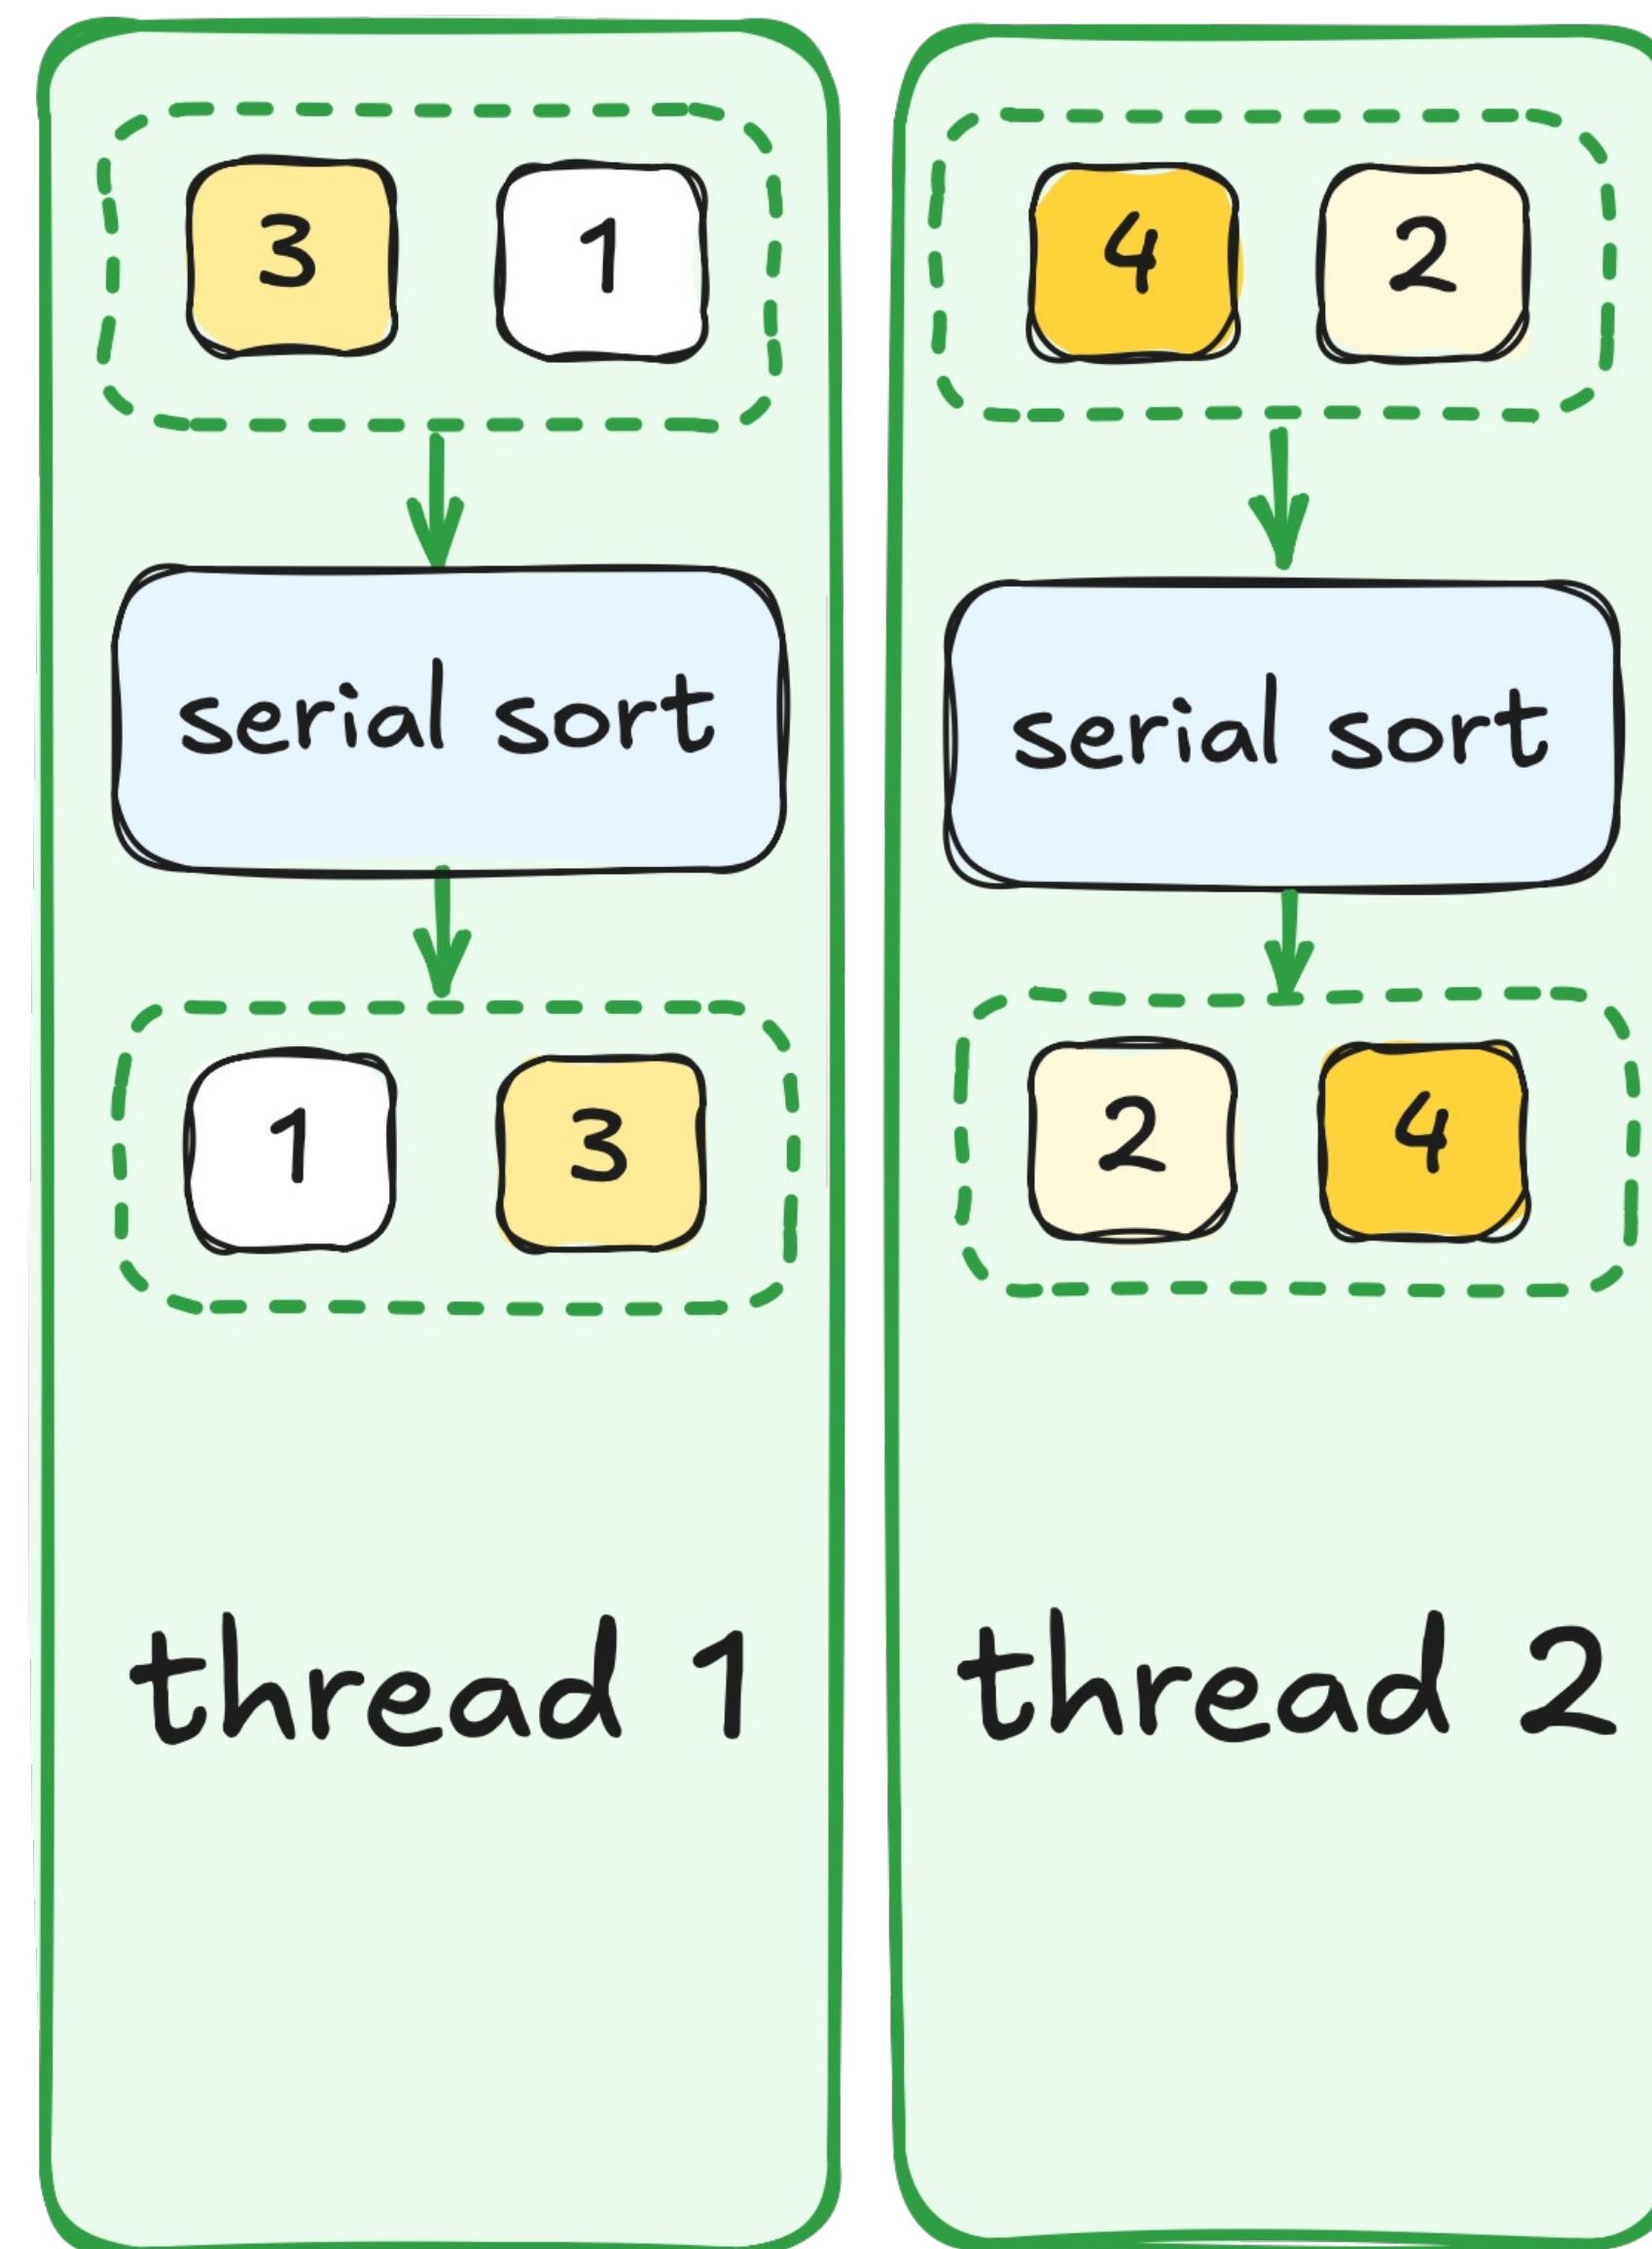
...

- Authoring a CUDA kernel doesn't mean implementing every algorithm from scratch
- CUDA provides variety of cooperative libraries that make your kernels faster and development times shorter
- But what exactly does cooperative mean?

# Cooperative Algorithms

## Serial

- **One thread** invokes algorithm
- **One thread** executes algorithm



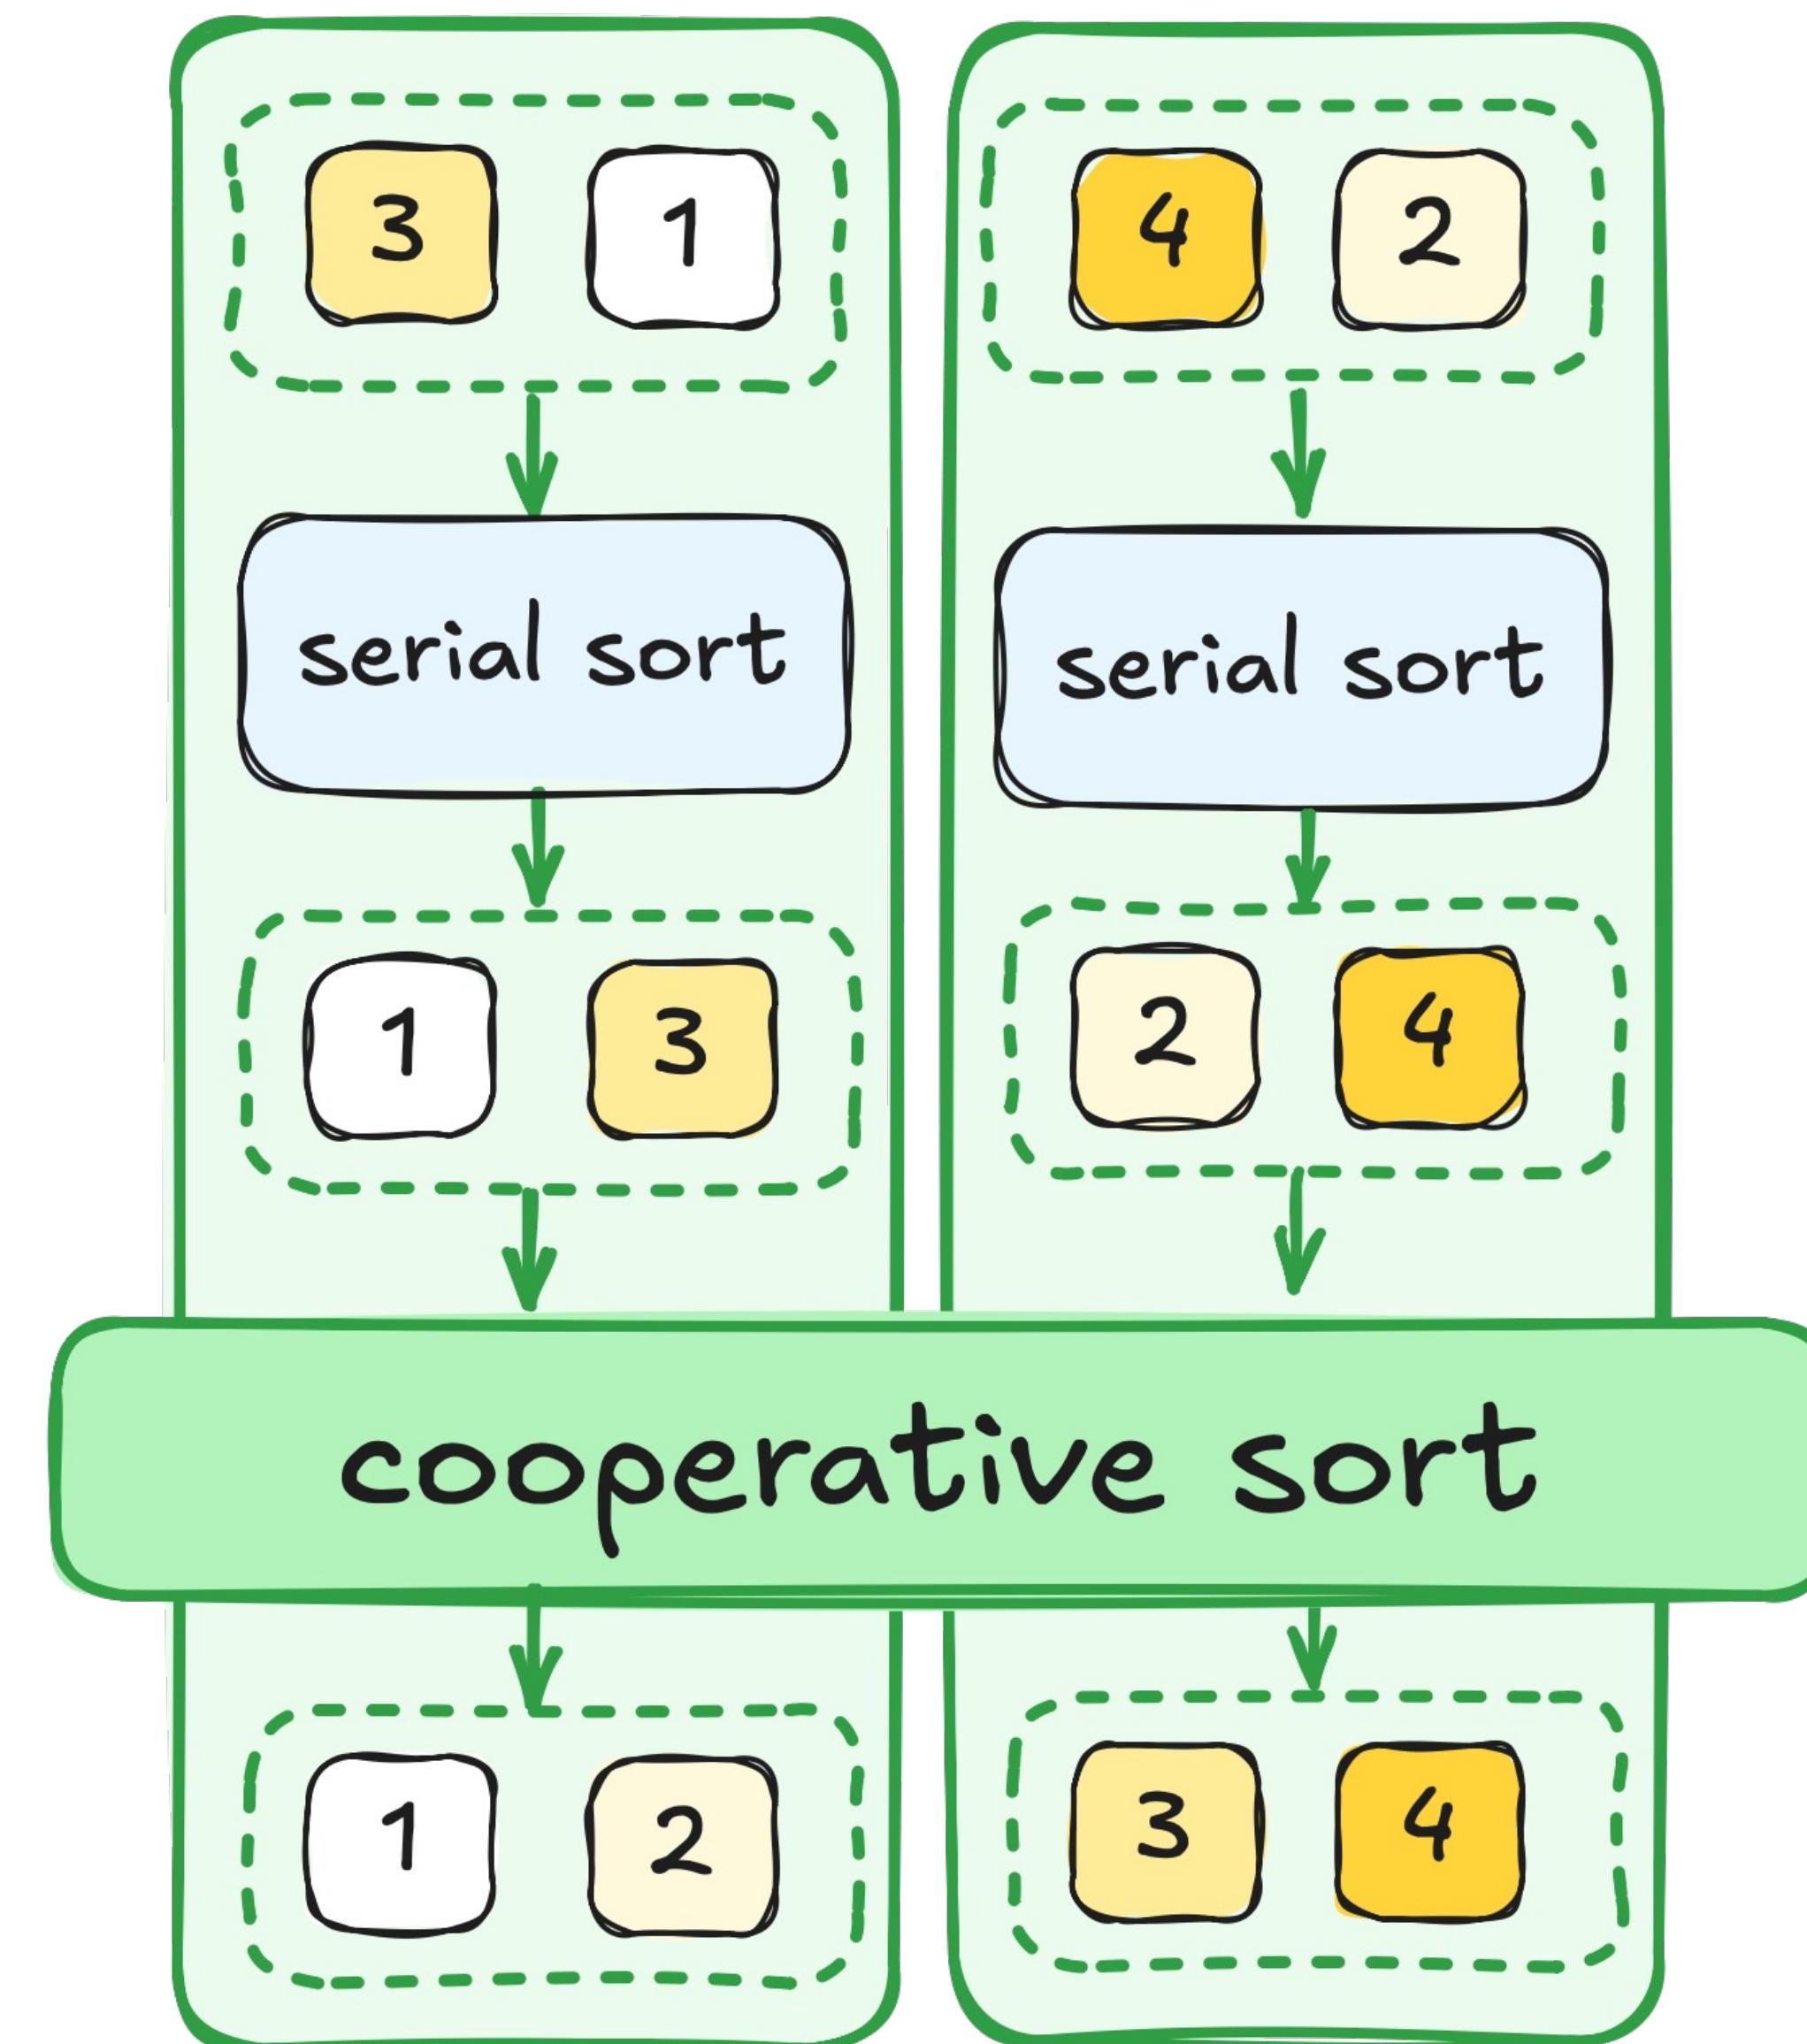
# Cooperative Algorithms

## Serial

- **One thread** invokes algorithm
- **One thread** executes algorithm

## Cooperative

- **Many threads** invoke algorithm
- **Many threads** execute algorithm



# Cooperative Algorithms

## Serial

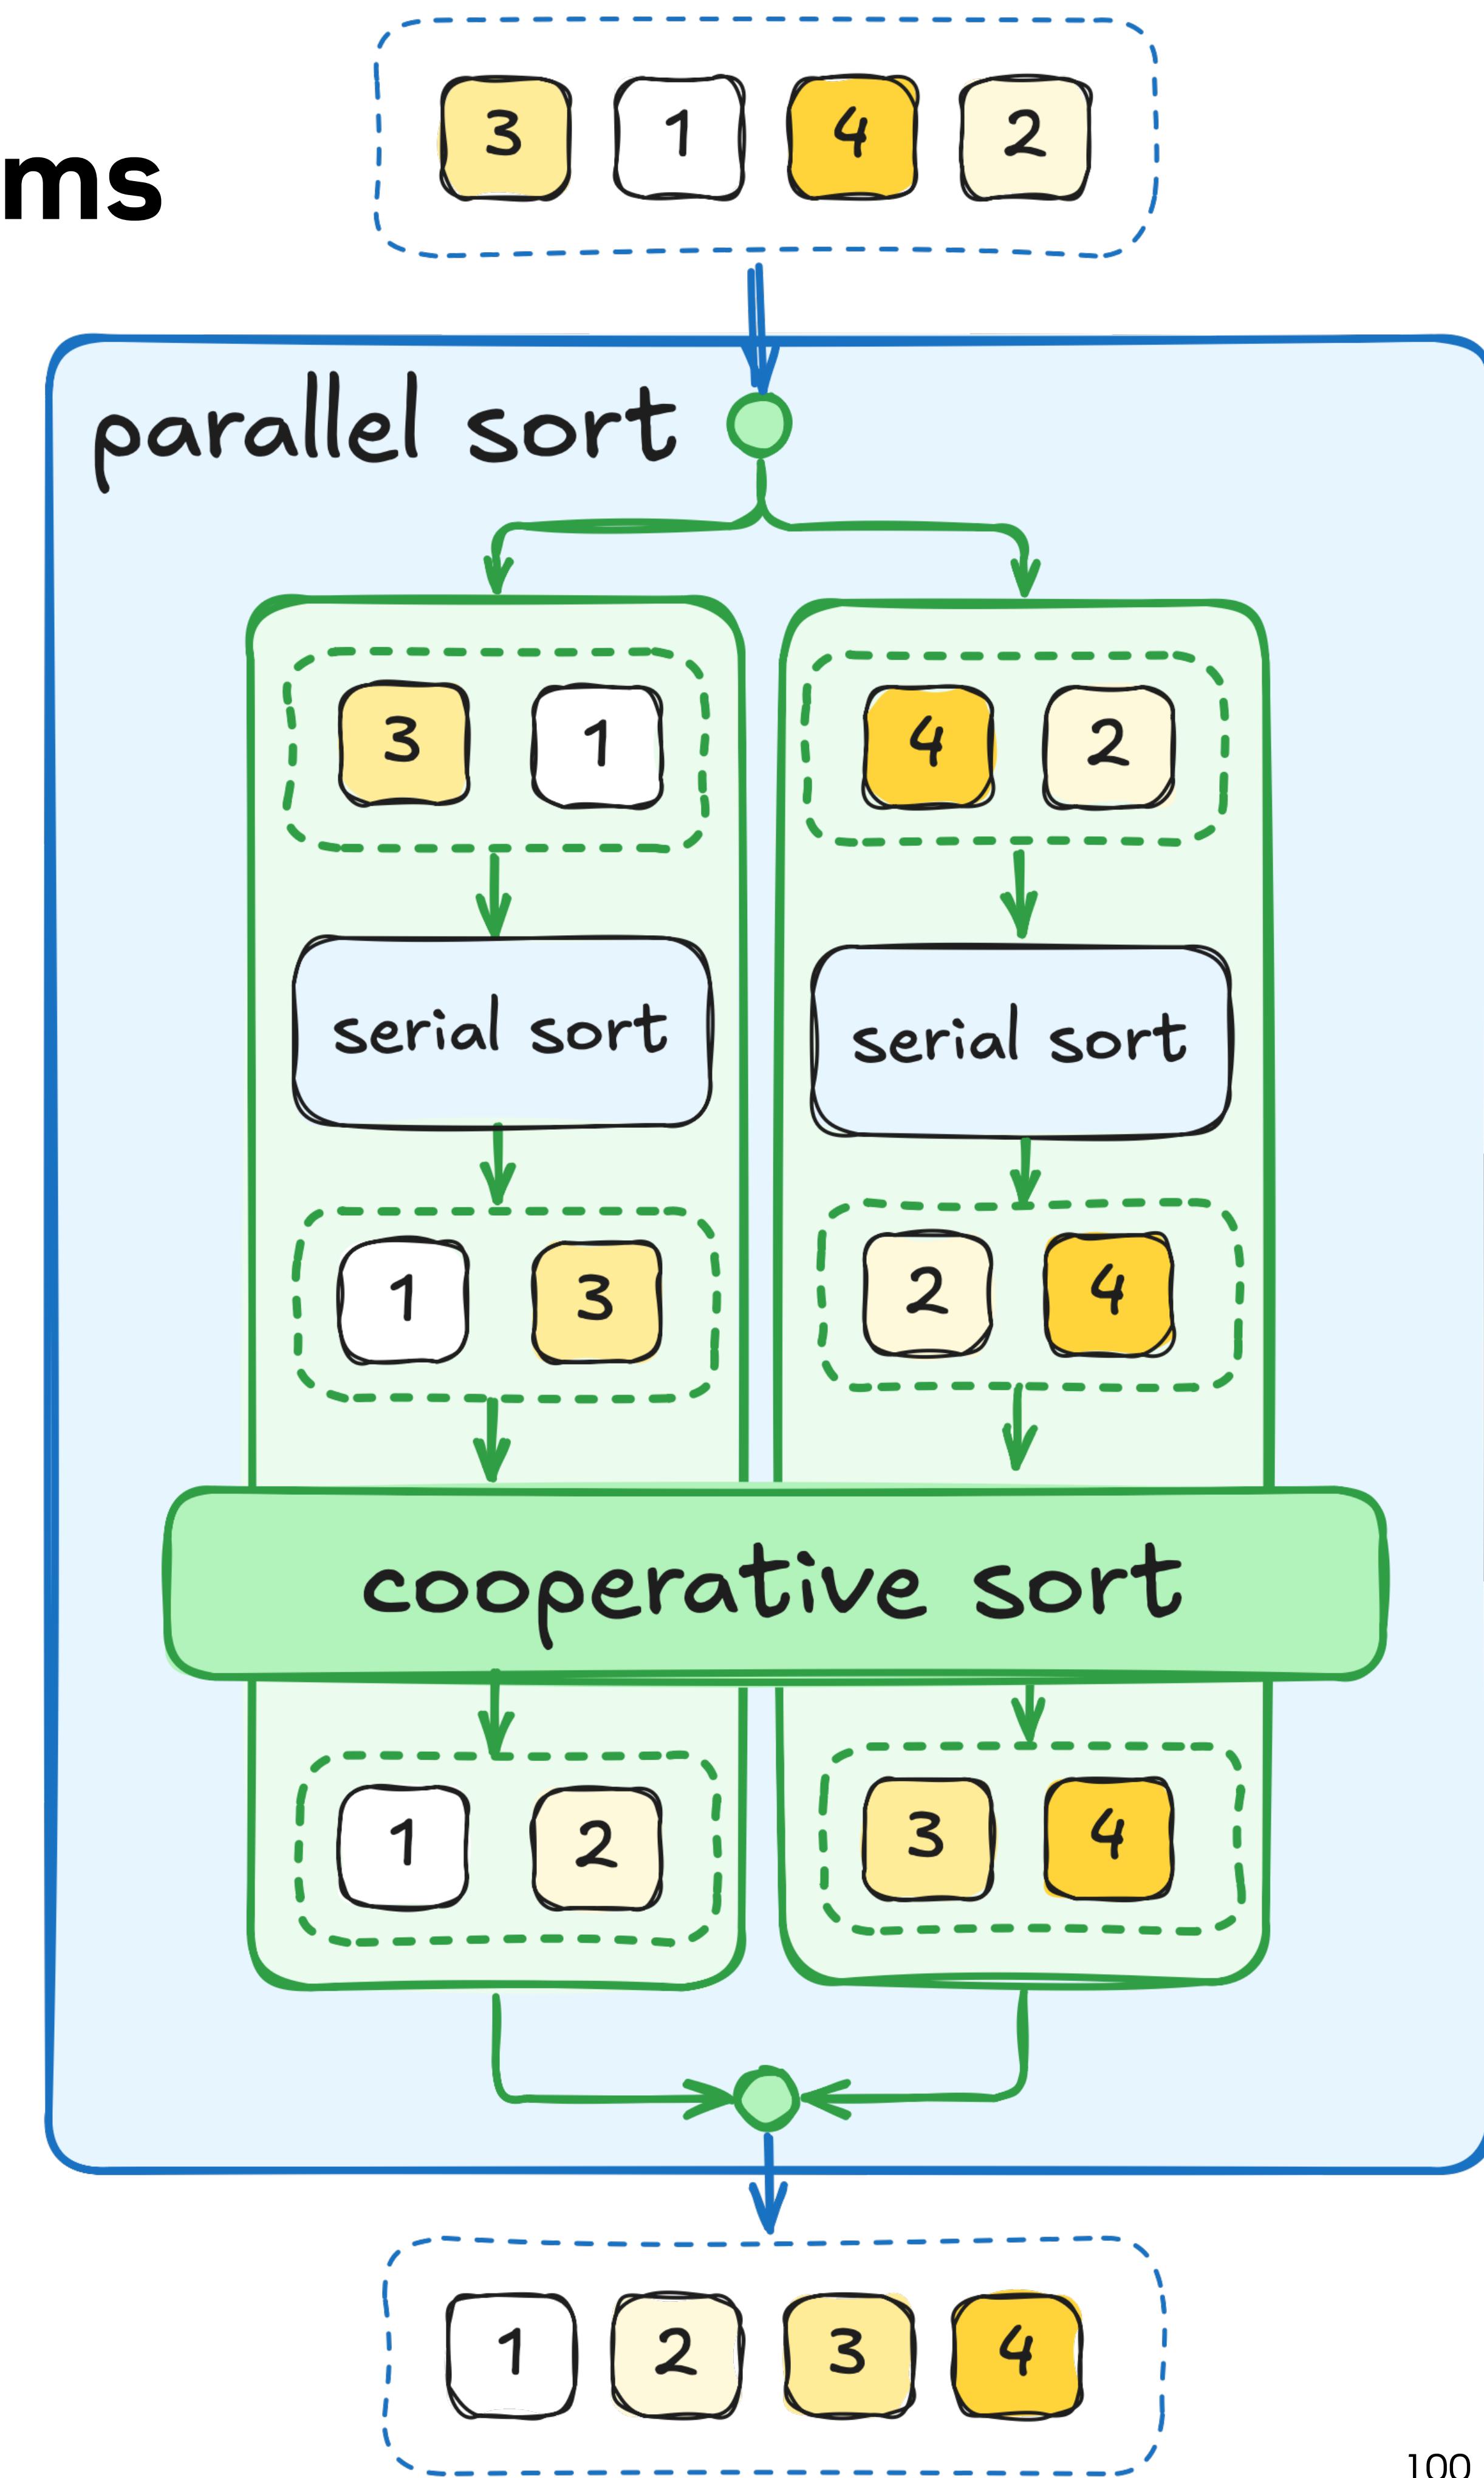
- **One thread** invokes algorithm
- **One thread** executes algorithm

## Cooperative

- **Many threads** invoke algorithm
- **Many threads** execute algorithm

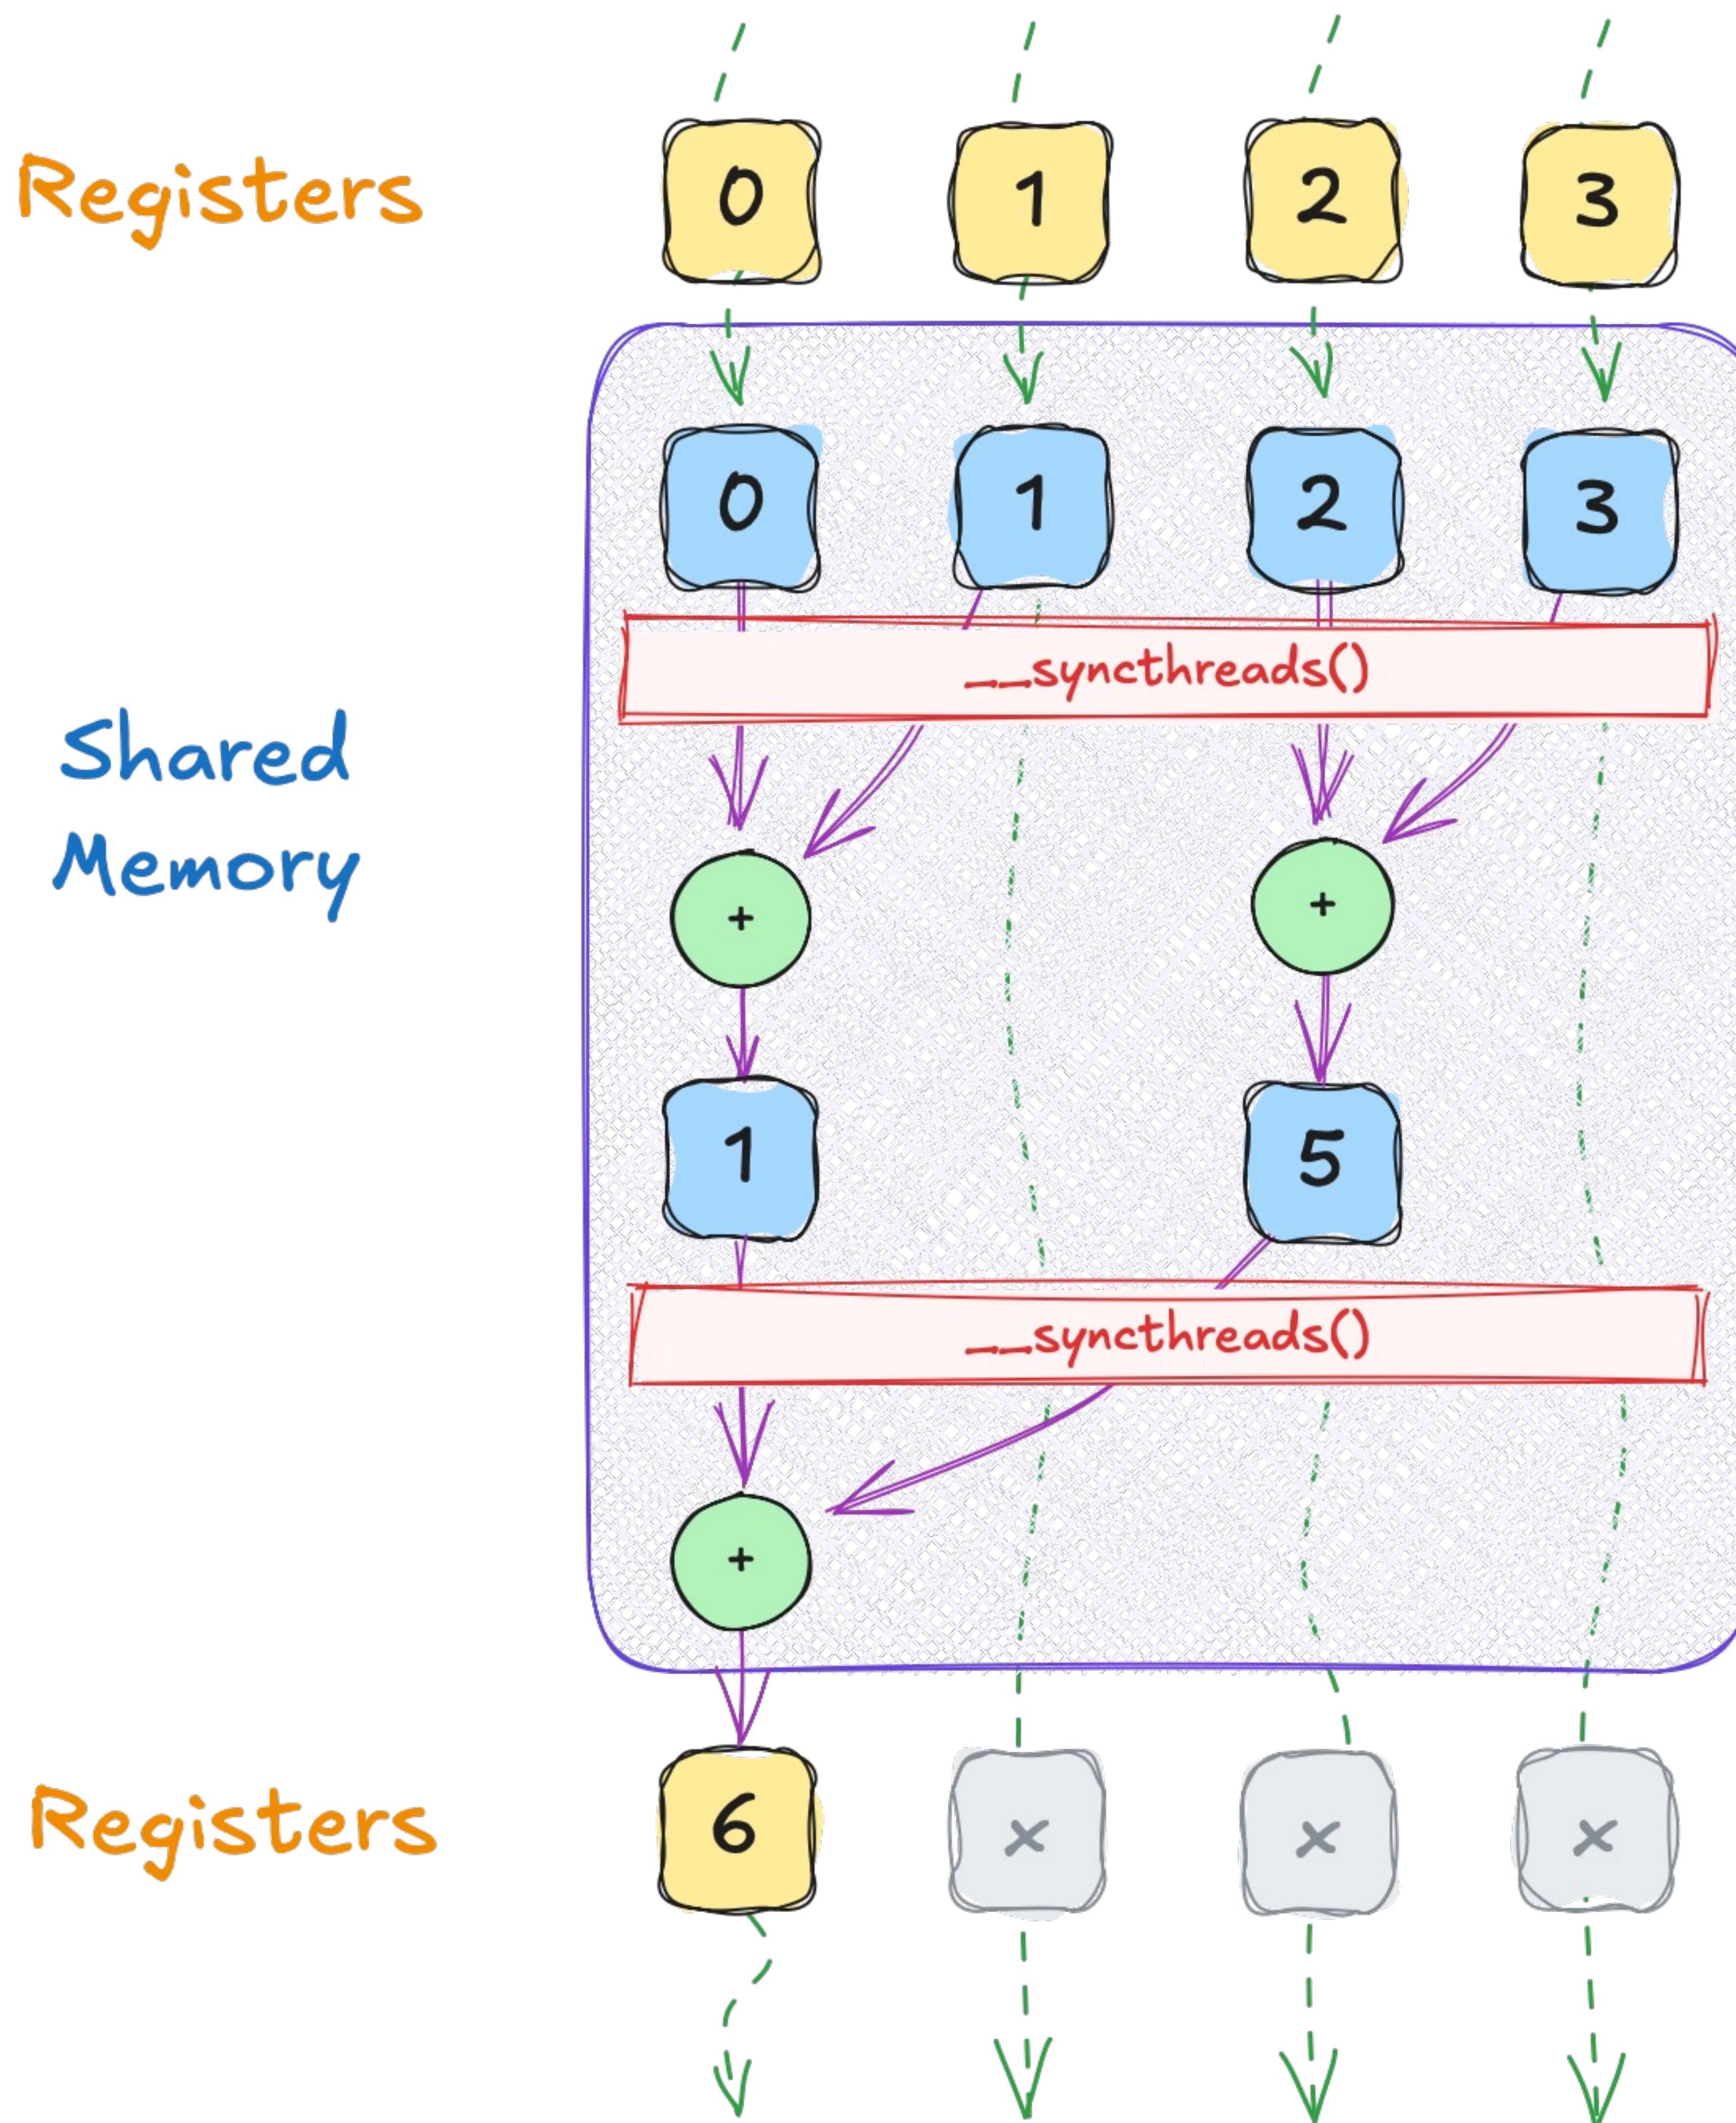
## Parallel

- **One thread** invokes algorithm
- **Many threads** execute algorithm



# Mental Model

- Cooperative algorithms frequently need shared memory for inter-thread communication
- All threads of a block are expected to participate in block-level cooperative algorithm
- Sometimes output is valid only in main thread



# CUB Interface

```
template <typename T, int BlockDimX>
struct cub::BlockReduce
{
    struct TempStorage { ... };

    __device__ BlockReduce(TempStorage& temp_storage) { ... }

    __device__ T Sum(T input) { return ...; }
}
```

- Unlike traditional function-oriented interfaces, CUB exposes its cooperative algorithms as templated structs
- Template parameters used to specialize algorithms for the problem at hand:
  - data type (int, float, etc.)
  - number of threads in the thread block
  - grain size (number of items per thread)
  - etc.

# CUB Interface

```
template <typename T, int BlockDimX>
struct cub::BlockReduce
{
    struct TempStorage { ... };

    __device__ BlockReduce(TempStorage& temp_storage) { ... }

    __device__ T Sum(T input) { return ...; }
}
```

- Nested **TempStorage** type provides type of temporary storage needed by cooperative algorithm for thread communication
- Instance of this type has to be allocated in shared memory
- **TempStorage** allocation must be provided to construct an instance of the algorithm

# CUB Interface

```
template <typename T, int BlockDimX>
struct cub::BlockReduce
{
    struct TempStorage { ... };

    __device__ BlockReduce(TempStorage& temp_storage) { ... }

    __device__ T Sum(T input) { return ...; }
}
```

- Member functions represent different flavors of a given cooperative algorithms

```
__shared__ cub::BlockReduce<int, 4>::TempStorage storage;

int block_sum = cub::BlockReduce<int, 4>(storage).Sum(threadIdx.x);
```

# CUB Interface

```
__shared__ cub::BlockReduce<int, 4>::TempStorage storage;    1. Allocate temporary storage in shared memory  
cub::BlockReduce<int, 4> reducer(storage);                  2. Construct an instance of the algorithm  
  
int block_sum = reducer.Sum(threadIdx.x);                   3. Invoke method of the algorithm instance
```

# CUB Block-Level Algorithms

- `cub::BlockDiscontinuity`
  - `cub::BlockExchange`
  - `cub::BlockHistogram`
  - `cub::BlockLoad`
  - `cub::BlockMergeSort`
  - `cub::BlockRadixSort`
  - `cub::BlockRadixRank`
  - `cub::BlockReduce`
  - `cub::BlockRunLengthDecode`
  - `cub::BlockScan`
  - `cub::BlockShuffle`
  - `cub::BlockStore`
- 
- CUB provides many general-purpose algorithms at thread block level

# CUB Block-Level Algorithms

- `cub::BlockDiscontinuity`
  - `cub::BlockExchange`
  - `cub::BlockHistogram`
  - `cub::BlockLoad`
  - `cub::BlockMergeSort`
  - `cub::BlockRadixSort`
  - `cub::BlockRadixRank`
  - `cub::BlockReduce`
  - `cub::BlockRunLengthDecode`
  - `cub::BlockScan`
  - `cub::BlockShuffle`
  - `cub::BlockStore`
- CUB provides many general-purpose algorithms at thread block level

# CUB Block-Level Algorithms

```
template <
    typename T,
    int BlockDimX,
    int ItemsPerThread,
    int Bins,
    cub::BlockHistogramAlgorithm
        Algorithm = BLOCK_HISTO_SORT>
class cub::BlockHistogram
{
    struct TempStorage
    { ... };

    __device__ BlockHistogram(TempStorage& temp_storage)
    { ... }

    template <typename CounterT>
    __device__ void Histogram(T (&items)[ITEMS_PER_THREAD], CounterT histogram[BINS])      - thread bins
    { ... }                                                               - block histogram (in shared memory)
};

};
```

- `cub::BlockHistogram::Histogram` initializes shared histogram and incorporates bin indices provided by each thread

# Exercise: Use Cooperative Histogram

- Use `cub::BlockHistogram` to compute block histogram

```
__shared__ int block_histogram[10];

int cell = blockIdx.x * blockDim.x + threadIdx.x;
int bins[1] = {static_cast<int>(temperatures[cell] / 10)};

if (threadIdx.x < 10)
{
    block_histogram[threadIdx.x] = 0;
}
__syncthreads();

cuda::atomic_ref<int, cuda::thread_scope_block>
    block_ref(block_histogram[bin]);
block_ref.fetch_add(1);
__syncthreads();
...
```

# Exercise: Use Cooperative Histogram

## Solution

```
__shared__ int block_histogram[10];

int cell = blockIdx.x * blockDim.x + threadIdx.x;
int bins[1] = {static_cast<int>(temperatures[cell] / 10)};

using histogram_t = cub::BlockHistogram<int, block_size, 1, 10>;
__shared__ typename histogram_t::TempStorage temp_storage;

histogram_t(temp_storage).Histogram(bins, block_histogram);

__syncthreads();
...
```

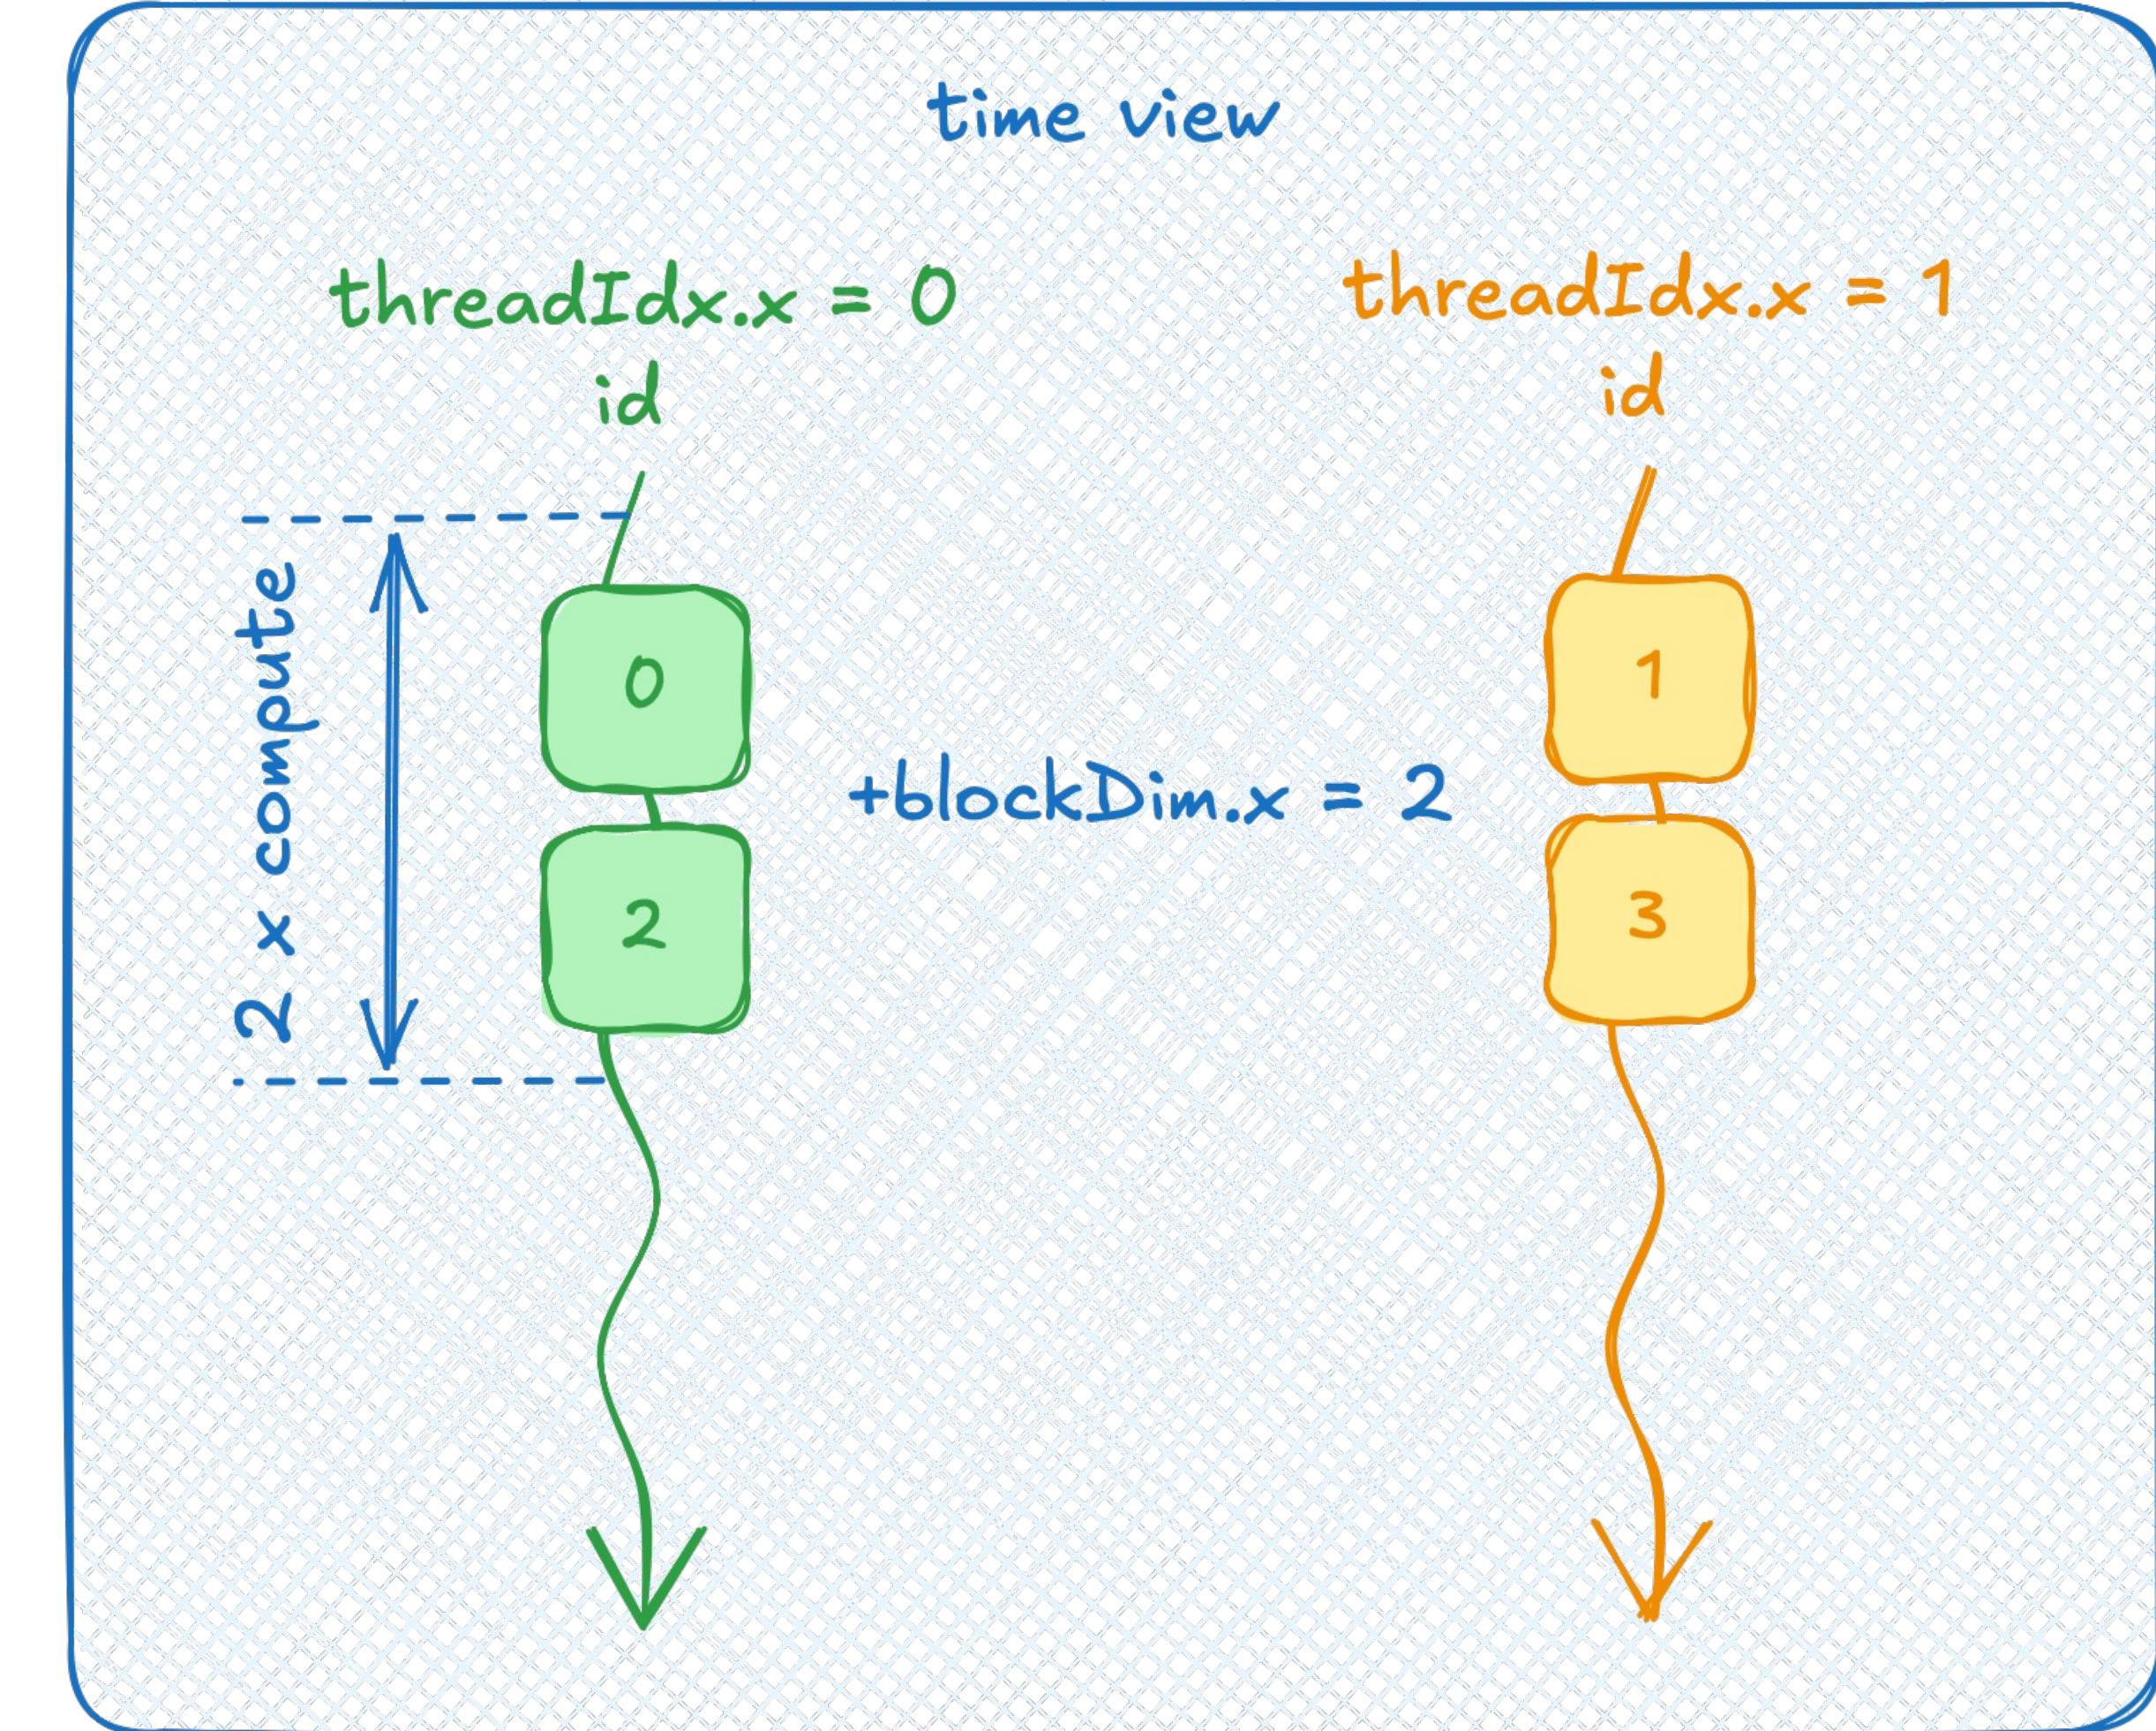
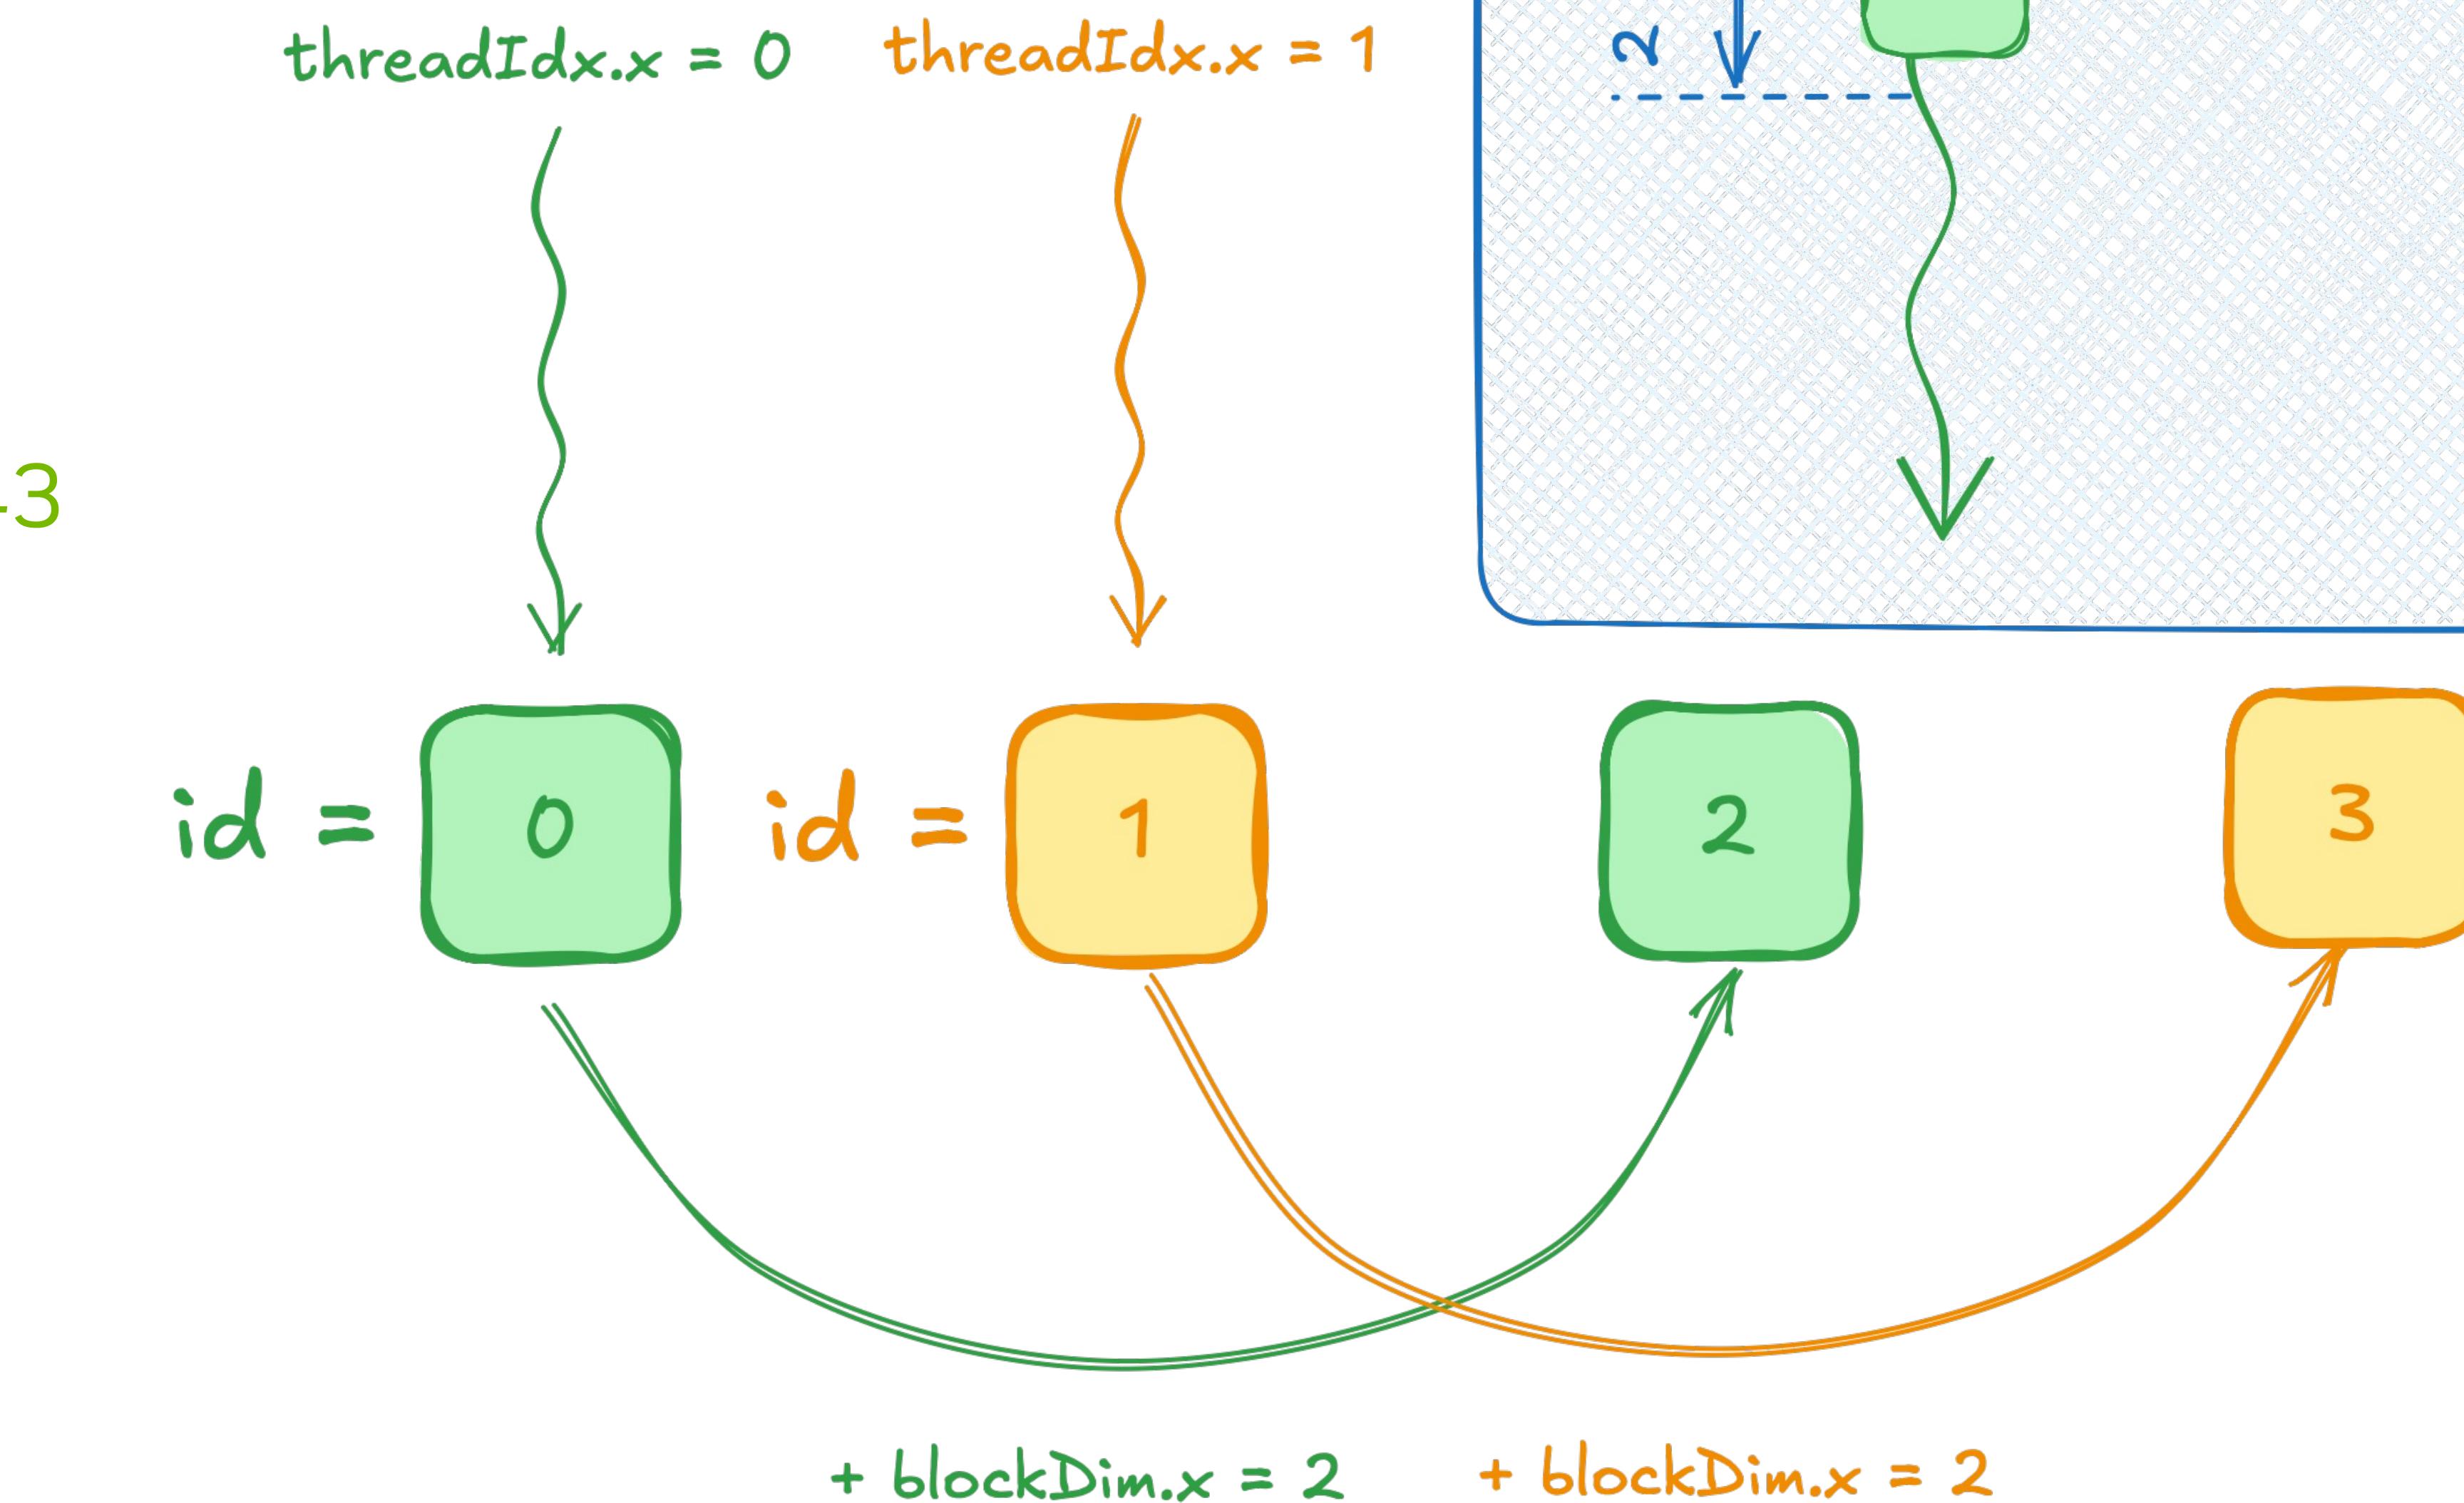
# Takeaways

Write custom CUDA kernels only when:

- Your use case is not covered by existing accelerated libraries

10 compute steps take:

- cub: 0.000043
- 1 thread: 4.15
- 2 threads: 2.6
- 256 threads: 0.037
- 5'120'000 threads: 0.0003



# Takeaways

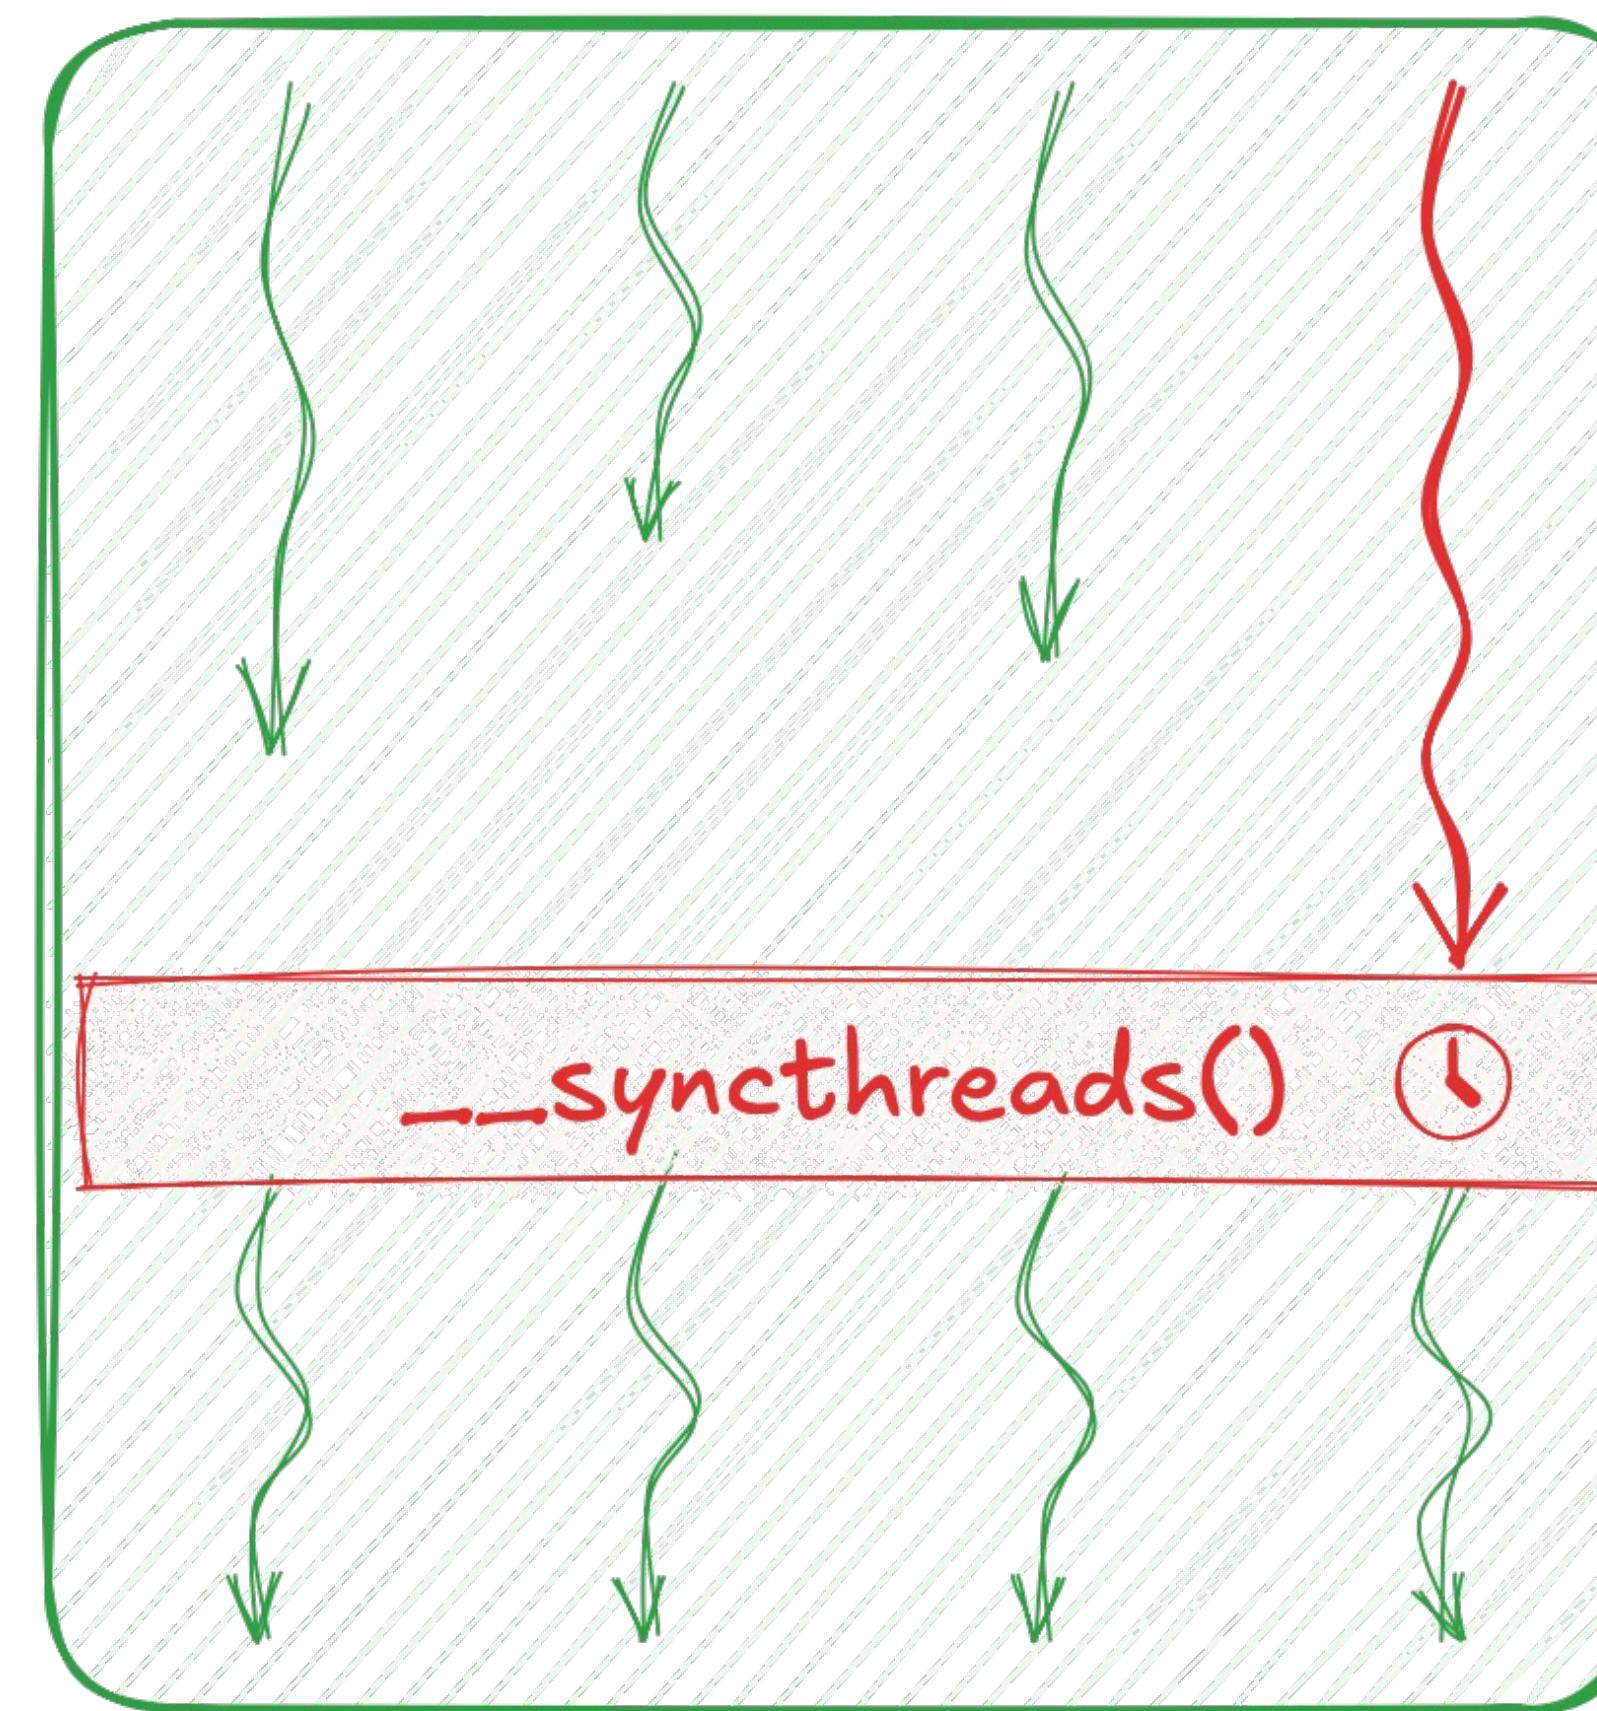
Write custom CUDA kernels only when:

- Your use case is not covered by existing accelerated libraries

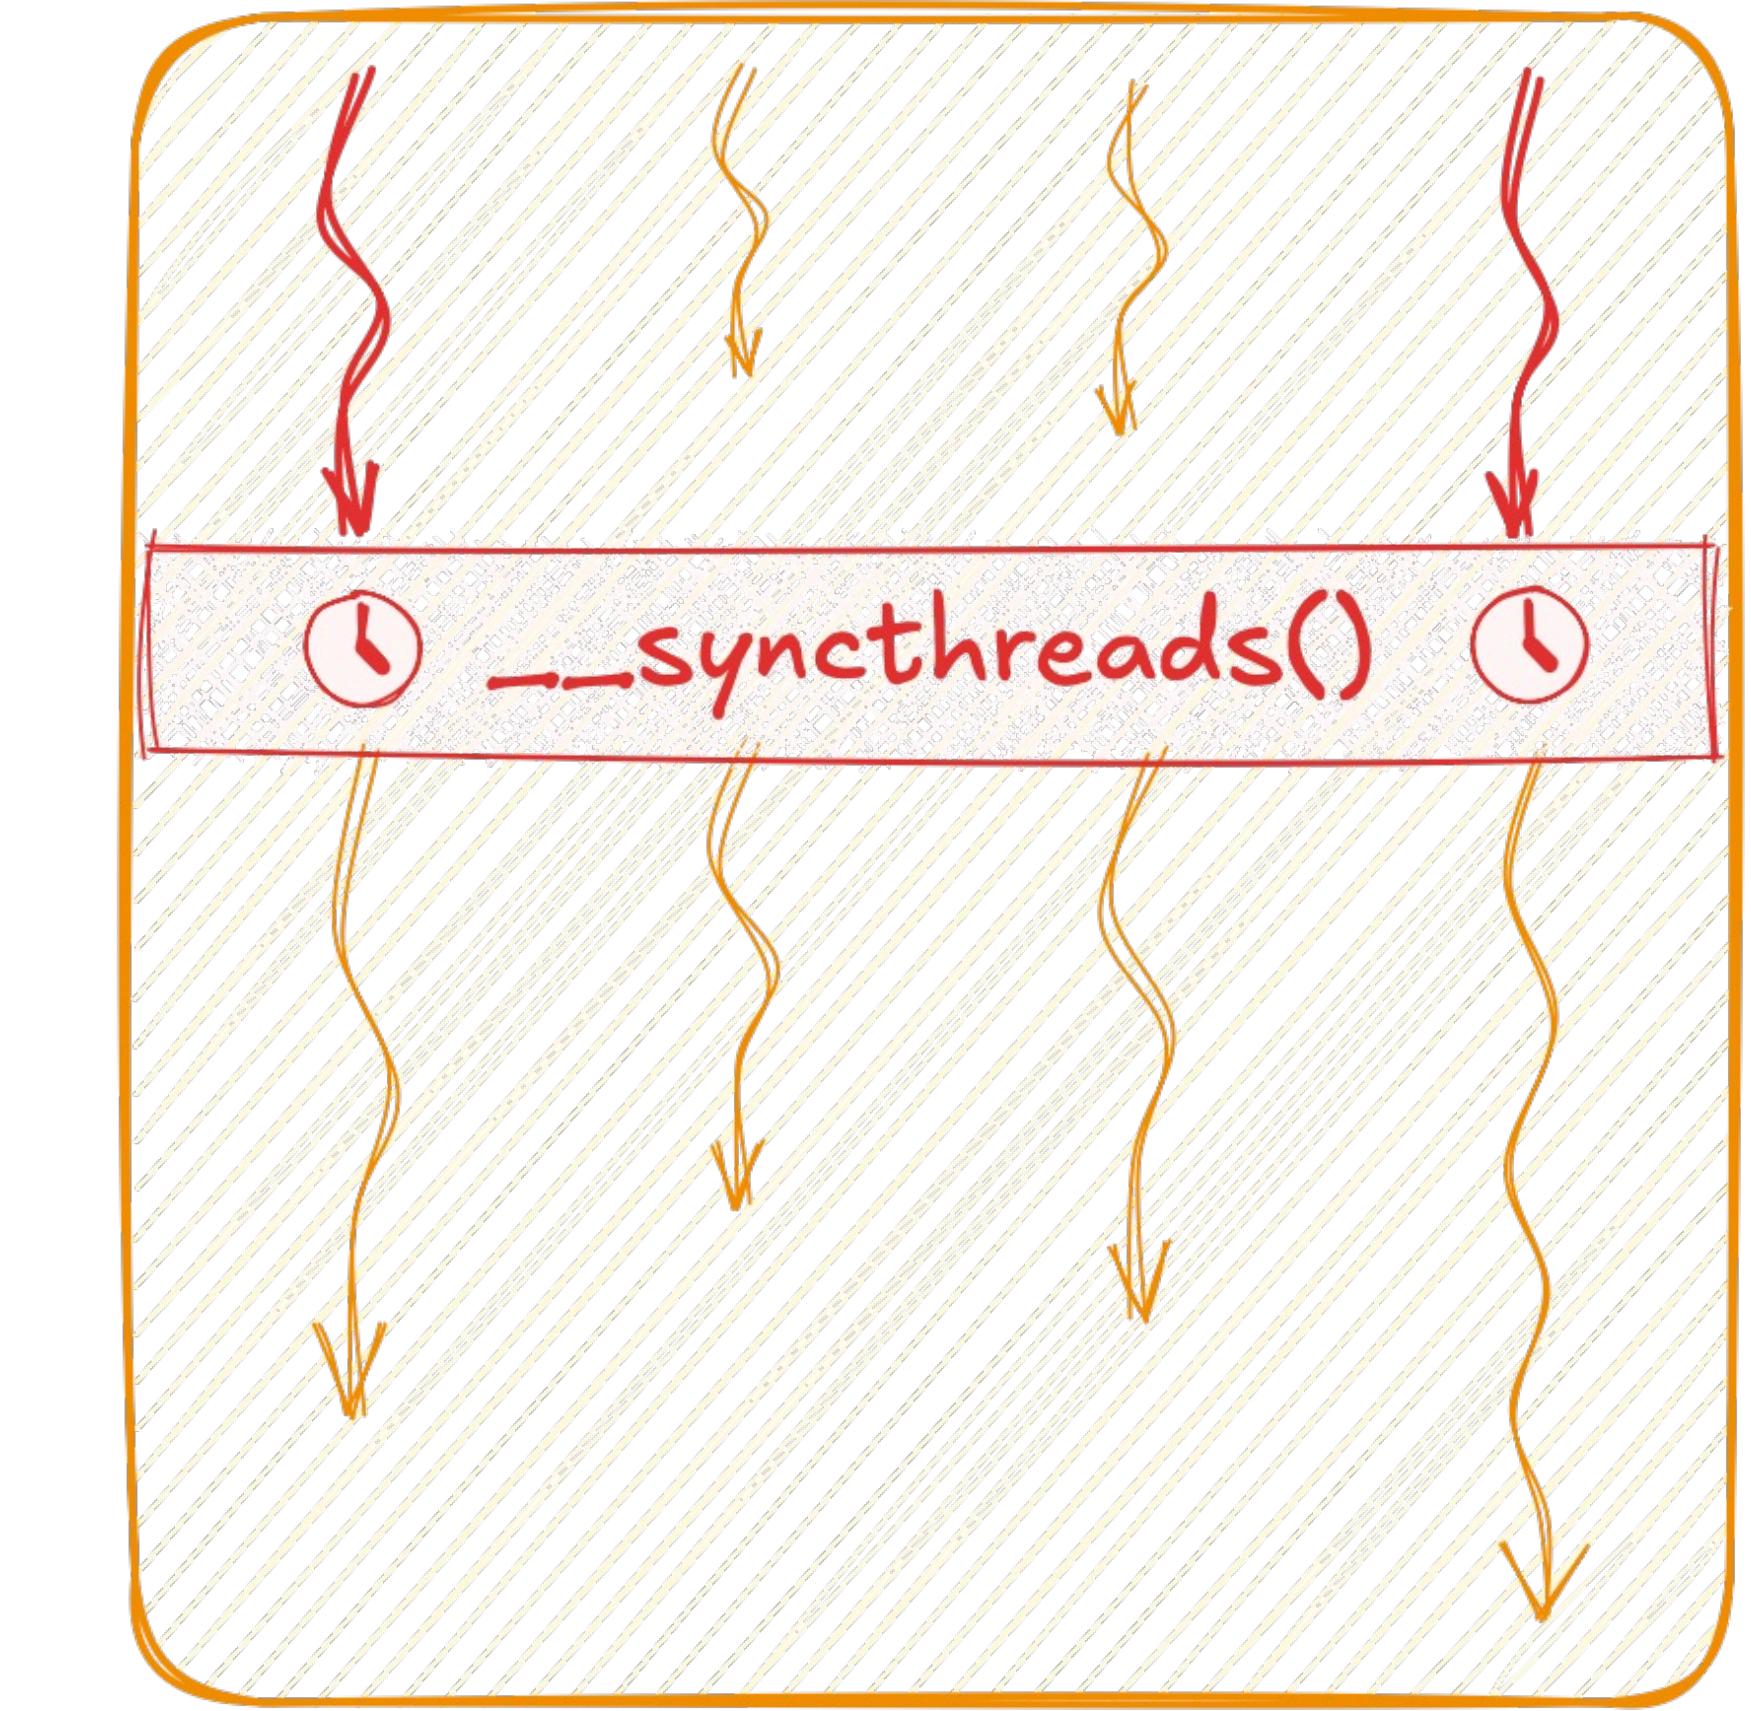
Use `__syncthreads`:

- To avoid data races within a thread block

`blockIdx.x = 0`



`blockIdx.x = 1`



# Takeaways

Write custom CUDA kernels only when:

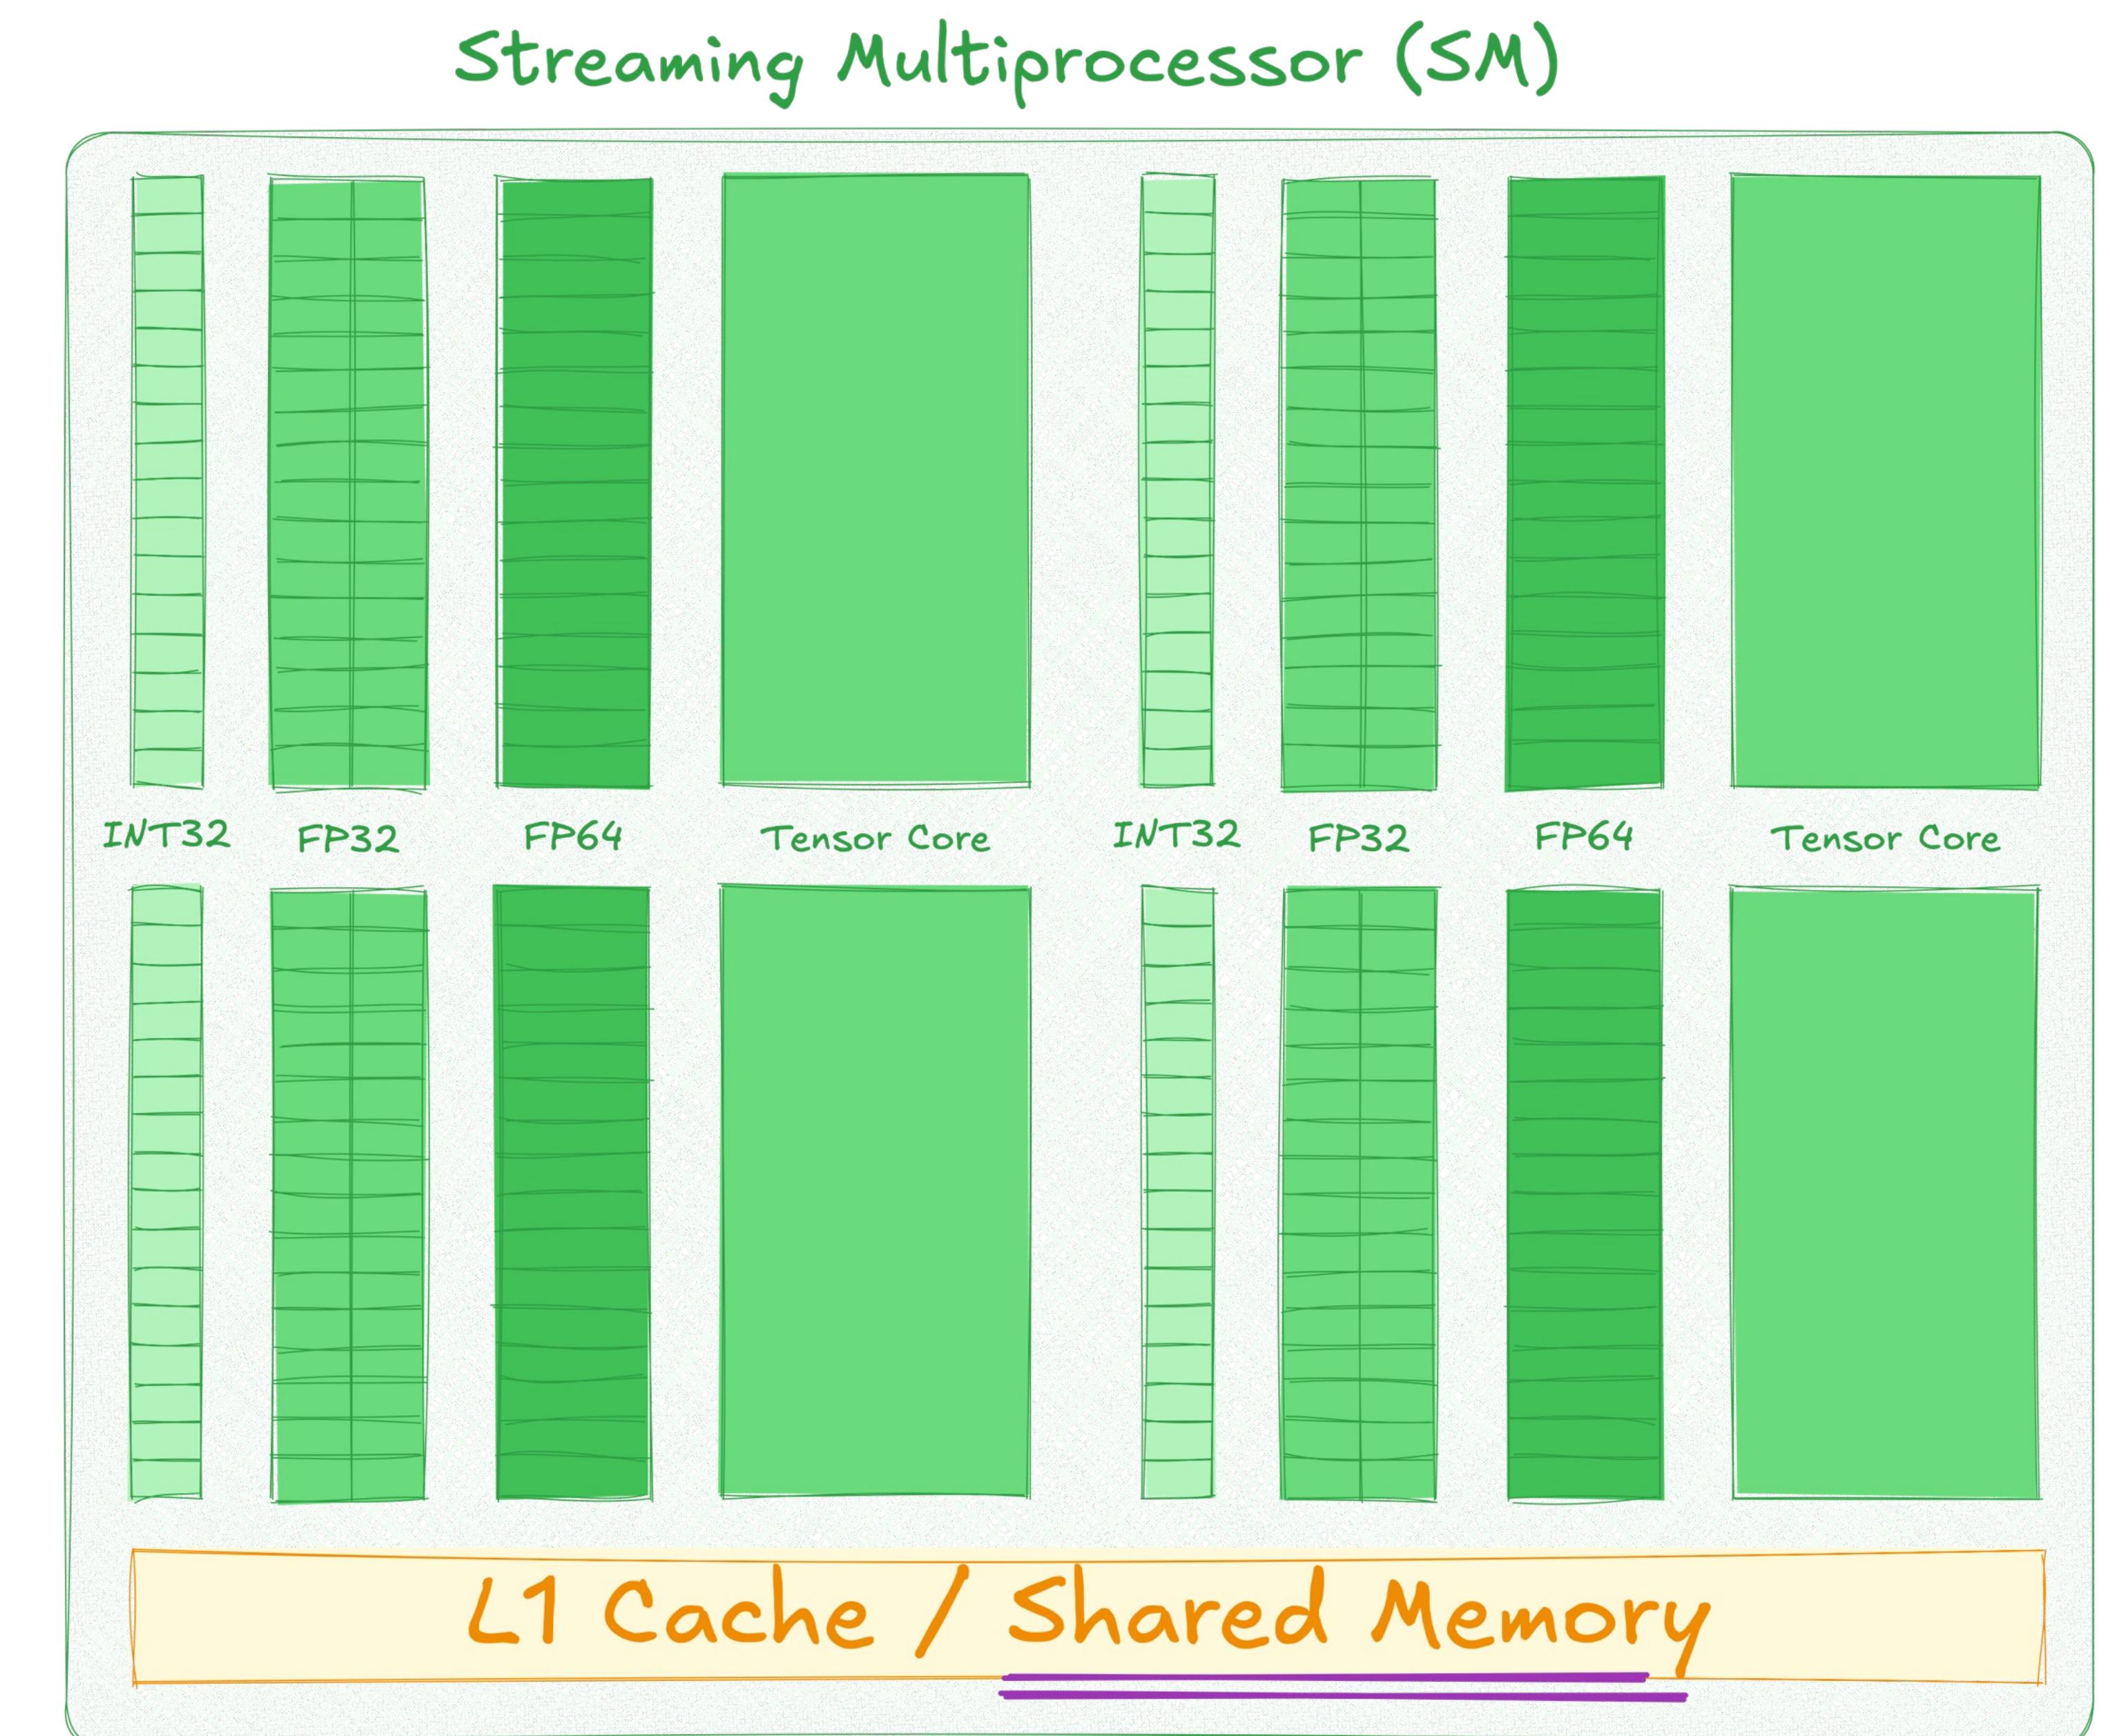
- Your use case is not covered by existing accelerated libraries

Use `__syncthreads`:

- To avoid data races within a thread block

Use `__shared__`:

- When you need a temporary storage inside a thread block
- When you need to exchange data between threads of a thread block
- When your data is not efficiently cached by L1



# Takeaways

Write custom CUDA kernels only when:

- Your use case is not covered by existing accelerated libraries

Use `__syncthreads`:

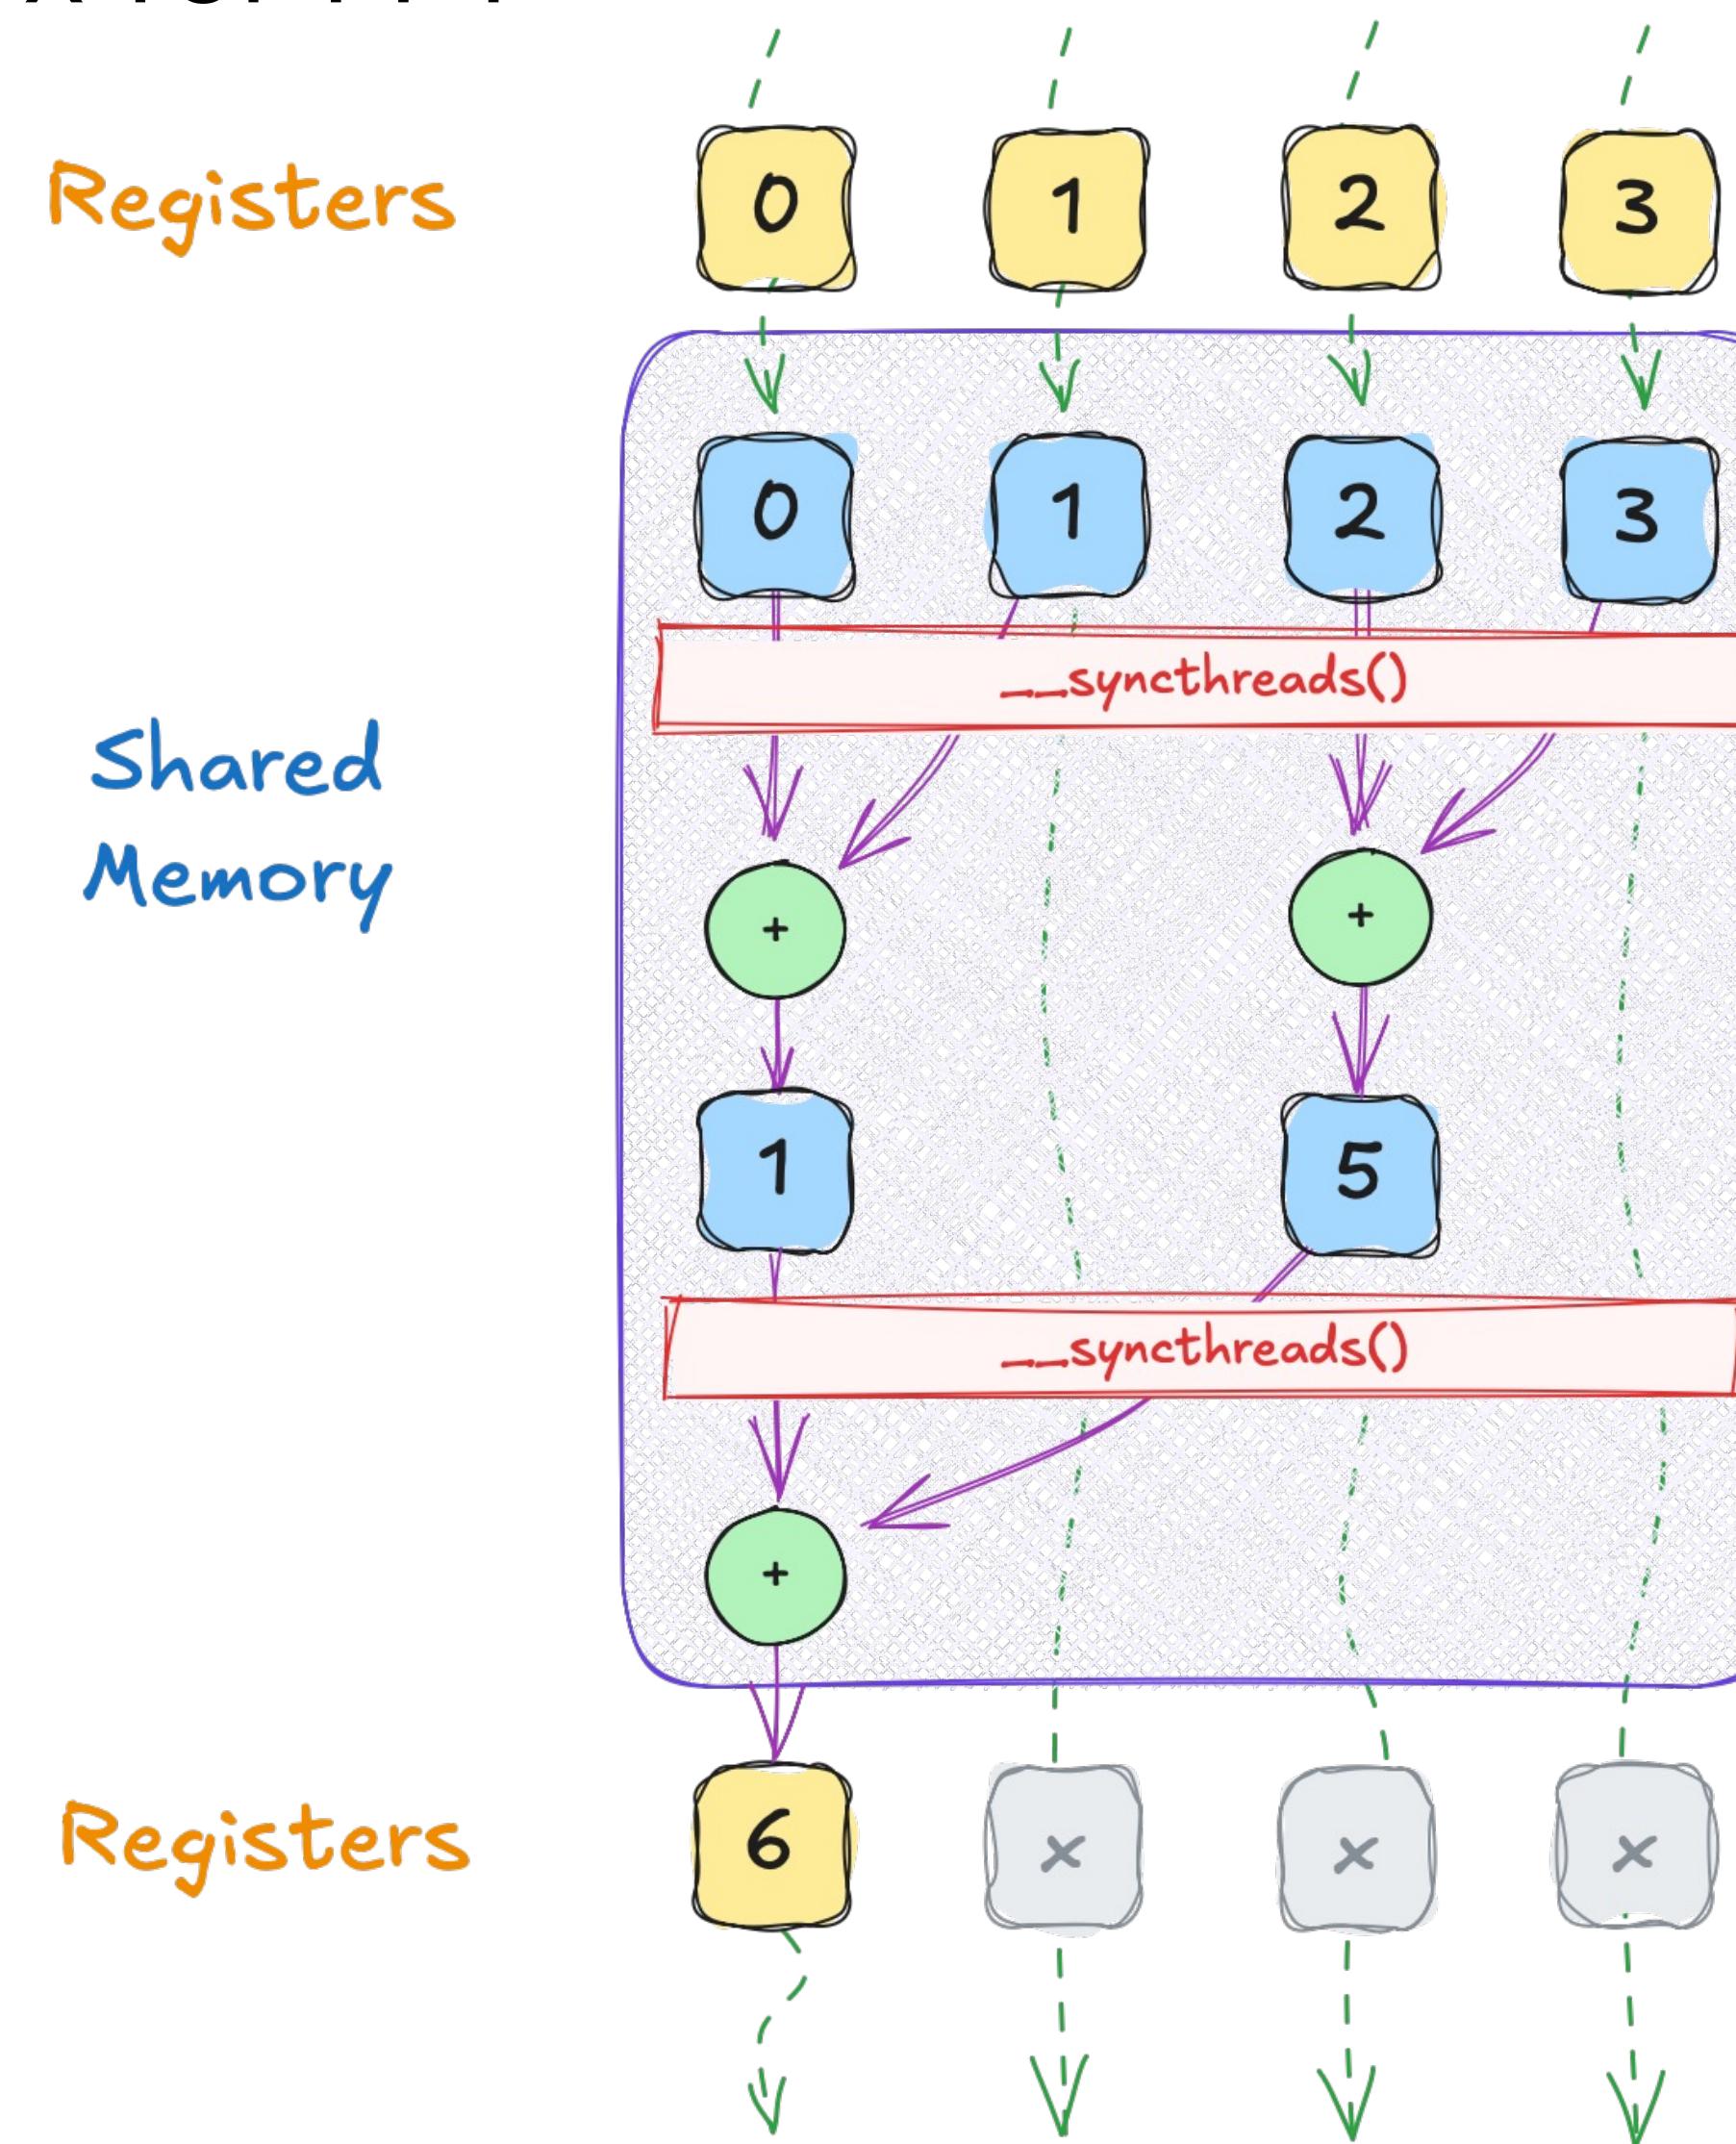
- To avoid data races within a thread block

Use `__shared__`:

- When you need a temporary storage inside a thread block
- When you need to exchange data between threads of a thread block
- When your data is not efficiently cached by L1

Use device-side libraries:

- CUB for block, and warp-scope cooperative general-purpose algorithms
- libcudacxx for vocabulary types
- cuBLASDx for linear algebra
- cuFFTDx for FFT
- etc.



# Takeaways

Write custom CUDA kernels only when:

- Your use case is not covered by existing accelerated libraries

Use `__syncthreads`:

- To avoid data races within a thread block

Use `__shared__`:

- When you need a temporary storage inside a thread block
- When you need to exchange data between threads of a thread block
- When your data is not efficiently cached by L1

Use device-side libraries:

- CUB for block, and warp-scope cooperative general-purpose algorithms
- libcudacxx for vocabulary types
- cuBLASDx for linear algebra
- cuFFTDx for FFT
- etc.

Use developer tools:

- compute-sanitizer for CUDA kernels correctness checking
- cuda-gdb for CUDA kernels debugging
- Use NVIDIA Nsight Systems to identify bottlenecks in your application
- NVIDIA Nsight Compute for in-depth kernel profiling

